



Data Structures: Static Hashing

Won Kim

(Lecture by Youngmin Oh)

Spring 2022



Static Hashing



Hashing

- Hash: “chop (meat or potatoes) into small pieces”
- Static Hashing (will learn today)
- Dynamic Hashing (will learn later)

Hash Brown



Corned Beef Hash





Static Hashing

- Technique for storing and searching a key in an array (hash table) by mapping the key to an index into the array
 - Each entry in the table is called a bucket.
- $O(1)$ **average** performance for search, insertion, and deletion of a key



Static Hashing: Example

(using the first letter of the key)

index bucket

0	Korea
1	Japan
2	Qatar
3	Vietnam
4	China
5	Taiwan
6	Singapore
7	France



Static Hashing: Example (cont'd)

0	Korea	
1	Japan	
2	Qatar	
3	Vietnam	
4	China	
5	Taiwan	Thailand
6	Singapore	Spain
7	France	

collision



Static Hashing: Example (cont'd)

0	Korea	
1	Japan	
2	Qatar	
3	Vietnam	
4	China	
5	Taiwan	Thailand
6	Singapore	Spain
7	France	

overflow

Sri Lanka



Hash Table

- n indexes
- k buckets per index
 - Total $n*k$ Keys Stored
- Total potential keys $\gg n$
- Collision
 - > 1 key mapped to the same index
- Overflow
 - $> k$ keys mapped to the same index



Ultimate Objectives of Hashing

- (1) minimize collision
- (2) minimize overflow
- (3) minimize hash table size
 - Keep n (indexes) and/or k (buckets) reasonably small



Two Key Elements in Hashing

- Selecting a hash function
 - Should have a **uniform distribution** of keys across all indexes
 - Should be easy to compute
- Collision resolution
 - What to do with the keys hashed to the same index
 - (Collision is unavoidable)



Hash Functions

- If the key is a string, convert it to a natural number.
 - e.g., “EACH” = (ASCII) 69 65 67 72
 - e.g., “EACH” = (alphabet position) 5 1 3 8
- Many possible functions.
 - e.g., Pick the first character of a string, and map it to a number. The number may be used as index into the hash table.



4 Commonly Used Hash Functions (1/4)

- Modulo (Division) Function
- Most commonly used
 - $H(\text{key}) = \text{key} \bmod \text{hash_table_size}$
 - e.g., $375 \bmod 101 = 72$
 - All keys are guaranteed to fit in the hash table.
 - It is best to use a prime number for `hash_table_size`.



4 Commonly Used Hash Functions (2/4)

- Digit Folding (“**hashing**”)
 - Add combinations of the elements of the key
 - e.g., from the key = 9010302051218, compute $(9+0+1+0+3+0+2+0+5+1+2+1+8) = 32$.
 - $\text{hash_table_size} = 9 \times 13 = 117$
 - (* Why is the hash_table_size NOT 10×13 ?
(answer: We are summing the highest number for each digit)
 - e.g., from the key = 9010302051218, compute $(90+10+30+20+51+21+8) = 230$
 - $\text{hash_table_size} = 99 \times 6 + 9 = 603$



4 Commonly Used Hash Functions (3/4)

- Digital Selection

- Select some elements of the key
 - e.g., from the key = 9010302051218, select only the even elements (000011).
- Need to be careful to prevent heavy collision
 - e.g., (do not) select only the first two elements of the key = 9010302051218.



4 Commonly Used Hash Functions (4/4)

- Mid-Square Function
 - Square the key, convert the result to ASCII equivalent, and select k bits from the middle of the square.
 - (e.g.) If the key is 11, its square is 121. The ASCII equivalent for 121 is 111001, and some bits from the middle may be selected as index into a hash table.
 - The number of bits to select depends on the hash table size. If k bits are selected, the range of the values is 2^k .



Collision Resolution

- In general, no hash function can prevent collision.
- Colliding keys must be stored somehow.
- Two approaches
 - Closed addressing
 - Stays with the computed index
 - Open addressing
 - Invades any available bucket for any index
- Tradeoff
 - Hash table size (memory space)
 - Time to store and search the overflow keys



Closed Addressing (1/2)

- “Full” 2-D array for the hash table
 - For each index, enough buckets for all colliding keys.
 - Fast access, no overflow
 - Potential big waste of space
 - Not practical



Hash Table as a “Full” 2-D Array

0	Korea		
1	Japan		
2	Qatar		
3	Vietnam		
4	China		
5	Taiwan	Thailand	
6	Singapore	Spain	Sri Lanka
7	France		



Closed Addressing (2/2)

- Overflow Chaining for each index
 - without collision buckets
 - with collision buckets for some colliding keys

Overflow Chaining: without Collision Buckets

0	Korea	
1	Japan	
2	Qatar	
3	Vietnam	
4	China	
5	Taiwan	
6	Singapore	
7	France	


overflow
chaining



Overflow Chaining: with Collision Buckets

0	Korea		
1	Japan		
2	Qatar		
3	Vietnam		
4	China		
5	Taiwan	Thailand	
6	Singapore	Spain	
7	France		

overflow
chaining





Search

- Use the same hash function used for insert.
- Search for the key in the indexed bucket.
- If there is no match in the bucket, continue along the overflow chain.



Search

- Use the same hash function used for insert.
- Search for the key in the indexed bucket.
- If no match, continue along the overflow chain.



Exercise: Hash Table & Overflow Chaining

hash table size: 7

no collision buckets

hash function: $\text{key} \bmod 7$

set of keys

7

9

44

13

140

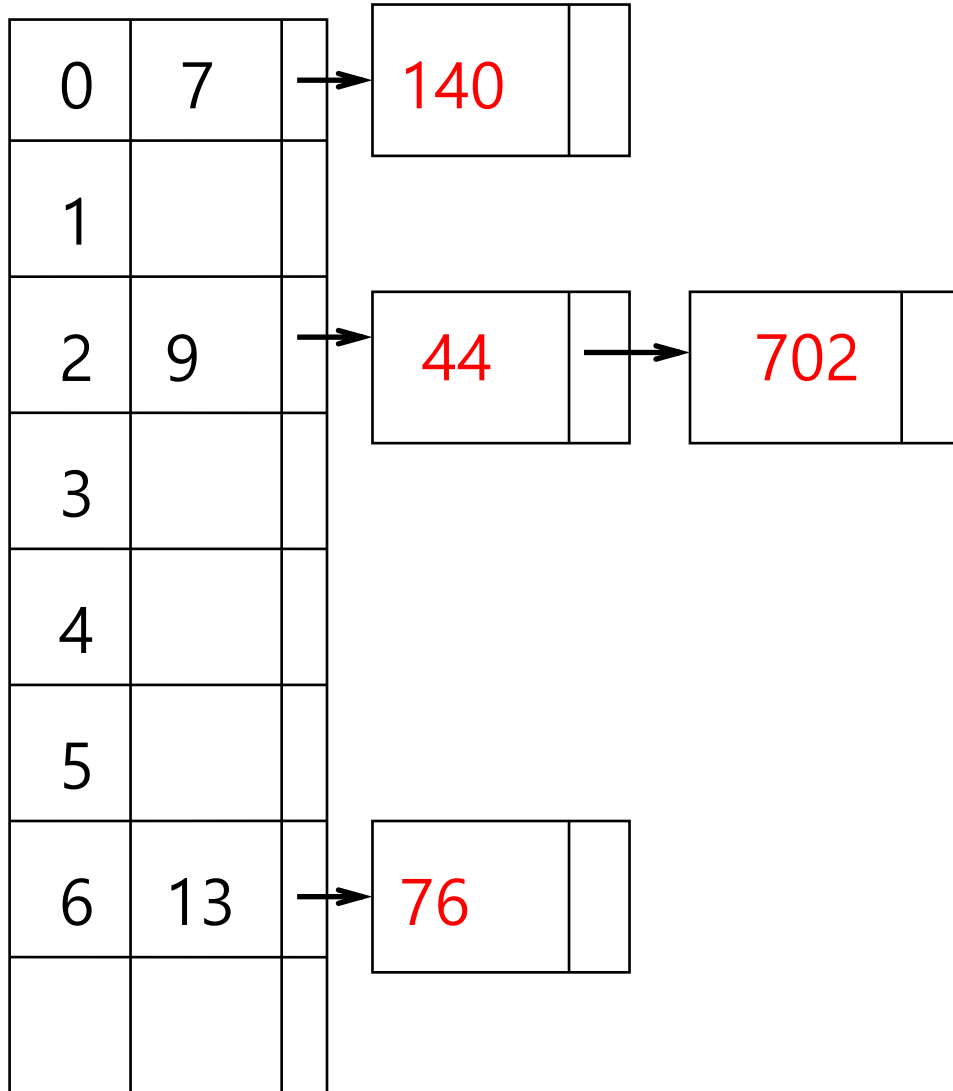
702

76



Solution

index bucket





Open Addressing

- Linear Probing
 - Quadratic Probing
 - Double Hashing
-
- Difference is the way to “determine where to look, if the indexed bucket is occupied”.



Linear Probing

- If collision occurs, sequentially (one bucket at a time) search for an empty bucket and store the “homeless” key there.



Linear Probing: Example

0	Korea
1	
2	Qatar
3	
4	China
5	Taiwan
6	Singapore
7	France

insert
Spain

Spain
↓



Linear Probing: Example (cont'd)

0	Korea
1	Spain
2	Qatar
3	
4	China
5	Taiwan
6	Singapore
7	France



insert
Spain



Linear Probing: Example (cont'd)

0	Korea
1	Spain
2	Qatar
3	
4	China
5	Taiwan
6	Singapore
7	France

insert
Sri Lanka

Sri Lanka



Linear Probing: Example (cont'd)

0	Korea
1	Spain
2	Qatar
3	Sri Lanka
4	China
5	Taiwan
6	Singapore
7	France



insert
Sri Lanka



Search

- Use the same hash function used for insert.
- If there is no match, use linear probing to find a match.



Exercise: Linear Probing

hash table size: 11

no collision buckets

hash function: $\text{key} \bmod 11$

set of keys

7

9

44

13

18

110





Solution

index	bucket
0	44
1	110
2	13
3	
4	
5	
6	
7	7
8	18
9	9
10	



Linear Probing: A Small Problem

- A “homeless” key may be deleted.
- Next search for another “homeless” key may prematurely stop there.



Linear Probing: Small Problem Example

0	Korea
1	Spain
2	Qatar
3	Sri Lanka
4	China
5	Taiwan
6	Singapore
7	France



delete
Spain



Linear Probing: Small Problem Example (cont'd)

0	Korea
1	
2	Qatar
3	Sri Lanka
4	China
5	Taiwan
6	Singapore
7	France



delete
Spain



Linear Probing: Small Problem Example (cont'd)

0	Korea
1	
2	Qatar
3	Sri Lanka
4	China
5	Taiwan
6	Singapore
7	France



search
Sri Lanka



Linear Probing: Small Problem Example (cont'd)

0	Korea
1	
2	Qatar
3	Sri Lanka
4	China
5	Taiwan
6	Singapore
7	France

↓
search ends here

search
Sri Lanka



Linear Probing: Solution to the Small Problem

- 3 states for each index
 - Occupied, empty, deleted
- If “occupied”, continue probing.
- If “empty”, insert key.
- If “deleted”,
 - (If the operation is to insert a key), insert.
 - (If the operation is to search for a key), continue probing.



Linear Probing: Big Problem

- Formation of “Cluster”s
 - “Homeless” keys occupy **somebody else**’s buckets, causing **somebody else** to become “homeless” in a continuous chain.
 - Search time complexity can approach $O(n)$, not $O(1)$.



Quadratic Probing

- Linear probing modified
 - Look for an empty bucket k^2 ($k = 1, 2, 3, \dots$) positions from the computed index.
 - Leaves room for legitimate keys.
- Partially solves the cluster formation problem of linear probing.



Quadratic Probing: Example

0	Korea
1	Qatar
2	
3	
4	China
5	Taiwan
6	Singapore
7	France

insert
Spain

Spain
↓



Quadratic Probing: Example (cont'd)

0	Korea
1	Qatar
2	
3	
4	China
5	Taiwan
6	Singapore
7	France

insert
Spain

Spain



$$k^2 = 1$$



Quadratic Probing: Example (cont'd)

0	Korea
1	Qatar
2	Spain
3	
4	China
5	Taiwan
6	Singapore
7	France

↓ $k^2=4$

insert
Spain



Exercise: Quadratic Probing

hash table size: 11

no collision buckets

hash function: $\text{key} \bmod 11$

set of keys

7

9

44

13

18

110





Solution

index	bucket
0	44
1	110
2	13
3	
4	
5	
6	
7	7
8	18
9	9
10	



Double Hashing

- Use of two separate hash functions
 - one to find the index
 - another to find the “index interval” for inserting a “homeless” key.
 - The same index interval is successively used.
 - e.g., $h1: \text{key} \% 13$,
 $h2: 1 + \text{key} \% 11$
 - The second hash function should
 - not compute to zero
 - not be the same as the first hash function.
 - major difference from linear and quadratic probing
 - Bucket index depends on the key value.



Double Hashing: Example

0	Korea
1	Qatar
2	
3	China
4	
5	Taiwan
6	Singapore
7	France

insert
Spain

Spain
↓



Double Hashing: Example (cont'd)

0	Korea
1	Qatar
2	
3	China
4	
5	Taiwan
6	Singapore
7	France

insert
Spain

Spain
↓

$h=3$



Double Hashing: Example (cont'd)

0	Korea
1	Qatar
2	
3	China
4	Spain
5	Taiwan
6	Singapore
7	France



occupied

$h=3$

insert
Spain



Exercise: Double Hashing

hash table size: 11

no collision buckets

hash function 1: $\text{key} \bmod 11$

hash function 2: $1 + \text{key} \bmod 7$

set of keys

7

9

44

13

18

110





Solution

index	bucket
0	44
1	18
2	13
3	
4	
5	
6	110
7	7
8	
9	9
10	



Note

- The collision scenarios shown in examples in this class are the result of using a very tiny array for a hash table just to convey the concepts to students.
- In practice, any reasonable software engineer will create a much larger array, and select hash functions carefully, and collisions do not happen too often.



Performance

- best case: $O(1)$
- avg case: $O(1)$
 - best possible search algorithm
- worst case: $O(n)$
 - (* as hash table gets nearly full *)



Open Addressing: Measures

- Loading density
 - $\alpha = \# \text{ of buckets occupied} / \text{total} \# \text{ of buckets in the hash table}$
- Avg. # of key comparisons
 - $(2 - \alpha) / (2 - 2\alpha)$
- Examples
 - for $\alpha = 0.2 \rightarrow (2 - 0.2) / (2 - 2 \times 0.2) = 1.125$
(almost 1)
 - for $\alpha = 0.99 \rightarrow (2 - 0.99) / (2 - 2 \times 0.99) = 50.5$



Problems with Hashing

- Cannot be used for a range query
 - (e.g.) all keys < 250
 - (e.g.) all keys between 30 and 45
- Uses an array (with all the problems of an array)
 - wastes space if the hash table has lots of unused buckets
 - requires the hash table to be recreated if the hash table is (nearly) full



Bloom Filters



Bloom Filter

- An application of hashing
- Invented by Burton Bloom in 1970
- It is used for quickly determining using a **small memory** whether new data **does not already exist** in a **large dataset**.
- Note the 3 key phrases above.
- Reference
 - <http://prakhar.me/articles/bloom-filters-for-dummies/>



2 Elements of the Bloom Filter

- An array of N bits (initially all set to 0)
- A collection of k hash functions
 - Each hash function maps a key to an index in the bit array.



Problem Bloom Filter Solves

- Does item “x” exist in a dataset?
- How to answer this fast using a small memory?
- Dataset may be very big or expensive to access.
- Filter negative results before accessing data.
- Allow **false positive** (“exist”) errors, as they only cost an extra data access.
- Don’t allow **false negative** (“not exist”) errors, because they result in wrong answers.



Application Example

- Web crawler (for a search engine) building up a large set of URLs
 - 1. Get a new URL (webpage).
 - 2. Extract all URLs on that page.
 - 3. For each URL on that page, check if this URL has not been crawled already. (* This step uses the Bloom Filter. *)
 - 4. If it has not been crawled already, add the new URL to the URL dataset, and return to Step 1.



Algorithm (1/2)

- An array of N bits (to represent the large dataset of the URLs)
 - Choose N to be much greater than the number of expected URLs.
 - Initialize all N bits to 0 (zero).
- Inserting new input data (new URL)
 - Apply each of the k hash functions to the input data. The result m ($m < N$) is an index in the bit array. There are k such results.
 - Set the m^{th} bit of the N -bit array to 1, for each of the k hashing results. The result is the N -bit array modified with k bits set to 1.



Algorithm (2/2)

- Check if the key is already in the dataset
 - Apply each of the k hash functions to the input data. The result m ($m < N$) is an index in the bit array. There are k such results.
 - Look up the k positions in the current bit array.
 - (Before setting bits in the k positions to 1) If at least one of the k positions in the current N -bit array is zero, then the key is DEFINITELY NOT in the dataset
 - Otherwise, the input data MAY BE in the dataset.
- Bloom Filter does not allow deletion.



Example: Bit Array and 2 Hash Functions

- The bit array has 11 bits, all of which are initially set to zero - 00000000000.
- $h(x)$ (* x is the new key to be inserted *)
 - Convert x to binary equivalent bx .
 - Take the odd numbered bits in bx and generate a new decimal number odx corresponding to these bits.
 - Apply the function $odx \text{ modulo } 11$.
- $g(x)$
 - the same as $h(x)$, except that even numbered bits are taken from bx



Illustration: First Input Key

- Assume the first input key is 25.
- Apply the key to the first hash function, $h(x)$.
 - $x = 25$. Then $bx = 11001$.
 - Taking odd bits from bx results in a binary number 101, which is binary for 5.
 - $5 \% 11 = 5$
- Apply the key to the second hash function, $g(x)$.
 - Taking even bits results in a binary number 10, which is binary for 2.
 - $2 \% 11 = 2$

First 3 Input Keys; and 2 Tests

x	bx	h(x)	g(x)	Bit array
25	0000011001	5	2	00100100000
159	0010011111	0	7	10100101000
585	1001001001	7	9	10100101010

118

162

5

3

2

0



Test Results

- 118
 - current array: 10100101010
 - $h(x)=5$: array[5] = 1
 - $g(x)=3$: array[3] = 0
 - conclusion: 118 is NOT in the dataset (correct)
- 162
 - current array: 10100101010
 - $h(x)=2$: array[2] = 1
 - $g(x)=0$: array[0] = 1
 - conclusion: 162 IS in the dataset (false positive)



Exercise

- query data: 317 \rightarrow 00100111101
 - current bit array: 10100101010



Solution

- test key: 317 \rightarrow 00100111101
 - current array: 10100101010
 - $h(x)$: $bx = 010111 = 23$; $23 \% 11 = 1$; $\text{array}[1] = 0$
 - $g(x)$: $bx = 00110 = 6$ $6 \% 11 = 6$; $\text{array}[6] = 0$
 - conclusion: 317 is NOT in the dataset (correct)



How to Reduce the False Positives

- Reducing the false positives means reducing the probability that a non-existing key will hash to 1 bits in the N-bit array
- Two Ways
 - One is to make the bit array larger
 - Another is to use additional hash functions
- In practice, false positive rate may be reduced to 15 to 1% of the tests.



End of Lecture
