


Chapter 3: Process Concept

School of Computing, Gachon Univ.
Jungchan Cho



Most slides from "Operating System Concepts – 10th Edition".
Many slides are taken from lecture notes of Prof. Joon Yoo.

Objectives

- To introduce the **notion of a process** -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore **interprocess communication** using shared memory and message passing

Chapter 3: Process Concept

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication


Process Concept

- **Process** (= job/task, #program, #processor)
 - a program *in execution*
 - ▶ Program is **passive** entity stored in disk (executable file), Process is **active** loaded into **main memory**
- One program can be several processes
 - e.g., same user executing multiple copies of a program
- These two terms are **different!!!!!!**
 - **Process**: a program in execution
 - **Processor**: processor chip (e.g., CPU)
 - Multiprocessor vs. Multiprocess(ing)




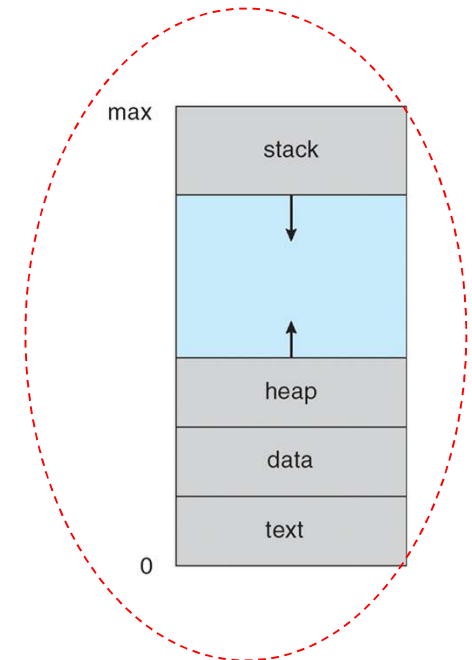
프로세스

프로세서



Process in Memory

- Process **in memory**
 - **Text**: The binary assembly **program code**
 - **Data**: global **variables**
 - **Stack**: temporary local **data**
 - ▶ Function parameters, return addresses, local variables
 - **Heap**: memory **dynamically** allocated **during run time** (e.g., C malloc(), Java objects)



Memory

Simple C-code example

Binary Assembly Code (text)

```
#include <stdio.h>
#include <stdlib.h>
```

```
int x;
int y = 15;
```

**Global variables
(data)**

```
int main()
{
```

```
    int *values;
    int i;
```


**Local variables
(stack)**

```
    values = (int *)malloc(sizeof(int)*5);
    for(i = 0; i < 5; i++)
        values[i] = i;
```

**C malloc
(heap)**

```
    return 0;
```

```
}
```



Simple C-code example

Binary Assembly Code (text)

```
#include <stdio.h>
FILE* fp;
int counter;

void main(int argc, char *argv[]) {
    char c;
    counter=0;
    fp = fopen(argv[1], "w");
    while(1) {
        c = char_read();
        if( c == 's' ) break;
        counter++;
        printf("%d th character is %c \n", counter, c);
        fprintf(fp, "%c", c);
    }
    fclose(fp);
}


char char_read()
{
    char c;
    scanf("%c", &c);
    return c;
}
```

Global variables (data)

Return address (stack)

Local variables (stack)

max



0

Program counter

Executing
here

Process in Memory: Stack

- Stack used for function calls

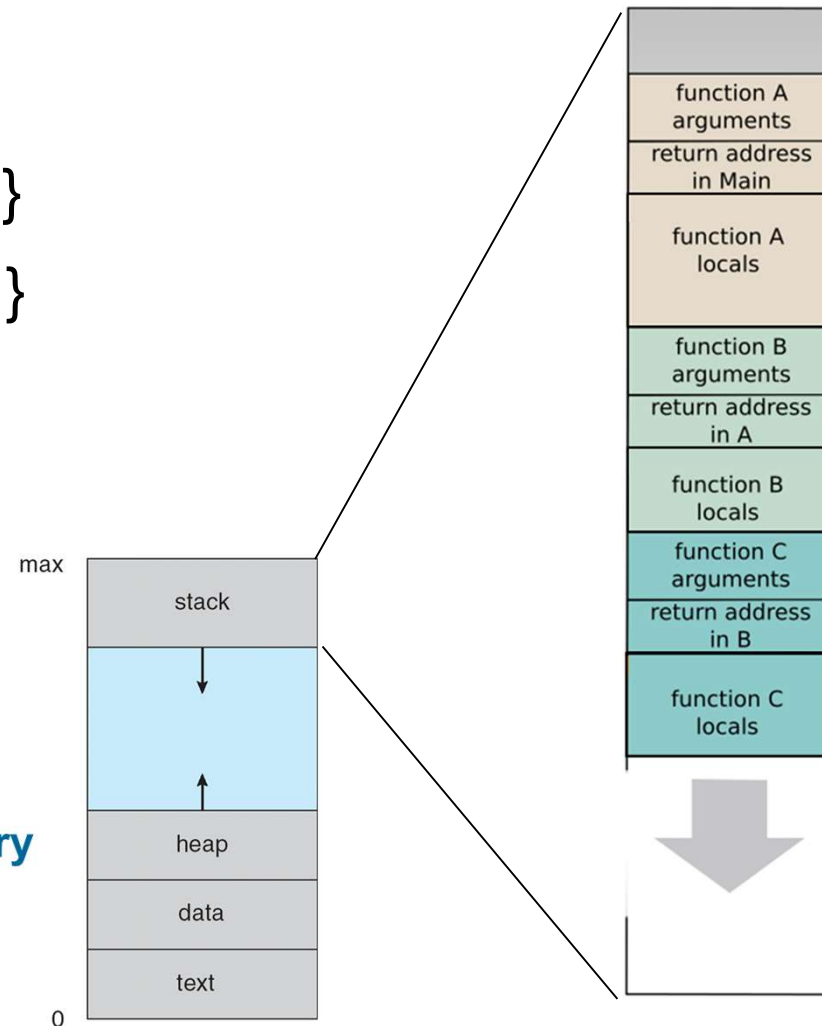
```
main() { func_A()... }
```

```
func_A() { func_B()... return }
```

```
func_B() { func_C()... return }
```

```
func_C() { ... return }
```

Process in Memory



Process in Memory: Text & Data


Process in Memory

```
.text
main:  la $t0,str          # put starting address of "hello world" into t0
loop:  lb $t1,0($t0)       # load byte in phrase
       beqz $t1,over      # if character null, we are finished
       beq $t1,0x6c,cnt   # if the character is an l, go to count
incr:  addi $t0,$t0,1      # add 1 to current byte address
       j loop            # get next byte to compare
cnt:    addi $t2,$t2,1      # add one to count of letter l's in phrase
       j incr           # go back into loop
over:  la $a0,rept         # Output report phrase
       li $v0,4
       syscall
       move $a0,$t2       # move total l-count to $a0 for output
       li $v0,1
       syscall            # output letter total
       li $v0,10
       syscall            # end program
```

Binary Assembly Code (text)


```
.data
str:    .asciiz "Hello, world!\n"
rept:   .asciiz "The total count of the letter l is "
```

Binary data (data)



Process Management

- Each process is registered to and managed by the OS
 - The OS manages and take care of all the processes.
 - Therefore, the OS needs to manage the current information (e.g., state) of each process – use a data structure called **PCB (Process Control Block)**



Process Control Block (PCB)


PCB: OS maintains the **information** for *each* process

- Process **state** – running, ready, waiting, etc (next slide)
- Process stops while running (e.g., interrupt) and now in ready state. Need to save current **running information** to resume running where we stopped...
 - Program counter (PC) – next instruction address, CPU registers – contents of all registers (in CPU)
 - CPU scheduling information
- Where is the process located in **memory**?
 - Memory-management information
 - memory allocated to the process
- Which **I/O devices/files** are the process using?
 - I/O status information
 - I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...

Process State

- As a process executes, it changes **state**
 - **new**: The process is *being* created
 - **ready**: The process (in memory) is ready to be assigned to CPU
 - **running**: Process instructions are being executed by CPU
 - **waiting**: The process is waiting for some event (e.g., I/O operation) to occur
 - **terminated**: The process has *finished* execution



Important!

- Only **one** process is **running** on any CPU core at any instant
- All the other processes are waiting in **ready** or **waiting** states

Chapter 3: Process Concept

- Process Concept
- **Process Scheduling**
- Operations on Processes
- Interprocess Communication

Questions


- How many processes can be executed *concurrently* in an **OS**?
 - hint: multi-tasking (programming, processing)
- How many processes can use a **CPU core** *in parallel* at the same time?
 - hint: CPU instruction execution

On a single CPU core...

- SUPPOSE: One process is currently using the CPU


```
long count = 0;  
while (count >= 0)  
    count++;
```

- What are other processes doing?
 - Usually nothing! (Just waiting for CPU in ____ state or doing I/O in ____ state)

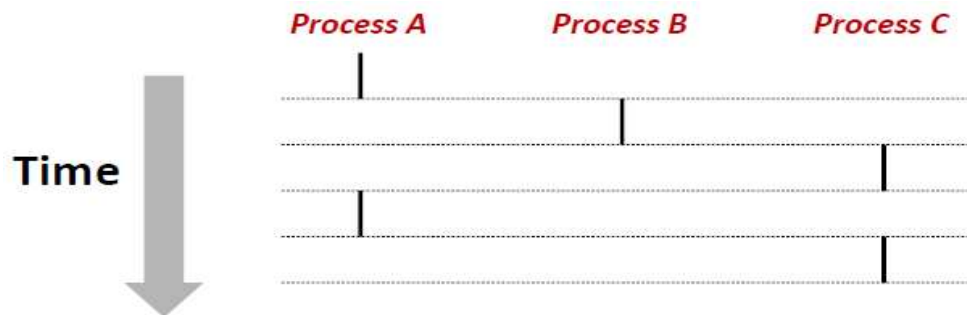


User View of Concurrent Processes

- Intuitively, we can think of **concurrent** processes as running **in parallel** with each other

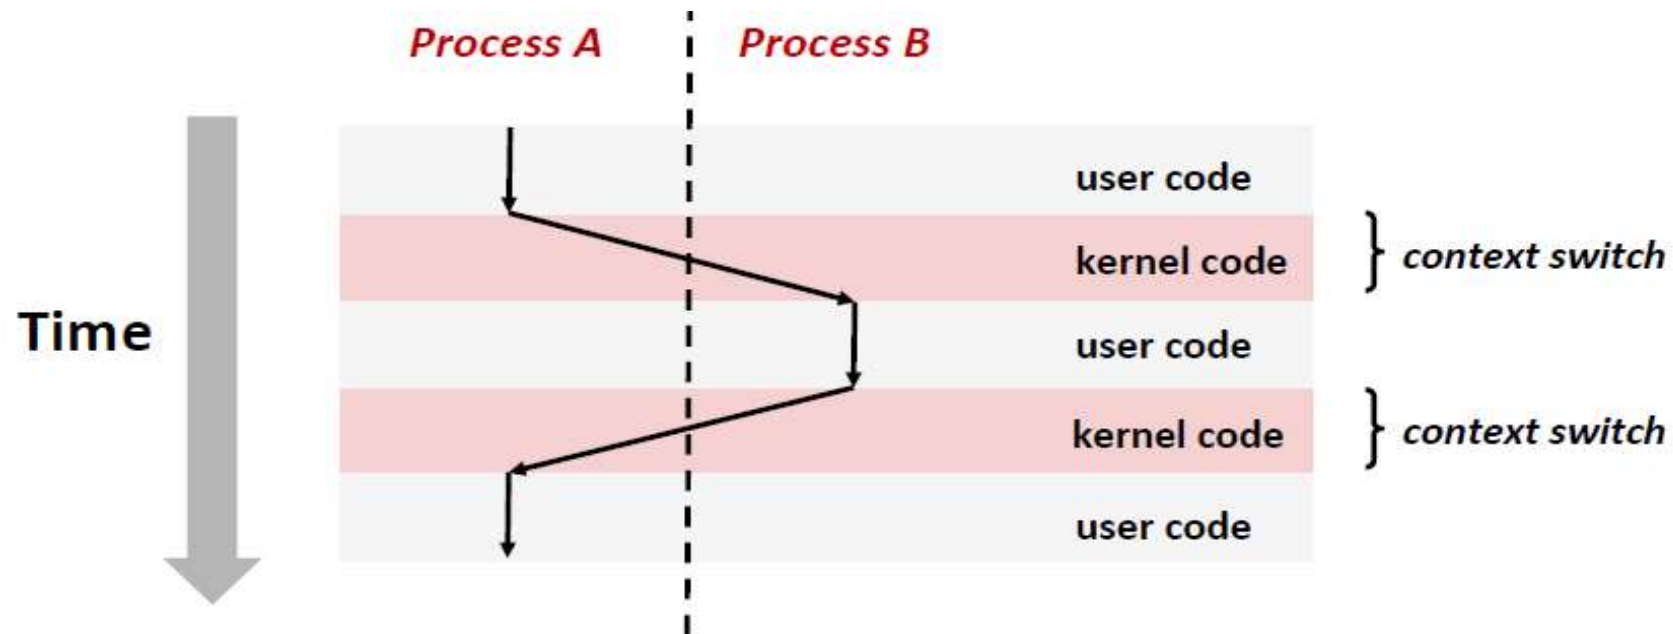


- However, in a single CPU core, only one process can run at a time



Concurrent Processes

- Process switching from one process to another by OS !
 - called **Context switch** (we will learn soon!)



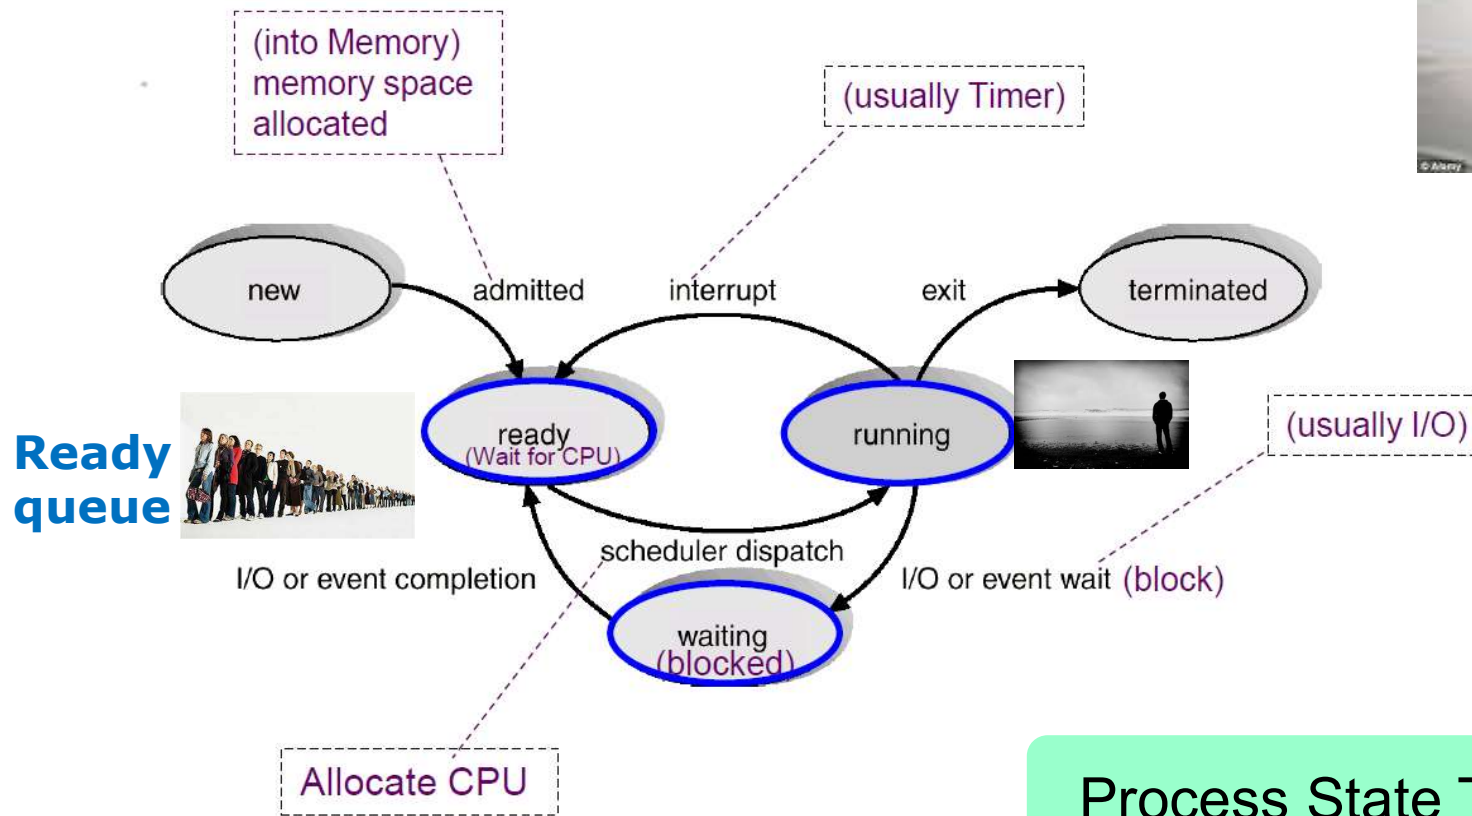
- Determining when to switch, to which process is very important to the system performance – Enter **process scheduling**!

Process Scheduling

- Objective of **Multiprogramming**
 - Maximize **CPU utilization**: Switch the CPU to another process when I/O happens
 - **Time-sharing**: Switch the CPU among processes so frequently that users can interact with each program while it is running
- How do we implement multi-programming?
 - **Process scheduler** selects among available processes (in Ready state) for next execution on CPU

Representation of Process State

- How to describe the state of a process ?
- A process's point of view




Process State Transition Diagram

Process Scheduling Queues



■ Ready queue

- The processes that are in the **main memory** and are ready to execute are kept on a list called **ready queue** (in Ready state; waiting for CPU)



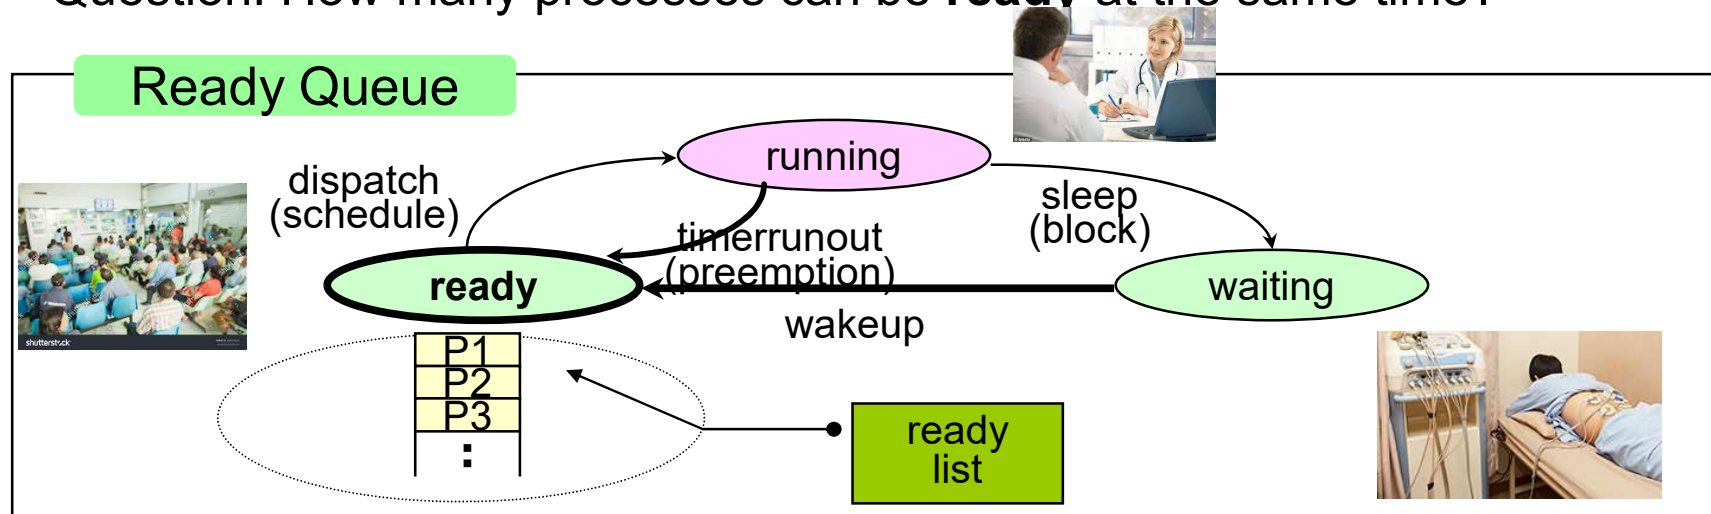
■ I/O (Device) queues

- The list of processes waiting for a particular **I/O** operation is called a device queue (waiting for I/O device)
 - ▶ Each device has its own device queue

Scheduling Queues

■ Ready queue

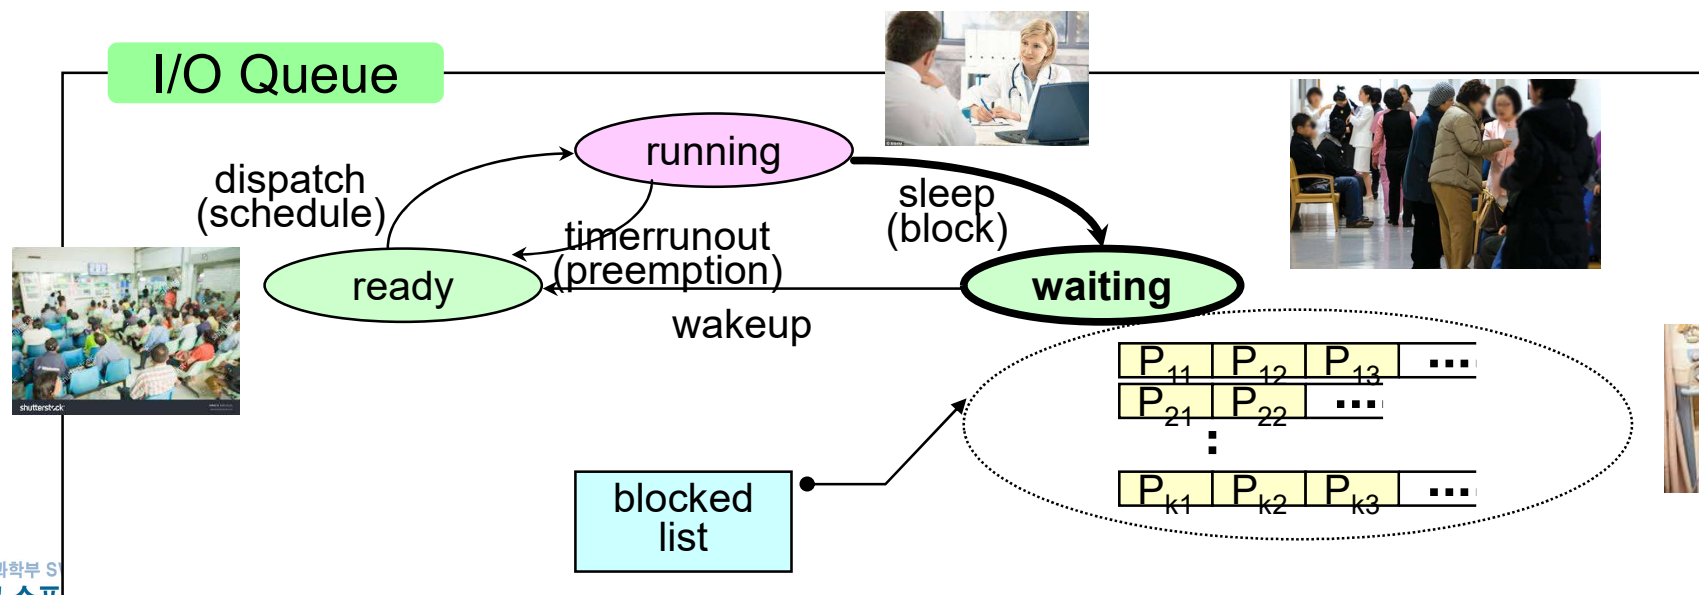
- Processes in _____ state
 - ▶ Awaiting for the processor (i.e. CPU)
- Process scheduling**
 - ▶ **Selecting** a process from the **ready queue** and **dispatch** it when the **processor** is available
- Question: How many processes can be **running** at the same time?
- Question: How many processes can be **ready** at the same time?



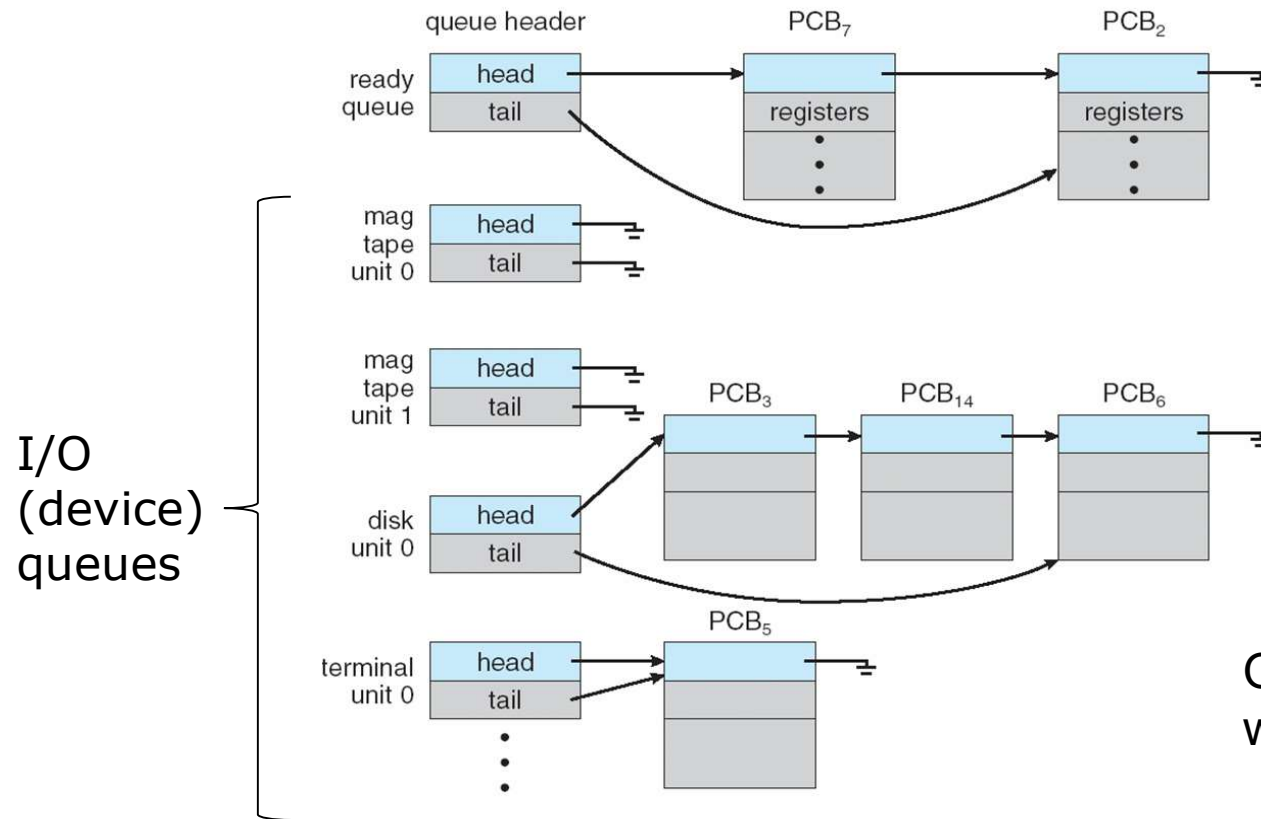
Scheduling Queues

I/O (Device) queue

- When a process is allocated the CPU, it **executes** for a while and eventually **quits**, is **interrupted**
- or **waits for an event** (e.g., I/O completion)
 - The process has to be **waiting** (or **blocked**) in the **I/O queue**
 - Each device as a I/O (device) queue – if I/O device is busy, multiple processes may be waiting in I/O queue
 - Also called Blocked list (block queue, I/O queue)



Ready Queue And Various I/O Device Queues



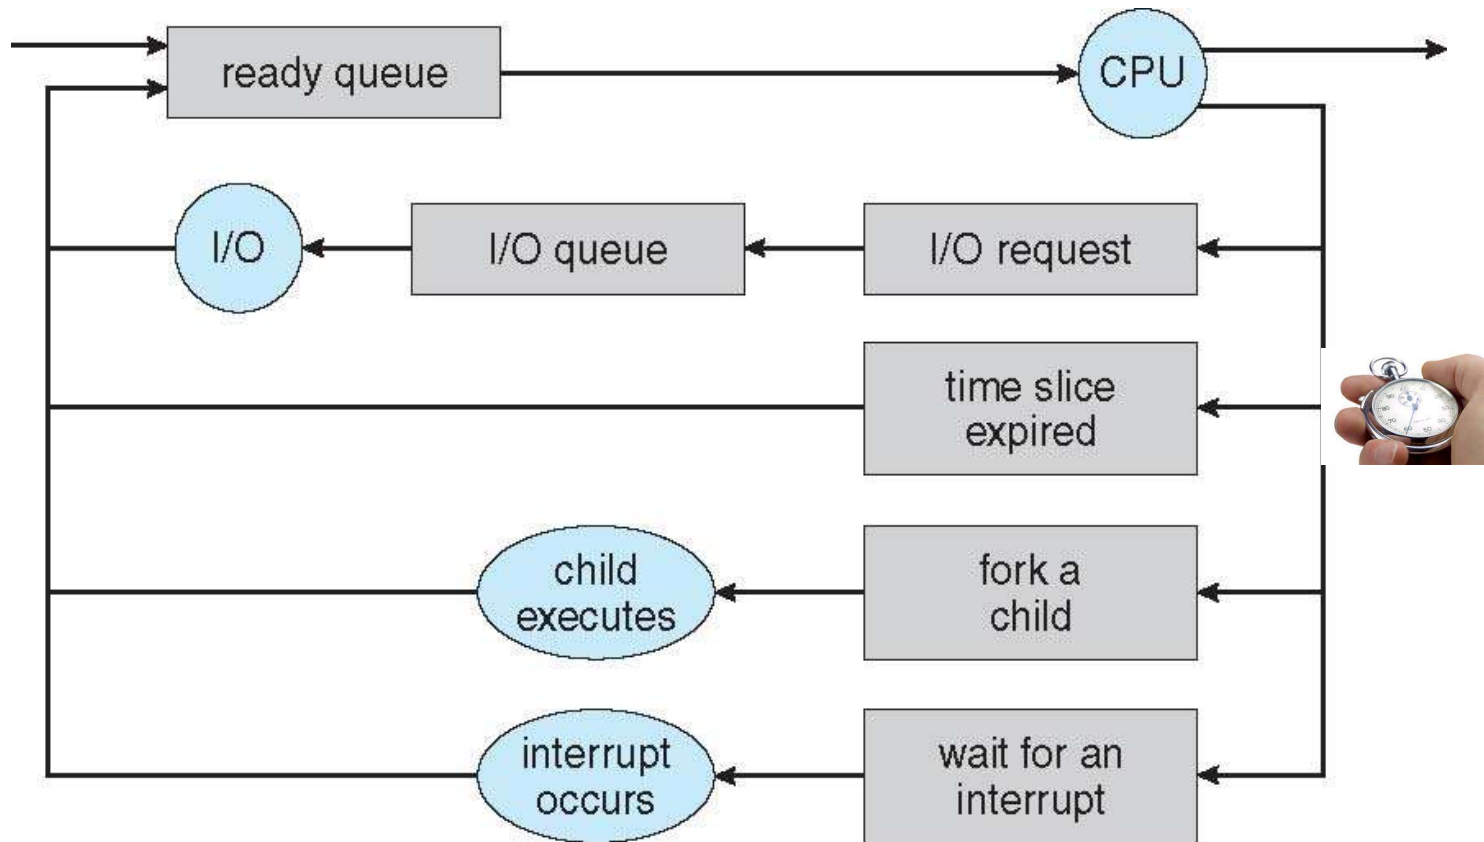
Queues are usually implemented with linked lists

Queue header → pointing the first and last PCB structures

Each PCB structure has the address of its next PCB structure

Representation of Process Scheduling

Queuing diagram



Process Scheduler

- Process Scheduler

- an **operating system module** that selects the next jobs to be admitted into the system and the next process to run

- Types

- Long-term scheduler (=job scheduler)
- Short-term scheduler (= CPU scheduler)

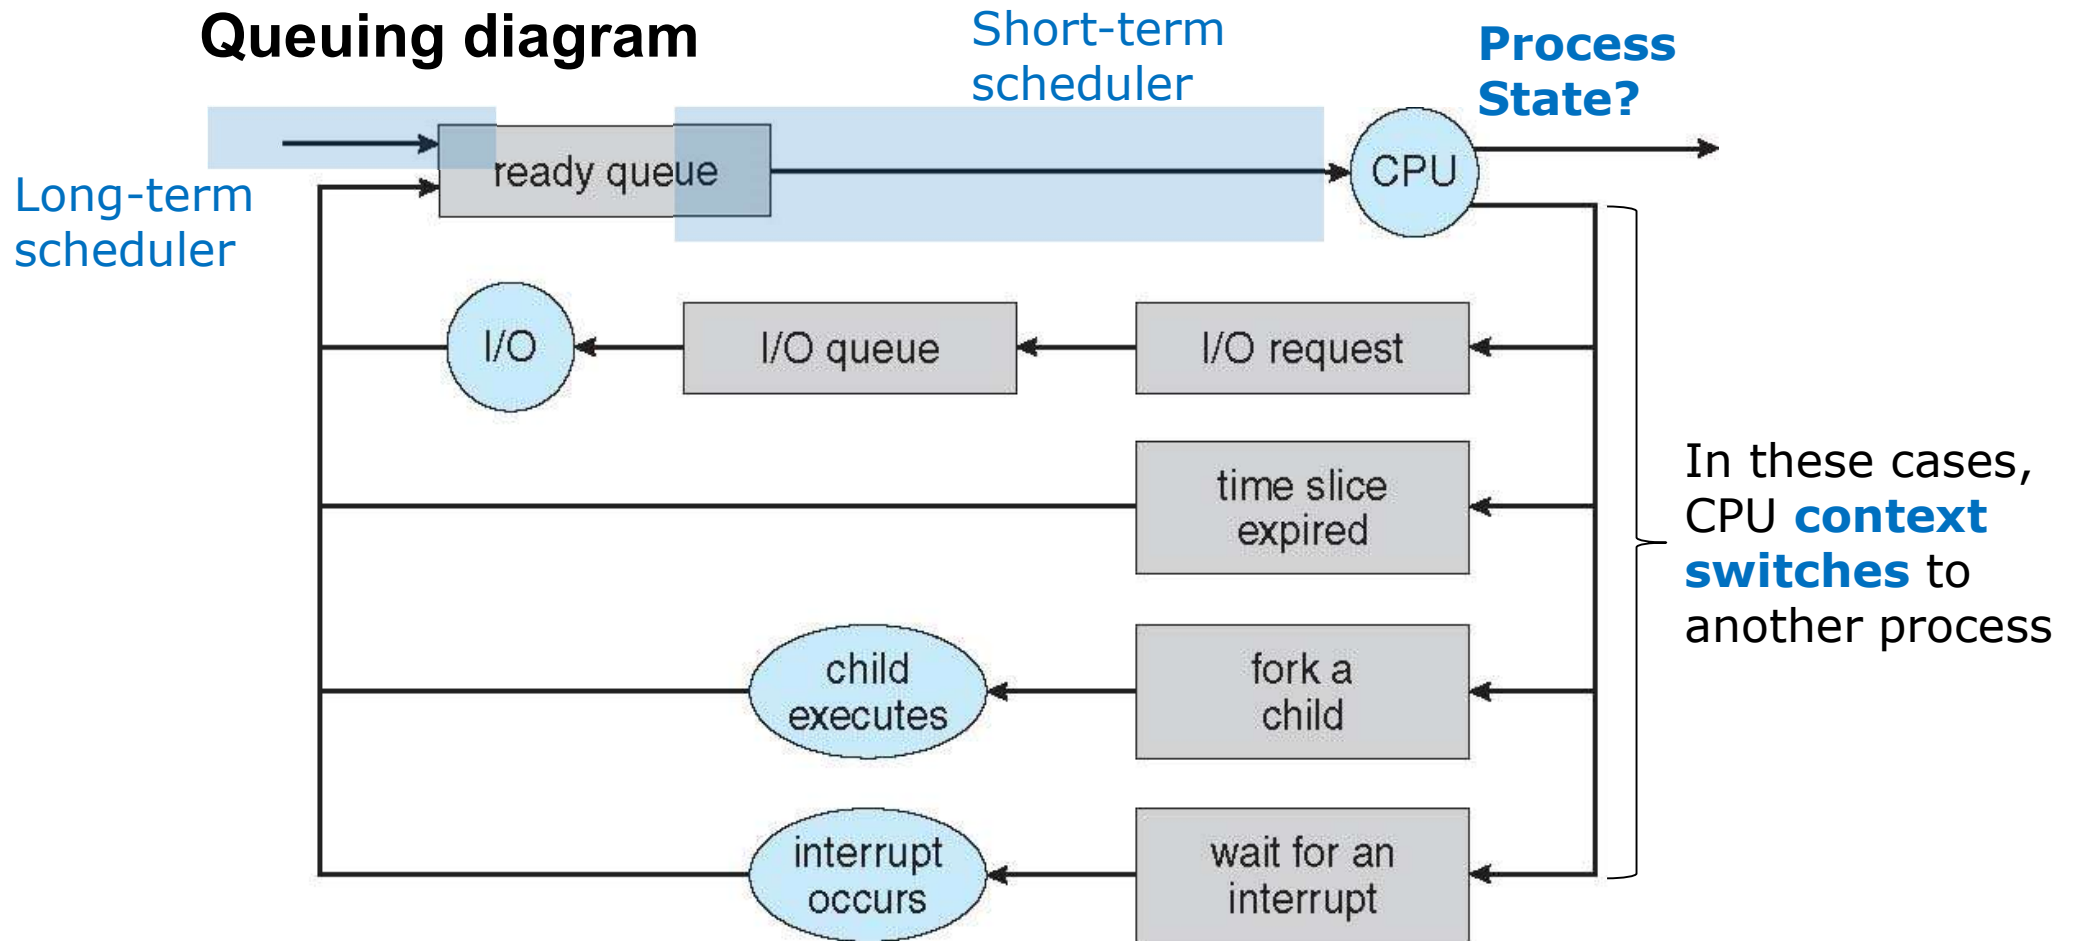
Schedulers

- Long-term scheduler (or job scheduler)
 - Decides which jobs or processes (in disk) are to be admitted into the ready queue (*i.e.* Main memory)
 - Giving Memory
 - invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
- Short-term scheduler (or **CPU scheduler**)
 - Among processes in ready queue, selects which process (memory resident) should be executed next and allocated CPU
 - Giving CPU
 - invoked very frequently (milliseconds) \Rightarrow (must be fast)

Long-term Schedulers


- The long-term scheduler controls the **degree of multiprogramming**
 - Degree of multi-programming : the number of processes in memory (determining the degree of contention for resources)
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Q: What happens if most processes in ready queue are I/O bound? CPU bound?
 - Long-term scheduler strives for good ***process mix***

Representation of Process Scheduling

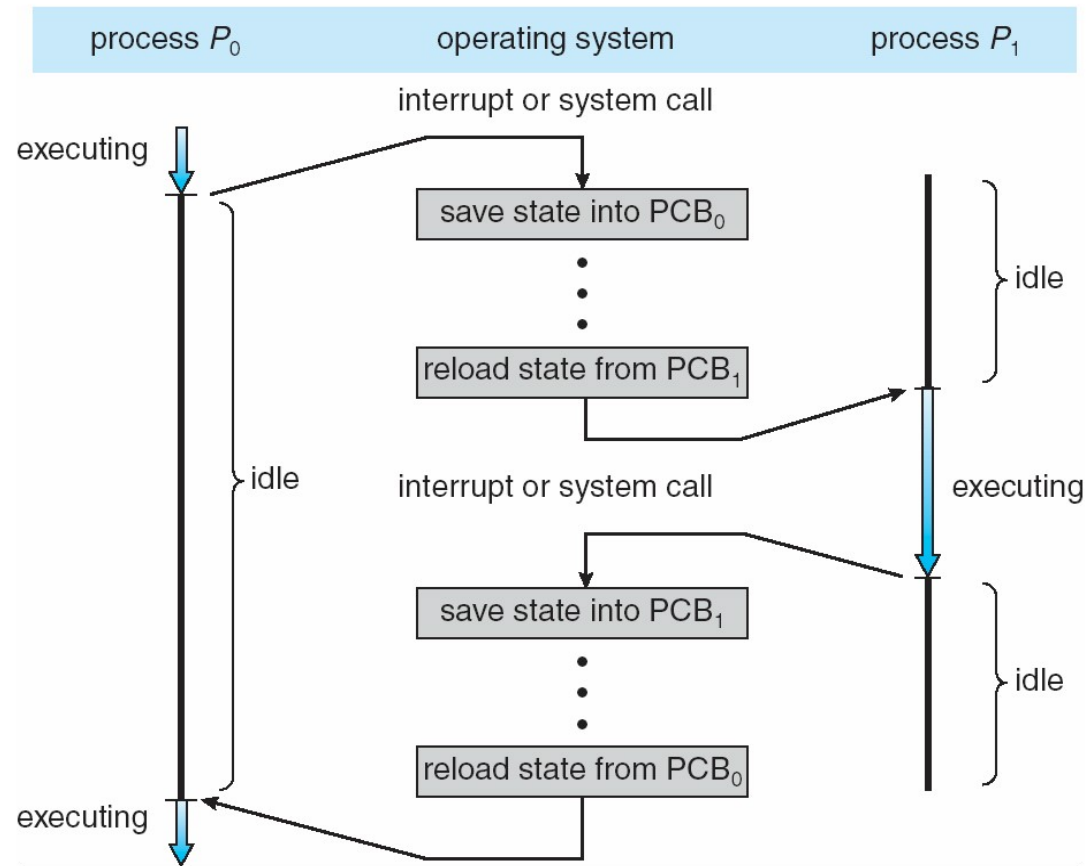


Context Switch

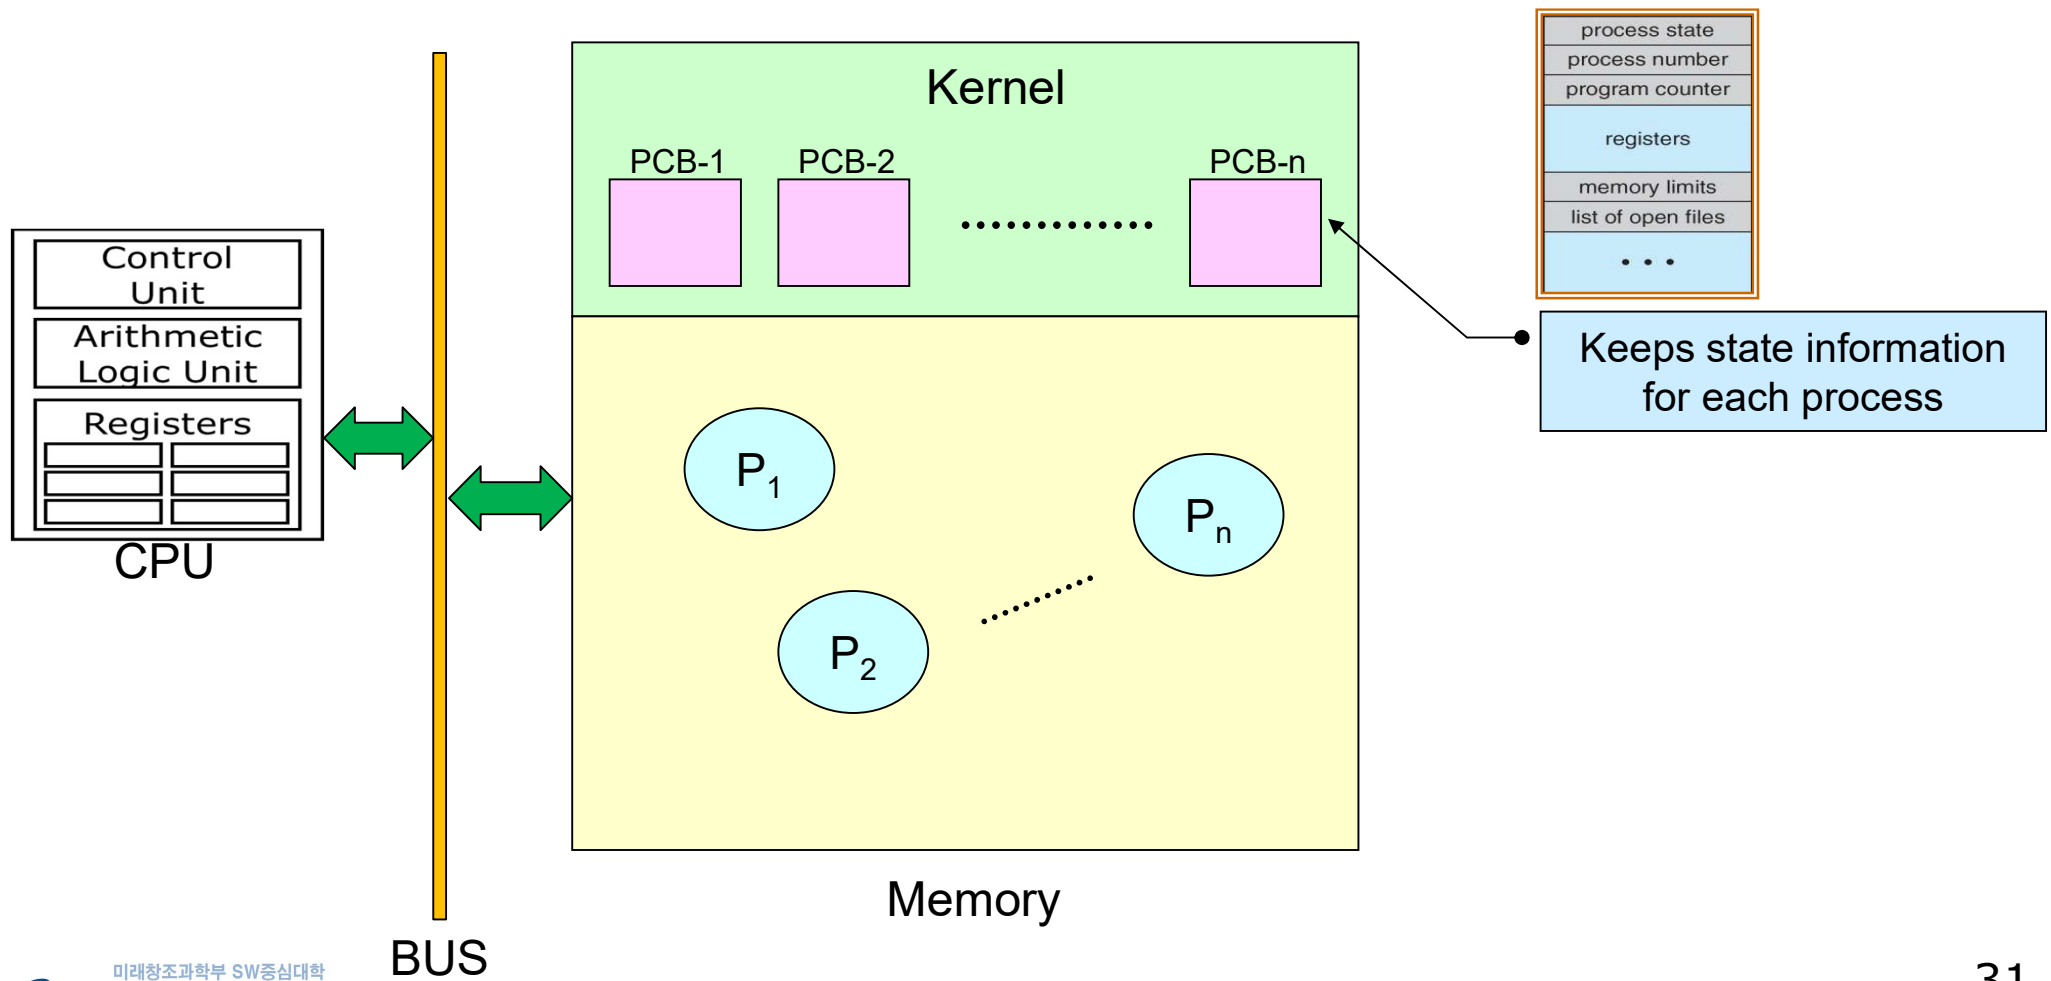
- When CPU switches to **run** another process
 - the system must **save the state** of the old process (to resume where we stopped) and
 - load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is **overhead**; the system (or CPU) cannot do any other work while switching
 - The more complex the PCB -> longer the context switch



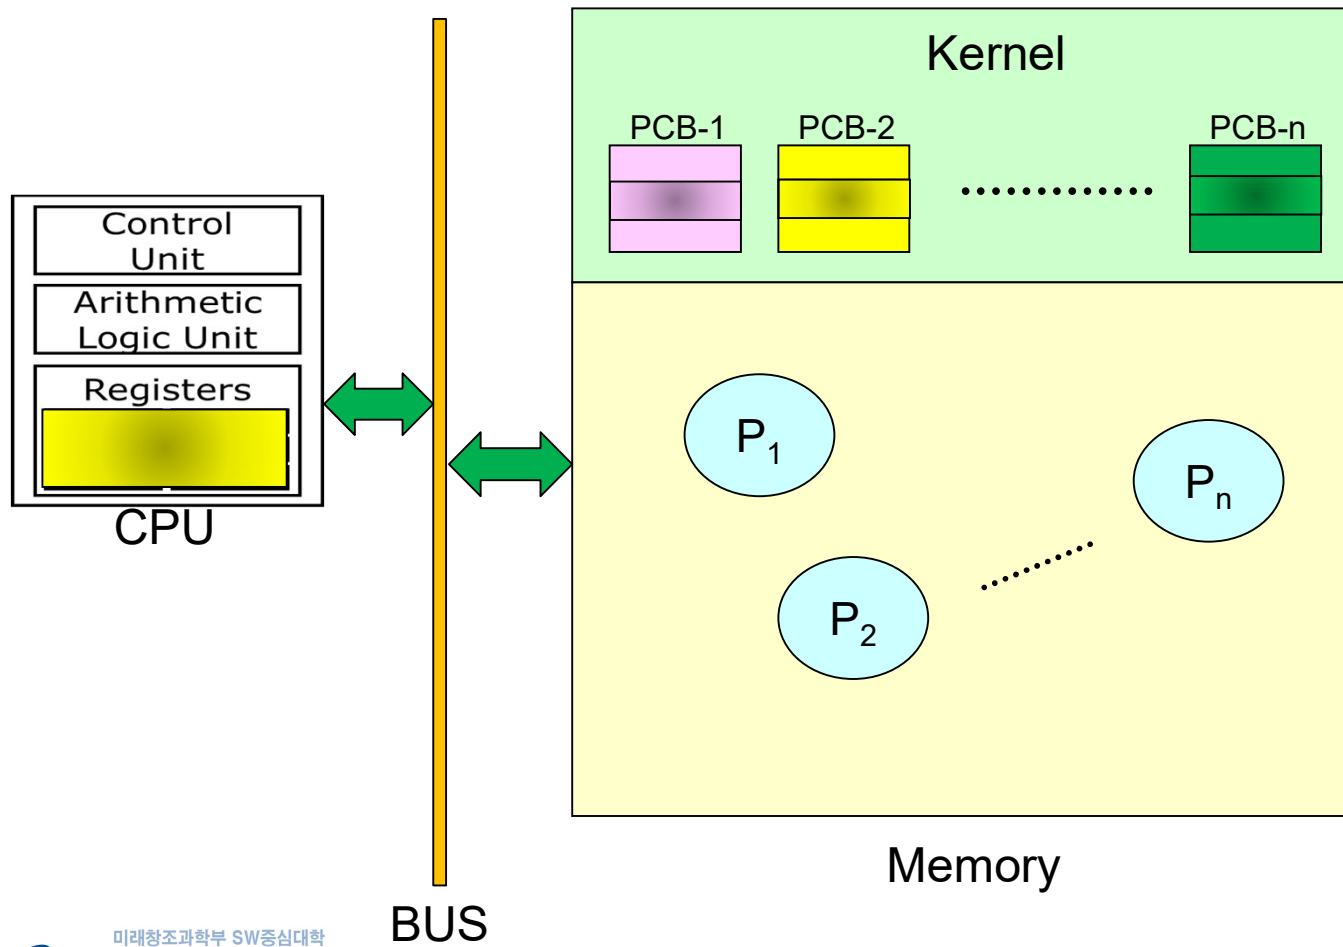
CPU Switch From Process to Process



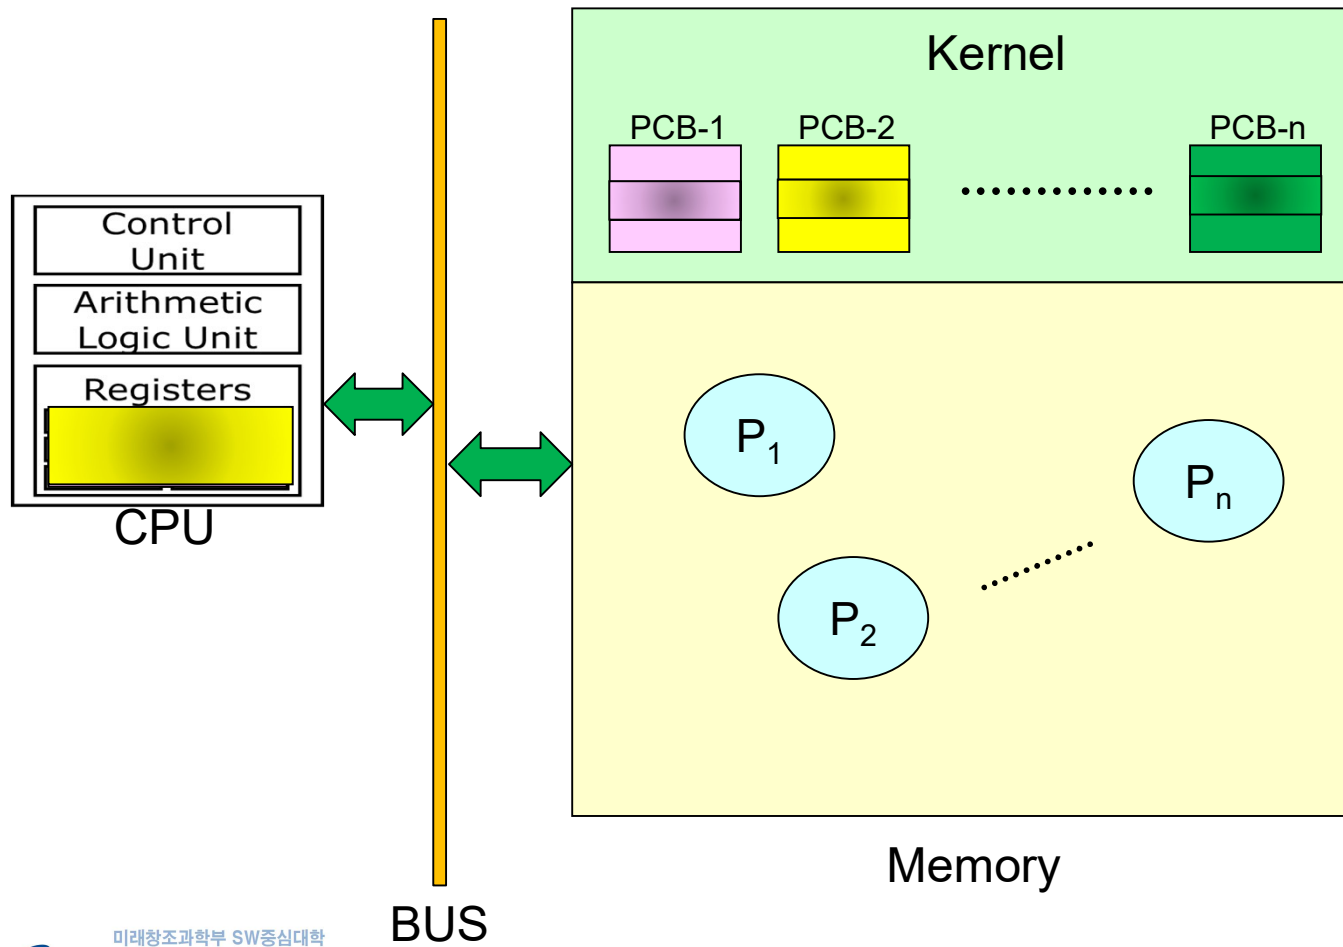
(Cont.)



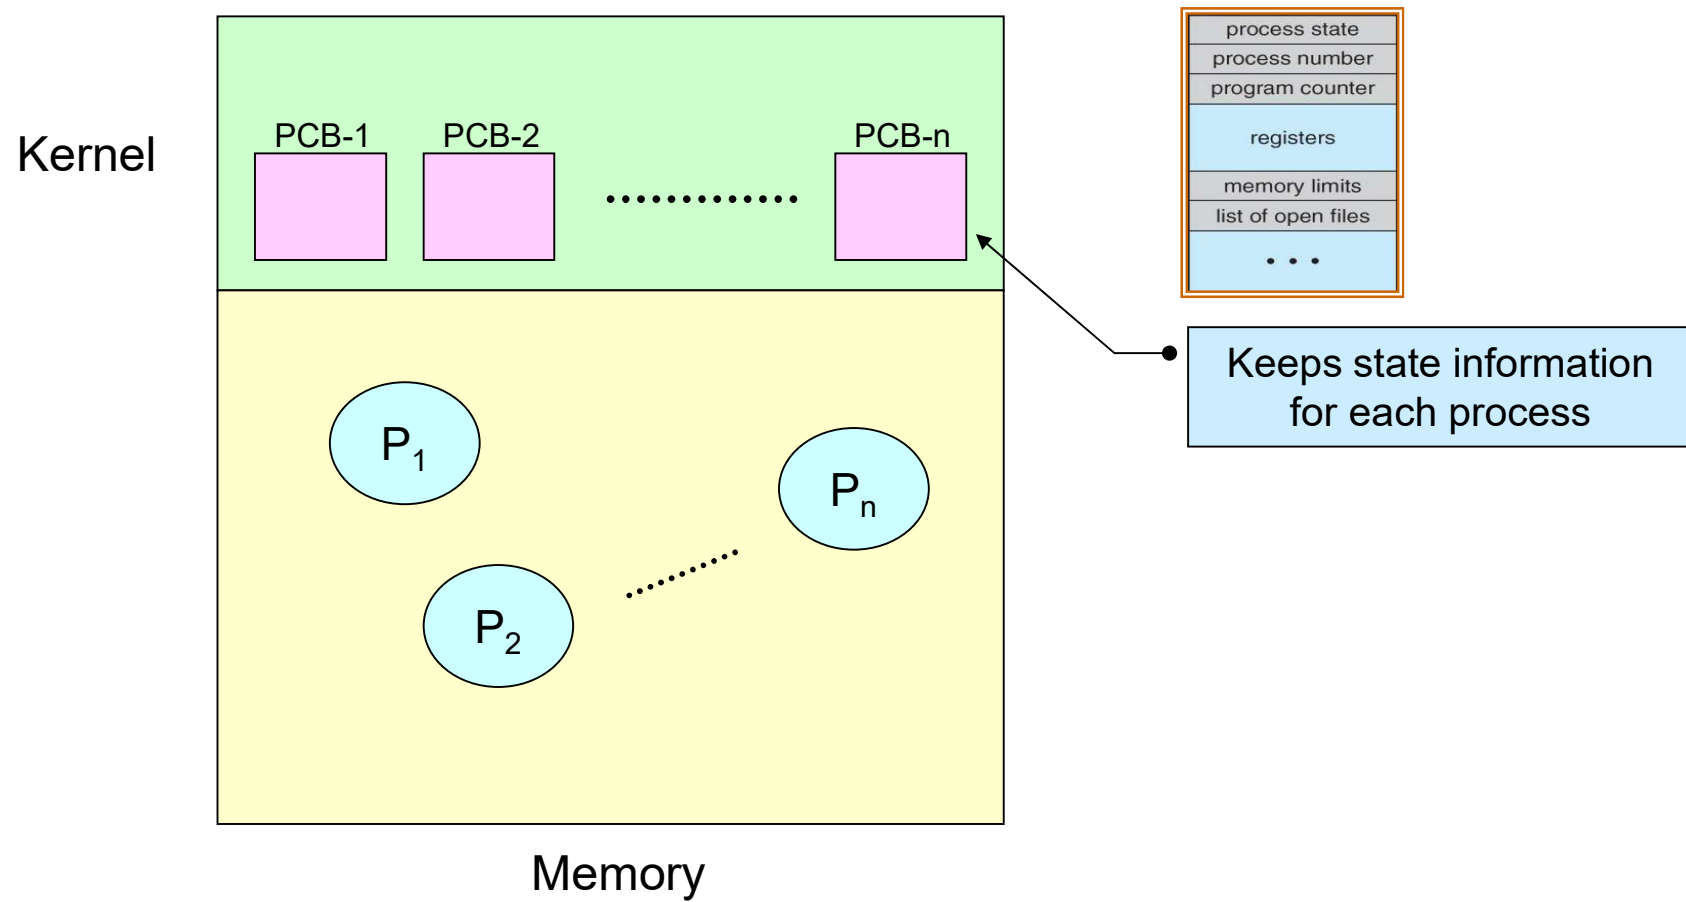
(Cont.)



(Cont.)



(Cont.)




Chapter 3: Operations on Processes

- Process Concept
- Process Scheduling
- **Operations on Processes**
- Interprocess Communication

Operations on Processes

- Processes can execute concurrently, thus they may be **created** and **deleted** dynamically.
 - Operating System must provide mechanisms for process creation, termination, and so on.
 - Generally, process identified and managed via a unique **process identifier (pid)**
- Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes

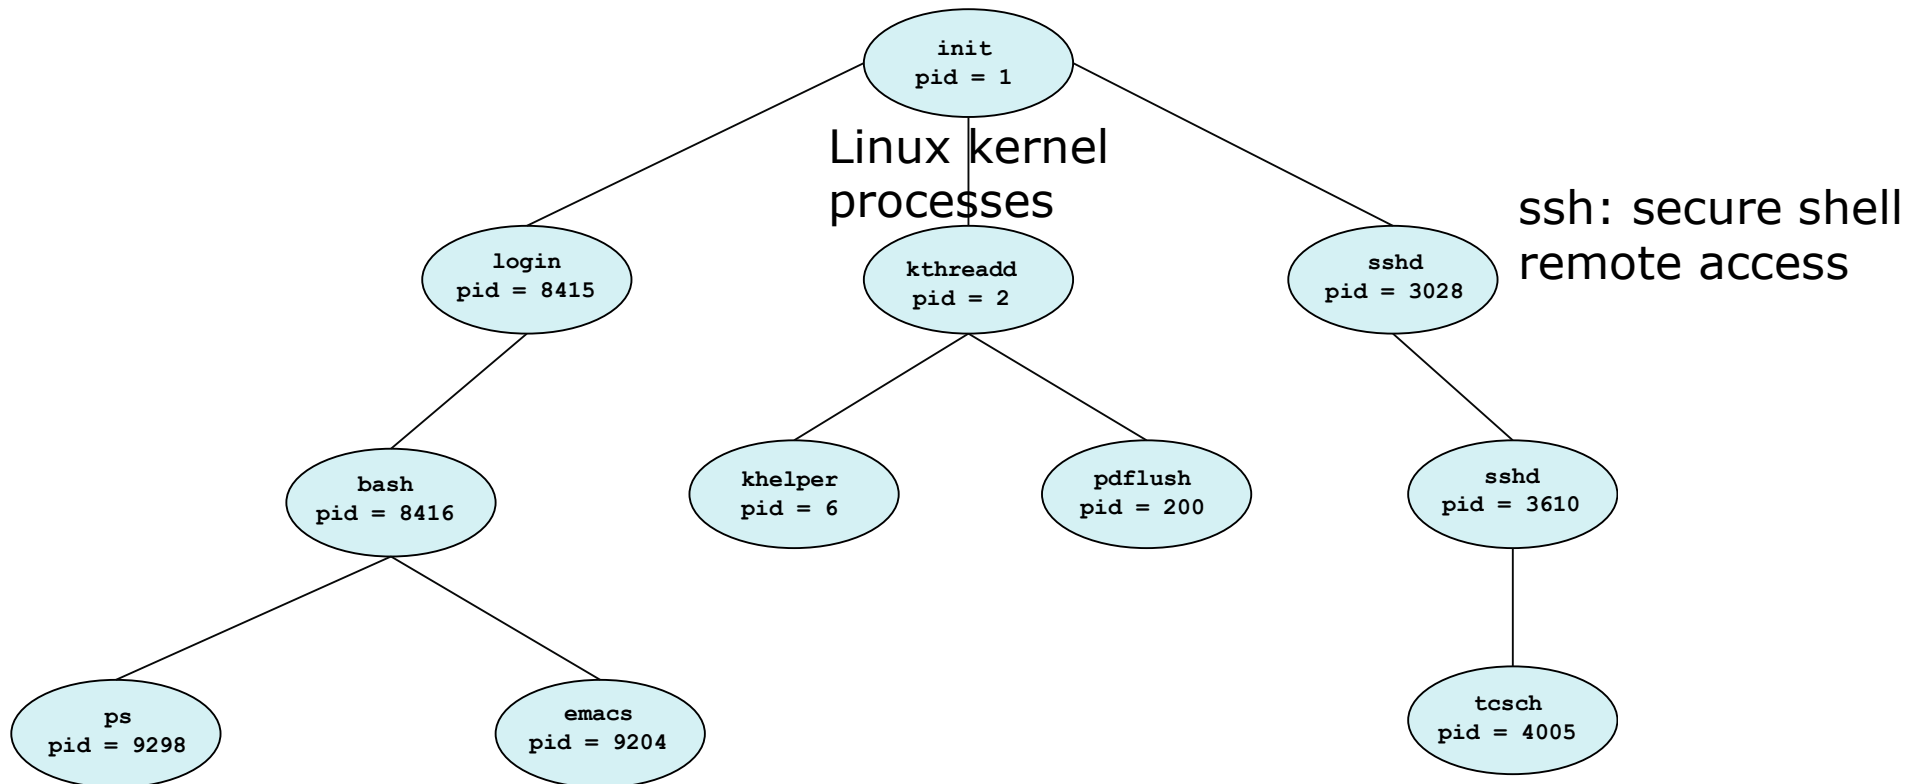


- So who created this universe (i.e., the first process)?



A Tree of Processes in Linux

Root process for all user processes



type `ps -el` in Linux shell


Process Creation

- Child process needs memory address space
 - Child duplicate of parent: **fork()** system call
 - Child has a program loaded into it: **fork()** → **exec()** system call

UNIX examples: Process Creation

- UNIX examples
 - `fork()` system call creates a new process
- Both processes (parent and child)
 - continue execution after `fork()`,
 - with one difference: `fork()` return value:
 - Child process: 0
 - Parent process: `pid` of child process (>0)

process identifier (`pid`)



```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
```

```
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
```

```
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
```

```
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
```

```
    return 0;
}
```


Parent

Child

Process Creation

■ Memory address space

- Child duplicate of parent: `fork()`
- Child has a program loaded into it: `exec()` system call
 - ▶ **`exec()`** system call used after a `fork()` to *replace* the process' memory space with a new program
 - loads a new binary file into memory (deletes original memory – copy of parent)



```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
```

```
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
```

```
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
```

```
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
```

```
    return 0;
}
```

Parent


Child

Process Creation

- Execution options
 - Parent and children: Which executes first?
 - Parent can
 - ▶ execute concurrently with child or
 - ▶ wait until children terminate: `wait()` system call returning the pid:

```
pid_t pid; int status;  
pid = wait(&status);
```

```
#include <sys/types.h>  
#include <stdio.h>  
#include <unistd.h>  
  
int main()  
{  
    pid_t pid;  
  
    /* fork a child process */  
    pid = fork();  
  
    if (pid < 0) { /* error occurred */  
        fprintf(stderr, "Fork Failed");  
        return 1;  
    }  
    else if (pid == 0) { /* child process */  
        execlp("/bin/ls", "ls", NULL);  
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        wait(NULL);  
        printf("Child Complete");  
    }  
  
    return 0;  
}
```



Process Termination



- Process executes last statement and asks the operating system to delete it: **exit()** system call or just **return()**
 - Return status value (integer) to parent process
 - Process' resources (memory, open files, I/O buffers) are deallocated by operating system
- Parent may **terminate** execution of children processes: **kill(pid, SIGKILL)**
 - E.g., child has exceeded allocated resources, task assigned to child is no longer required
 - If parent is exiting
 - ▶ Some operating systems do not allow child to continue if its parent terminates
 - ▶ All children terminated - **cascading termination**



Question 1

- How many processes are created ? (including the parent process)

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    /* fork a child process */
    fork();

    /* fork another child process */
    fork();

    /* and fork another */
    fork();
    printf("process forked\n");
    return 0;
}
```

Question 2

- Using the program in Figure, identify the values of `pid` at lines A, B, C, and D. (Assume that the actual `pids` of the parent and child are 2600 and 2603, respectively.)
- recall
 - **fork()** return value:
 - Child process: 0
 - Parent process: `pid` of child process (>0)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid, pid1;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        pid1 = getpid();
        printf("child: pid = %d",pid); /* A */
        printf("child: pid1 = %d",pid1); /* B */
    }
    else { /* parent process */
        pid1 = getpid();
        printf("parent: pid = %d",pid); /* C */
        printf("parent: pid1 = %d",pid1); /* D */
        wait(NULL);
    }

    return 0;
}
```

Figure 3.29 What are the pid values?

Question 3

- What is the output of line A?

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int value = 5;

int main()
{
    pid_t pid;
    pid = fork();

    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE A */
        return 0;
    }
}
```

fork() return value:

- Child process: 0
- Parent process: pid of child process (>0)

Figure 3.30 What output will be at Line A?

Chapter 3: Process Concept

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication

Interprocess Communication



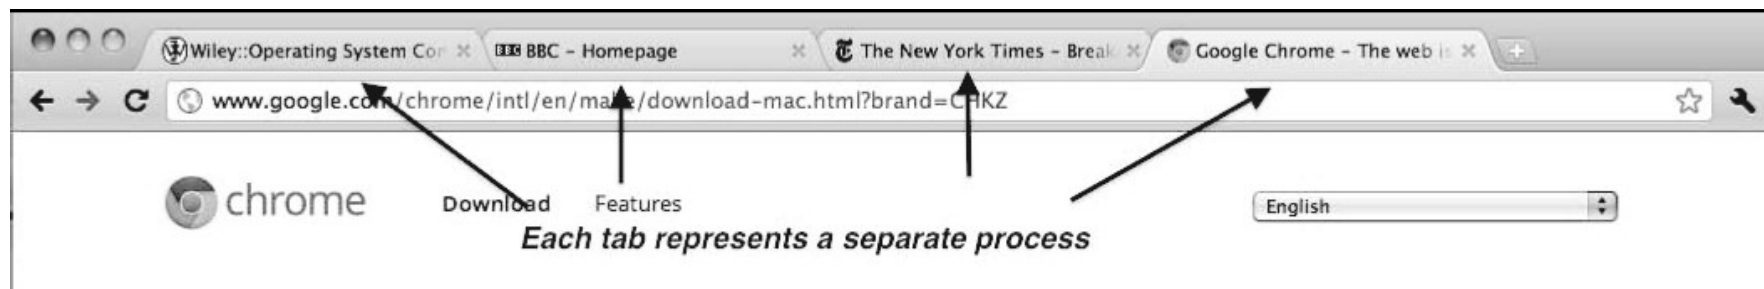
- In many cases, processes can **cooperate** with each others
 - Reasons for cooperating processes:
 - ▶ Information sharing – e.g., shared file
 - ▶ Computation speedup – e.g., parallel computing in multi-core environment
 - Example: PPT (100page) → Print
 - ▶ The **PPT process** needs to communicate to a **printer-process**, so that it sends data to the printer-process

Interprocess Communication

- Why can't processes just communicate with each other?
 - A process cannot directly access other process's **memory** – why not?
 - ▶ A: For system protection
- Cooperating processes need **interprocess communication (IPC)** provided by OS via system call

Example: Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble (e.g., JavaScript, Flash, HTML5), entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 categories
 - **Browser** process manages user interface, disk and network I/O (1 process created)
 - **Renderer** process renders web pages (deals with HTML, Javascript), new process for each website opened in a new tab
 - **Plug-in** process for each type of plug-in (e.g. Flash, QuickTime)



Communications Models

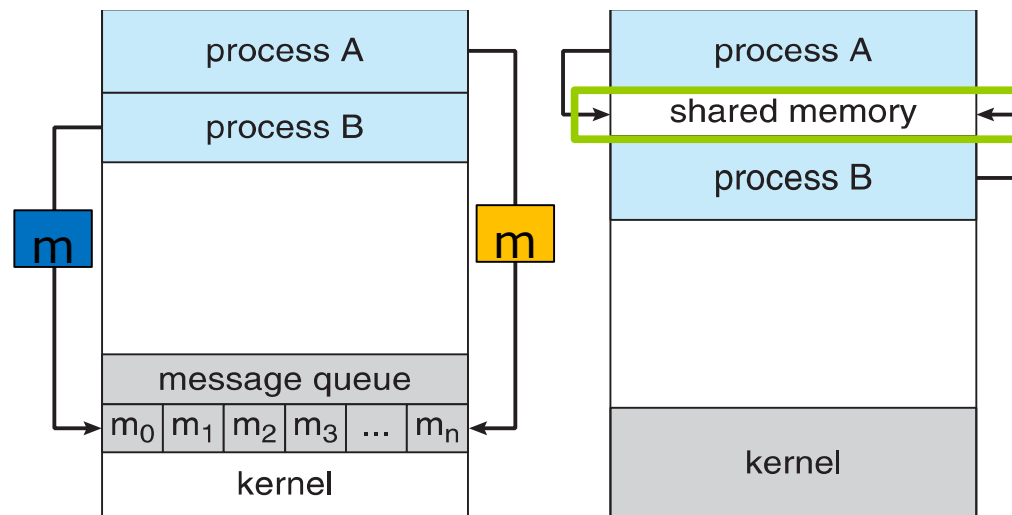
Mechanism

for processes to **communicate** & to **synchronize** their actions

Two fundamental models of IPC

Shared memory

Message passing




(a) Message passing

(b): Shared memory

Shared-Memory Systems


- OS prevents one process from accessing another process's memory – Protection
 - **Shared memory**: Two or more processes agree to remove this restriction
- A region of memory that is shared is created by system call
 - Processes exchange information by reading and writing data on shared area
 - All access to shared memory are treated as routine memory access
- The shared-memory can be modeled as **Producer-Consumer Problem**



Shared-Memory Systems

■ Producer-Consumer Problem

- A **producer process** produces information
- that is consumed by **consumer process**.
 - ▶ e.g. producer : PPT – send to printer;
consumer : printer program
 - ▶ e.g., producer: a web server produces -- HTML files and images ;
consumer -- a client web browser
- The producer and consumer must be **synchronized** so that
 - ▶ **Producer** does not try to produce an item when **buffer** (shared memory) is *full*
 - ▶ **Consumer** does not try to consume an item that has not yet been produced (i.e., buffer is *empty*)



Shared Memory


```
#define BUFFER_SIZE 6
```

```
typedef struct  
{  
    . . .  
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0; //next free position in buffer
```

```
int out = 0; // first full position in buffer
```




Shared Memory (producer view)

```
while (true)
{
    /* Produce an item */
    while (((in + 1) % BUFFER SIZE) == out)
        ; /* do nothing --no free buffers */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER SIZE;
}
```


Make sure buffer
is not full



Shared Memory (Consumer view)

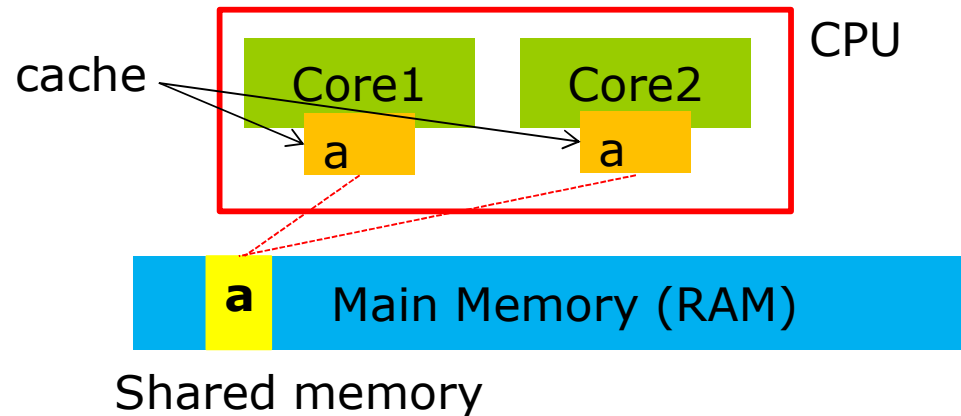
```
item nextConsumed;  
while (true)  
{  
    while (in == out)  
        ; // do nothing --nothing to consume  
        // until remove an item from the buffer  
  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
}
```

Make sure buffer
is not empty




Shared Memory Systems

- **Bigger problem:** What happens if both producer process and the consumer process attempt to access the shared buffer concurrently?
 - Process synchronization
- **Another problem:** Multicore processors
 - Each core have separate cache – cache coherence problem



Interprocess Communication – Message Passing

- Message system – processes communicate with each other by **messages** without resorting to shared variables
- IPC facility provides two operations:
 - **send(message)** **receive(message)**
- If P and Q wish to communicate, they need to:
 - establish a **communication link** between them
 - exchange messages via send/receive
 - ▶ need **system call**
 - ▶ More time-consuming compared to shared memory
 - e.g., **Microkernel structure**
 - e.g., **Sockets** (networking), RPC (distributed systems)



Message Passing Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous** and **Non-blocking** is considered **asynchronous**
- **Send**
 - **Blocking** send has the sender block until the message is received
 - **Non-blocking** send has the sender send the message and continue
- **Receive**
 - **Blocking** receive has the receiver block until a message is available
 - **Non-blocking** receive has the receiver receive a valid message or null