


Chapter 6: Synchronization

School of Computing, Gachon Univ.
Jungchan Cho



Most slides from "Operating System Concepts – 10th Edition".
Many slides are taken from lecture notes of Prof. Joon Yoo.

Chapter 6: Synchronization

- Background
- The Critical-Section Problem
- Software C.S.: Peterson's Solution
- Hardware C.S.: Mutex Locks
- Semaphores
- Classic Problems of Synchronization


Objectives

- To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- To present both software and hardware solutions of the critical-section problem
- To examine several classical process-synchronization problems

Background

■ Resource sharing

- memory, files
 - ▶ interprocess communication (IPC)
 - ▶ multi-threads



■ Why process synchronization?

- Concurrent or parallel access to shared data may result in data inconsistency

- ▶ Recall the multithread example in Ch. 4:

- **Solution to the problems**

: Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

- ▶ **Process synchronization** mechanisms

Output:

x is 2
x is 3

Output:

x is 3
x is 2

Output:

x is 3
x is 3

Output:

x is 2
x is 2

Possible output?

- Shared code:

```
int x = 1; //global variable
void* func(void* p) {
    x = x + 1;
    printf("x is %d\n", x);
    return NULL;
}
```

- fork version:

```
main(...) {
    fork();
    func(NULL);
}
```

- threads version:

```
main(...) {
    pthread_t tid;

    pthread_create(&tid, NULL, func, NULL);
    func(NULL);
}
```

Possible output: threads case 1

```
int x = 1; //global variable
```


```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Parent thread

Child thread

time



Possible output: threads case 2

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;
```

```
    printf("x is %d\n", x);  
    return NULL;
```

```
}
```

Parent thread

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Child thread

time



Possible output: threads case 3

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
  
    // interrupted during printf()  
  
    printf("x is %d\n", x);  
  
    return NULL;  
}
```


Parent thread

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Child thread

time

Example: ATM Bank Server



- ATM (현금인출기)
 - Balance (잔고)
 - Withdrawal (인출)
 - Deposit (입금)

Example: ATM Bank Server


- Single transaction

- Withdraw



1. Check balance : x

ATM Bank Server



2. Withdraw : m



4. Update: $x \leftarrow x'$

3. Update balance : $x' \leftarrow x - m$

- Deposit

- $x \leftarrow x + m$

Example: ATM Bank Server

Concurrent transactions

Withdraw

W1. Check balance : x

W2. Withdraw : m

W3. Update balance : $x' \leftarrow x - m$

ATM Bank Server

Update: $x \leftarrow x'$

Deposit

D1. Check balance : x

D2. deposit : n

D3. Update balance : $x' \leftarrow x + n$

Order of execution

W1 → W2 → W3 → D1 → D2 → D3 : final $x = ?$


vs.

W1 → W2 → D1 → D2 → W3 → D3 : final $x = ?$

Example2: Producer-Consumer problem

■ Producer-consumer problem

- Producer threads
 - ▶ Set of threads that generates messages
- Consumer threads
 - ▶ Set of threads that consumes messages



Example3: Producer-Consumer problem


`int counter=0; // global variable` ← Number of items in buffer

PRODUCER thread

```
item  nextProduced;  
  
while (TRUE) {  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

CONSUMER thread

```
item  nextConsumed;  
  
while (TRUE) {  
    while (counter == 0);  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```



C code → Assembly code

- Note that `counter++;` ← Actually, this C code is compiled into 3 assembly codes:

<code>register₁ = counter</code>	<code>#load from memory</code>
<code>register₁ = register₁ + 1</code>	
<code>counter = register₁</code>	<code>#store in memory</code>
- `counter--;` ← is compiled into:

<code>register₂ = counter</code>	<code>#load from memory</code>
<code>register₂ = register₂ - 1</code>	
<code>count = register₂</code>	<code>#store in memory</code>
- Consider this execution with “count = 5” initially. **Producer** adds item (counter++) then **Consumer** deletes an item (counter--). What is the result?

T ₀ :	producer execute	<code>register₁ = counter</code>	{register ₁ = 5}
T ₁ :	producer execute	<code>register₁ = register₁ + 1</code>	{register ₁ = 6}
T ₄ :	producer execute	<code>counter = register₁</code>	{counter = 6}
T ₂ :	consumer execute	<code>register₂ = counter</code>	{register ₂ = 6}
T ₃ :	consumer execute	<code>register₂ = register₂ - 1</code>	{register ₂ = 5}
T ₅ :	consumer execute	<code>count = register₂</code>	{ <u>counter = 5</u> }

Race Condition

- But not always!
- Consider this execution with “count = 5” initially. **Producer** adds item (count++) then **Consumer** deletes an item (count--). What is the result?
- Interleaving instructions

Context switch
P1→P2

T₀: producer execute $\text{register}_1 = \text{counter}$ {register₁ = 5}

Context switch
P2→P1

T₁: producer execute $\text{register}_1 = \text{register}_1 + 1$ {register₁ = 6}

T₂: consumer execute $\text{register}_2 = \text{counter}$ {register₂ = 5}

T₃: consumer execute $\text{register}_2 = \text{register}_2 - 1$ {register₂ = 4}

Context switch
P1→P2

T₄: producer execute $\text{counter} = \text{register}_1$ {counter = 6}

T₅: consumer execute $\text{counter} = \text{register}_2$ {counter = 4} ↕ ?

- Result (counter) can be either 4, 5 or 6!!
 - This is called race condition

Race Condition

- Although both the producer and consumer routines are correct separately, they may not function correctly when executed concurrently (or in parallel)
- **Race Condition:**
 - Several processes access and manipulate the same data concurrently
 - The outcome of the execution depends on the particular order in which the access takes place
- To prevent race conditions, concurrent processes must be synchronized




But why does Race Condition happen?

- Processes can execute **concurrently** or in **parallel**


- Concurrent** execution: thread scheduling

- CPU scheduler switches rapidly between threads for concurrent execution
- One thread may only partially complete execution before another thread is scheduled – interrupt



- Parallel** execution: multicore

- Two threads execute simultaneously on separate processing cores
- Each thread can share data



- Race condition!**

Enforcing mutual exclusion

- How can we avoid the race condition?
- Answer: We must **synchronize** the execution of the threads
 - i.e., need to guarantee **mutually exclusive access** to **critical sections**
- Classic solution
 - Software C.S.: Peterson's Solution
 - Hardware C.S.: Mutexes Locks
 - Semaphores (Edsger Dijkstra)

Chapter 6: Synchronization

- Background
- The Critical-Section Problem
- Software C.S.: Peterson's Solution
- Hardware C.S.: Mutex Locks
- Semaphores
- Classic Problems of Synchronization

Critical-Section Problem

- N processes all competing to use some shared data
- Each process has a code segment, called **critical section (C.S.)**, in which the **shared data** *is accessed*.
 - Only one process execute in its critical section at a time
 - ▶ ensure that when one process is executing in its **critical section**, no other processes are allowed to execute in its critical section
 - **Critical section**
 - ▶ a piece of code that accesses a shared resource (e.g., data structure or device)
- The critical-section problem
 - Design a protocol that the processes can use to cooperate

Critical-Section Problem

- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

do {

entry section

critical section

exit section

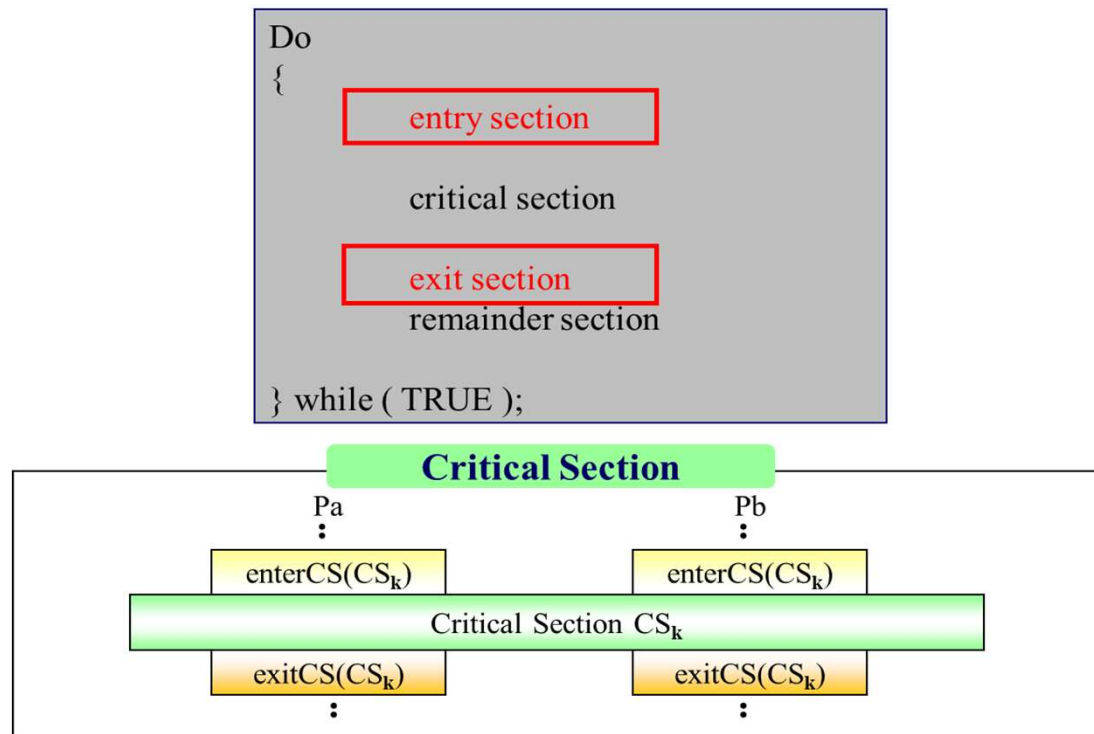
remainder section

} while (true);

No one else
can use the
CS

Critical-Section Problem

- General Structure of Critical Section



Example

- In the producer-consumer problem,

PRODUCER

```
item  nextProduced;

while (TRUE) {
    while (counter == BUFFER_SIZE);
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;

    counter++;

}
```

CONSUMER

```
item  nextConsumed;

while (TRUE) {
    while (counter == 0);
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    counter--;

}
```

Critical section

Example – Lock & Unlock

- In the producer-consumer problem,

```
static pthread_mutex_t cs_mutex = PTHREAD_MUTEX_INITIALIZER;
```

PRODUCER

```
item  nextProduced;

while (TRUE) {
    while (counter == BUFFER_SIZE);
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
```

```
    pthread_mutex_lock( &cs_mutex );
```

```
    counter++;
```

```
    pthread_mutex_unlock( &cs_mutex );
```

```
}
```

CONSUMER

```
item  nextConsumed;

while (TRUE) {
    while (counter == 0);
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
```

```
    pthread_mutex_lock( &cs_mutex );
```

```
    counter--;
```

```
    pthread_mutex_unlock( &cs_mutex );
```

```
}
```

entry

exit

Critical section

No one else
can use the
CS

C.S. Requirements *

1. Mutual Exclusion

If a process P_i is executing in its C.S., then no other processes can be executing in their C.S.

2. Progress

If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the process that will enter the critical section next cannot be postponed indefinitely – **deadlock free**



3. Bounded Waiting

Each process should be able to enter its C.S. after a finite number of trials – **starvation free**

Chapter 6: Synchronization

- Background
- The Critical-Section Problem
- **Software C.S.: Peterson's Solution**
- Hardware C.S.: Mutex Locks
- Semaphores
- Classic Problems of Synchronization

C.S. Algorithm

- Shared variables
 - array wants [] : true or **false**
 - ▶ wants[i] indicates if P_i wants to enter C.S.

do {

wants[*i*] = true;

while (wants[*j*]) {;}

CRITICAL SECTION

wants[*i*] = false;

REMAINDER SECTION

} while (TRUE);

Algorithm (cont.)

- Consider two P_i where $i = \{1, 2\}$, i.e., P_1, P_2

```
do {  
    wants[ i ] = true;  
    while (wants[ j ]) {;}  
    CRITICAL SECTION  
    wants[ i ] = false;  
    REMAINDER SECTION  
} while (TRUE);
```

• Process P_1

```
do {  
    wants[ 1 ] = true;  
    while (wants[ 2 ]) {;}  
    CRITICAL SECTION  
    wants[ 1 ] = false;  
    REMAINDER SECTION  
} while (TRUE);
```

• Process P_2

```
do {  
    wants[ 2 ] = true;  
    while (wants[ 1 ]) {;}  
    CRITICAL SECTION  
    wants[ 2 ] = false;  
    REMAINDER SECTION  
} while (TRUE);
```

- What happens if both $wants[1]$ and $wants[2]$ are true?

Peterson's Solution

- A **software-based** solution
 - Solution to CS problem w/o H/W support
- Suppose two processes : $P_i (=P_0)$, $P_j (=P_1)$
The two processes share two variables:
 - bool **wants**[i]; – wants[i] indicates if P_i wants to enter C.S.
 - int **not_turn**; – not this thread's turn to enter C.S. Other threads can enter C.S. if they want to...

```
for (;;) { /* assume i is thread number (0 or 1) */
    wants[i] = true;
    not_turn = i;          /* "Not my turn, you go first!" – yield (양보) */
    while (wants[j] && not_turn == i)
        ; /* other thread wants in and not our turn, so loop */
    CRITICAL SECTION
    wants[i] = false;      /* I'm finished, so others can enter now */
    REMAINDER SECTION
}
```

Algorithm for Process P_i

```
while (true) {  
    // entry section  
    wants[i] = TRUE;  
    not_turn = i;  
    while ( wants[j] && not_turn == i);
```

CRITICAL SECTION

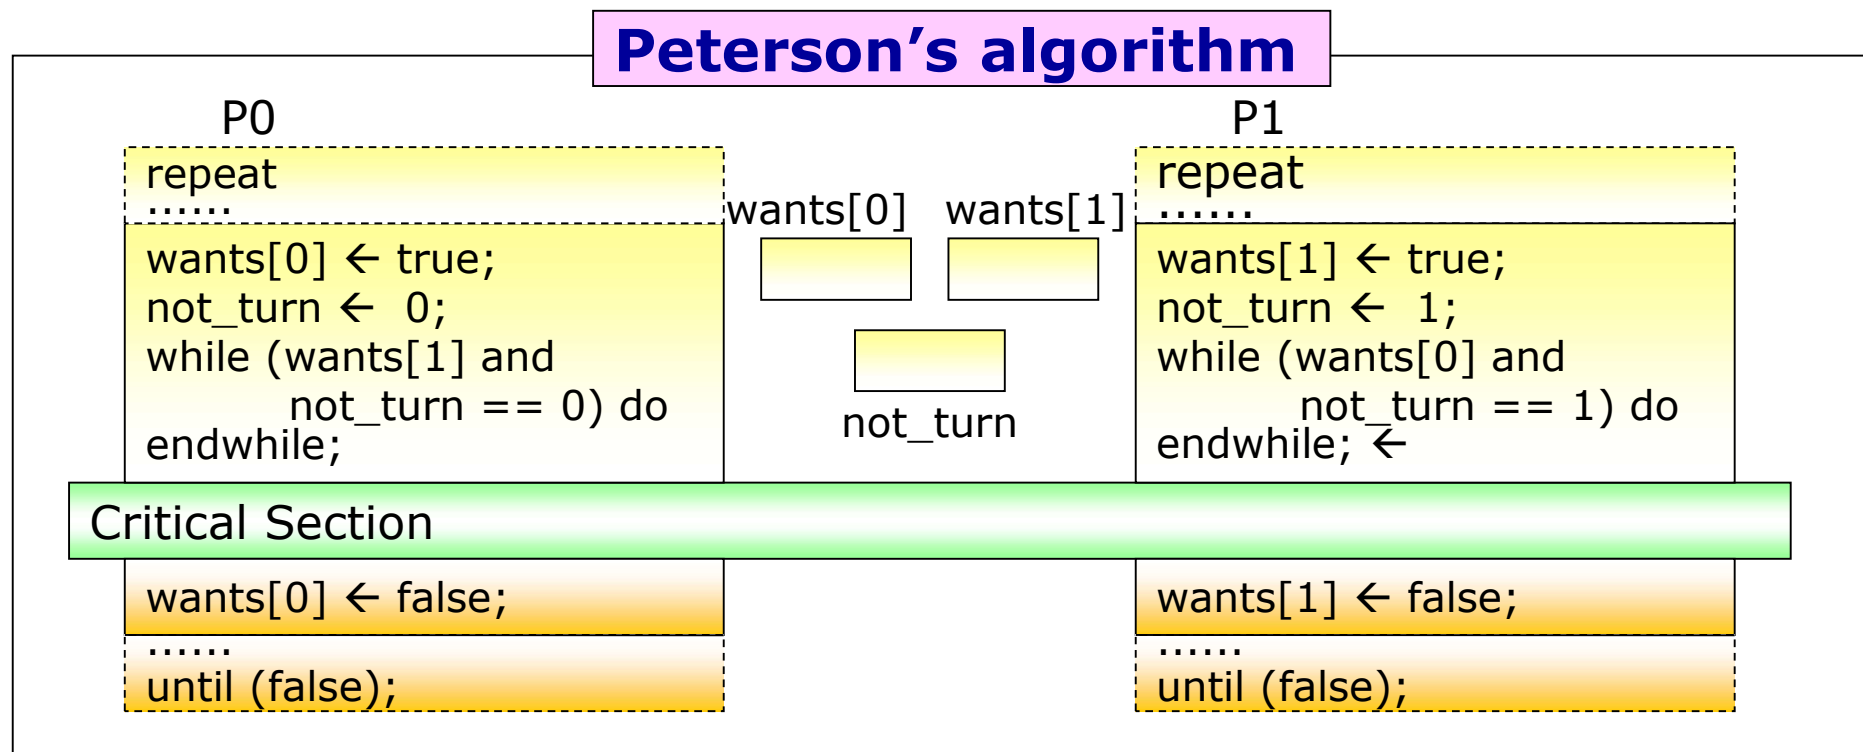
```
    // exit section  
    wants[i] = FALSE;
```

REMAINDER SECTION

```
}
```

Peterson's algorithm

- Peterson's algorithm



Does Peterson's solution work?

```
for (;;) { /* assume i is thread number (0 or 1) */
    wants[i] = true;
    not_turn = i;
    while (wants[j] && not_turn == i)
        ; /* other thread wants in and not our turn, so loop */
    CRITICAL SECTION
    wants[i] = false;
    REMAINDER SECTION
}
```

■ Mutual exclusion?

- can't both be in C.S. - Would mean `wants[0] == wants[1] == true` and – both cannot be in critical section!
- Also, **not_turn** would have blocked one thread from C.S.

Does Peterson's solution work?

```
for (;;) { /* assume i is thread number (0 or 1) */
    wants[i] = true;
    not_turn = i;
    while (wants[j] && not_turn == i)
        ; /* other thread wants in and not our turn, so loop */
    CRITICAL SECTION
    wants[i] = false;
    REMAINDER SECTION
}
```

■ Progress (Deadlock free)? Bounded–waiting?

- Can P_i be stuck in while-loop ($wants[j] == true \ \&\& \ not_turn == i$) forever?
 - ▶ Case1: P_j does not want to enter C.S then $wants[j] == false$
 - ▶ Case2: P_j wants to enter C.S then $wants[j] == true$, but $not_turn == i$ or j
 - if $not_turn == i$ then j enters C.S, if $not_turn == j$ then i enters C.S.
- Since P_i does not change the value of the 'not_turn' while executing the while statement, P_i will enter the CS (progress) after at most one entry by P_j (bounded waiting).
- Peterson's solution considers strict alternation so, alternatively P_i and P_j will get access to critical section.

Chapter 6: Synchronization

- Background
- The Critical-Section Problem
- Software C.S.: Peterson's Solution
- **Hardware C.S.: Mutex Locks**
- Semaphores
- Classic Problems of Synchronization

Hardware Approach

- Peterson's algorithm is expensive, only works for 2 processes
- Peterson's algorithm is a software solution
 - Low speed
- Many systems provide **hardware support** for C.S.
 - All solutions below based on idea of protecting critical regions via **locks**
- Modern machines provide special atomic hardware instructions
 - ▶ **Atomic** = non-interruptible

Note

Characteristics of a machine instruction

- Atomicity, indivisibility

(No interrupt during the execution of a machine instruction)

Synchronization Hardware: TestAndSet

- Mutual exclusion with **TestAndSet()** instruction

Initially
lock = FALSE;

```
do {  
    while ( TestAndSet (&lock ) )  
        ; // do nothing
```

CRITICAL SECTION

```
    lock = FALSE;
```

REMAINDER SECTION

```
} while ( TRUE);
```

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

**Atomic (indivisible) operation
Supported by H/W**

Uninterruptible
Machine Instructions

Mutex Locks

- Hardware based solutions are complex and unavailable to the application programmer
- **Mutex lock**
 - A simple software tool to solve the C.S problem
- Product critical regions with it by first **acquire()** a lock then **release()** it

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
release() {  
    available = true;  
}
```

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

Mutex Locks are
provided as **API**

Hardware Approach: Mutex Locks

- Calls to **acquire()** and **release()** must be atomic
 - Usually implemented via hardware atomic instructions

▶ e.g., TestAndSet

```
// initialization
mutex->available = 0;

// acquire using test_and_set()
void acquire(lock *mutex) {
    while (test_and_set(&mutex->available) != 0)
        ;
    return;
}

void release(lock *mutex) {
    mutex->available = 0;

    return;
}
```

- But this solution requires busy waiting

***Busy waiting** is a process synchronization technique in which a process/task waits and constantly checks for a condition to be satisfied before proceeding with its execution.

Example Code (POSIX pthread)

- Example Code For Critical Sections with POSIX pthread library

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

Busy waiting in Mutex

■ Busy waiting

- loop continuously while waiting
- (=spinlock): spins while waiting for the lock to become available



```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

■ Cons

- Can do nothing while waiting
- Waste CPU cycles – some other process could have used it

■ Alternative: Process goes to “waiting” state and context switch to another process

■ Pros

- No context switch is required during spinlock
- Useful when locks are expected to be held for short times
 - ▶ tradeoff: context switch time (waiting) vs. spinlock time

Chapter 6: Synchronization

- Background
- The Critical-Section Problem
- Software C.S.: Peterson's Solution
- Hardware C.S.: Mutex Locks
- **Semaphores**
- Classic Problems of Synchronization

Semaphore



- Proposed by Dijkstra in 1965
- Semaphore **S** –integer variable
- Two atomic standard operations modify S: **wait()** and **signal()**

- 1972 Turing Award
- ACM Dijkstra Prize
- Known for Dijkstra Algorithm, Semaphore

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal (S) {  
    S++;  
}
```

will see other
version later

Compare this
with Mutex Lock

- Only one process can modify the semaphore value (S)
 - $S \leq 0$, $S--$ and $S++$ are atomic operations – hardware support

Semaphore

- Semaphore provides more sophisticated ways than mutex lock for synchronization
- **Binary semaphore**
 - The semaphore can be set to 0 or 1, i.e., $S = 0$ or 1
 - Same as **Mutex locks**
- **Counting semaphore**
 - The semaphore can initially have nonnegative integer values, i.e., the ***number of resources available***
 - Used for solving producer-consumer problems and etc.

Binary Semaphore

- Provides mutual exclusion

- Semaphore S – shared by all processes

- Process P_i

Semaphore S ; // initialized to 1

```
do {
    wait (S);

    //Critical Section

    signal (S);

    //Remainder Section

} while (true);
```

- N processes shares semaphore S
- S is initialized to 1 (# of resources?)

P1	S	P2
wait(S) {	1	
while(S<=0) ; //pass		
S--; }	0	
//Critical Section		wait(S) {
		while(S<=0) ; //wait
signal(S) {		
S++; }	1	//pass
//Remainder Section	0	S--; }
		//Critical Section
	1	signal(S) {
		S++; }
		//Remainder Section

Q: what happens if there are 2 resources?

counting semaphore

Problem 1: Deadlocks and Starvation

■ Deadlock

- two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0

wait (S);
wait (Q);
.
.
.
signal (S);
signal (Q);


P_1

wait (Q);
wait (S);
.
.
.
signal (Q);
signal (S);

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal (S) {  
    S++;  
}
```

Problem 2: Priority Inversion

- Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Example: 3 process **L**, **M**, **H** with scheduling priorities $L < M < H$ (**L** has lowest priority, **H** has highest)




- Resource **R** is protected by a Semaphore
- Indirectly, Process **M** has affected Process **H**'s waiting time for resource **R**! – **Priority Inversion!**

▶ **M** → **L** → **H**

Here, we can see that running of **M** has delayed the running of both **L** and **H**. Precisely speaking, **H** is of higher priority and doesn't share **R** with **M**; but **H** had to wait for **M**.

Problem 2: Priority Inversion

- Priority Inversion solved via **priority-inheritance protocol**
 - A process (process L) that is accessing resources needed by a higher-priority process (process H) inherits the higher priority (H) until they are finished with the resources



Problem 3: Incorrect Use of Semaphores

- Correct use of semaphore operations may not be easy.
- Suppose semaphore variable called mutex is initialized to 1.
- Incorrect use of semaphore operations:
 - 1) signal (mutex) wait (mutex)
 - ▶ What happens?
 - 2) wait (mutex) ... wait (mutex)
 - ▶ What happens?
 - 3) Omitting of wait (mutex) or signal (mutex) (or both)
 - ▶ **Mutual exclusion is violated (1) or deadlock will occur (2) or both (3)**

```
do {  
    wait (S);  
    //Critical Section  
    signal (S);  
    //Remainder Section  
} while (true);
```

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal (S) {  
    S++;  
}
```


Semaphore with Block Operation

- Avoid **busy waiting** (spinlock)
- **block()**
 - If $S \leq 0$ then wait using block
 - Instead of using busy waiting, the process is placed into **waiting queue** (process state?)
 - CPU scheduler selects another process (in ready queue) to execute
- **wakeup()**
 - The blocked process restarts when some other process executes a **signal()** operation
 - The blocked process is moved from waiting queue to ready queue

Counting Semaphore with Block Operation

- Implementation of *wait()*:

```
wait (S) {  
    S.value--;  
    if (S.value < 0) {  
        add this process to waiting queue  
        block();  
    }  
}
```

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore
```

- Implementation of *signal()*:

```
signal (S) {  
    S.value++;  
    if (S.value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P);  
    }  
}
```

Pthreads Semaphore

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```

Chapter 6: Synchronization

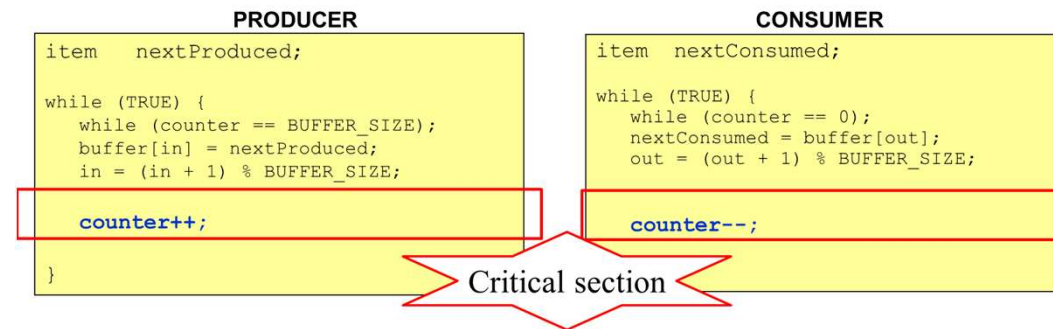
- Background
- The Critical-Section Problem
- Software C.S.: Peterson's Solution
- Hardware C.S.: Mutex Locks
- Semaphores
- Classic Problems of Synchronization

Classical Problems of Synchronization


- Semaphore-based concurrent programming
 - Bounded-Buffer Problem
 - Reader-Writer problem
 - Dining philosopher problem
 - Etc.

The Bounded-Buffer Problem

- Producer processes
- Consumers processes
- **N** size buffer (Limited buffer)
 - Buffer can hold up to **N** items



- Solutions with semaphore – use 3 semaphores
 - Semaphore **mutex** = 1 // critical section
 - Semaphore **full** = 0 // number of items in buffer
 - Semaphore **empty** = N // number of empty slots in buffer
 - ▶ full + empty = N



Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do
{
    // Produce an Item

    wait (empty);           // Check if buffer is full
    wait (mutex);           // Enter into Critical Section


    // Critical Section
    // add an item to buffer

    signal (mutex);         // Leave C.S.
    signal (full);          // Produce an Item

} while (TRUE);
```

```
wait (S) {
    while (S <= 0)
        ; // busy wait
    S--;
}

signal (S) {
    S++;
}
```



Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
do  
{
```


```
    wait (full);           // Check if buffer is empty  
    wait (mutex);         // Enter into Critical Section
```

```
    // Critical Section  
    // remove an item from buffer
```

```
    signal (mutex);        // Leave C.S.  
    signal (empty);        // Consume an Item
```

```
} while (TRUE);
```

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal (S) {  
    S++;  
}
```



Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write.
- Problem
 - allow **multiple readers** to read at the same time.
 - **One single writer** can access the shared data at the same time.
- Shared Data
 - Data set
 - Semaphore **mutex** initialized to 1 // critical section
 - Semaphore **wrt** initialized to 1. // mutual exclusion for writers
 - Integer **readcount** initialized to 0. // number of readers

To protect readcount!

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do
{
    wait (wrt);

    // writing is performed

    signal (wrt);

} while (TRUE);
```

```
wait (S) {
    while (S <= 0)
        ; // busy wait
    S--;
}

signal (S) {
    S++;
}
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait (mutex);  
    readcount ++;  
    if ( readcount == 1 )  
        wait(wrt);  
    signal(mutex)  
  
    // reading is performed  
  
    wait(mutex);  
    readcount --;  
    if ( readcount == 0 )  
        signal(wrt);  
    signal (mutex);  
} while (TRUE);
```

writer process

```
do  
{  
    wait (wrt);  
  
    // writing is performed  
  
    signal (wrt);  
  
} while (TRUE);
```

Conclusion

- Processes in multiprogramming systems
 - Asynchronous, concurrent
 - Needs mutual exclusion (ME) and process synchronization mechanisms
- SW solution for ME and synchronization
 - Peterson's algorithm
- HW solution for ME and synchronization
 - Test-and-Set instruction (atomic)
- Semaphore with no busy waiting