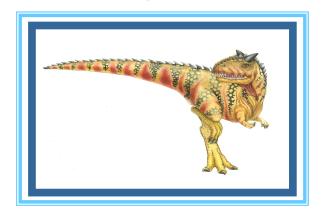
Chapter 7: Deadlocks_2

School of Computing, Gachon Univ. Jungchan Cho



Most slides from "Operating System Concepts – 10th Edition". Many slides are taken from lecture notes of Prof. Joon Yoo.



Chapter 7: Deadlocks

- Deadlock Concept
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Avoidance



Methods for Handling Deadlocks

- Ensure that the system will never enter a deadlock state
 - Deadlock Prevention
 - Deadlock Avoidance
- Allow the system to enter a deadlock state and then recover
 - Deadlock Detection
 - Deadlock Recovery
- Ignore the problem and pretend that deadlocks never occur in the system (Ostrich Algorithm)
 - actually used by most operating systems, including Linux and Windows
 - up to the application developer to handle deadlocks





Chapter 7: Deadlocks

- Deadlock Concept
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Avoidance



Deadlock Prevention

 For deadlock to occur, each of the four necessary conditions (mutual exclusion, hold and wait, no preemption, circular wait) must hold

Deadlock prevention

 Ensure at least one of above conditions cannot hold – deadlock is impossible!!

1. No Mutual Exclusion



Practically impossible. Some resources are non-sharable (e.g., mutex lock, printers)



OK Cancel



- 4 necessary conditions for Deadlock
 - 1. Mutual exclusion
 - 2. Hold and wait
- 3. No preemption
- 4. Circular wait

2. No Hold and Wait



- must guarantee that whenever a process requests a resource, it does not hold any other resources
- Total allocation: process requests and be allocated all its resources before it begins execution
 - e.g., A process copies data from <u>DVD drive</u> to a file on <u>disk</u>, sort the file, and then prints the results to a <u>printer</u>
 - ▶ The process must initially request the DVD drive, disk file, and printer
 - If all resources are not available must wait
 - hold all resources until process terminates; then release all resources

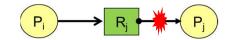


- 4 necessary conditions for Deadlock
 - 1. Mutual exclusion
 - 2. Hold and wait
- 3. No preemption
- 4. Circular wait
- Alternative protocol: process can request resources only when it has none
- Problems of No hold and wait
 - Low resource utilization
 - Resources may be allocated but unused for a long period total allocation
 - Starvation
 - A process that needs several popular resources may have to wait indefinitely –
 at least one popular resource it needs is always allocated to some other process



- 4 necessary conditions for Deadlock
- 1. Mutual exclusion
- 2. Hold and wait
- 3. No preemption
- 4. Circular wait

3. Allow Preemption



- A lock can be taken away (preempted) from current owner
- Often applied to resources whose state can be easily saved and restored
- Example: preemptive CPU scheduling RR
 - ▶ P1 is using (holding) CPU, P2 is waiting in ready queue
 - ▶ P1's time quantum has expired preempt resource (CPU) from P1.
 - A preempted process can restart only if it can restore previous state: Save P1's register to PCB1, restore P2's register from PCB2
- Problem: OK for CPU but Basically impossible for some resources
 - ▶ E.g., Printers, mutex locks, semaphores



- 4 necessary conditions for Deadlock
- 1. Mutual exclusion
- 2. Hold and wait
- 3. No preemption
- 4. Circular wait

4. No Circular Wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration
- Example

$$F(disk drive) = 10$$

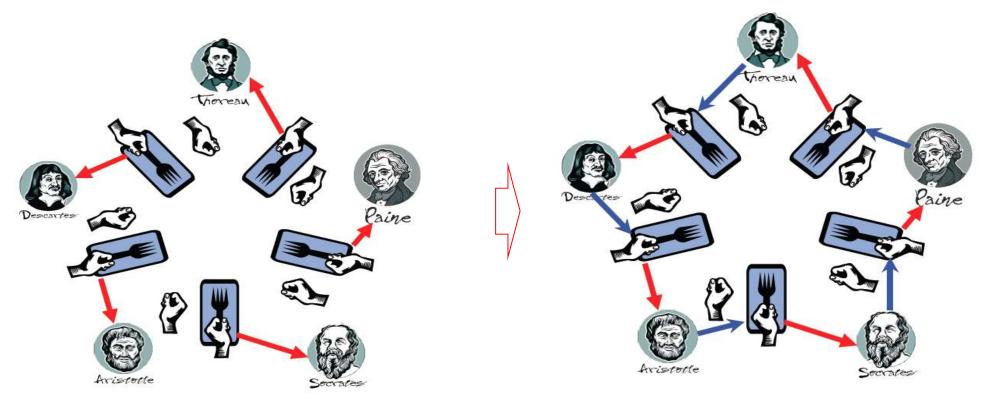
 $F(printer) = 20$
 $F(tape drive) = 30$

- Each process can request resources only in an increasing order of enumeration a process requests R_i . Then, it request instances of resource type R_i if and only if $F(R_i) > F(R_i)$.
- Circular wait can never happen



Example of Denying Circular Wait

Dining Philosophers



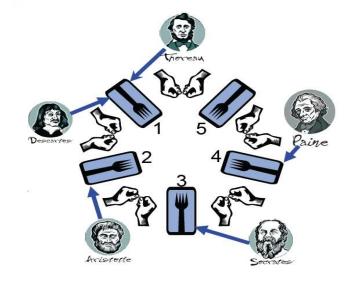


Example of Denying Circular Wait

Dining Philosophers

```
void philosopher (int i) {
   while (TRUE) {
      think();
      take fork(i);
      take fork((i+1)%N);
      eat(); /* yummy */
      put fork(i);
      put fork((i+1)%N);
```

```
void philosopher (int i) {
   while (TRUE) {
      think();
      take_fork(LOWER(i));
      take_fork(HIGHER(i));
      eat(); /* yummy */
      put_fork(LOWER(i));
      put_fork(HIGHER(i));
   }
}
```





Deadlock Prevention

- Summary of prevention schemes
 - Deny a (one) necessary condition for deadlocks among the 4 necessary conditions
 - Serious resource waste
 - ▶ Low device utilization, Reduced system throughput
 - May incur high cost to redesign the system
- Deadlock avoidance!



Chapter 7: Deadlocks

- Deadlock Concept
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Avoidance



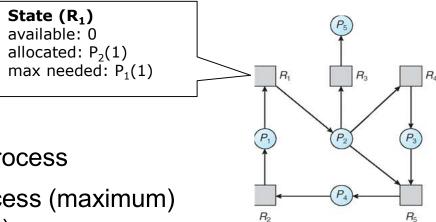
Deadlock Avoidance

- Require additional information about how resources are to be requested
 - System knowledge of complete sequence of requests and releases of resources
 - Example: P1 requests first hard disk then printer. Then release both resources,
 P2 requests first printer then hard drive.
 - System can decide whether the process should wait in order to avoid possible deadlock
- Construct algorithm so that circular-wait can <u>never</u> exist
- What Information (state) is required from each process?
 - Process declares maximum number of resources of each type it needs
 - ▶ P1: I want to use 2 hard disk instances



System Model

- What information in state?
 - For each resource type state,
 - 1. Current amount available instances
 - ▶ 2. Current amount allocated to each process
 - 3. Future amount needed by each process (maximum) (declared by each process – prev. slide)



Assumptions

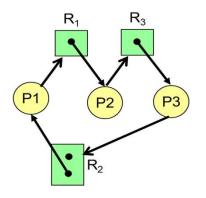
- Assume OS knows: current state of the system!, i.e.,
 - 1. Number of available instances of each resource
 - ▶ 2. For each process, *current* amount of each resource it owns
 - ▶ 3. For each process, *maximum* amount of each resource it needs in future

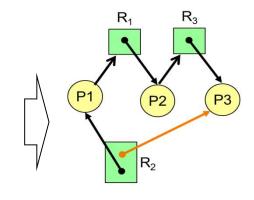


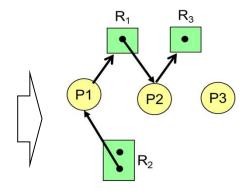
System Model (cont.)

- Assume worst case...
 - A process will never release its resources until it's done! (non-preemptive) processes are independent
 - While one blocks, others can finish if they have "enough resources"
 - After finishing, each process releases the resources it owns within finite time
 - request use release!!

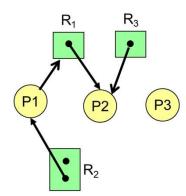
Example











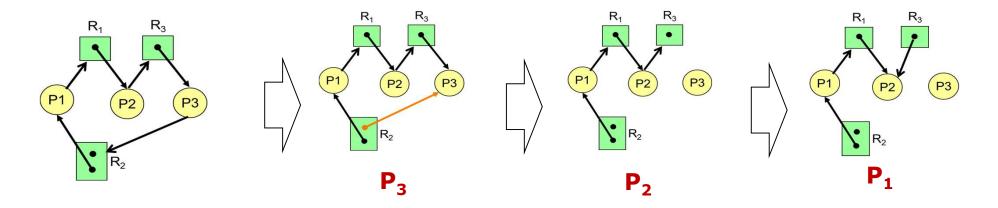


Safe Sequence

Safe = there's definitely a way to finish the processes without deadlock

Safe Sequence

 Sequence (order) with which the system can allocate resources to each process and avoid a deadlock





Safe sequence

- Sequence <P₁, P₂, ..., P_n> is safe if for each P_i, the resources that P_i can request can be satisfied by currently available resources + resources held by other P_i, with j < i.</p>
 - $P_1, ..., P_j, ..., P_i, P_{i+1}, ... >$
 - If P_i resource needs are not immediately available, then P_i can wait until all P_i have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on.
- The system is in safe state only if there exists a safe sequence



Example: Safe State

- We consider <u>a system that has 12 tape instances</u> and 3 processes: P₀, P₁, P₂.
 - P_0 needs 10 tapes, P_1 needs 5 and P_2 needs 9.
 - Currently (at time t₀), P₀ has 5, P₁ has 2 and P₂ has 2

How many tape instances are available?

Q : Is the system in a safe state at t_0 ?

= Is there **any** execution order that can finish without deadlock?

A process might request its maximum resources at any time

	Max needs	Current holding
$\overline{P_0}$	10	5
P_1	5	2
P_2	9	2
P ₁ P ₂	5 9	2 2

A process will never release its resources until it's done



Example: Safe State (cont.)

Q : Is the system in a safe state at t_0 ?

= Is there **any** execution order that can finish without deadlock?

12 tape instances	Max needs	Current holdings
$\overline{P_0}$	10	5
P_1°	5	2
P_2	9	2

Solution ?:

Search for an order P_i, P_{i+1}, P_{i+2}, ... such that:

for each P_i , the resources that P_i can request can be satisfied by currently available resources + resources held by other P_j , with j < i.



Example: Safe State (cont.)

Same example as before: <u>12 tape instances</u> (available = 3)

	Max needs	Current holdings
P_0	10	5
P_1	5	2
P_2	9	2

- Q : Is the system in a safe state at t₀?
 - A: Yes, since <P₁,P₀,P₂> satisfied safe state condition
- Q : At t₁, if P2 requires and is allocated one more tape by the system, is the system is in a safe state?
 - A: No, since there is no safe sequence
 - See how many instances are available.



Deadlock Avoidance

Safe state vs. unsafe state

Definition: safe state

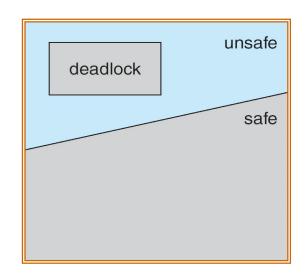
A state is *safe* if the system can allocate resources to each process in some order and still avoid a deadlock. A system is in a safe sate only if there exist one or more **safe sequences**Unsafe state: a state that is not safe

- Meaning of safe state
 - Can guarantee that the system does not get into the deadlock state
- Meaning of unsafe state
 - ✓ May not avoid getting into the deadlock state
 - ✓ conversely, Deadlock → unsafe state
 - ✓ Does not necessarily mean that deadlock will occur in the future



Basic Facts; Safe, Unsafe, Deadlock

- If a system is <u>in safe state</u> ⇒ no deadlocks
- If a system is <u>in unsafe state</u> ⇒ possibility of deadlock.



■ Avoidance ⇒ ensure that a system will never enter an unsafe state.



Deadlock Avoidance Methods

Single instance of a resource type
 ⇒ Use a Resource-Allocation Graph

- Multiple instances of a resource type
 - ⇒ Use the Banker's algorithm



Resource-Allocation Graph Algorithm

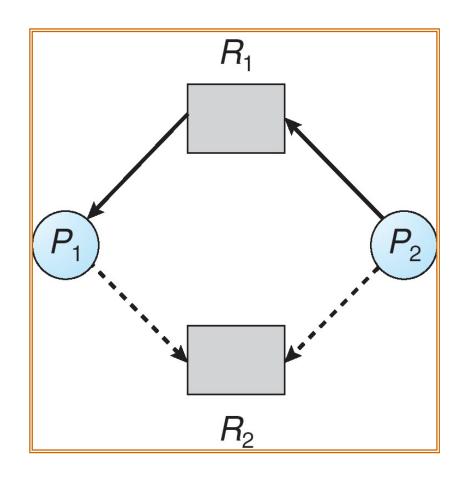
Claim edge P_i → R_j indicates that process P_i may request resource R_j
(at some time in the future); represented by a dashed line.



- Claim edge is converted to a request edge when a process actually requests a resource.
- An <u>assignment</u> edge P_i ← R_j is reconverted to a <u>claim</u> edge when a resource is released by a process.



Resource-Allocation Graph Algorithm



P_i requests resource R_j - request is granted *only if*

converting request edge $P_i \rightarrow R_j$ to assignment edge $R_i \rightarrow P_i$ does not result in **cycle**

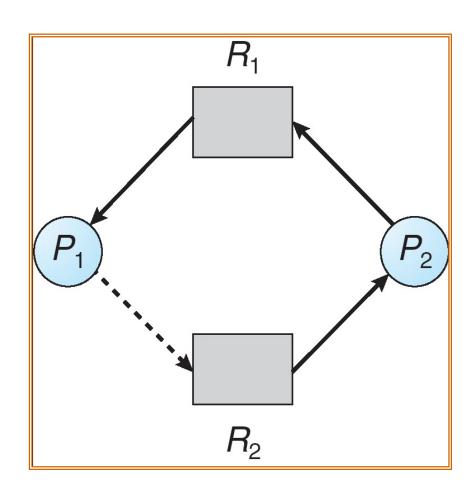
Example: allocation sequence

$$P_1, P_2 \rightarrow YES$$

Safe state!



Resource-Allocation Graph Algorithm



- ➤ Allocation sequence
 - \blacksquare P₂, P₁ \rightarrow NO
 - Unsafe state!

Banker's Algorithm

- A resource allocation and deadlock avoidance algorithm
- <u>Multiple</u> instances
- Assumptions
 - Each process must a priori claim maximum use.
 - When a process requests a resource, it may have to wait.
 - When a process gets all its resources, it must return them in a finite amount of time.
- Consists of two sub-algorithms
 - Safety Algorithm
 - Resource-Request Algorithm



Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- Available: Vector of length m.
 - If Available [j] = k, there are k instances of resource type R_j available.



- Max: n x m matrix.
 - If Max[i, j] = k, then process P_i may request at most k instances of resource type R_{i-1}
- Allocation: n x m matrix.
 - If Allocation[i, j] = k then P_i is currently allocated k instances of R_j.
- Need: n x m matrix.
 - If Need[i, j] = k, then P_i may need k more instances of R_i to complete its task.

Need
$$[i, j] = Max[i, j] - Allocation [i, j]$$
.



- Algorithm for finding out whether a system is in a safe state
 - 1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

Work = Available
Finish
$$[i]$$
 = false for i = 0, 1, 2, ..., n -1.

- 2. Find and *i* such that both:
 - (a) Finish [i] == false
 - (b) $Need_i \leq Work$

If no such *i* exists, go to step 4.

- 3. Work = Work + Allocation; Finish[i] = true go to step 2.
- 4. If *Finish* [*i*] == true for all *i*, then the system is in a safe state.



- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	ABC	ABC	ABC
P_0	010	753	3 3 2
P_1	200	3 2 2	
P_2	302	902	
P_3	211	222	
P_4	002	4 3 3	



The content of the matrix.

Need is defined to be *Max* – *Allocation*

5 processes P_0 through P_4 ; 3 resource types A (10 instances),

B (5 instances), and C (7 instances).

• Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
P_0	010	753	743	3 3 2
P_1	200	3 2 2	122	
P_2	302	902	600	
P_3	211	222	0 1 1	
P_4	002	4 3 3	4 3 1	



- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
P_0	010	753	743	3 3 2
$\overline{P_1}$	200	322	122	5 3 2
P_2	302	902	600	
P_3	2 1 1	222	0 1 1	
P_4	002	4 3 3	4 3 1	



- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
P_0	010	753	743	3 3 2
P 1	200	322	122	5 3 2
P_2	302	902	600	7 4 3
$\frac{1}{\sqrt{3}}$	211	222	011	
' 3				
P_4	0 0 2	4 3 3	4 3 1	



- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
P_0	010	753	743	3 3 2
P 1	200	322	122	5 3 2
$P_2^{'}$	302	902	600	7 4 3
$\frac{1}{\sqrt{3}}$	211	222	011	7 4 5
7/	002	433	431	



- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
\mathcal{F}_0	0 1 0	753	743	3 3 2
P ₁		322	122	5 3 2
P_2		902	600	7 4 3
$\frac{1}{2}$	211	222	011	7 4 5
$\frac{3}{P_4}$		433	431	7 5 5



Example of Bankers Algorithm

- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Work</u>
	ABC	ABC	ABC	ABC
\mathcal{P}_0	0 1 0	753	743	3 3 2
P 1	200	322	122	5 3 2
P_2	302	902	600	7 4 3
$\frac{1}{7}$	211	222	011	7 4 5
$\frac{1}{2}$	002	433	431	7 5 5
•				10 5 7

Q: The system is in a safe state?

Yes, since the sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety criteria.



Example: P_1 requests (1,0,2)

- 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5 instances), and C (7 instances).
- Snapshot at time T₀:

	<u>Allocation</u>	<u>Max</u>	Need	<u>Available</u>
	ABC	ABC	ABC	ABC
P_0	010	753	743	3 3 2
P_1	200	3 2 2	122	
P_2	302	902	600	
P_3	2 1 1	222	0 1 1	
P_4	002	4 3 3	4 3 1	

Question. What will happen if process P₁ requests one additional instance of resource type A and two instances of resource type C?



To decide whether the request is granted we use Resource-Request algorithm



Resource-Request Algorithm for Process Pi

Can we remain in **safe state** after **request**?

Request = request vector for process P_i . Request_i[j] = $k : P_i$ wants k instances of resource type R_i .

- If Request_i ≤ Need_i go to step 2.
 Otherwise, raise error condition, since process has exceeded its maximum claim.
- 2. If $Request_i \le Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
 - 3. Pretend to allocate requested resources to P_i by modifying the state as follows:

```
Available = Available - Request<sub>i</sub>;

Allocation<sub>i</sub> = Allocation<sub>i</sub> + Request<sub>i</sub>;

Need<sub>i</sub> = Need<sub>i</sub> - Request<sub>i</sub>;
```

If safe \Rightarrow the resources are allocated to Pi.

If unsafe => Pi must wait, and the old resource-allocation state is restored



Example: P_1 requests (1,0,2) (Cont.)

Request 1 = (1,0,2)



1. Check that Request $1 \le \text{Need } 1$ that is, $(1,0,2) \le (1,2,2) \Rightarrow \text{true}$.

	Allocation	Max	Need	Available
	ABC	ABC	ABC	ABC
P_0	010	753	743	332
P_1	200	322	122	No. Okenie
P_2	302	902	600	
P_3	211	222	011	
P_4	002	433	431	

- 2. Check that Request $1 \le$ Available that is, $(1,0,2) \le (3,3,2) \Rightarrow$ true.
- 3. Pretend to allocate requested resources to P_i by modifying the state as:

Available = Available - Request_i; Allocation_i = Allocation_i + Request_i; Need_i = Need_i - Request_i;



Allocation	Max	Need	Available
ABC	ABC	ABC	ABC
010	753	743	2 3 0
302	322	020	
302	902	600	
211	222	011	
002	433	431	
	ABC 010 302 302 211	ABC ABC 010 753 302 322 302 902 211 222	ABC ABC ABC 010 753 743 302 322 020 302 902 600 211 222 011



Example: P_1 requests (1,0,2) (Cont.)

Pretend to allocate requested resources to P_i by modifying the state as:

	<u>Allocation</u>	<u>Need</u>	<u>Work</u>
	ABC	ABC	ABC
Po	010	743	230
-	302	020	5 3 2
	302	600	7 4 3
	244	0 0 0	7 4 5
7 3	211	4 2 4	7 5 5
F	002	431	10 5 7

- Executing safety algorithm shows that sequence <P1, P3, P4, P0, P2> satisfies safety requirement.
 - → The new system state is safe.
 - → We can immediately grant the request for process P1.
- Q: Can request for (3,3,0) by P4 be granted?
- Q: Can request for (0,2,0) by P0 be granted?



Conclusion

- Deadlock Concept: A set of blocked processes each holding a shared resource and waiting to acquire a resource held by another process in the set. – Dining Philosopher's Problem
- 4 necessary conditions for Deadlock: Mutual Exclusion, Hold and Wait,
 No Preemption, Circular Wait
- Deadlock Prevention: Make sure 4 necessary conditions for Deadlock does not hold
- Deadlock Avoidance: Always keeps safe state
 - Resource allocation graph
 - Banker's Algorithm



- Appendix
 - Methods for Handling Deadlocks
 - Deadlock Detection
 - Recovery from Deadlock



Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
 - Examine the state of the system to determine whether a deadlock has o ccurred
 - (i.e., Check to see if a deadlock has occurred!)
 - Cases
 - single instance per resource type
 - multiple instance per resource type
- Recovery scheme
 - Recover from the deadlock

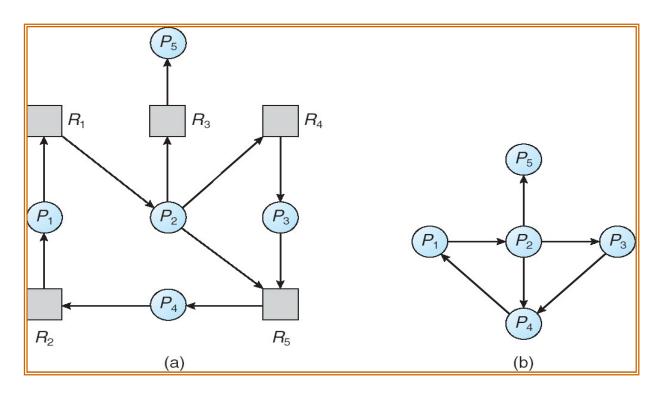


Single Instance of Each Resource Type

- Maintain <u>wait-for graph</u>
 - Nodes are processes.
 - Remove the in-between resource nodes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j .
- Periodically invoke an algorithm that searches for a cycle in the graph.



Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

We can see dependencies between process



Detection-Algorithm Usage

- When, and how often, to invoke depends on:
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - one for each disjoint cycle

- Invoking detection-algorithm for every resource request vs.
 Invoking detection-algorithm at defined interval
 - High overhead vs. detection accuracy
 - In a periodic detection, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes "caused" the deadlock.



Deadlock Recovery

- Introduction
 - Recovery after detection of deadlocked processes
 - Eliminates detected deadlocks in the system
- Deadlock recovery methods
 - Process termination
 - Terminate (abnormally) part of the deadlocked processes
 - Terminated processes are restarted or rolled back afterwards
 - Resource preemption
 - Elect resources to be preempted to eliminate the deadlock
 - Preemption of the resources from the processes that currently owns them



Deadlock Recovery: Process Termination

- Abort all deadlocked processes
 - Simple, but at great expense
 - May terminate unnecessary processes
- Partial Termination: Abort one process at a time until the deadlock cycle is eliminated. (Based on which process is the min cost)
 - Choose optimal set of processes to be terminated
 - Minimum overall termination cost
 - Complex: a deadlock-detection algorithm must be invoked

