


# Chapter 8: Memory-Management Strategies

School of Computing, Gachon Univ.  
Jungchan Cho



Most slides from "Operating System Concepts – 10th Edition".  
Many slides are taken from lecture notes of Prof. Joon Yoo.

# Objectives

---

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- Important keywords
  - Logical (=Virtual) Memory, Paging (TLB)
  - We use logical address and virtual address interchangeably in this text.

# Chapter 8: Memory Management Strategies

---

- **Background**
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table


# Memory Technology (2012)

- SRAM semiconductor memory - **Cache**
  - 0.5ns – 2.5ns, \$500 – \$1000 per GB
- **DRAM semiconductor** memory – **Main memory**
  - 50ns – 70ns, \$10 – \$20 per GB
- Flash semiconductor memory – **Flash disk/memory**
  - 5,000 – 50,000ns, \$0.75 - \$1.00 per GB
- Magnetic disk – **Hard disk**
  - 5ms – 20ms, \$0.05 – \$0.1 per GB
- Ideal memory: Fast, large and cheap?
  - Access time of SRAM
  - Capacity and cost/GB of disk



# DRAM: main memory


---



<http://zeptobars.com/en/read/how-to-open-microchip-asic-what-inside>

# DRAM: main memory


- **Dynamic random access memory**
  - Storage cells: One capacitor and one transistor = one data bit
- As software engineers,
  - Let the Electrical Engineers worry about these complexities...
  - We are only interested in memory address and byte data



[https://en.wikipedia.org/wiki/Dynamic\\_random-access\\_memory](https://en.wikipedia.org/wiki/Dynamic_random-access_memory)

# DRAM: main memory

- Dynamic random access memory
  - Give address to DRAM
    - ▶ Read/write data
  - **Byte address**
    - ▶ Each data is 1-byte

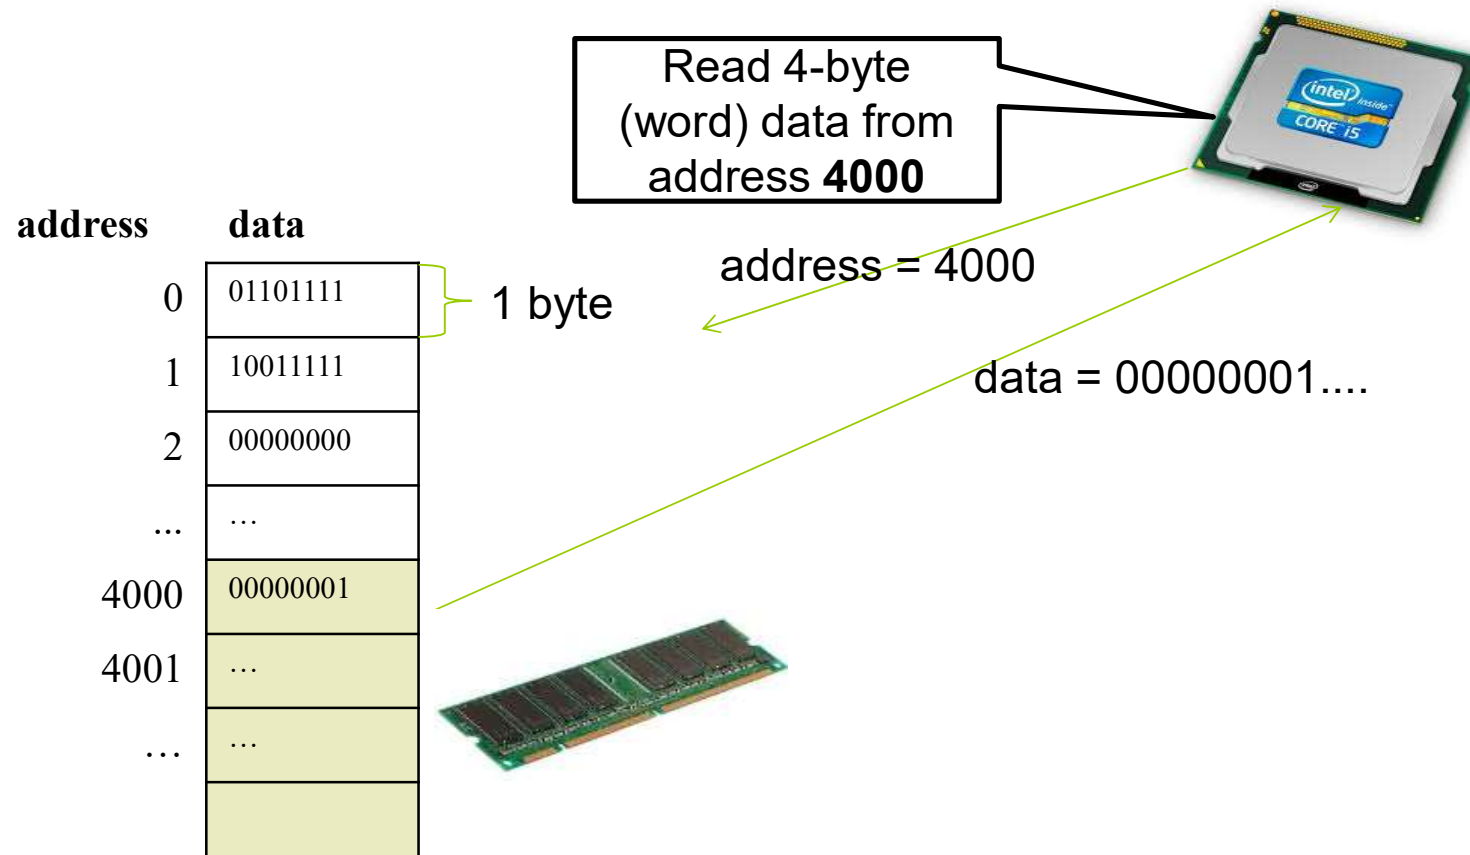


<https://chortle.ccsu.edu/AssemblyTutorial/>



# DRAM: main memory


- Dynamic random access memory
  - Byte address



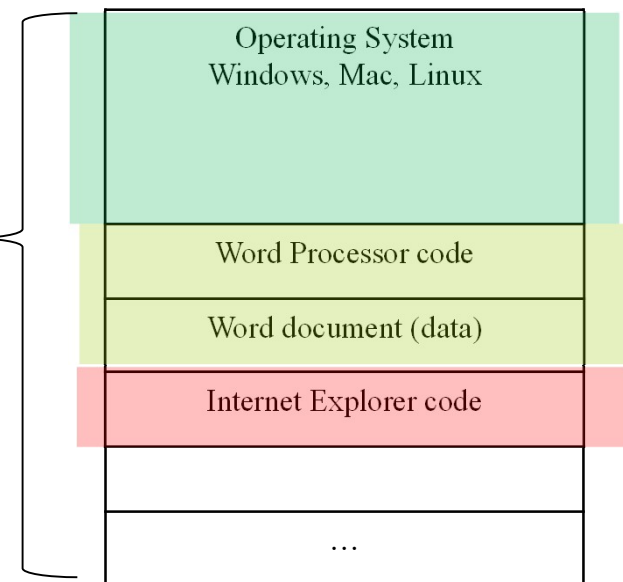


# DRAM: main memory

- Dynamic random access memory
  - **Programs** are usually stored in the **hard disk**
    - ▶ When programs are executed, it is loaded on the **main memory** – faster than disk
  - Similarly, when **file** is **open**, it is loaded on the main memory




## Memory




# DRAM: main memory

- Why load program/data from disk to memory?
  - DRAM access time  $\approx 50\text{ns}$
  - Disk access time  $\approx 5,000,000\text{ns}$
  - If we use disks to run programs it will be **100,000 times** slower
- But DRAM is still much slower than CPU
  - It takes about **100~400 CPU cycles** to access DRAM
  - Use a small but very fast memory = **Cache!!**



# Background

- Memory is **central** to the operation of a computer system
- A typical **CPU instruction-execution cycle**
  - 1. The CPU **fetches** an **instruction** from **memory**
  - 2. The instruction is then decoded and may cause operands (data) to be fetched from memory
  - 3. the result may be **stored back in memory**



# Background Summary

---

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory is the only storage that CPU can access directly
  - CPU cannot directly access disk, network, printers, ... – **must** be copied to memory first
- Speed?
  - Register access in one CPU clock (or less)
  - Main memory can take many cycles, causing a **stall**
    - ▶ Note: Faster **cache memory** sits between main memory and CPU registers. **OS view of cache is same as main memory!!**


# Sharing/Protecting the Memory

---

- Kernel, System/user Processes, Data share the main memory
- **Q1:** How can the processes share the main memory?
- **Q2:** How can the processes be protected from each other?
- Answer: Each process uses **Logical (or Virtual) address** instead of actual physical memory address!

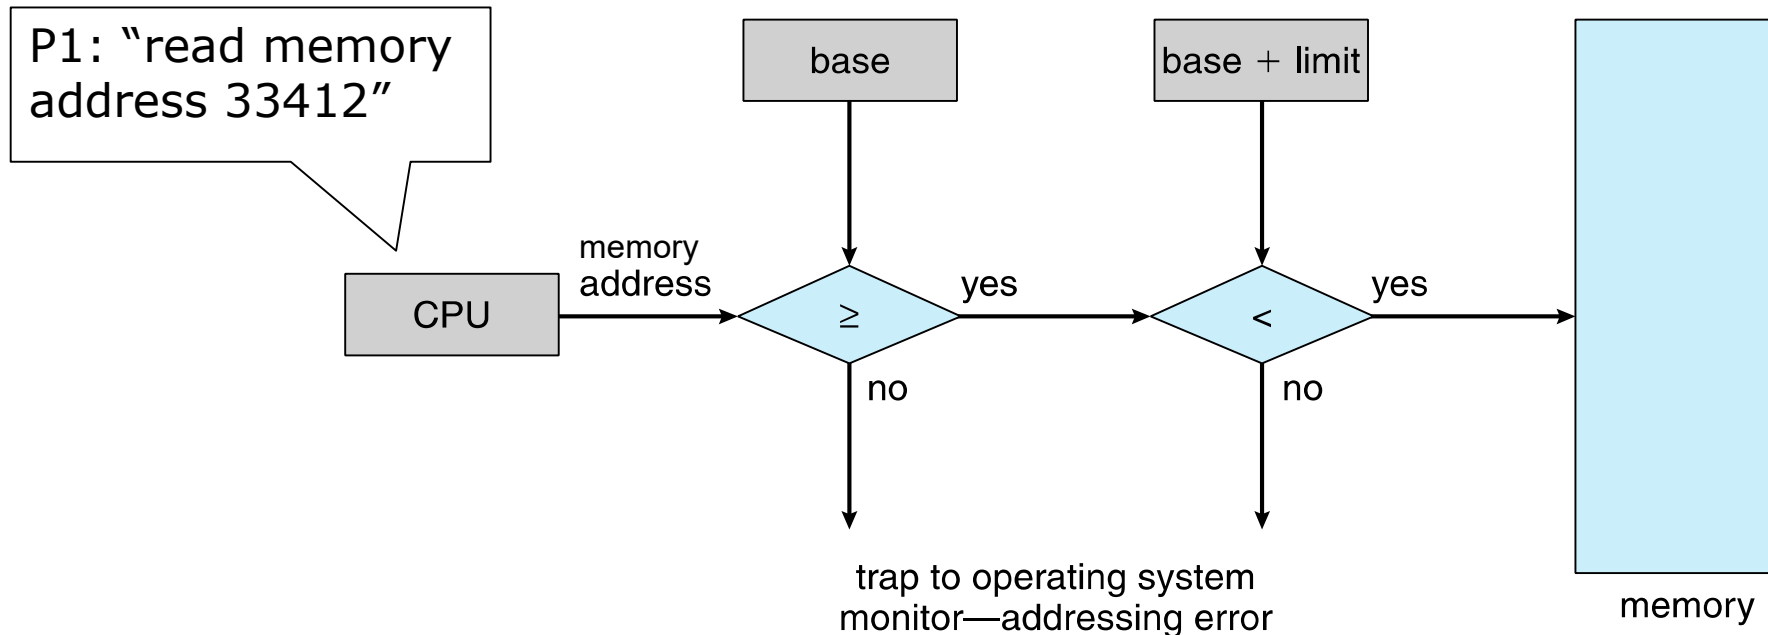
# Base and Limit Registers

- **Protection** of memory required to ensure correct operation
  - Protect OS from user programs
  - Protect user programs from one another
- Each process must have separate memory space
- Hardware protection: a pair of **base register** and **limit register** define the **logical address space** of each process
  - Memory access outside the logical space?
    - ▶ Trap to OS



## Hardware Address Protection with Base and Limit Registers

CPU must check every memory access generated in user mode to be sure it is between **base** and **base+limit** for that user

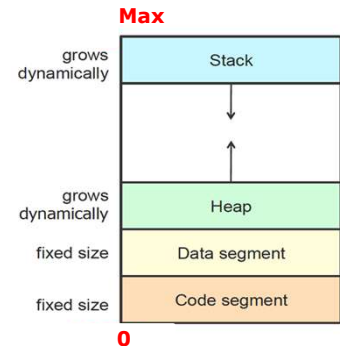




# Logical vs. Physical Address Space

## ■ Logical address (= virtual address)

- An address generated by **CPU**: meaningful only to the **user process**
- Always starts from **address 0**
- **Logical Address space**
  - ▶ **Set** of all logical addresses generated by a program




## ■ Physical address

- Address seen by the **physical memory unit**
- **Physical Address space**
  - ▶ Set of all physical addresses

A d d r e s s	0xFFFFFFFF	1000 0000
		.....
		.....
	0x00000008	0100 1001
	0x00000007	1100 1100
	0x00000006	0110 1110
	0x00000005	0110 1110
	0x00000004	0000 0000
	0x00000003	0110 1011
	0x00000002	0101 0001
	0x00000001	1100 1001
	0x00000000	0100 1111
		Main Memory

# Address Space

- A **logical (virtual) address** is a memory address that a process uses to access its own memory
  - Virtual address  $\neq$  actual physical RAM address
  - When a process accesses a virtual address, the **memory-management unit (MMU)** translates the virtual address into a physical address
  - The OS determines the mapping from virtual address to physical address



# Memory-Management Unit (MMU)

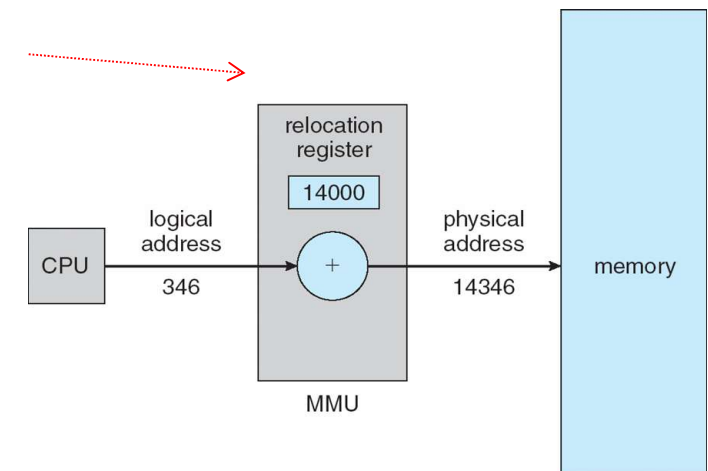
## ■ Memory-Management Unit (MMU)

- Maps virtual to physical address at **run time** with **hardware support**.
- Many methods possible, covered in this chapter (e.g., relocation register, paging, segmentation)

- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory

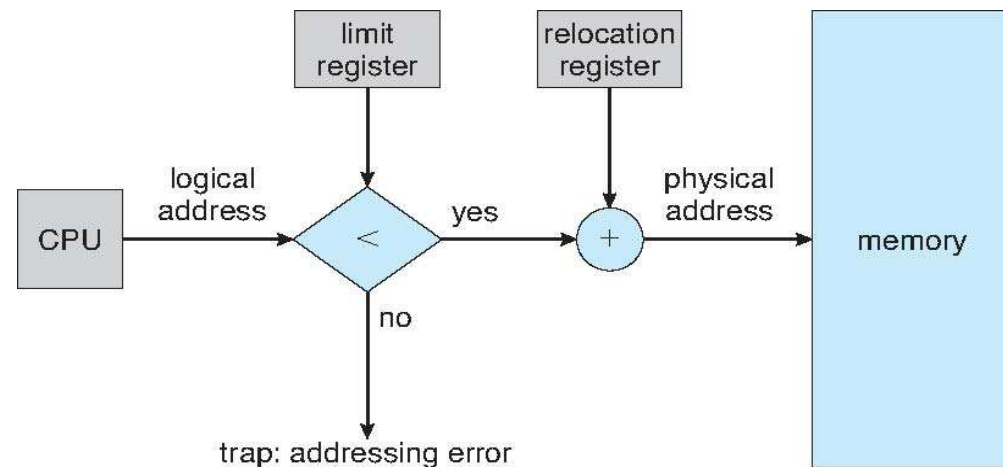
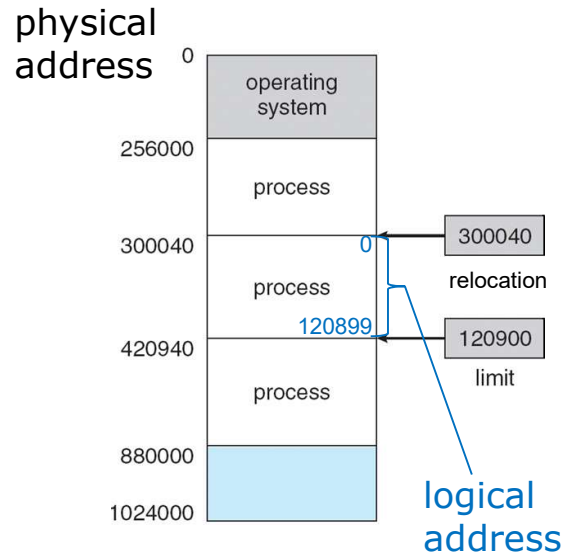
- Base register now called **relocation register**

- The user program deals only with **logical** addresses; it never sees the real **physical** addresses



# Memory Protection


- Need to **protect** user processes from each other, and from changing operating-system code and data
- **Memory Protection: Relocation-register scheme**
  - A relocation register and a limit register define a logical address space
  - Address protection with relocation and limit registers



# Address Binding

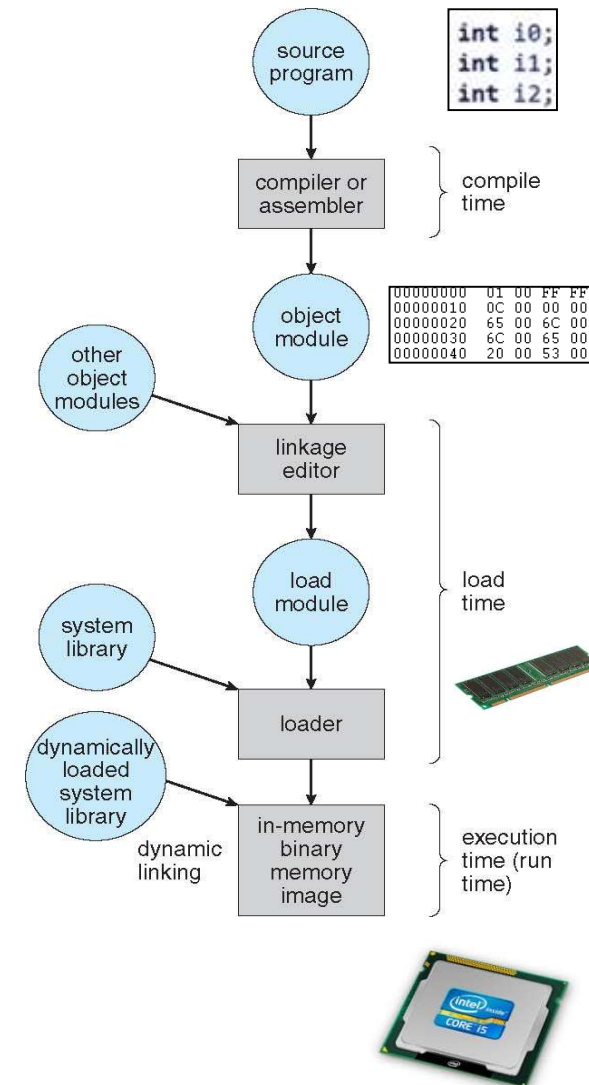
- Program on disk as binary executable file
  - ready to be brought into memory to execute as process
  - And then system **loads** the process in physical memory
- **Memory addresses** represented in different ways at different stages of a program's life
  - Source code addresses usually **symbolic**
    - ▶ e.g., “int **count**”
  - Binary code needs **address binding** between *logical* and *physical* addresses:
    - ▶ e.g., “14 bytes from beginning of this program” (logical) vs. 74014 (physical)
- When should we do address binding?

```
1 #include "myMult.h"
2
3 void myMult(const double a[12], const double b[20], double c[15])
4 {
5     int i0;
6     int i1;
7     int i2;
8     for (i0 = 0; i0 < 3; i0++) {
9         for (i1 = 0; i1 < 5; i1++) {
10             c[i0 + 3 * i1] = 0.0;
11             for (i2 = 0; i2 < 4; i2++) {
12                 c[i0 + 3 * i1] += a[i0 + 3 * i2] * b[i2 + (i1 << 2)];
13             }
14         }
15     }
16 }
```

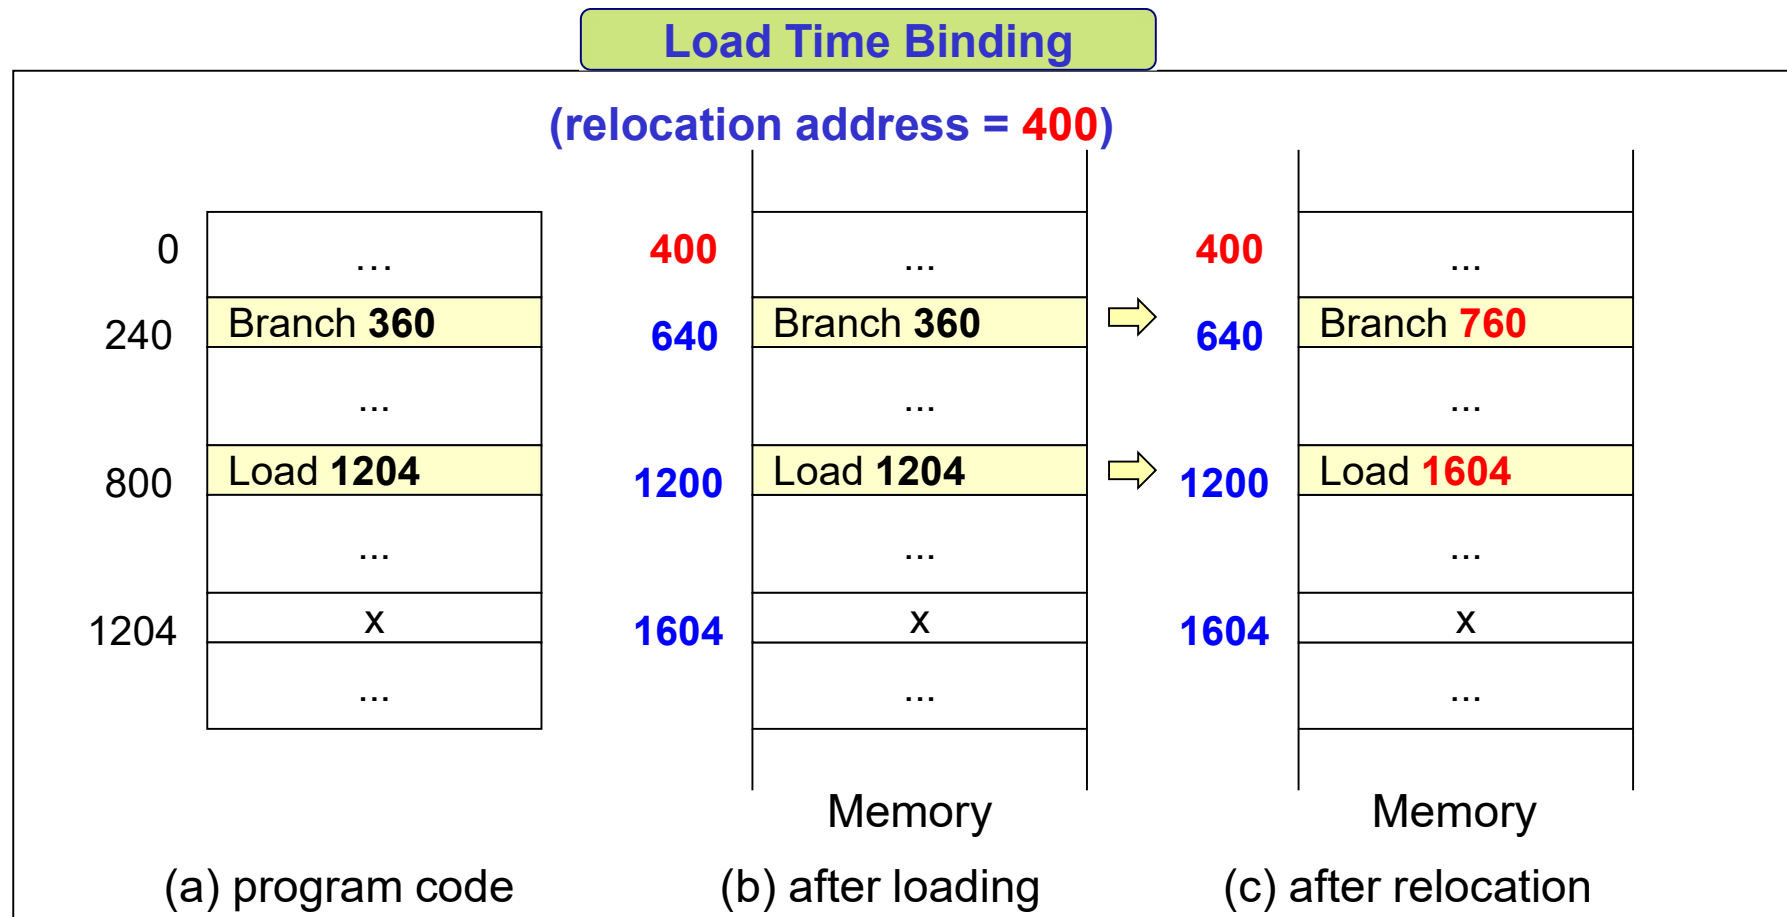


# Binding of Instructions and Data to Memory

- **Address binding** can happen at three different stages: compile time, load time, execution time
- **Compile time**
  - If memory location known at compile time, **absolute code** can be generated
    - ▶ e.g., user process will reside starting at location  $R_1$ , compiled code will start at  $R_1$
  - If starting location changes (e.g.,  $R_2$ ), must **recompile** code
- **Load time**
  - If memory location is not known at compile time, compiler must generate **relocatable code**
  - Address binding is delayed until **load time**
  - If starting address changes, need to **reload** user code



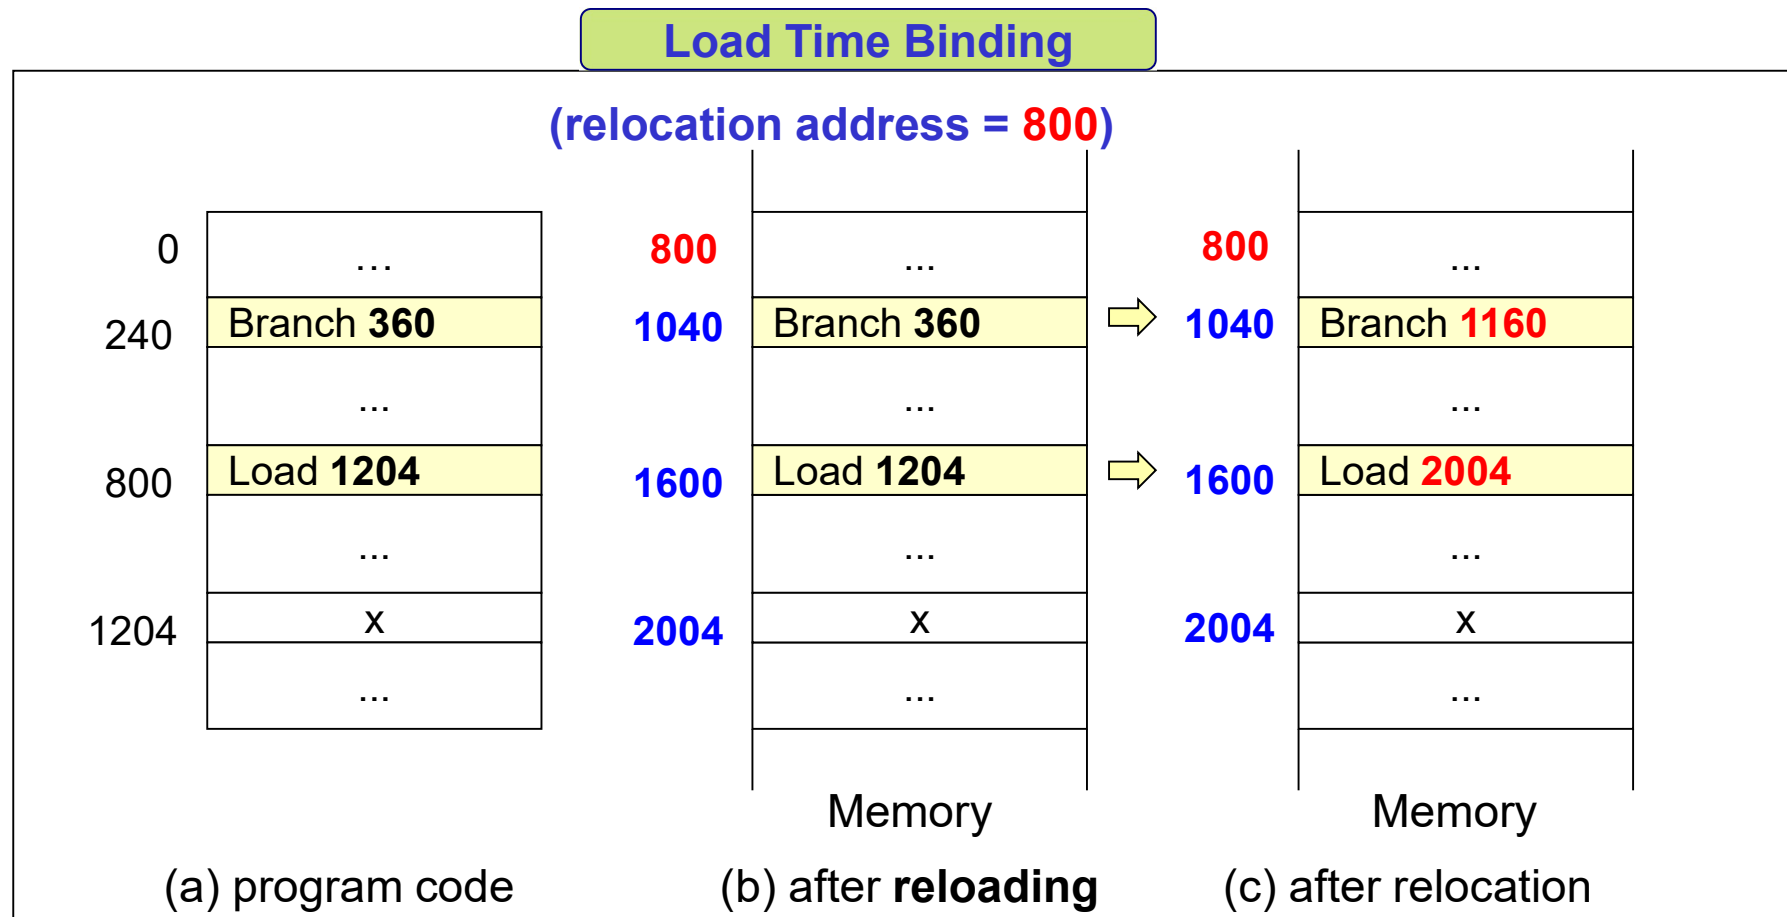
# Example: Load time binding





# Example: Load time binding

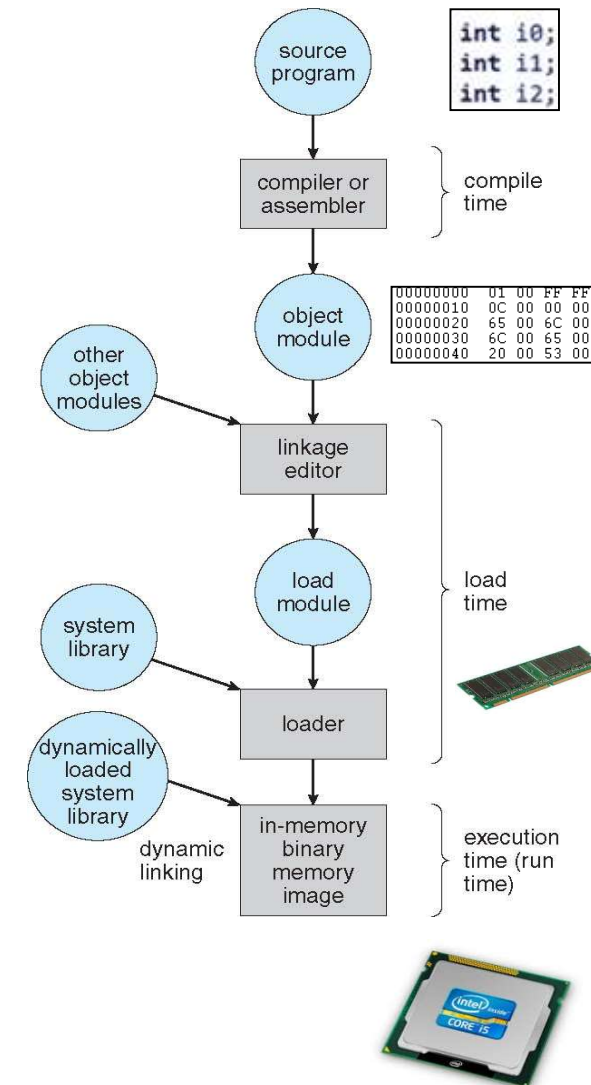
Q: What if relocation address changes to 800? Need to reload



# Binding of Instructions and Data to Memory

- **Execution time**


- ▶ process memory location can be moved during execution
- ▶ Address binding delayed until **run time**
- ▶ Need hardware support for address maps (e.g., MMU such as base and limit registers)
- ▶ Most OSs uses this method




# Example: Execution (Run) time binding

## Execution Time Binding


(relocation address = 1400)



(a) program code



(b) after loading



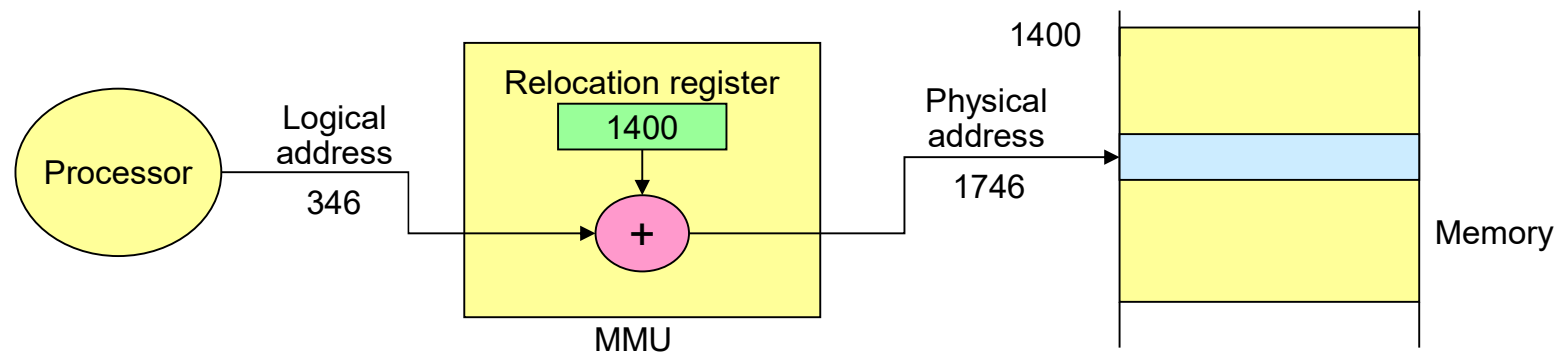
(c) On every memory access



# Dynamic Relocation

## ■ Run-time binding

- Run-time mapping from logical(virtual) address to physical address is performed by **MMU**(Memory Management Unit)



Q: What if relocation address changes to 1800 in previous slide?

# Chapter 8: Memory Management Strategies

---

- Background
- **Contiguous Memory Allocation**
- Paging
- Structure of the Page Table


# Memory Protection

## ■ Memory Allocation

- Allocate memory to OS and user processes in an efficient way possible

## ■ Main memory usually into two partitions:

- Resident **operating system**
  - ▶ usually held in low memory with interrupt vector
- User **processes**
  - ▶ held in high memory

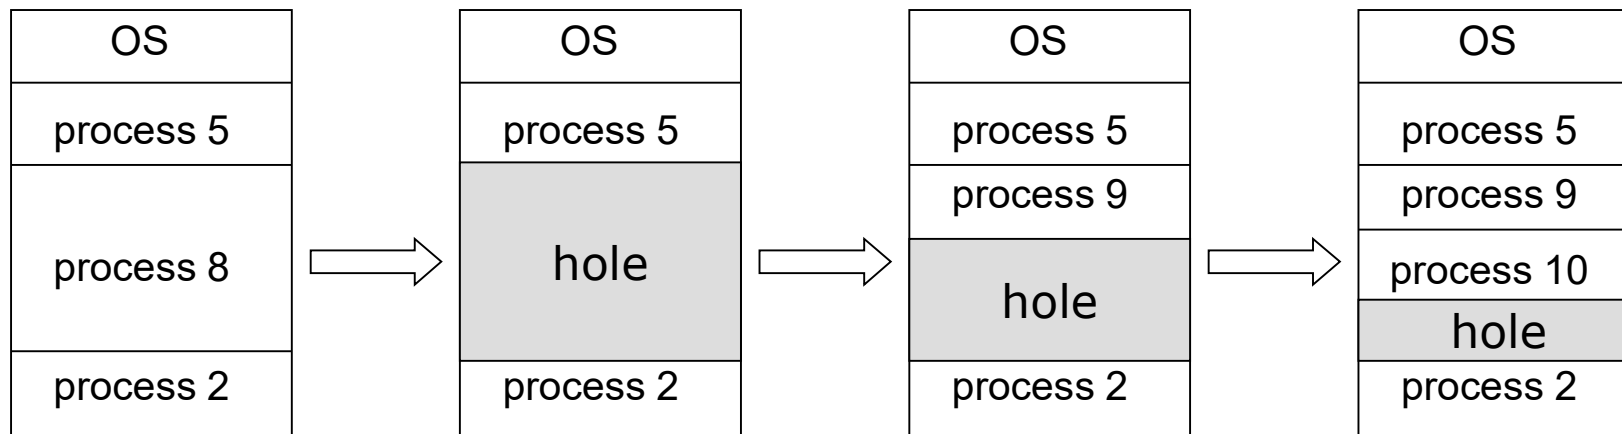


## ■ Contiguous Memory Allocation

- Allocate available memory to processes
- Each process is contained in a single section of contiguous memory

# Memory Allocation: Contiguous Allocation

- Multiple-partition allocation
  - Degree of multiprogramming limited by number of partitions
  - **Variable-partition** sizes for efficiency (sized to a given process' needs)
  - **Hole** – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it
  - Process exiting frees its partition, adjacent free partitions combined
  - Operating system maintains information about:
    - a) allocated partitions
    - b) free partitions (hole)





# Contiguous Allocation: Example

## ■ Example

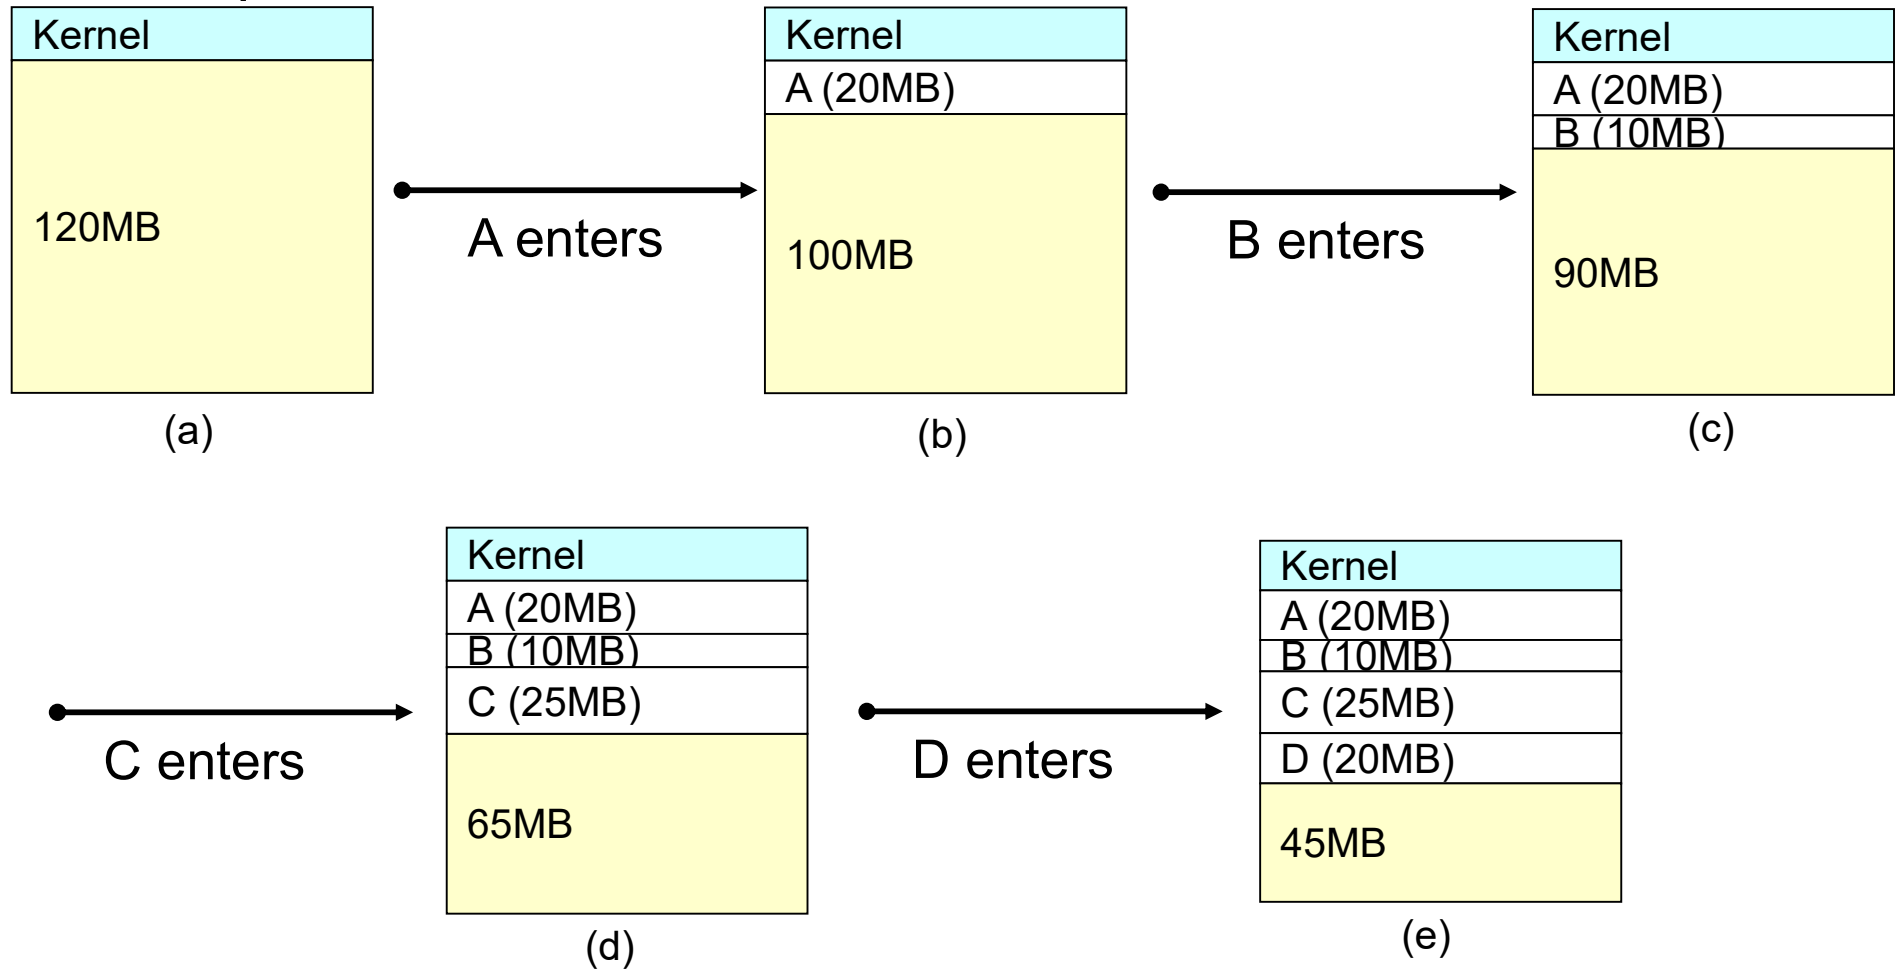
### Memory allocation and partition scenario

#### ◆ Assumption

- Memory space: 120 MB
- (a) : Initial state
- (b) : After loading process A(20MB)
- (c) : After loading process B(10MB)
- (d) : After loading process C(25MB)
- (e) : After loading process D(20MB)
- (f) : After process B releases memory
- (g) : After loading process E(15MB)
- (h) : After process D releases memory

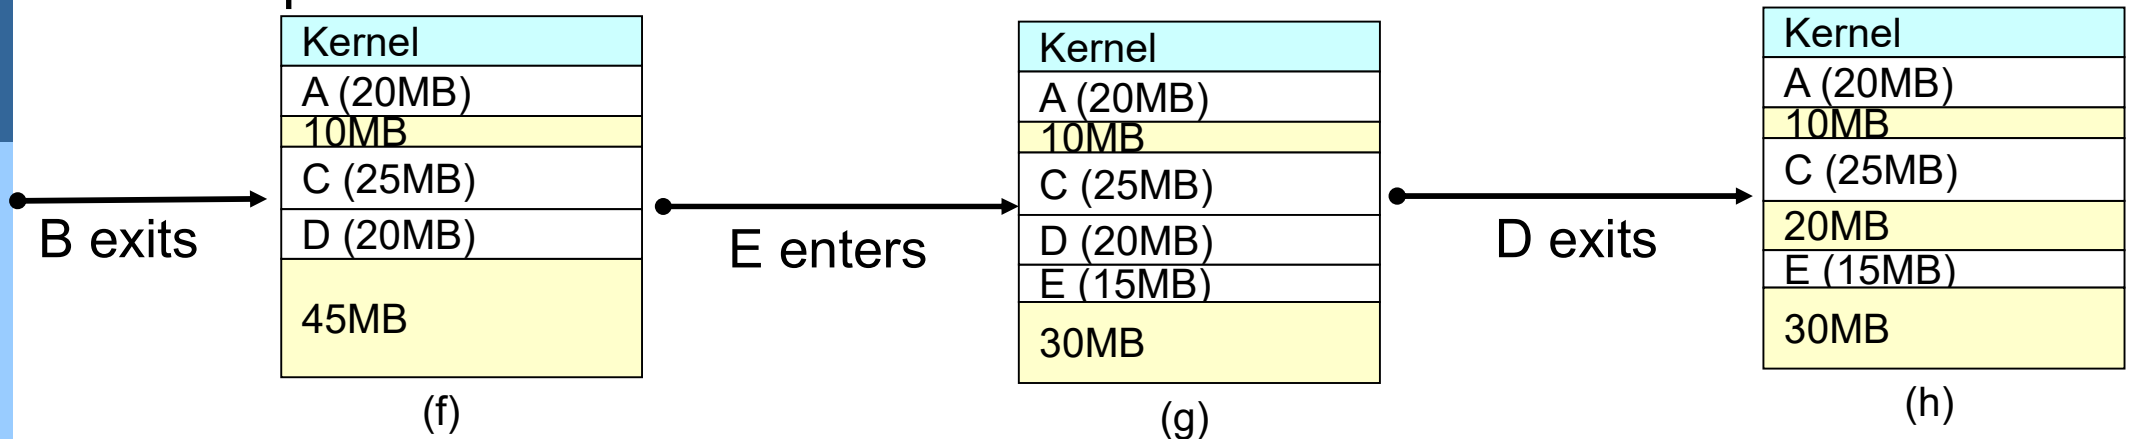
# Contiguous Allocation: Example

## ■ Example



## Contiguous Allocation: Example (Cont'd)

### ■ Example



□ **Hole** – block of available memory; holes of various size are scattered throughout memory

# External Fragmentation

F (40MB)?

Kernel
A (20MB)
10MB
C (25MB)
20MB
E (15MB)
30MB

- **External Fragmentation**
  - Total memory space exists to satisfy a request, but it is not contiguous
- Analysis reveals that given  $N$  blocks allocated,  $0.5 N$  blocks lost to external fragmentation
  - $1/3$  may be unusable -> **50-percent rule**
- We will discuss **internal fragmentation** later

# External Fragmentation Solutions?

---

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
    - ▶ **NOTE:** We cannot preempt memory space if logical memory is determined in load time
- Another possible solution to the external-fragmentation problem
  - Permit the logical address space of the processes to be **noncontiguous**
  - Allowing a process to be allocated physical memory wherever such memory is available

# Chapter 8: Memory Management Strategies

---

- Background
- Contiguous Memory Allocation
- **Paging**
- Structure of the Page Table


# Paging : Basic Method

## ■ Paging

- **Noncontiguous memory management** scheme
  - ▶ No external fragmentation – no need for compaction

## ■ Basic method

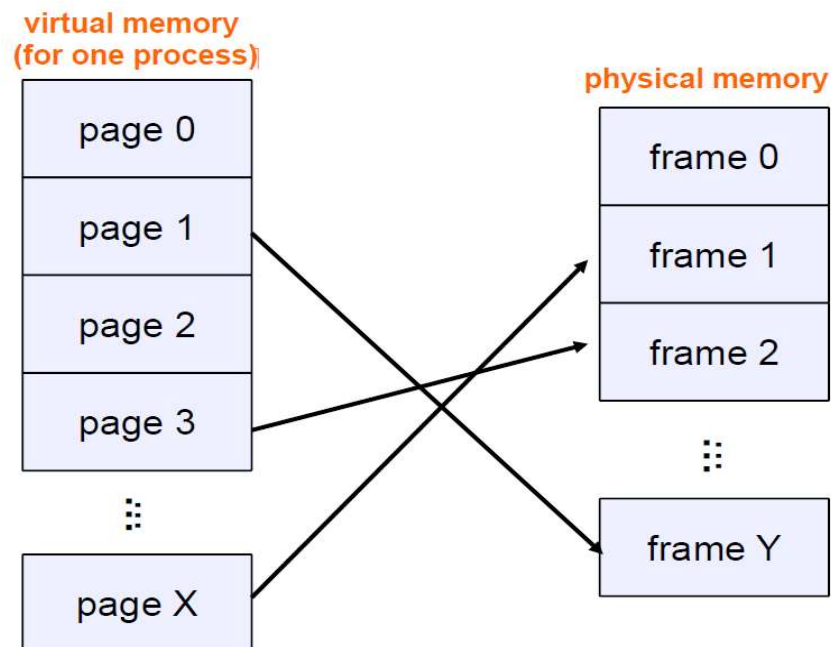
- Divide
  - ▶ **physical memory** → *fixed-sized blocks* : **frames**
  - ▶ **logical memory** → *blocks of same size* : **pages**
  - ▶ size of one frame = size of one page
- Set up a **page table** to translate logical to physical addresses



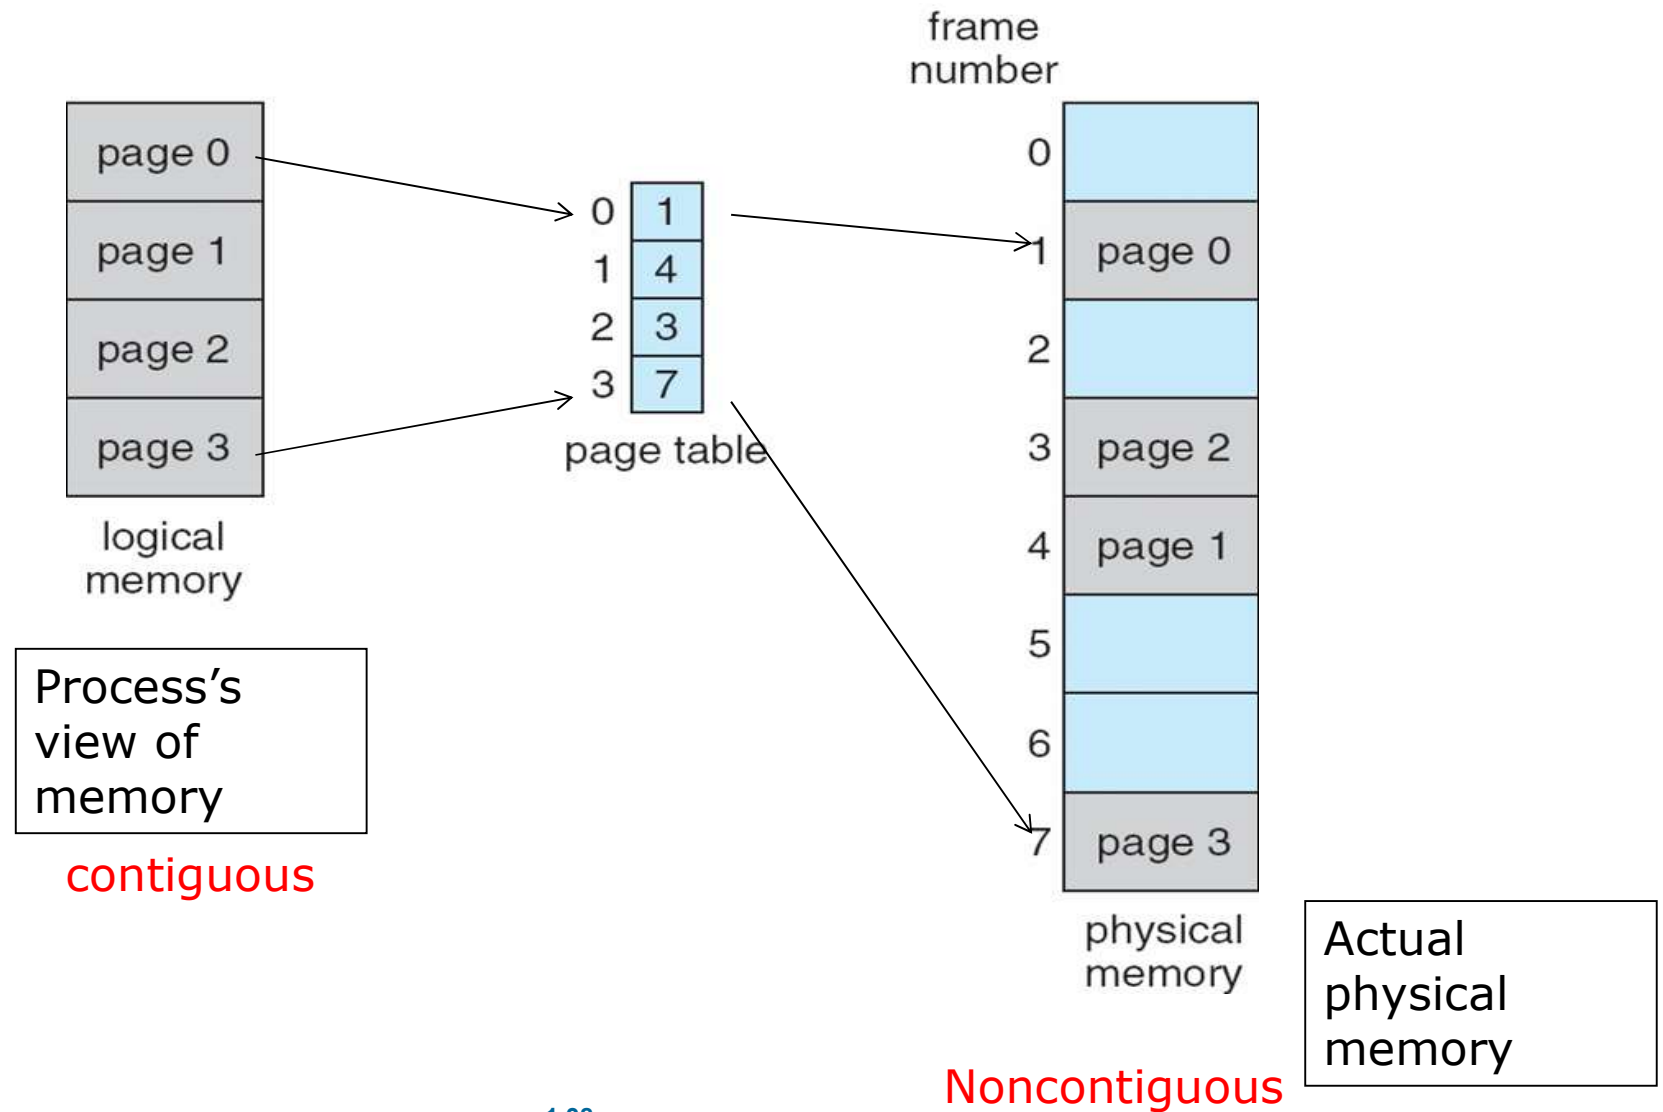


## Cont'd.

- Solve the external fragmentation problem by using **fixed size chunks** of virtual (**page**) and physical memory (**frame**)

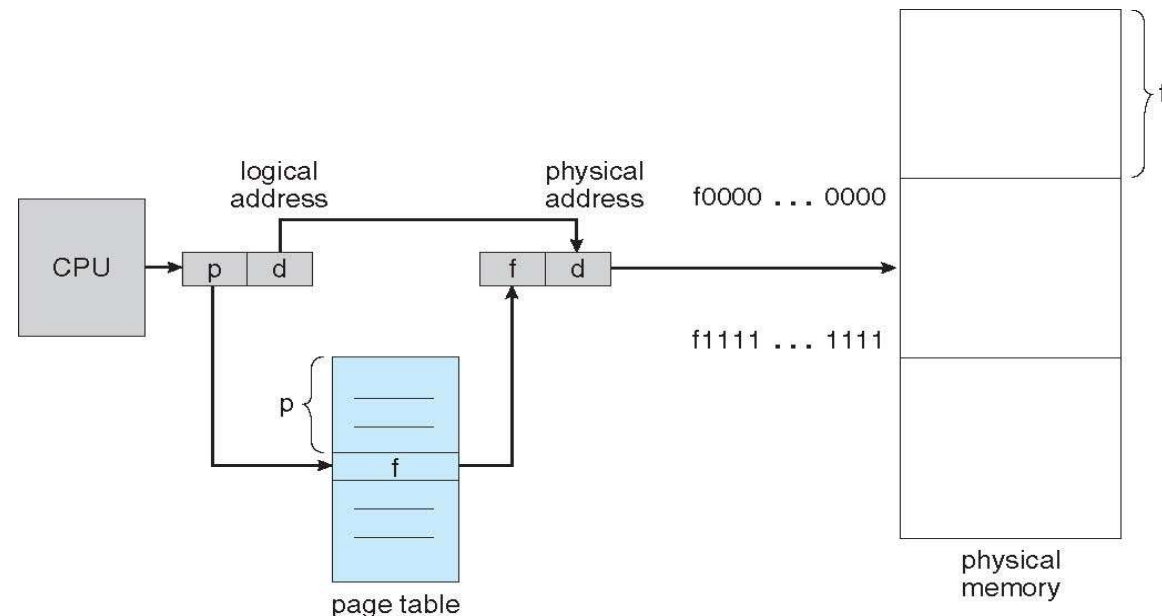


# Paging Example




# Page Number and Offset

- When a process is executed, its pages are loaded into any available memory frames
- Every address generated by CPU is divided into two parts: **page number (p)** and **page offset (d)**
  - **page number (p)**: index into a **page table**
  - **page table**: contains **base address** of each page in physical memory
  - **page offset (d)**: page number and page offset combines to define the physical memory address



# Page Size

- Fixed sized chunk: **page size** (=frame size) defined by hardware
  - vary between 512B ( $=2^9\text{B}$ ) and 1GB ( $=2^{30}\text{B}$ ) per page – depending on computer architecture
- If the size of **logical address space** is  $2^m$  and **page size** is  $2^n$ 
  - **page number(p)**: high-order  $m-n$  bits of a logical address: **index** into a *page table*
  - **page offset(d)**: low-order  $n$  bits of a logical address



page size?

# Paging in bits Example

page#

0	0	a
	1	b
	2	c
	3	d
1	4	e
	5	f
	6	g
	7	h
2	8	i
	9	j
	10	k
	11	l
3	12	m
	13	n
	14	o
	15	p

logical memory

**n = 2**  
**m = 4**

- **logical address space**  
=  $2^4 = 16$
- **page size** is  $2^2 = 4$

page table

0	5
1	6
2	1
3	2

Frame#

0		0
4	i	1
	j	
	k	
8	m	2
	n	
	o	
	p	
12		
16		
20	a	5
	b	
	c	
	d	
24	e	6
	f	
	g	
	h	
28		

physical memory

□ **m = 4, n = 2**

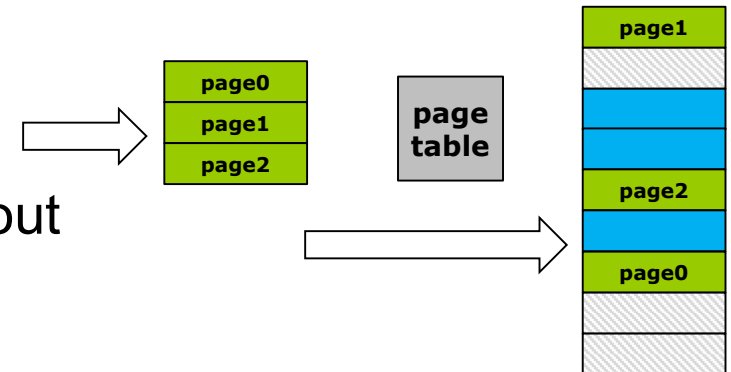
- Logical address space:  $2^m = 16$  bytes
- Page size :  $2^n = 4$  bytes
- Page number:  $2^{m-n} = 4$  bytes

□ **Example**

- Logical address 0
  - ✓ Page #0 → Frame #5, offset# 0
  - ✓ physical address =  $(5 \times 4) + 0 = 20$
  - data "a"**
- Logical address 1, 11?
  - ✓ Page# ? → Frame #? offset?
  - ✓ physical address? data?

# Why Paging? (1)


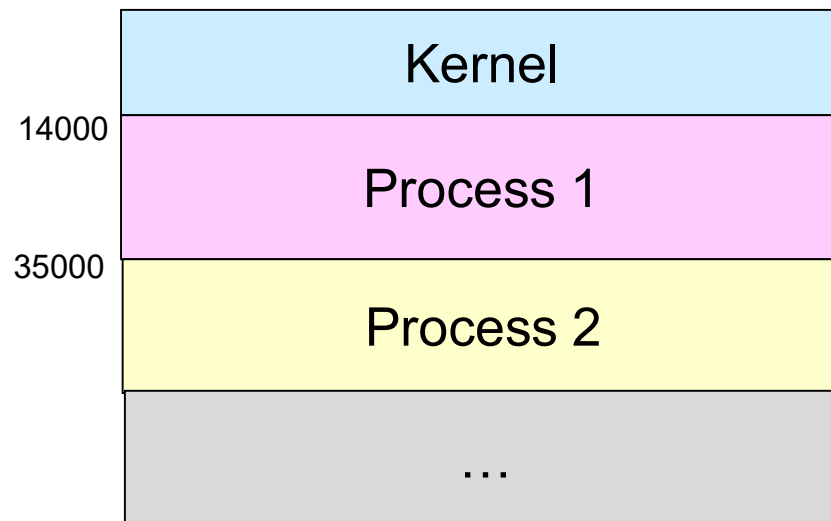
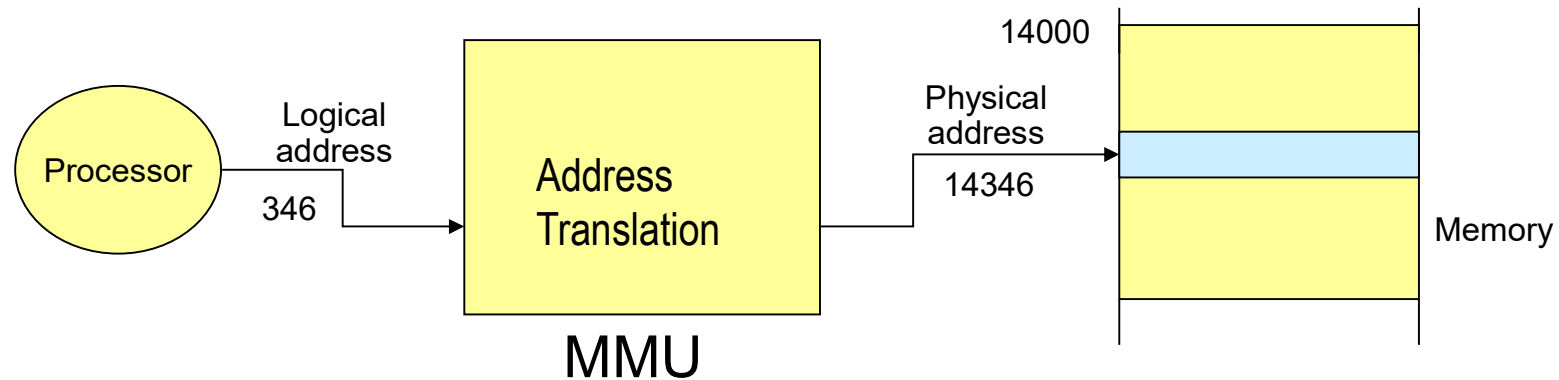
- Clear separation between the programmer's view of memory (logical memory) and physical memory
  - Programmer: One contiguous single space
  - Actual physical memory: scattered throughout the physical memory
  - Needs **page table** for address translation



## 1. Provides **Transparency**



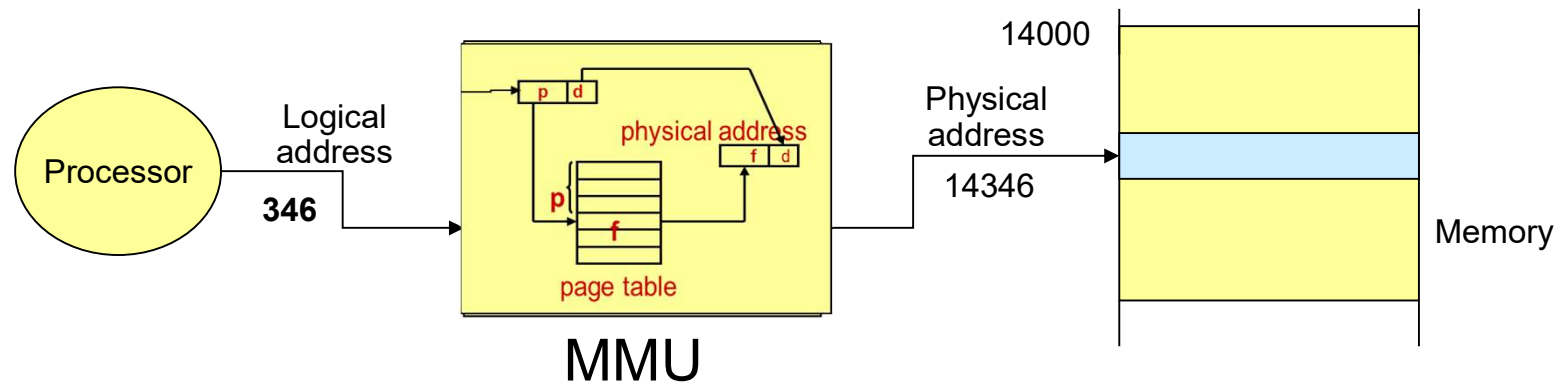
- ▶ Programmer only sees *logical address*, never sees actual *physical address* – *physical address* is transparent to programmer
- ▶ This mapping is *hidden* from programmer and controlled by OS

# Contiguous vs. Paging



MMU for Contiguous Allocation


# Contiguous vs. Paging





# Free-frame management

- To execute a process whose size requires  $n$  pages, at least  $n$  free frames are required – How do we find the **free frames**?
  - The first page of the process is loaded into the first frame listed on **free-frame list**
  - and the allocated frame number is put into **page table**
  - The next page is loaded into another frame ...



A new process with 4 pages

## Why Paging? (2)

---

### 2. No external fragmentation

- Noncontiguous memory allocation - ANY free frame can be allocated to a process that needs it.
- However, we may have **internal fragmentation**
  - ▶ e.g.) if a page size is 2048 bytes, a process of 72766 bytes would need 35 pages plus 1086 bytes
  - ▶ In worst case, a process would need  $n$  pages plus 1 byte, and it would be allocated  $n+1$  frames – internal fragmentation in last frame

# Internal Fragmentation

## External Fragmentation


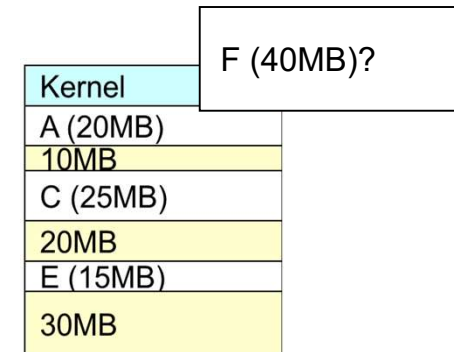
- Contiguous allocation: Total memory space exists to satisfy a request, but it is not contiguous

## Internal Fragmentation

- Memory is allocated in fixed-sized **pages**
- Allocated **pages** may be larger than requested memory;
  - page size 4kB,
  - process needs 9kB – allocated 3 pages (12kB)
- this size difference is memory internal to a partition, but not being used

## How do we determine the Page Size?

- Consider **tradeoff**: Internal fragmentation, page table size, disk I/O efficiency
- Practically 4 – 8KB in size



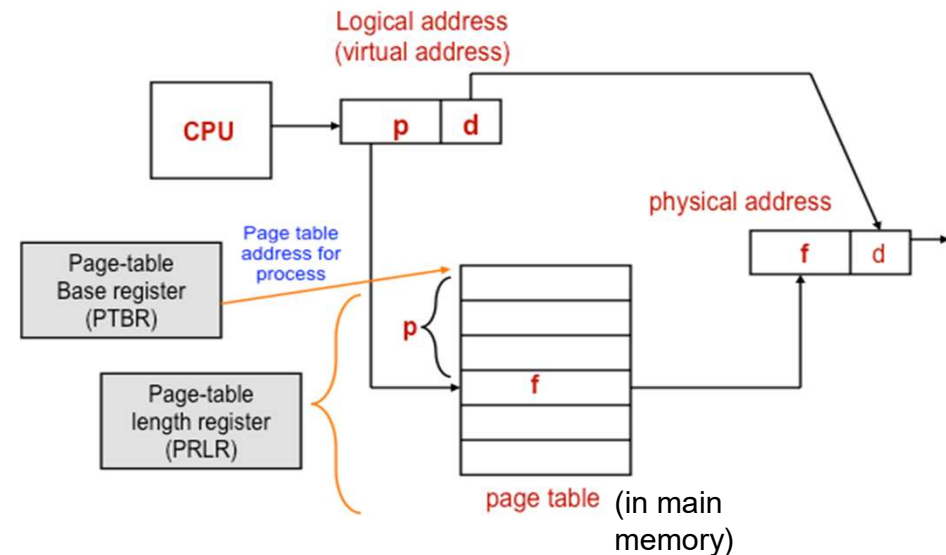
# Implementation of Page Table

---

- There is a **page table** for each process!!
  - e.g., process P1's page table, P2's page table, ...
- **Option 1**: Keep in registers and make it fast
  - Implement page table as **registers**: translation very fast but expensive – need to keep it small (e.g., 256 entries)
  - But page table can be very large! (e.g., 1 million entries)
- **Option 2**: Keep in memory with register help
  - **Page table** is kept in **main memory**
  - Use **CPU registers** to store location and size of page table – hardware support

# Implementation of Page Table

- **Page table** is kept in main memory with **hardware support**:
  - **Page-table base register (PTBR)**
    - ▶ points to the page table for its process
  - **Page-table length register (PTLR)**
    - ▶ size of the page table (for the process)



- When context switch happens,
  - Need to change the page table changing page tables requires changing the **PTBR** and **PTLR**

process state
process number
program counter
registers
memory limits
list of open files
...

# Page Table in Main Memory

- **Page table** is kept in **main memory**: Problems?

- Every data/instruction memory access *actually* requires two memory accesses


- ▶ 1. page table 2. data/instruction

- Example

- ▶ Process wants to read memory at logical address 010 (page 0, offset 10)

- 1. First read page table to look for physical frame number of page 0 (=14)

- 2. Read frame number 14



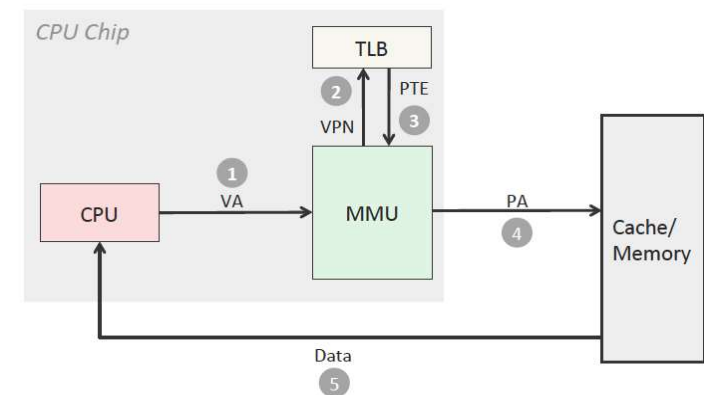
# Page Table in Main Memory

---

- Isn't it slow to have to go to memory twice every time?
  - It takes about 100~400 CPU cycles to access DRAM
  - Yes, it would be... so, real MMUs don't
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
  - **TLB** is like a cache to the page table
    - ▶ e.g., cache memory is a cache to the main memory (RAM)
  - Store recently referenced page table entries in the TLB

# TLB (Translation Look-aside Buffer)

- **Small, dedicated, super-fast hardware cache in MMU**
  - High speed memory technology, approaching register speeds
    - ▶ Implement using the dedicated registers or cache memories
    - ▶ Parallel search
  - Each TLB entry contains a page number and the corresponding page-table entry
    - ▶ **Search key** = logical address, **result** = frame number
- Memory references that can be translated using TLB entries are much faster to resolve
  - high speed, but Expensive hardware



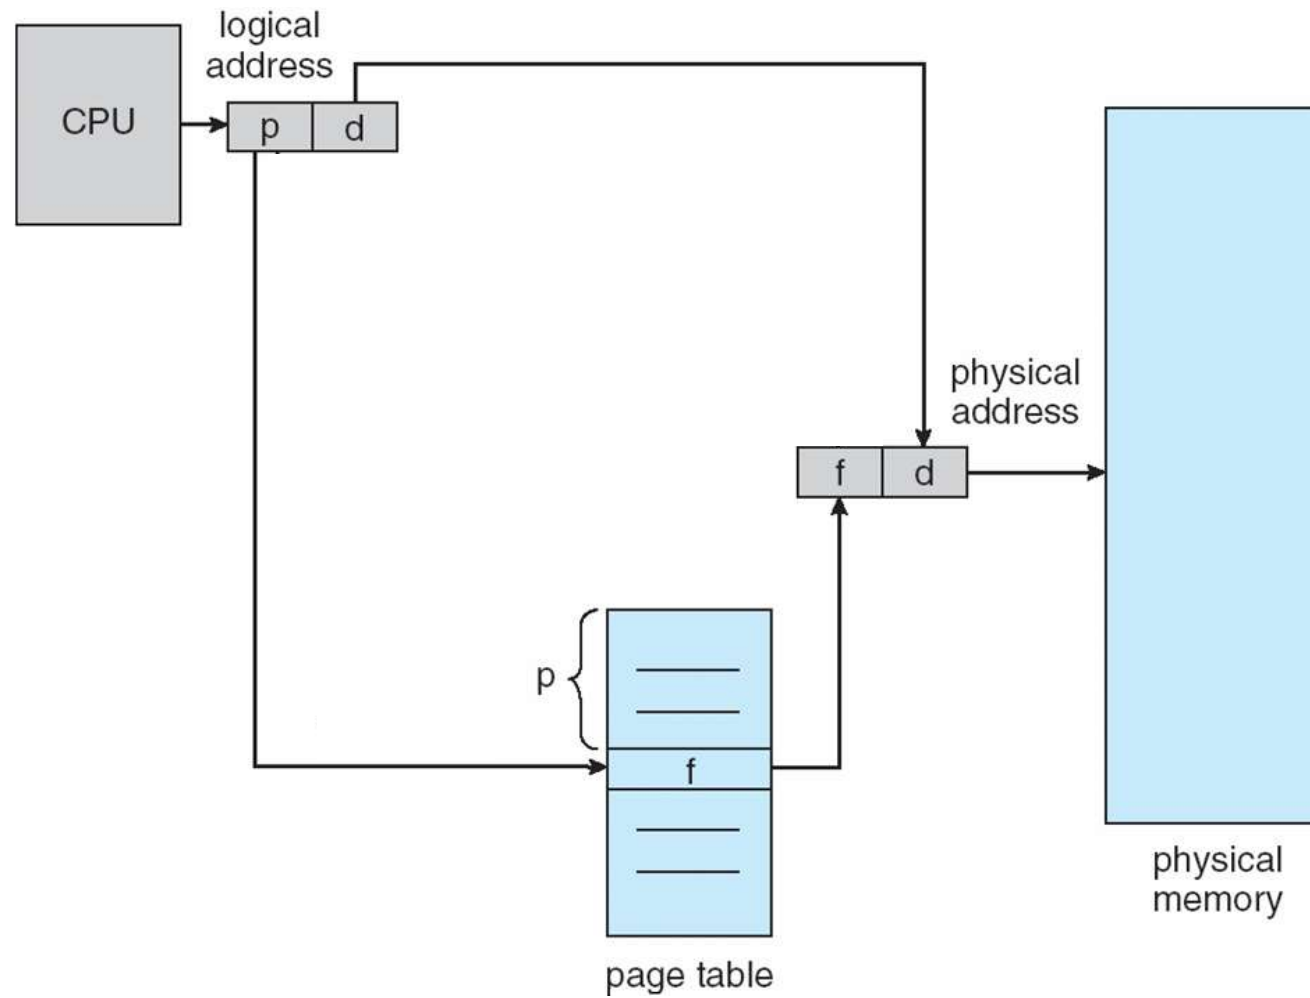


# TLB

---


- **Recently-referenced** Page Table Entries (PTE's) are stored in the TLB
  - only small part of the page table is kept in TLB, typically 32 ~ 1,024 entries
- **TLB hit:** memory reference is in TLB
- **TLB miss:** memory reference is not in TLB
  - A TLB miss incurs an additional memory access to the page table
  - In this case, the page number/frame number is *added* to TLB for next page reference

# TLB



# TLB with cache

- Memory access order
  - **Logical (virtual) address**
  - Check **TLB**
    - **TLB hit**: get address from TLB
    - **TLB miss**: check page table in memory
  - Acquire **physical address**
  - Check **cache**
    - **cache hit**: get data from cache
    - **cache miss**: get data from main memory
  - Acquire data



# Effective Access Time (EAT)

## ■ Hit ratio

- percentage of times that a page number is found in **TLB**
- e.g., 80% hit ratio: find the page in TLB 80% of the time

## ■ Assume

- memory access time = 100 ns, TLB access time = 1ns
- Hit ratio =  $\alpha$

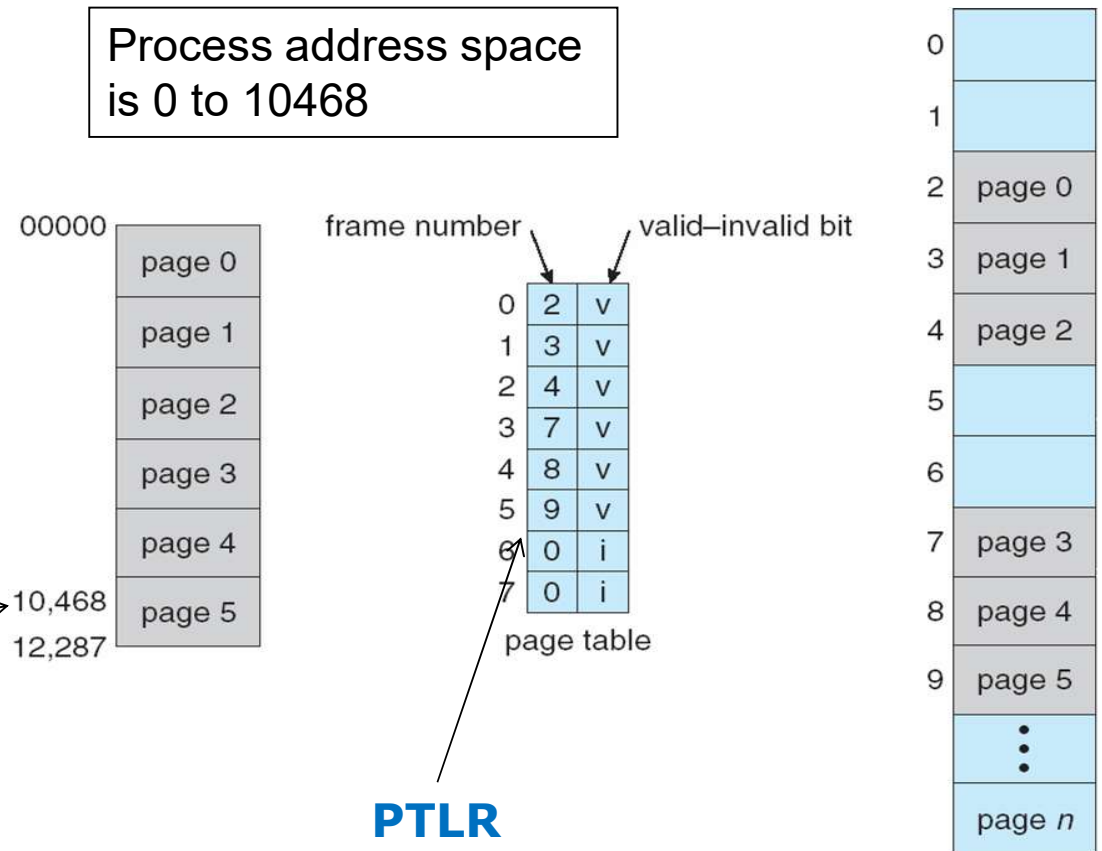
## ■ Effective Access Time (EAT)

$$\text{EAT} = \alpha \times (1 + \underbrace{100}_{\text{data/instruction}})\text{ns} + (1 - \alpha) \times (1 + \underbrace{100}_{\text{TLB}} + \underbrace{100}_{\text{page table}} + \underbrace{100}_{\text{data/instruction}})\text{ns} = 201 - 100\alpha$$

if  $\alpha = 0.8$  then EAT = 121ns, if  $\alpha = 0.99$  then EAT = 102ns

# Memory Protection

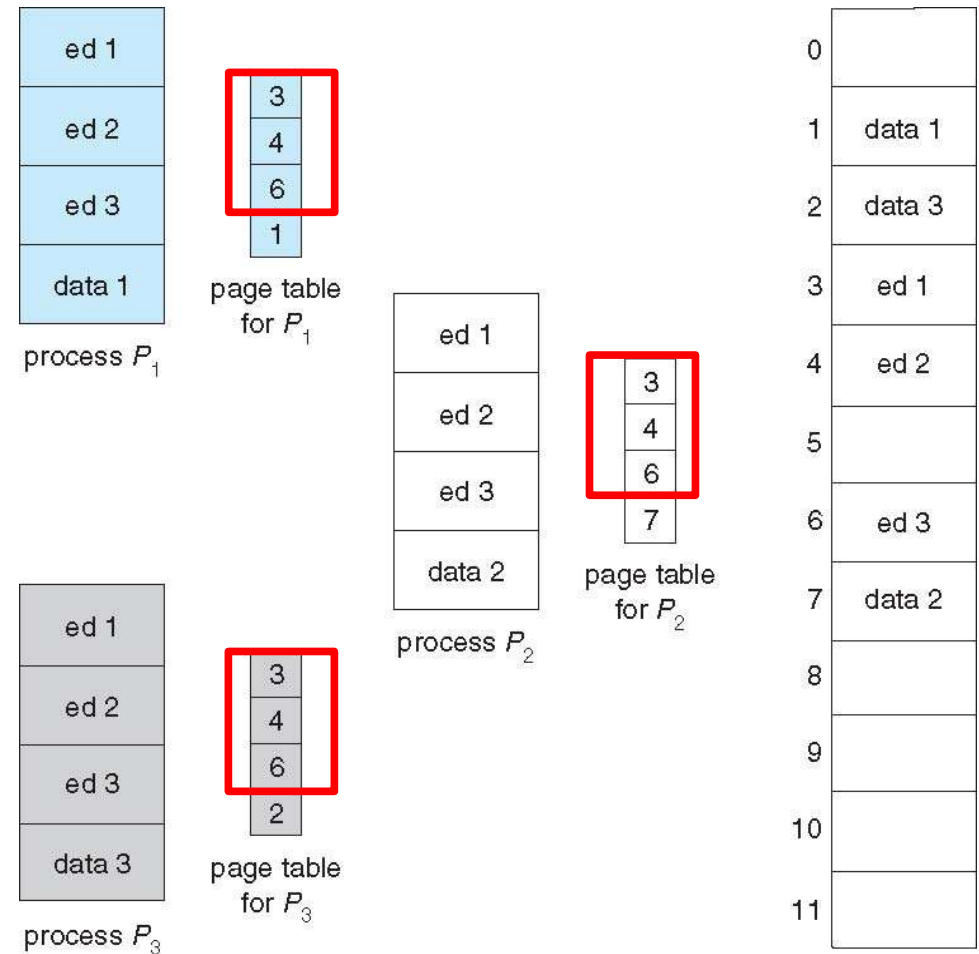
- Valid-invalid bit attached to each entry in the page table:
  - valid bit** indicates that the associated page is in the process' logical address space, and is thus a *legal* page
  - invalid bit** indicates that the page is *not* in the process' logical address space
- Internal fragmentation
  - page 5 is valid, but some parts of the page 5 are invalid



# Shared Page (Shared Memory)

## ■ Shared code

- One copy of code shared among processes
  - ▶ Example: 40 users run same program using 150KB code and 50KB data
    - need total 8,000KB physical memory
  - ▶ Share the code!
    - need total  $2,000 + 150 = 2,150$ KB physical memory
- IPC, Multiple threads sharing the same memory space



## Why Paging? (3)

---

- Provides **Protection** to memory
  - Programmer can never access memory outside the page table.
  - Programmer can never access memory outside the page table length defined in **page-table length register (PTLR)**
- Provides **Shared memory**
  - Use shared pages

# Chapter 8: Memory Management Strategies

---

- Background
- Contiguous Memory Allocation
- Paging
- **Structure of the Page Table**
  - Hierarchical Paging (Two-level page table)
  - Hashed Page Table
  - Inverted Page Tables




# How to Store Page Table

---

- Consider a system with 32-bit logical address space
  - 32-bit (4GB) logical address space
  - If page size: 4KB ( $=2^{12}\text{B}$ )
  - Each page table needs up to 1 million ( $2^{20}$ ) entries  $\rightarrow$  4MB contiguous physical memory space for page table
- Solution? **Hierarchical paging**
  - Multi-level paging, in which the page table is also paged

# Hierarchical Page Tables

- Break up the logical address space into multiple page tables (logically)
- A simple technique is a **two-level page table**




## Two-Level Paging Example

---

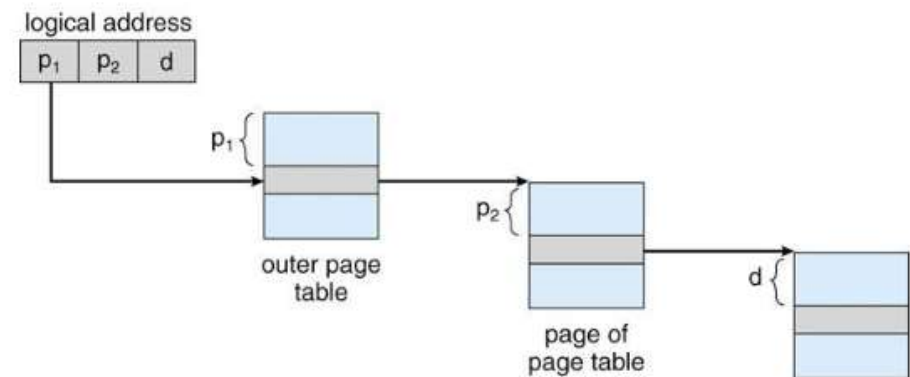
- A logical address (on 32-bit machine with 4K page size) is divided into:
  - a **page number** consisting of 20 bits, a **page offset** consisting of 12 bits.
  - There should be  $2^{20} \approx 10^6 = 1$  million page entries ( $4 \times 10^6 = 4\text{MB}$ )
    - ▶ Page table becomes too large!
- Since **the page table is paged**, the page number is further divided into:
  - a 10-bit page number.
  - a 10-bit page offset.

## Two-Level Paging Example

- Thus, a logical address is as follows:

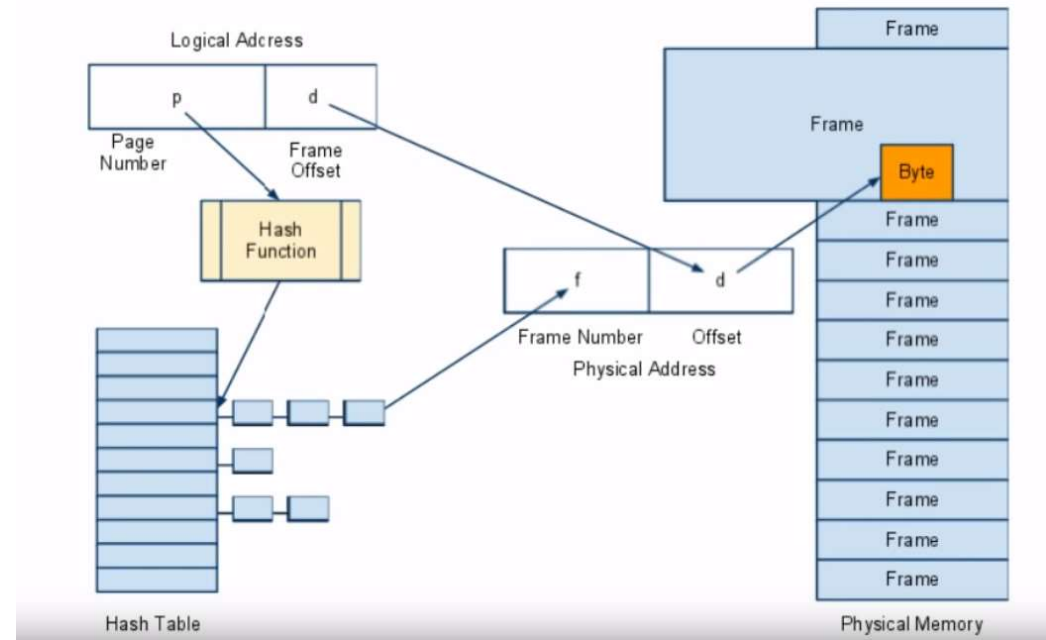


- where  $p_1$  is an index into the **outer page table** ( $2^{10} = 1,024$  entries: 4KB), and  $p_2$  is the displacement within the page of the **inner page table** ( $2^{10} = 1,024$  entries: 4KB).



# Hashed Page Table

- Address spaces on 64-bit architectures are large
  - Page table search times can become excessive
- Hash table improves search times
- Page table stores linked list of page-to-frame mappings



# Inverted Page Table

- Usually, each process has a page table
  - Translate logical address (page) into physical address (frame)
  - Each page table can be very large! (as seen in last few slides)

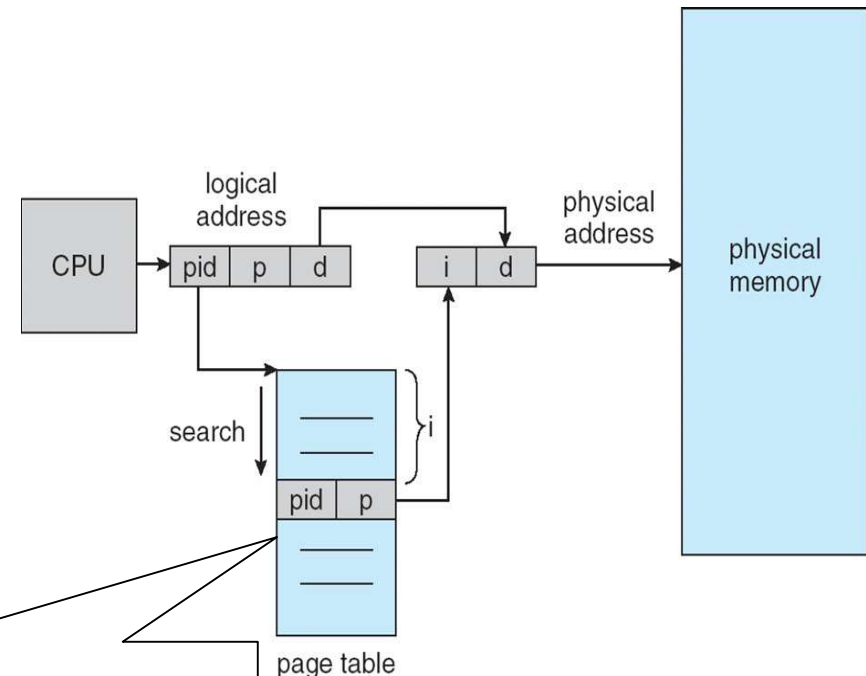
- **Inverted page table**

- Only one page table in OS!
- One entry for each real frame of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page

- Decreases memory needed to store each page table

- Total table size = 1M entries (4MB size)

- Any disadvantages?



PT: page# → frame#  
inverted PT: frame# → pid, page#

# Inverted Page Table

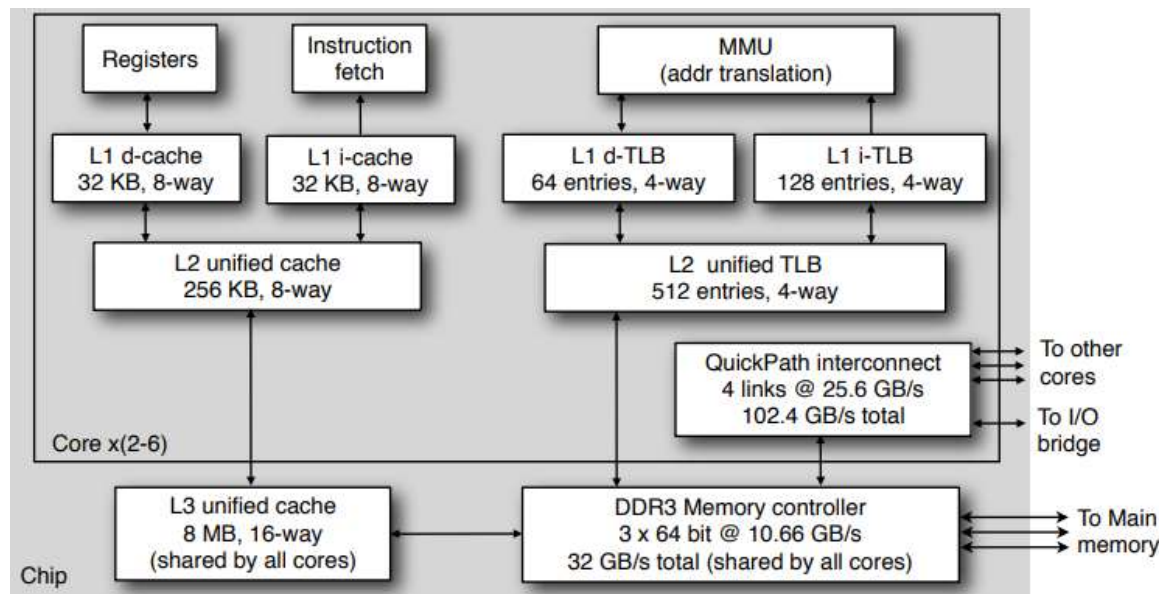
---

- Increases time needed to search the table when a page reference occurs
  - If there are  $N$  entries, may need up to  $N$  searches!
- Solution
  - Use **hash table** to limit the search page-table entries
  - **TLB** can accelerate access
- Difficult to implement shared memory?
  - One mapping of a virtual address to the shared physical address

# Example: Intel Core i7

## ■ Nehalem microarchitecture

- Allows for full 64-bit virtual and physical address
- Current implementations (Bloomfield, Lynnfield, etc.)
  - ▶ 48-bit virtual address (256 TB)
  - ▶ 52-bit physical address (4 PB)
- Page size is either 4KB or 4MB (Linux uses 4KB)





# Summary

---

- Various memory management schemes
  - Contiguous memory allocation
  - Noncontiguous memory allocation – paging
- H/W support for memory protection
  - Base and limit registers
  - TLB – associative mapping table
- Fragmentation problems
  - Internal and external fragmentation
- Paging system