


Chapter 9: Virtual-Memory Management

School of Computing, Gachon Univ.
Jungchan Cho



Most slides from "Operating System Concepts – 10th Edition".
Many slides are taken from lecture notes of Prof. Joon Yoo.

Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging
- To explain the page-replacement algorithms


Chapter 9: Virtual-Memory Management

- **Background**
- Demand Paging
- Page Replacement
- Page Replacement Algorithms

Chapter 8: Swapping

- Q: What if memory needed by process exceeds physical memory size?
 - In other words, can process size > physical memory?
- **Input queue** – collection of processes on the disk that are waiting to be brought into memory to run the program

- **Swapping**
 - A whole process can be swapped temporarily out of memory to a backing store (swap device – usually HDD, SSD, or Flash) then brought back into memory for continued execution



Chapter 8: Swapping

- Context switching time of swapping system is fairly high
 - e.g., user process 100MB, hard disk 50MB per second: then, $100\text{MB} / 50\text{MB per second} = 2 \text{ seconds}$. Swap-in + swap-out = $2 \text{ seconds} \times 2 = 4 \text{ seconds}$
- Swapping takes too much time! Modern OSs try to avoid swapping – only enable if available physical memory is small
- Can we do better?
 - Swap only portions of process (rather than entire process) to decrease swap time: **Demand Paging**

Background


- Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures (arrays)
- Entire program code not needed at same time
- Consider ability to execute **partially-loaded program**
 - e.g., 500MB program. Only need to load 100MB to start program. The rest can be loaded on-demand!

Chapter 9: Virtual-Memory Management

- Background
- **Demand Paging**
 - Basic Concepts
 - Valid-Invalid Bit
 - Page Fault
 - Performance of Demand Paging
- Page Replacement
- Page Replacement Algorithms

Demand Paging Concept

- We can view a process as
 - a sequence of pages (rather than large contiguous address spaces)
- How can an executable program be loaded from disk into memory?
 - **Option 1**
 - ▶ Load the entire program into memory
 - **Option 2**
 - ▶ **Load pages only as they needed**




Demand Paging

- Concept:
 - Instead of loading the whole program into memory, bring pages into memory only when it is needed

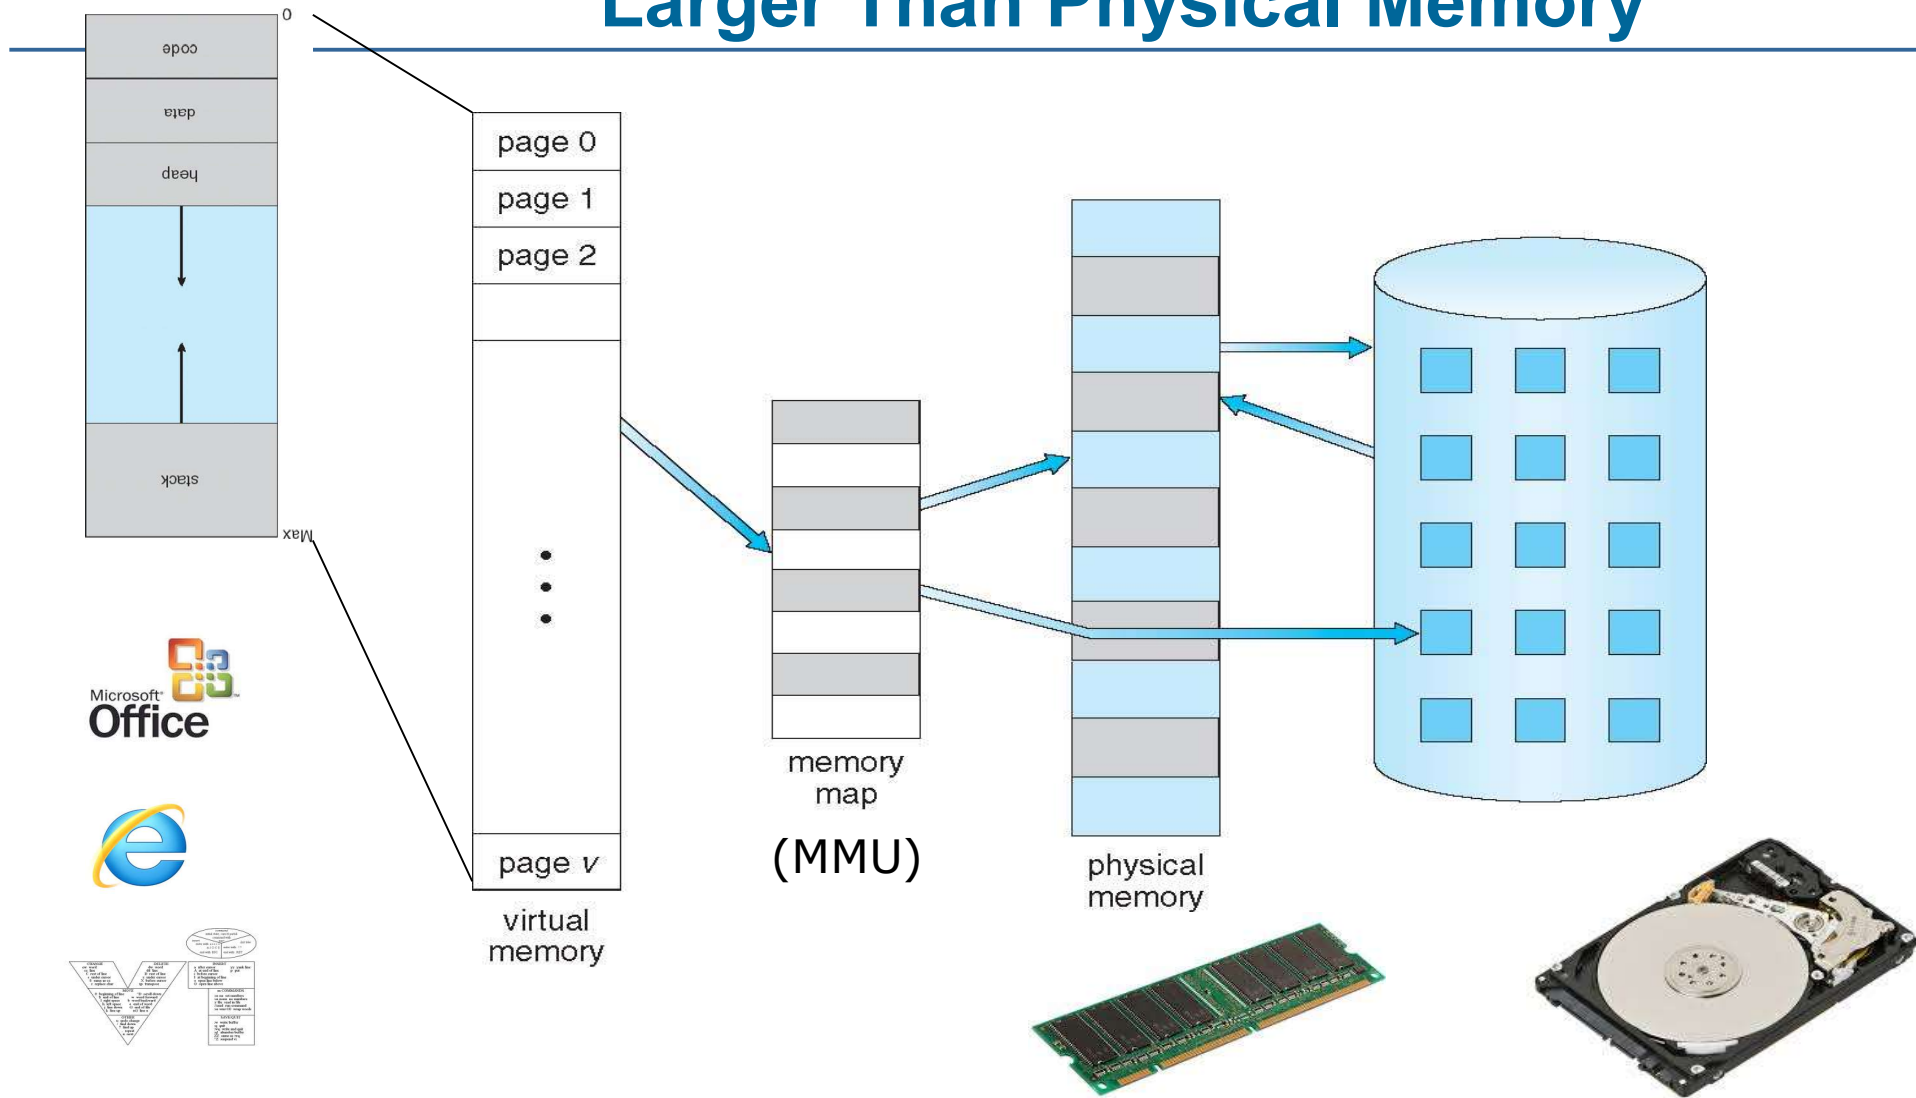
- Why? (Advantages)
 - Less I/O needed from disk to load memory
 - ▶ Faster response to load programs
 - Less memory needed
 - ▶ Larger programs
 - ▶ Better multi-program – more users for server

Demand Paging

- Pages are loaded only when they are demanded during execution
 - Pages that are never accessed are never loaded into physical memory
- **Pager**
 - Conducts **demand paging**
 - c.f., swapper: swaps in/out the whole process




Virtual Memory That is Larger Than Physical Memory



Demand paging usage: Heap/Stack

- Heap: grow upward in memory as it is used for dynamic memory allocation
- Stack: grow downward in memory through successive function calls
- Large blank space is part of virtual address space, but will require actual physical pages only if the heap and stack grows
- Using **sparse** address spaces with holes left for growth can be filled during execution via **demand paging**
 - as the stack or heap segments grow



Demand paging usage: Dynamically linked libraries (DLL)

- **Dynamically linked libraries (DLL)**
 - Library linking is postponed until execution time
- Now some pages are in memory but some are still in disk (not in memory)
- Q: How can distinguish this? - **Need H/W support**
 - **Valid-Invalid Bit**

Valid-Invalid Bit

- With each page table entry, a **valid-invalid bit** is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
 - or it can be represented as **1** / **0** for valid/invalid.
- Initially valid-invalid bit is set to **i** on all entries
- Example of a page table snapshot:
- During address translation, if valid-invalid bit in page table entry is **i** \Rightarrow **page fault**

Frame #	valid-invalid bit
	v
	i
	v
	v
	i
....	
	v
	i

page table



Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

	valid-invalid bit	
frame		
0	4 v	
1	i	
2	6 v	
3	i	
4	i	
5	9 v	
6	i	
7	i	
	page table	


0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory



Page Fault


■ Page fault



- If there is a reference to an **invalid** page
- Reference to that page will **trap (software interrupt)** to operating system

■ OS handling page fault

1. Operating system looks at page table to decide:
 - ▶ Illegal reference (address outside of virtual memory space) \Rightarrow abort
 - ▶ **validation bit = i** \Rightarrow Just not in memory
2. Get empty frame from **physical memory**
3. Swap page from **disk** into **physical memory frame**
4. Reset **page table** to indicate **page** now in memory
Set **validation bit = v**
5. **Restart** the instruction that caused the page fault



Aspects of Demand Paging

- Hardware support needed for demand paging
 - **Page table** with valid / invalid bit
 - **Secondary storage**
 - ▶ Swap device with **swap space** (usually disk)
 - Restart Instruction (after page fault)

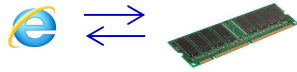
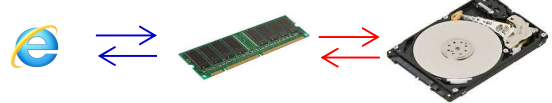
Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- **Effective Access Time (EAT)** for *demand paging*

$$\text{EAT} = (1 - p) \times \underline{\text{memory access}} + p \times \underline{\text{page fault time}}$$

page fault time = page fault overhead
+ swap page out
+ swap page in
+ restart overhead

Demand Paging Example

- Memory access time = **200 ns** 
- Average page-fault service time = **8 ms** (=8,000,000 ns) 
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 200 \text{ ns} + p (8 \text{ ms}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \text{ (ns)} \\ &= 200 + p \times 7,999,800 \text{ (ns)} \end{aligned}$$
- If one access out of 1,000 causes a page fault ($p=0.001$), then
$$\text{EAT} = 8,200 \text{ ns} (= 200 \text{ ns} \times 41)$$

This is a slowdown by a factor of 41x!! (or 97.6% degradation)
- If we want performance degradation < 10 % (or $\text{EAT} < 220 \text{ ns}$)
 - $220 > \text{EAT} = 200 + 7,999,800 \times p$
 $20 > 7,999,800 \times p$
 - $p < .0000025$
 - < one page fault in every 400,000 memory accesses
- It is important to keep the page-fault rate low in a demand-paging system!**

Chapter 9: Virtual-Memory Management


- Background
- Demand Paging
- **Page Replacement**
 - Basic Page Replacement
 - Modify (Dirty) Bit
- Page Replacement Algorithms

What Happens if There are no Free Frames?

- All memory may be in use
 - Multiprogramming – multiple processes share the memory
 - Page fault → OS finds the desired page on the disk → find the free-frame
 - If NO free frames on the **free-frame list** ? – i.e., physical memory is full!
- **Page replacement**
 - Find some page in memory, but not really in use, **page out**
 - Performance – want an algorithm which will result in minimum number of page faults
 - ▶ **Page replacement Algorithms**





Need For Page Replacement



Basic Page Replacement

- Find the location of the desired page on disk
- Find a **free frame**:
 - If there is a free frame, use it
 - If there is no free frame, use a **page replacement algorithm** to select a **victim frame**
- Read the desired page into the (newly) free frame. Update the page and frame tables.
- Restart the process



Modify Bit (= Dirty Bit)

- If no frames are free → **TWO** page transfers are required
 - One page out and one page in (see previous slide)
 - Can we reduce the number of disk writes (page-outs)?
 - Solution: **Modify bit (dirty bit)**
 - Initially set to zero
 - Set the bit (to 1) whenever any word or byte in the page is modified
 - ▶ only modified pages are written to disk – unmodified pages can just be deleted without writing on disk (it has **not** changed!)
- can reduce the page-transfer overhead

Benefits of Virtual Memory (using Demand Paging)

■ Protection

- A bug in one process can corrupt memory in another
- A process can only access the virtual address (pages)

■ Transparency

- A process shouldn't require particular physical memory
- A process sees a contiguous virtual memory space (pages)

■ Resource exhaustion

- Sum of sizes of all processes often greater than physical memory
- **Demand paging**
 - ▶ e.g., If a process is 20 pages, we can execute with 10 frames and the rest in the swap space (disk) – use demand paging and page-replacement

Chapter 9: Virtual-Memory Management

- Background
- Demand Paging
- Page Replacement
- **Page Replacement Algorithms**
 - FIFO Page Replacement
 - Optimal Page Replacement
 - LRU Page Replacement


Page and Frame Replacement Algorithms

■ Page-replacement algorithm

- Which frame should be replaced?
- Want lowest page-fault rate

■ Evaluate algorithm by computing the number of page faults

- The number of page faults will decrease as the number of frames available increases



Page Replacement Algorithms

- FIFO page replacement
- Optimal page replacement
- LRU (Least-Recently-Used) algorithm
 - clock, stack, reference-bit

First-In-First-Out (FIFO) Algorithm

- When a page must be replaced, the oldest page is chosen

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2	2	4	4	4	0			0	0			7	7	7
	0	0	0		3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

page frames

15 Page Replacements!!

- How to track ages of pages?
 - Just use a FIFO queue

Problem of FIFO Algorithm

■ Example 2

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

-3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

-4 frames:


1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

Why did this happen???

FIFO Illustrating Belady's Anomaly

- FIFO replaces the **oldest** page – is this always good?
 - An old page may contain a **frequently used variable**
- **Belady's anomaly**
 - Sometimes, page-fault may increase as the number of allocated frames increase!



Optimal Algorithm

- Algorithm that has the lowest page-fault rate
- Replace page that **will not** be used for longest period of time
 - This is a design to guarantee the lowest page-fault rate for a fixed number of frames

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2			2			2				7		
	0	0	0		0		4			0			0				0		
		1	1		3		3			3			1				1		

page frames

9 Page Replacements!!

Optimal Algorithm

- Unfortunately, the optimal page-replacement algorithm is difficult to implement
 - Why?
 - ▶ Require **future knowledge** of the reference string – We do not know the future!
- Mainly used for comparison studies
 - Evaluating a new algorithm
 - e.g. “the new algorithm is within 12.3% of optimal at worst and within 4.7% on average”

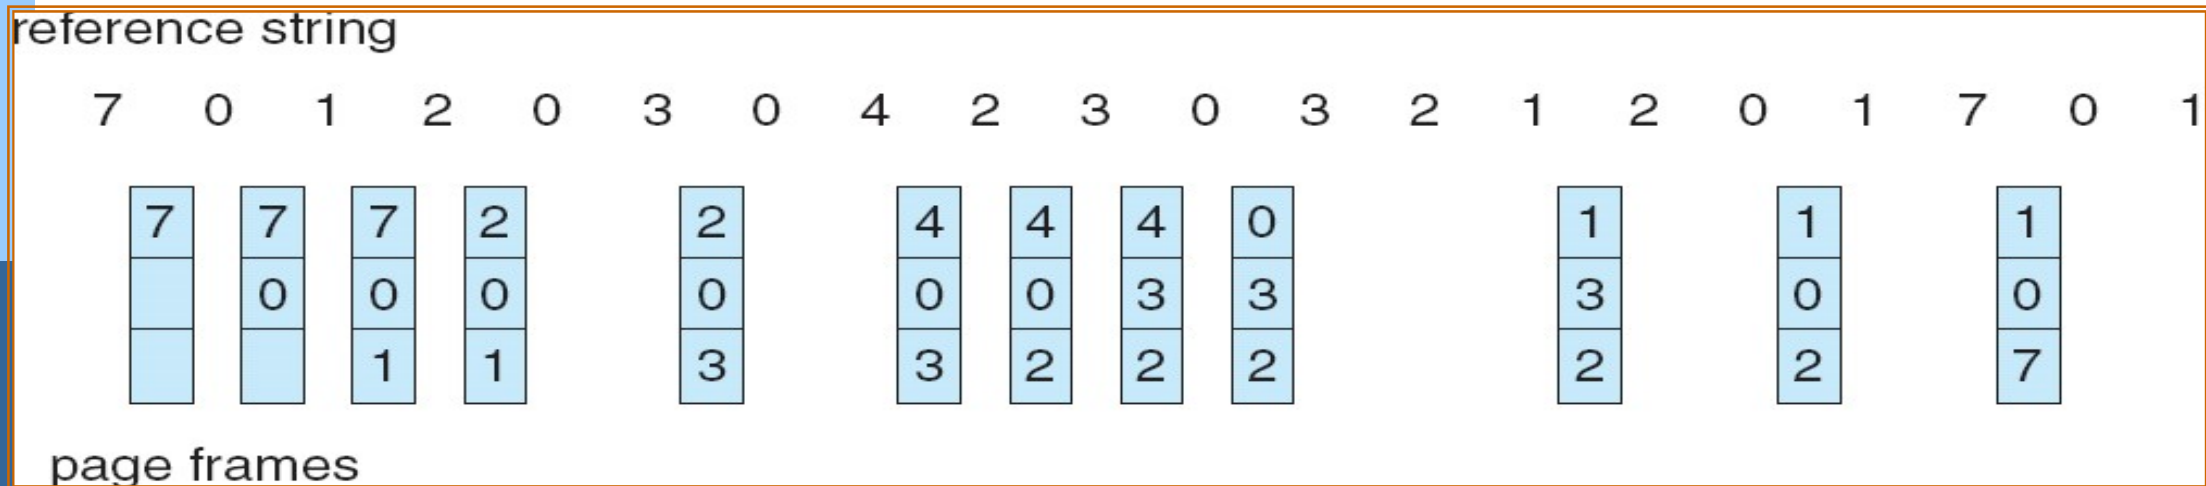
Least Recently Used (LRU) Algorithm

- FIFO: Use time when a page **was brought** into the memory
- OPT: Use time when a page is to be **used**
- Let us predict the future by using the past!

- **Least Recently Used (LRU) Algorithm**
 - Use past knowledge rather than future
 - **Replace** page that has **not been used** in the most amount of time
 - Generally good algorithm and frequently used

Least Recently Used (LRU) Algorithm

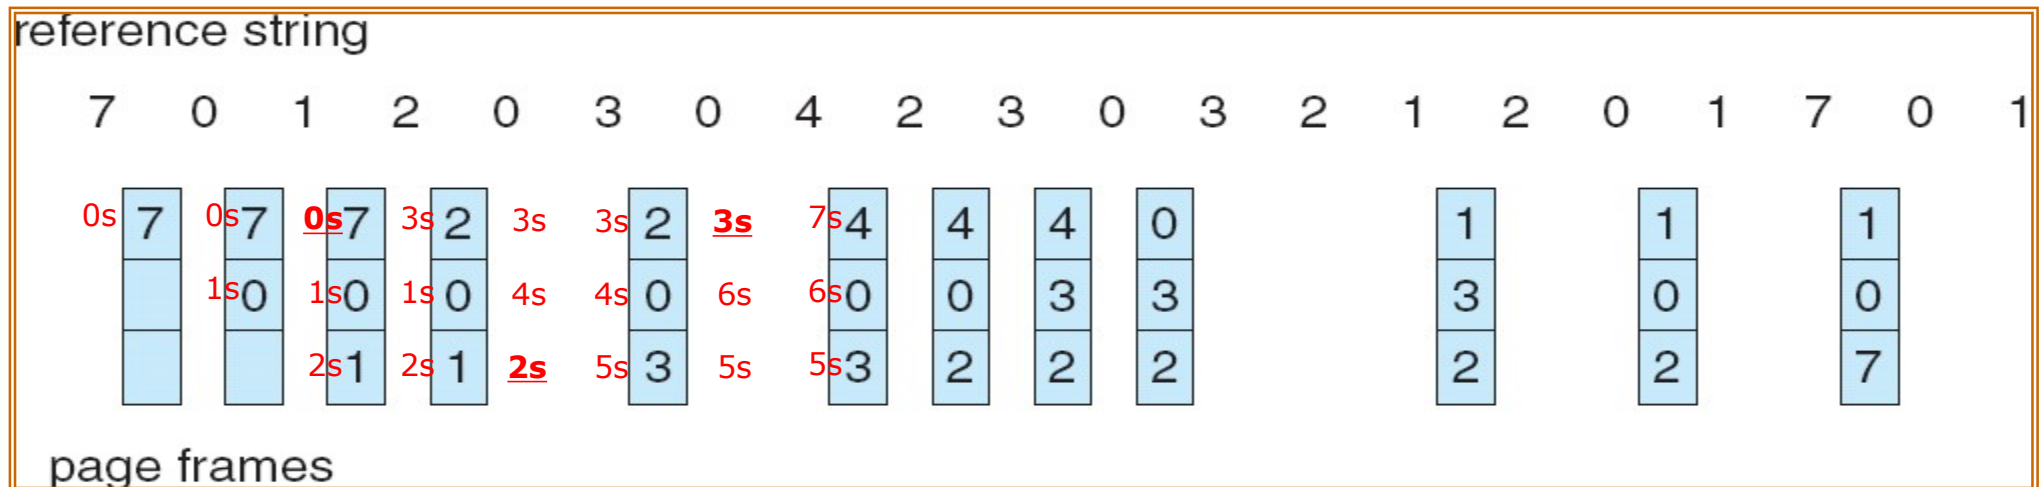
- LRU replacement associates with each page the time of that page's last use
- Replace the page that has not been used for the longest period of time



12 Page Replacements!!

Implementation of LRU Algorithm

- The major problem is **how to implement LRU replacement**
 - Use *counter* or *stack*
 - LRU implementation requires H/W support (due to speed)
- Counter implementation example




Implementation of LRU Algorithm

- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters – replace page with **smallest** timer value
 - ▶ Con1: Need to **search** to replace page
 - ▶ Con2: Timer has to be updated for each memory reference

Implementation of LRU Algorithm

- Stack implementation
 - Page referenced:
 - ▶ move it to the top
 - Use doubly-linked list
 - ▶ Why not array?
 - Con: Replacement requires 6 pointers to be changed at worst
 - No search for replacement



LRU Implementation

- Updating of **clock fields** or **stack** must be done for **every** memory reference
- LRU needs special hardware and still slow
 - Timer and stack approach use too much resources
- LRU Approximation: **Reference bit**
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace any with reference bit = 0 (if one exists)
 - In this method, although we can not know the order of page reference, we can know that whether the page is referenced or not

Summary

- Virtual memory is how we stuff large programs into small physical memories.
- We perform this magic by using demand paging, to bring in pages only when they are needed.
- But to bring pages into memory, means kicking other pages out, so we need to worry about paging algorithms.

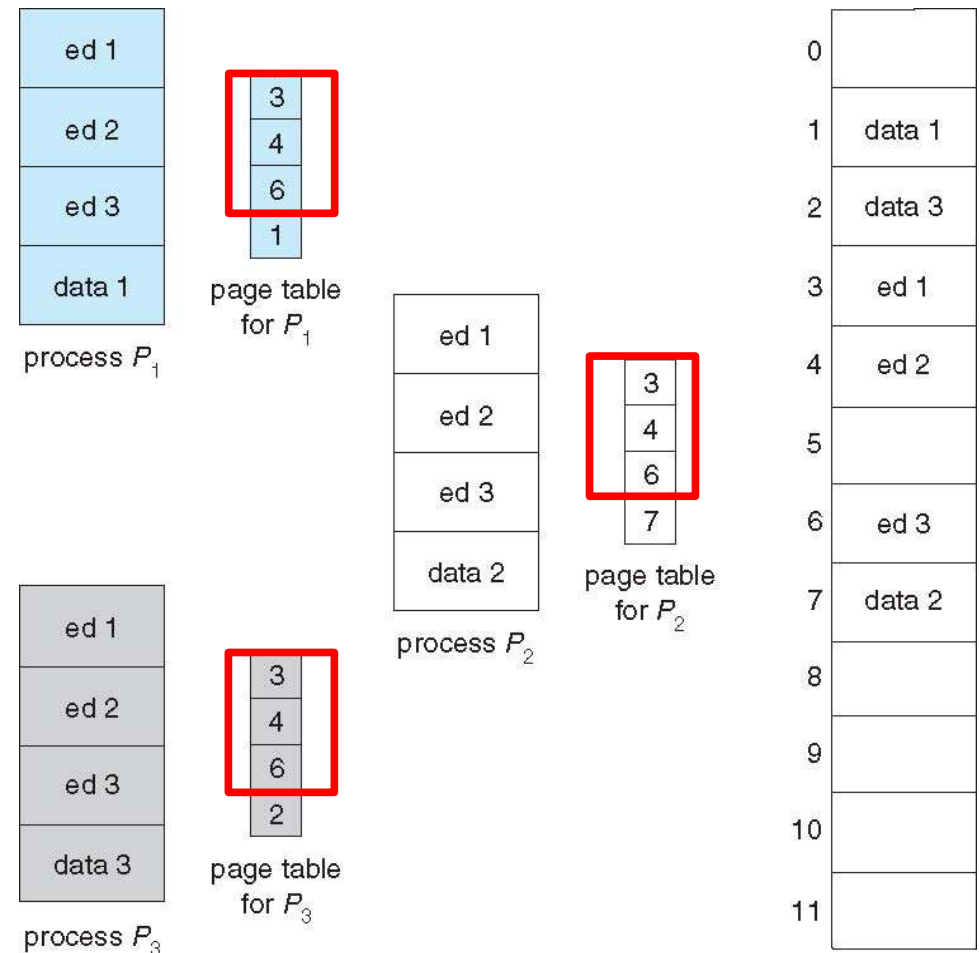
Chapter 9: Virtual-Memory Management

- Appendix
 - Sharing and Memory-Mapped Files

Shared Page


■ Shared code

- One copy of code shared among processes
 - ▶ Example: 40 users run same program using 150KB code and 50KB data
 - need total $200 \times 40 = 8,000\text{KB}$ physical memory
 - ▶ Share the code!
 - need total $150 + 50 \times 40 = 2,150\text{KB}$ physical memory
- IPC, Multiple threads sharing the same memory space



Virtual Address Space

- System libraries shared via mapping into virtual address space
- Shared memory by mapping pages read-write into virtual address space
 - System libraries can be shared
 - ▶ part of its virtual address space, but actual physical memory are shared
 - Shared memory by processes (IPC)
 - ▶ each process considers the shared memory as part of their virtual address space



fork(): Copy-on-Write

- **Copy-on-Write** (COW) allows both parent and child processes to initially *share* the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- In general, free pages are allocated
 - when the stack or heap for a process must expand
 - or when there are copy-on-write pages

