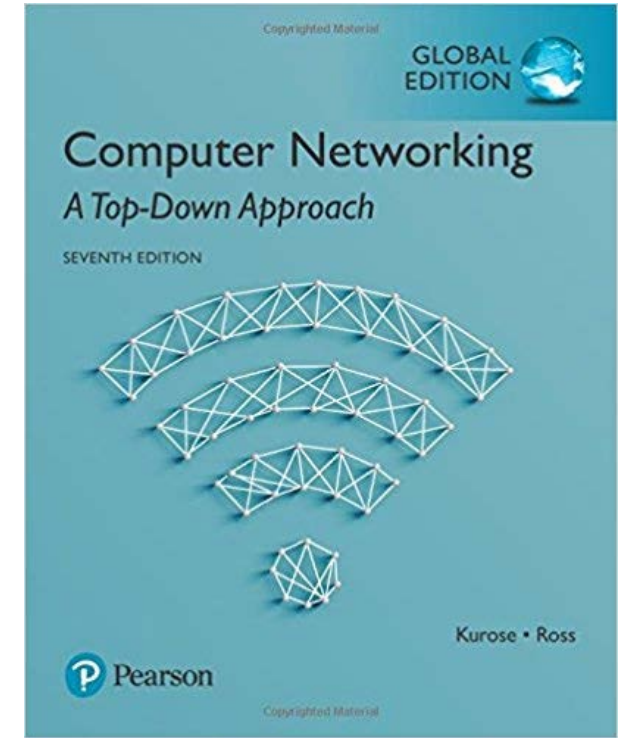


Chapter 2

Application Layer part 2

School of Computing
Gachon Univ.
Joohyung Lee

Most of slides from J.F Kurose and K.W. Ross. And, some slides from Prof. Joon Yoo



*Computer
Networking: A Top
Down Approach*

7th edition

Jim Kurose, Keith Ross
Pearson, 2017

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

Electronic Mail

- ❖ **Electronic Mail (E-mail)** has been around since the beginning of the Internet – and still remains most important and utilized application
- ❖ E-mail is an *asynchronous* communication
 - People can send out messages when it is convenient for them, no need to coordinate other people's schedules
 - Faster, easier to distribute, and inexpensive compared to postal mail

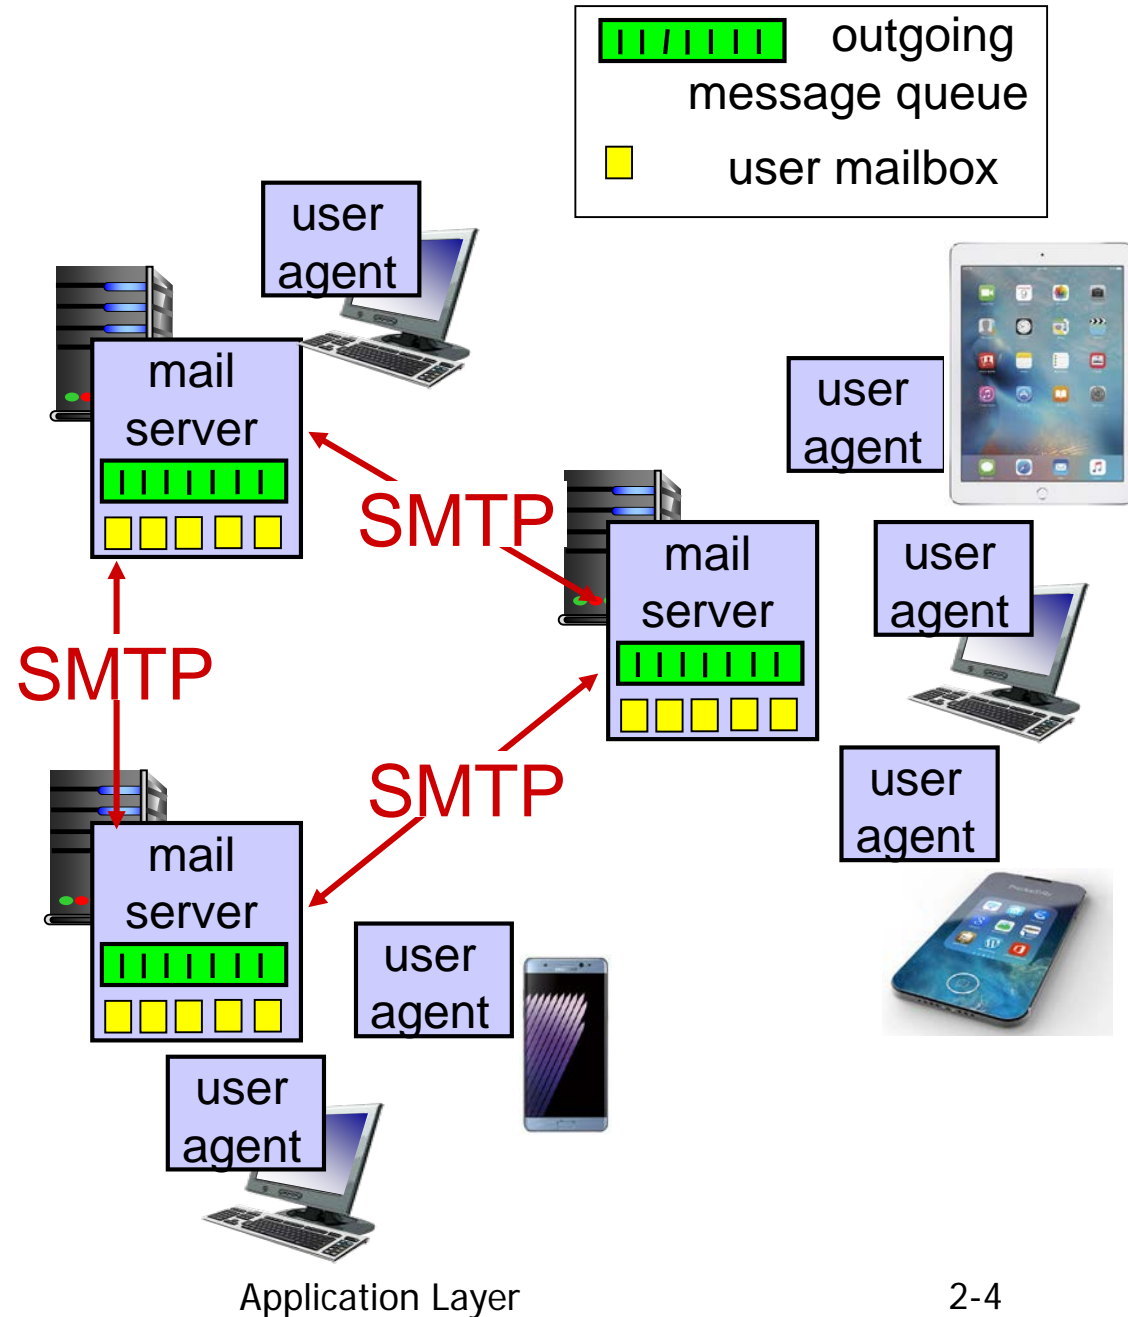
Electronic mail

Three major components:

- ❖ user agents
- ❖ mail servers
- ❖ simple mail transfer protocol: SMTP

User Agent

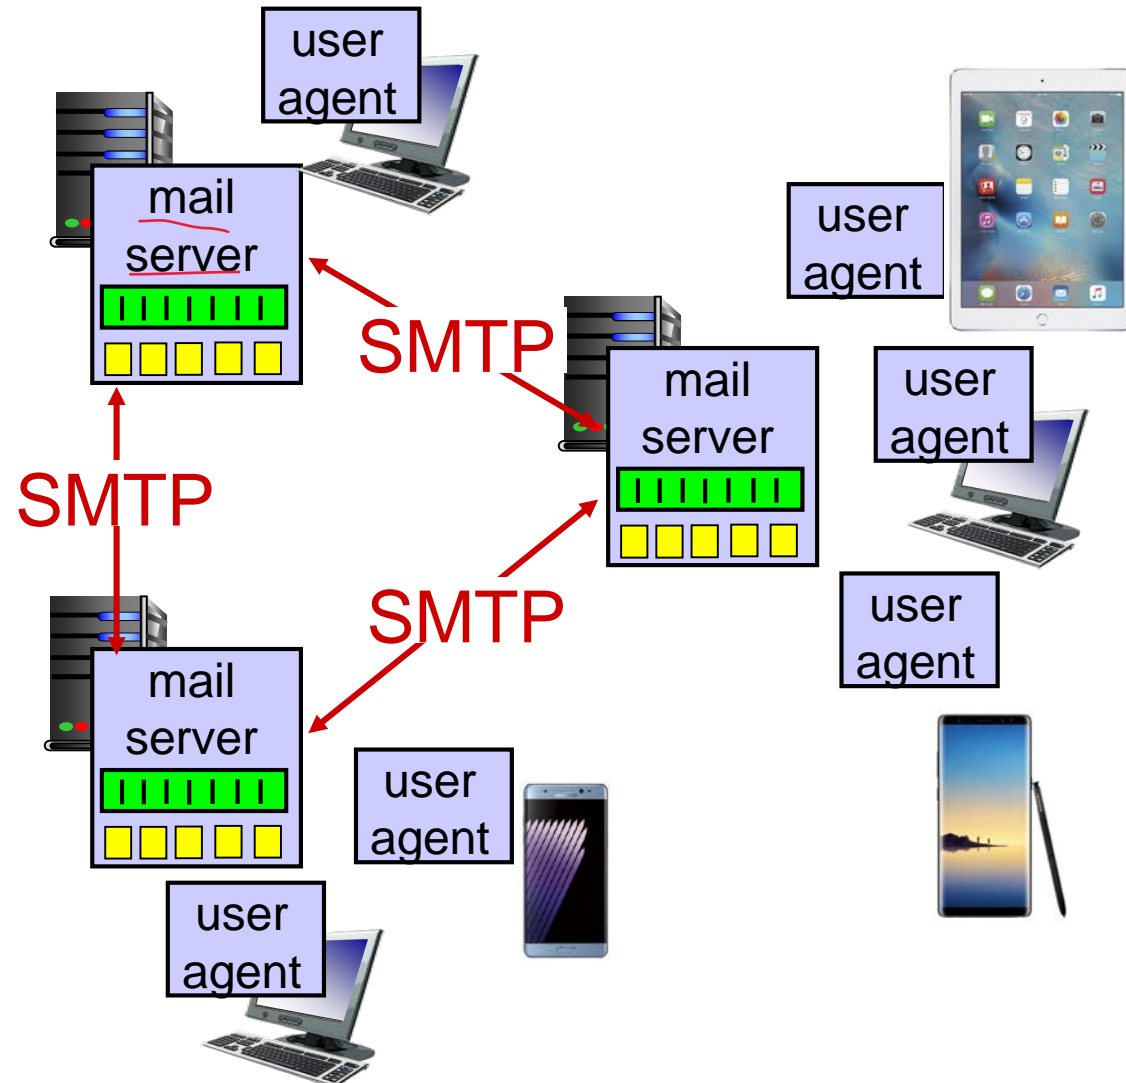
- ❖ a.k.a. “mail reader”
- ❖ composing, editing, reading mail messages
- ❖ e.g., Outlook, Naver web client, Android google web client, iPhone mail client



Electronic mail: mail servers

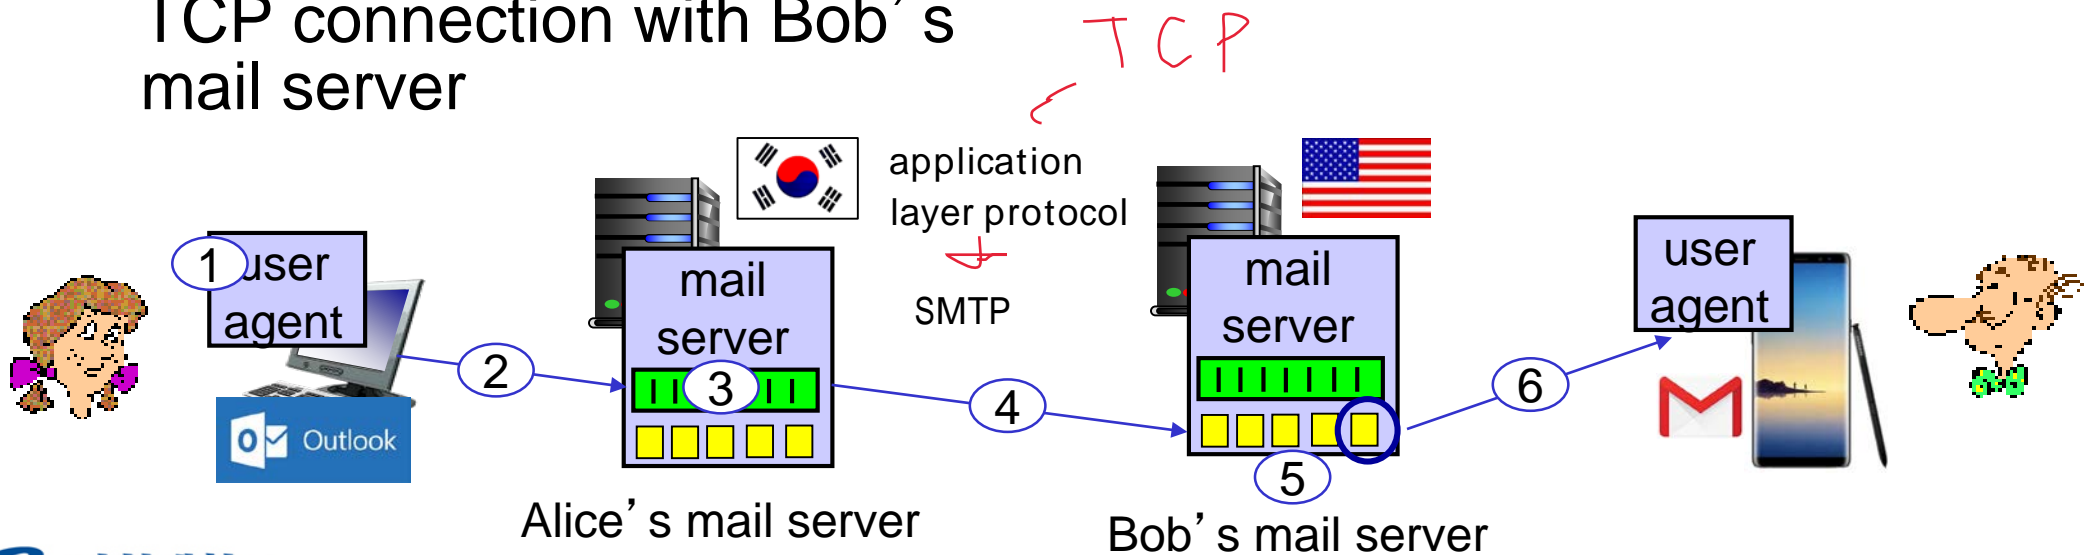
mail servers:

- ❖ *mailbox* contains incoming messages for user
- ❖ *message queue* of outgoing (to be sent) mail messages
- ❖ *SMTP protocol* between mail servers to send/receive email messages
 - Simple Mail Transfer Protocol (SMTP)



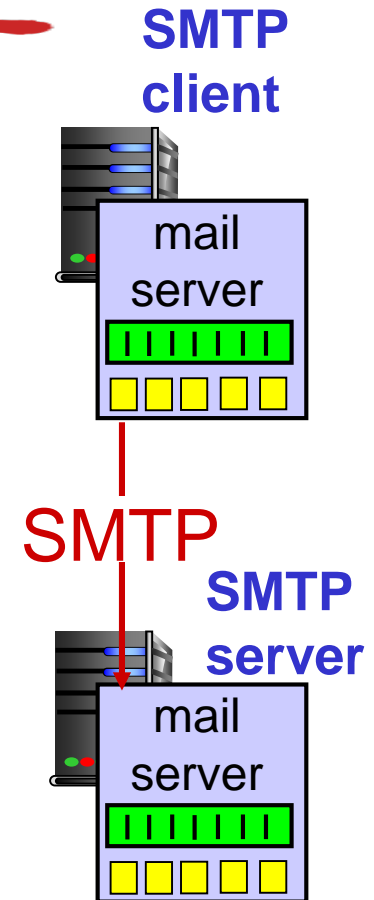
Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message "to" bob@someschool.edu
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message

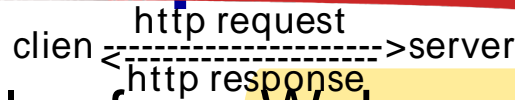


Mail Server-to-server : SMTP

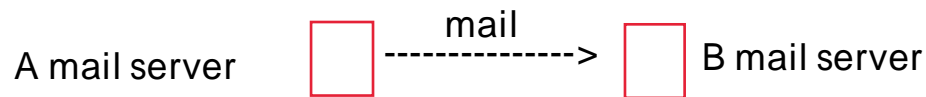
- ❖ direct transfer: sending server (SMTP client) to receiving server (SMTP server)
- ❖ uses TCP transport protocol to reliably transfer email message from client to server, port 25
- ❖ three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure



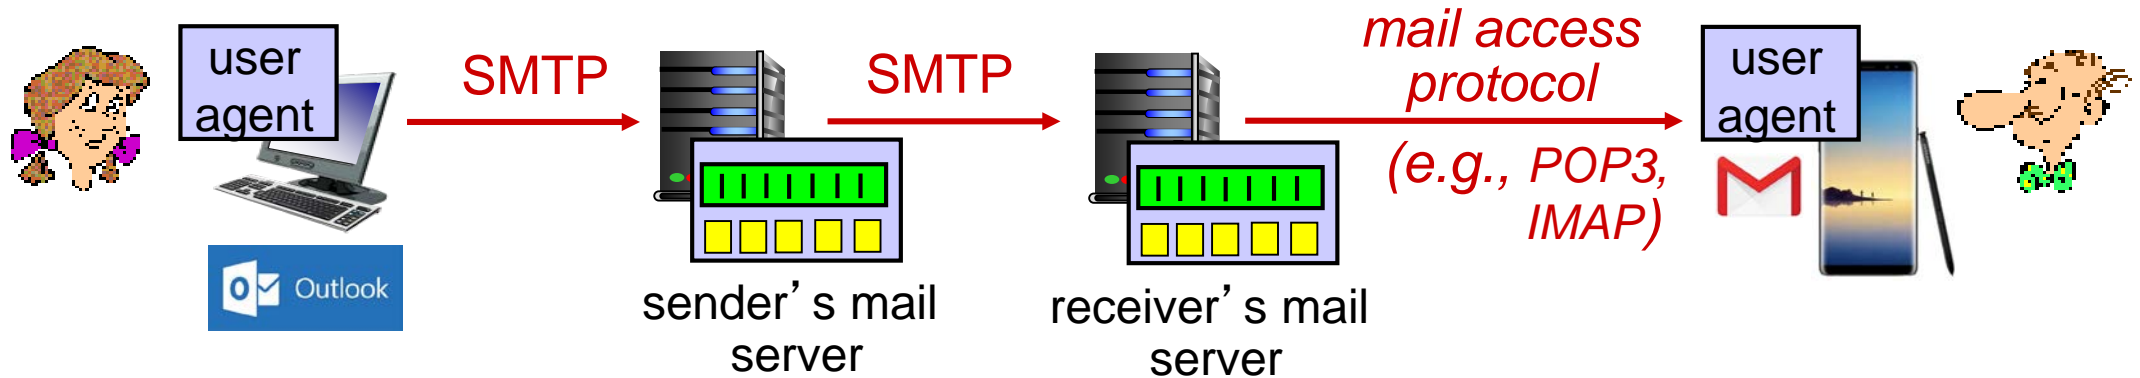
SMTP: Comparison with HTTP




- ❖ HTTP: transfer files from Web server to Web client (browser)
 - **Pull protocol:** HTTP client pulls the information from server
 - TCP connection is initiated by machine that wants to receive
- ❖ SMTP: transfer files from one mail server to another mail server
 - **Push protocol:** sending mail server pushes the file to the receiving mail server
 - TCP connection is initiated by machine that wants to send
- ❖ SMTP requires message (header & body) to be in **7-bit ASCII**
 - ❌ Need to encode all binary multimedia data into ASCII before sending over SMTP (No such restriction in HTTP) decoding
 - Image (sender) → 7-bit ASCII text (in SMTP msg) → Image (receiver)
 - This made sense in early 80s when transmission capacity was scarce, so all messages were text – but now it is archaic



Mail access protocols



- ❖ **SMTP**: delivery/storage to receiver's server
- ❖ mail access protocol: retrieval from server
 - **POP**: Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP**: Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
 - **HTTP**: gmail, Hotmail, Yahoo! Mail, etc. 

IMAP & Web mail

POP3

- ❖ Transfer mail from recipient's mail server to user agent (client)
- ❖ POP3 uses “download and delete” mode

IMAP

- ❖ keeps all messages in one place: at server – doesn't delete

Web-based E-mail

- ❖ User agent is Web browser and communicate with mailbox via HTTP (rather than SMTP, POP3, or IMAP)
- ❖ The mail server still uses SMTP to send/receive messages to/from other mail servers

push

SMTP

HTTP

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

Host name



But, Google server's actual address is **IP address**!

```
C:\Users\jyoo>ping google.co.kr  
Ping google.co.kr [59.18.46.113] 32바이트 데이터 사용:
```

host

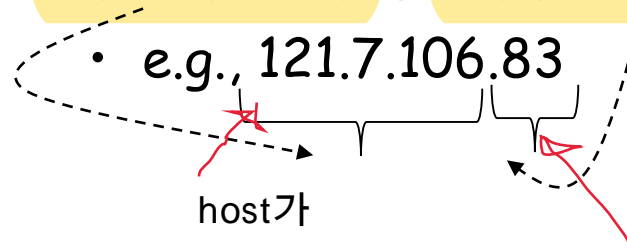
ip

IP address

The network ID identifies the network segment to which the host belongs.
The host ID identifies an individual host on some specific network segment.

❖ IP address is used to identify Internet hosts (Ch. 4)

- 4-bytes (=32 bits), *hierarchical* structure
 - e.g., 121.7.106.83
 - each period separates one byte (0~255)
 - try finding the current IP address of your machine - use commands such as ipconfig, ifconfig, ...
- It is an **ID** of the host – unique IP address for each Internet Host
- It shows the **network** of the host - IP address is composed of **network id & host id**



network id is used to locate the network

DNS: domain name system

people: many identifiers:

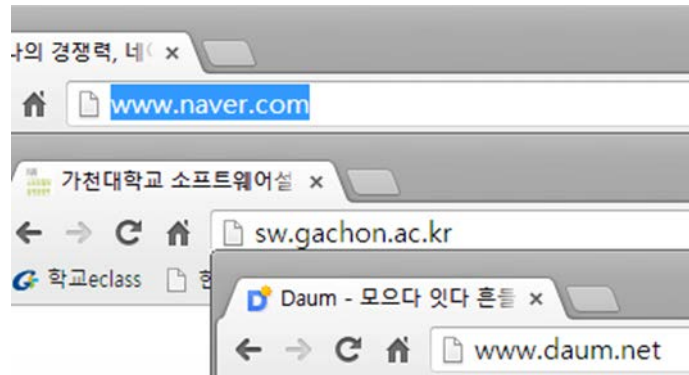
- SSN (≈주민번호), driver's license#, passport #, ...

Internet hosts, routers:

- **IP address** or **host name** (e.g., www.yahoo.com)
- People prefer _____ and routers prefer _____

Q: how to map between IP address and name, and vice versa ?

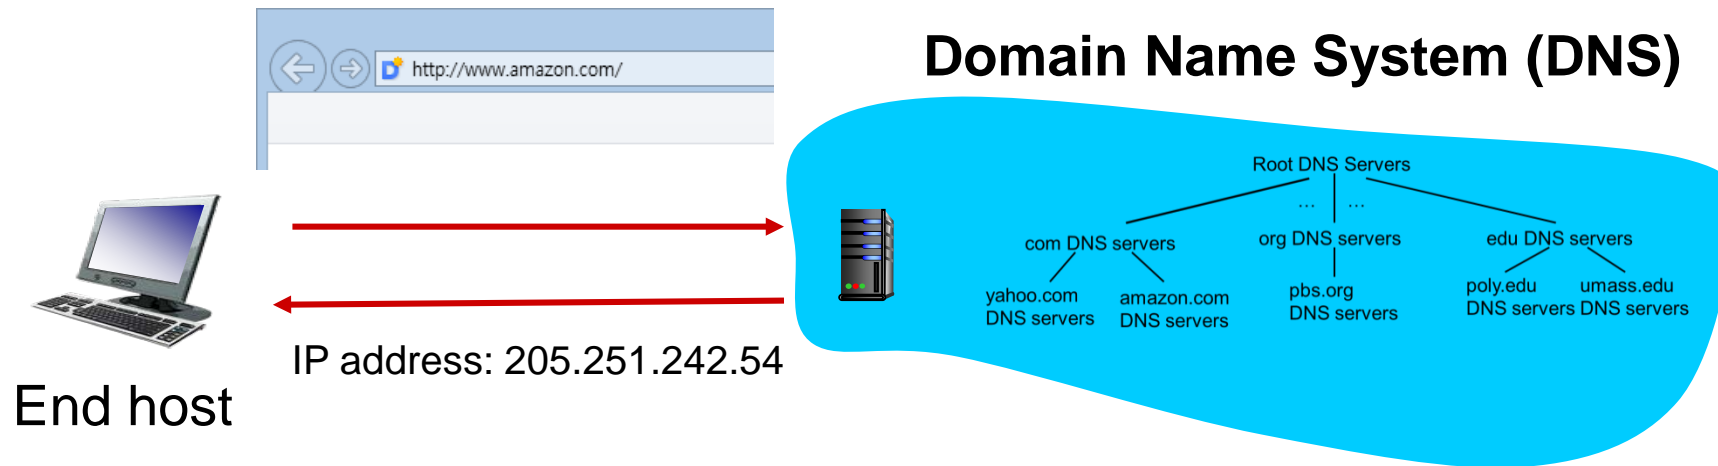
A: DNS



DNS: domain name system

Domain Name System (DNS):

- ❖ *distributed database* implemented in hierarchy of many *name servers*
- ❖ *application-layer protocol*: hosts, name servers employ DNS to *translate* host names into IP addresses



DNS: distributed vs. central

Single centralized server

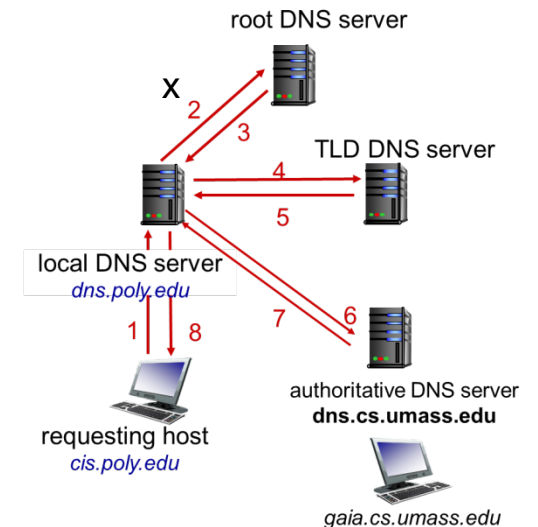
- ❖ Ask one DNS server to translate name to IP address
- ❖ Very simple!

Then, why not centralize DNS?

- ❖ single point of failure 가 1
- ❖ traffic volume
- ❖ distant centralized database
- ❖ maintenance

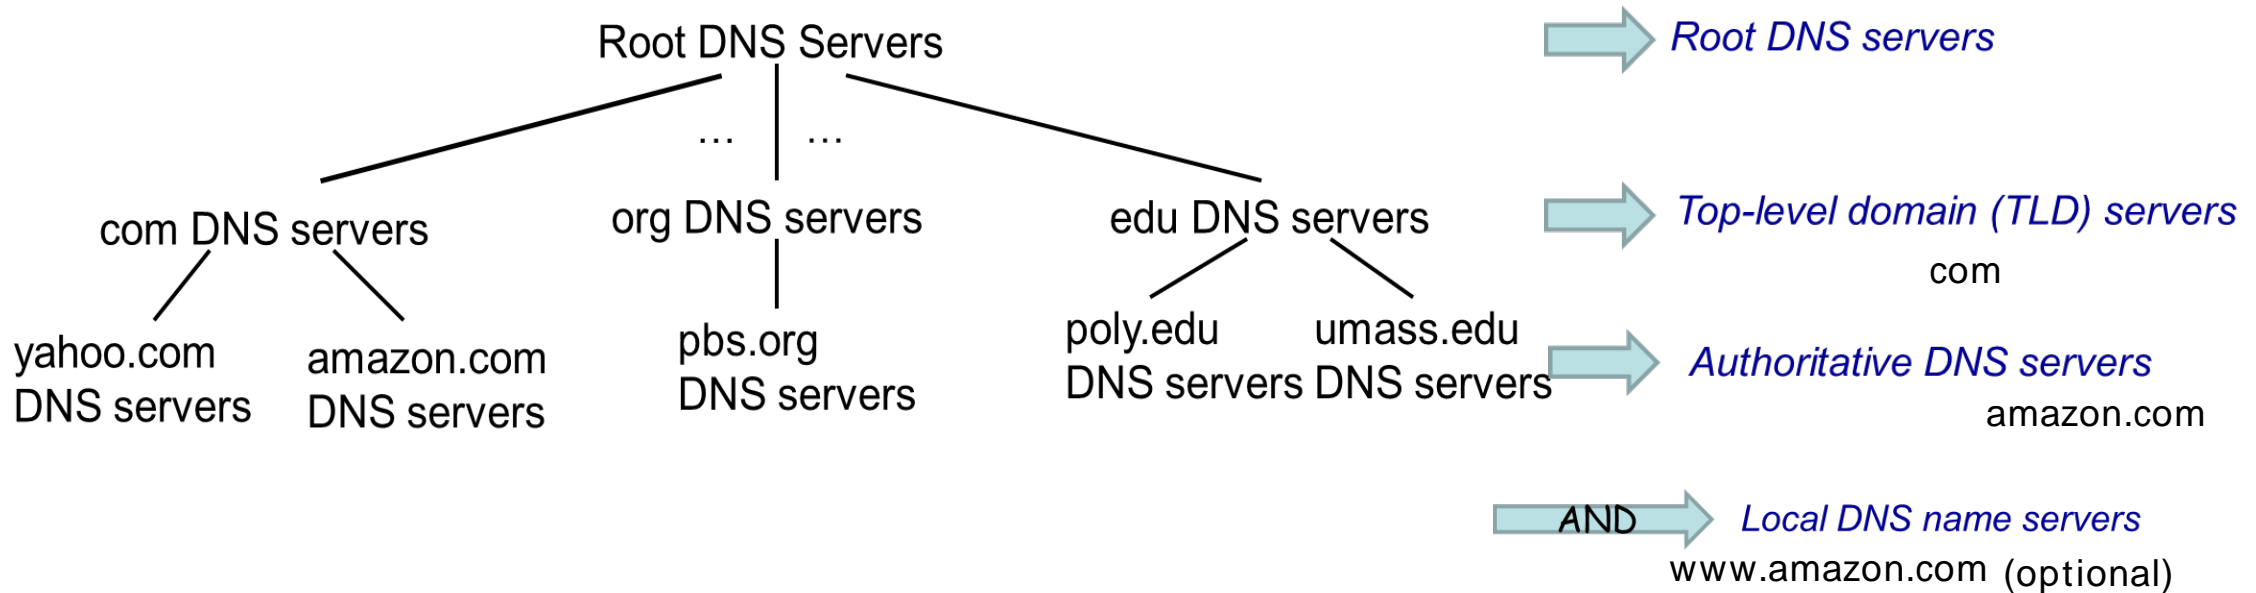


vs.



A: doesn't scale!

DNS: a distributed, hierarchical database

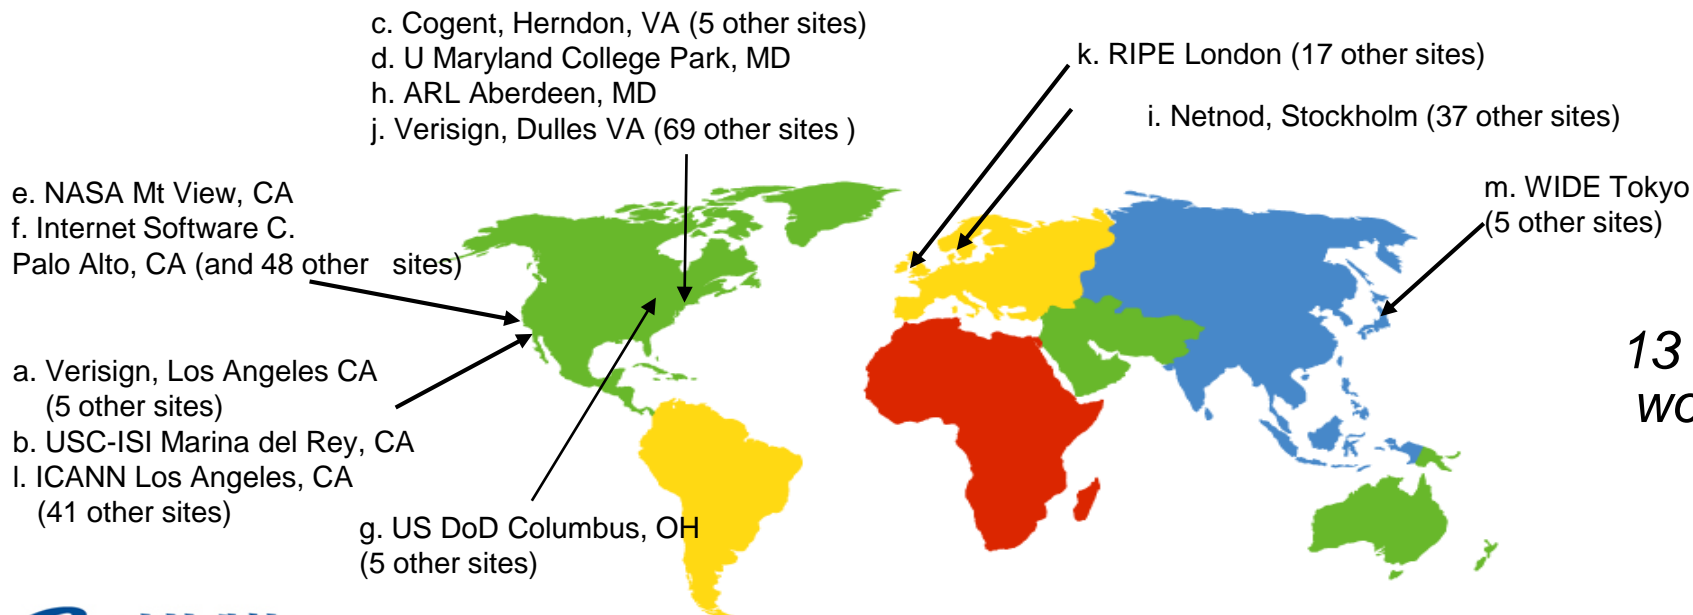


client wants IP address for `www.amazon.com`; 1st approx:

- ❖ client queries **root server** to find TLD (`.com`) DNS server TLD: `.com`
- ❖ client queries `.com` DNS server to get authoritative (`amazon.com`) DNS server
amazon.com authoritative DNS server가
- ❖ client queries `amazon.com` DNS server to get IP address for `www.amazon.com`

DNS: root name servers

- ❖ contacted by local name server that can not resolve name
- ❖ **root name server:**
 - contacts authoritative name server if name mapping not known
 - returns mapping to local name server



13 root name “servers” worldwide

TLD, authoritative servers

top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp

authoritative DNS servers:

ex) amazon.com

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's name
- can be maintained by organization or service provider



Local DNS name server

- ❖ each ISP (residential ISP, company, university) has one

- also called “default name server”

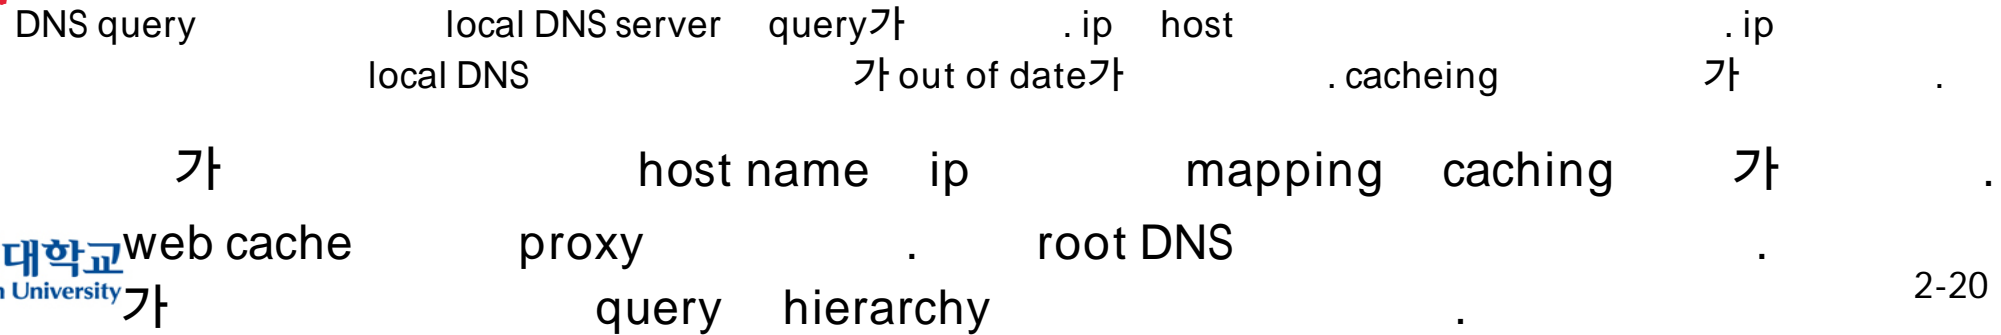
ipconfig/all

```

IPv4 주소 . . . . . : 121.135.107.150<기본 설정>
서브넷 마스크 . . . . . : 255.255.255.0
임대 시작 날짜 . . . . . : 2014년 9월 21일 일요일 오후 1:59:47
임대 만료 날짜 . . . . . : 2014년 9월 21일 일요일 오후 11:11:20
기본 게이트웨이 . . . . . : 121.135.107.1
DHCP 서버 . . . . . : 121.137.7.58
DHCPv6 IAID . . . . . : 199754650
DHCPv6 클라이언트 DUID. . . : 00-01-00-01-18-39-90-22-E8-03-9A-66-87-44
DNS 서버 . . . . . : 168.126.63.1
                  168.126.63.2
Tcpip를 통한 NetBIOS. . . . : 사용
    
```

- ❖ when host makes DNS query, query is sent to its local DNS server

- has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

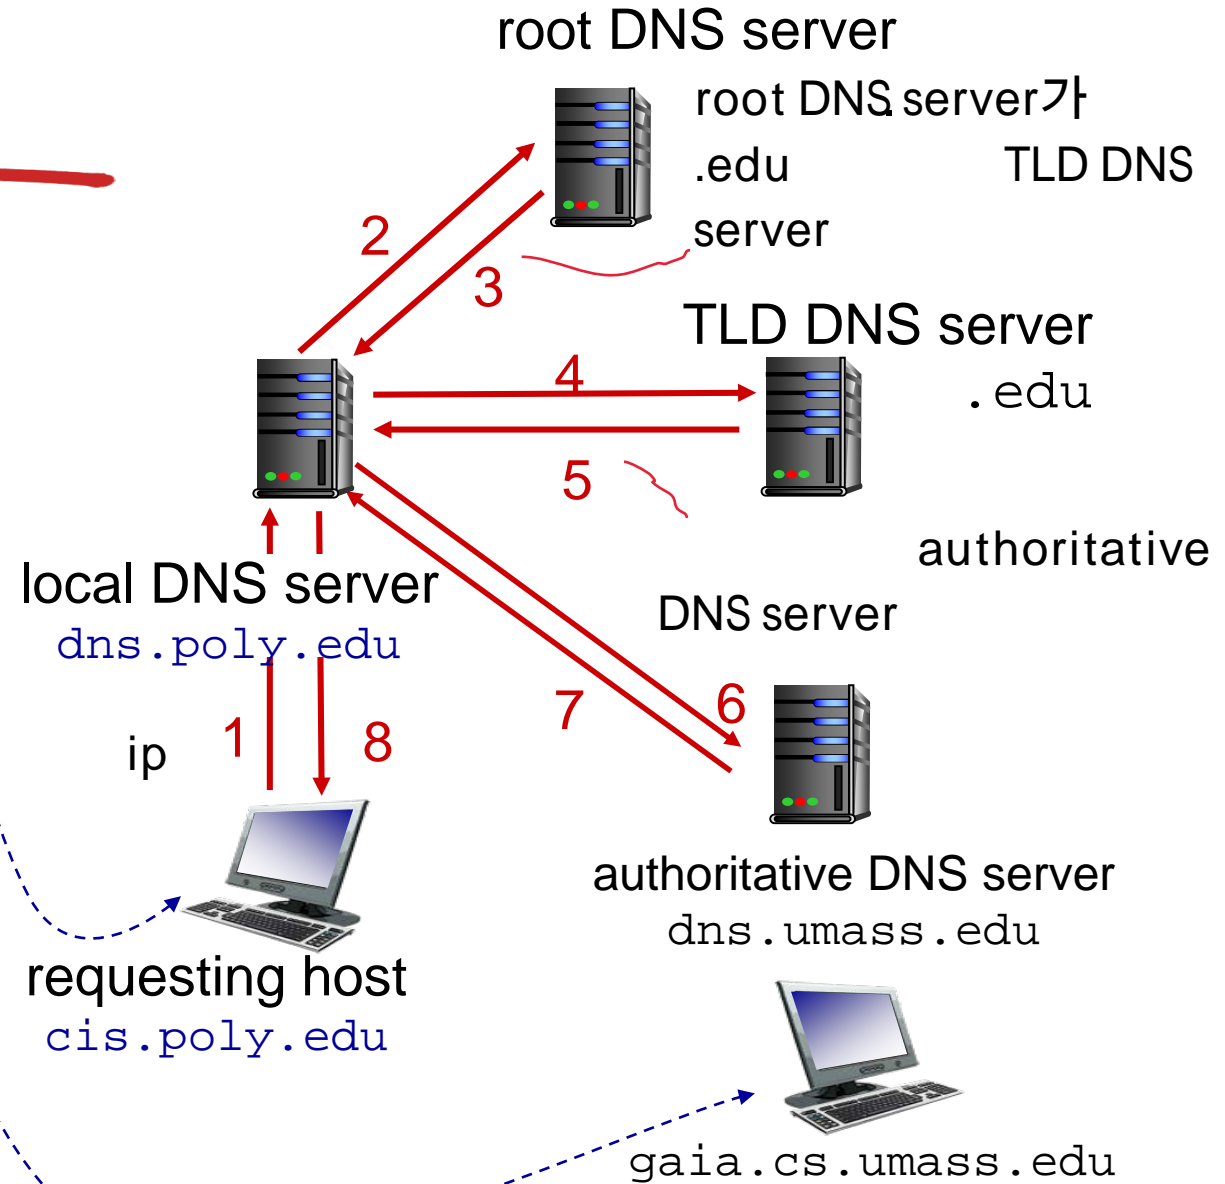


DNS name resolution example

- ❖ host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

iterated query:

- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”



DNS name resolution example

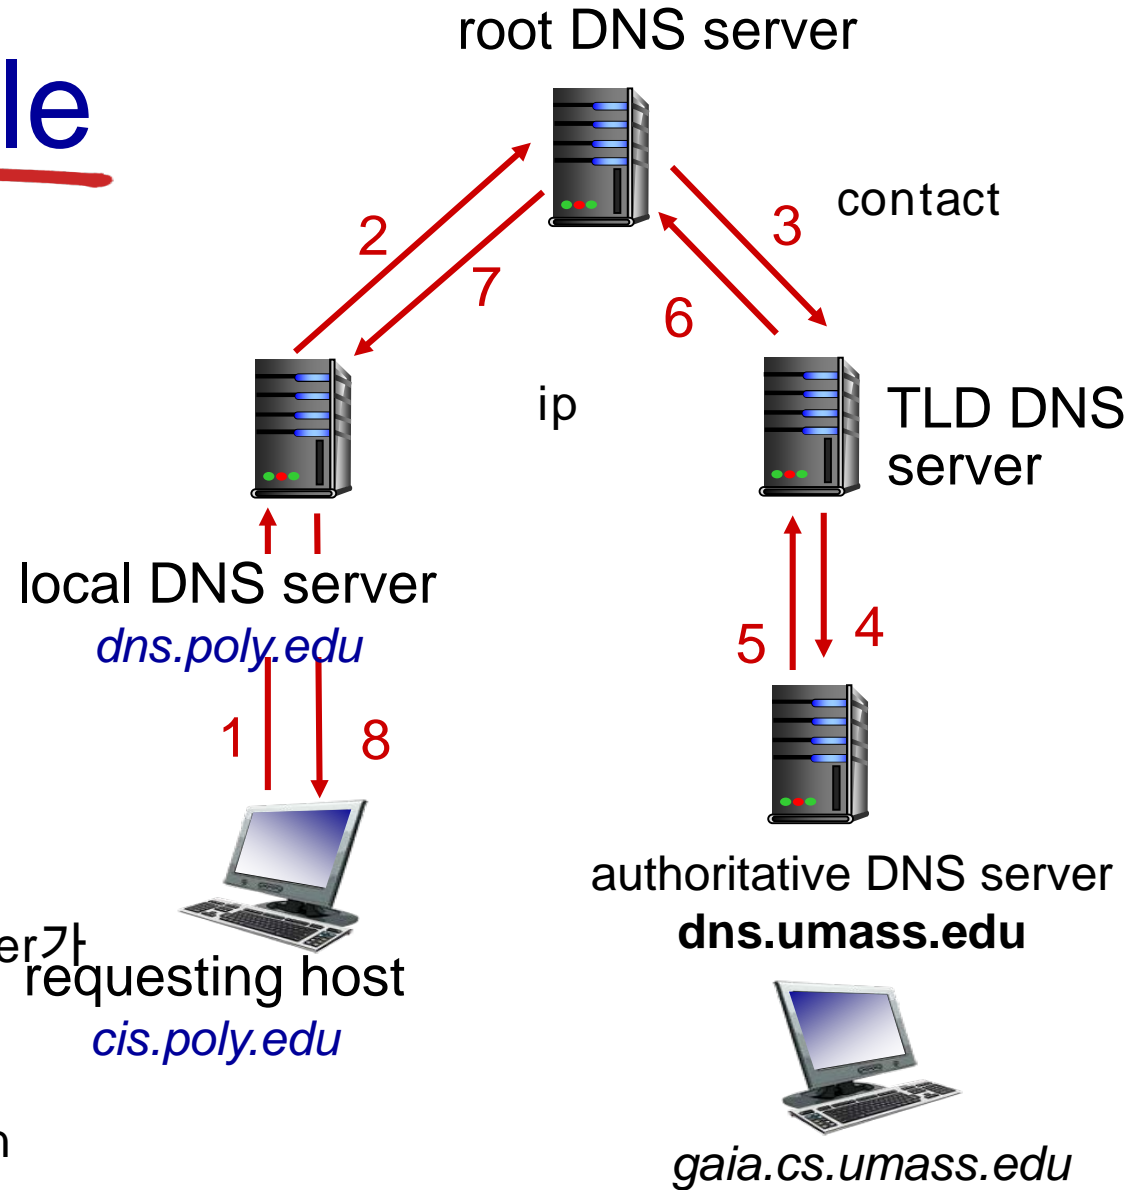
recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?

contact name server가 root DNS server

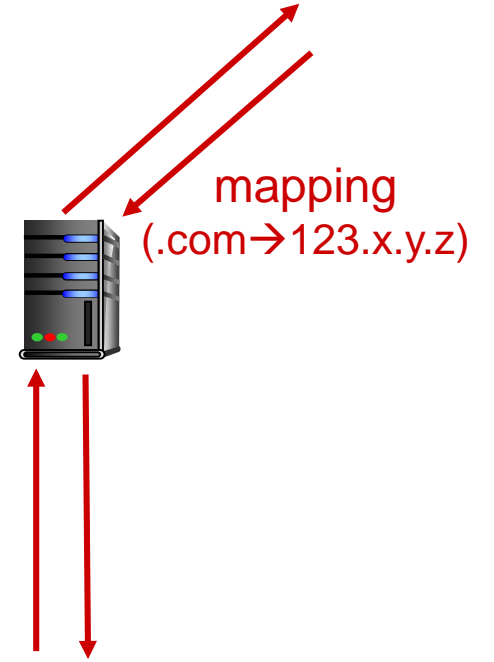
contact name server burden

local DNS server burden



Questions

Is it appropriate that (any) name server *caches* mapping after it learns the mapping ?



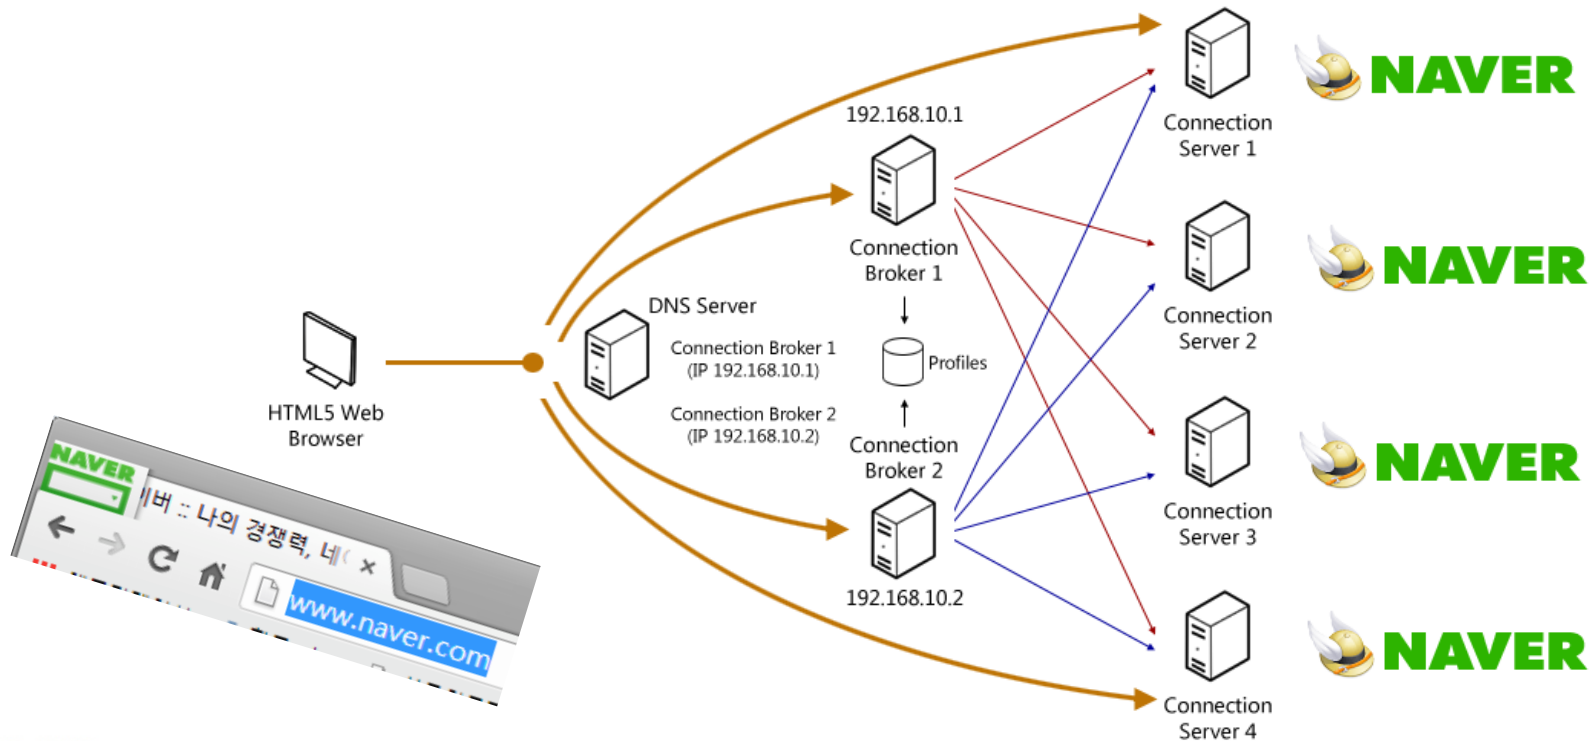
If so, what is a possible problem occurred from DNS caching?

DNS: caching, updating records

- ❖ once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - thus root name servers not often visited
- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!) ()
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire

DNS load distribution

- ❖ Help web servers do **load distribution (or load balancing)**
 - replicated Web servers: set of IP addresses for one canonical name (정식 이름)



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

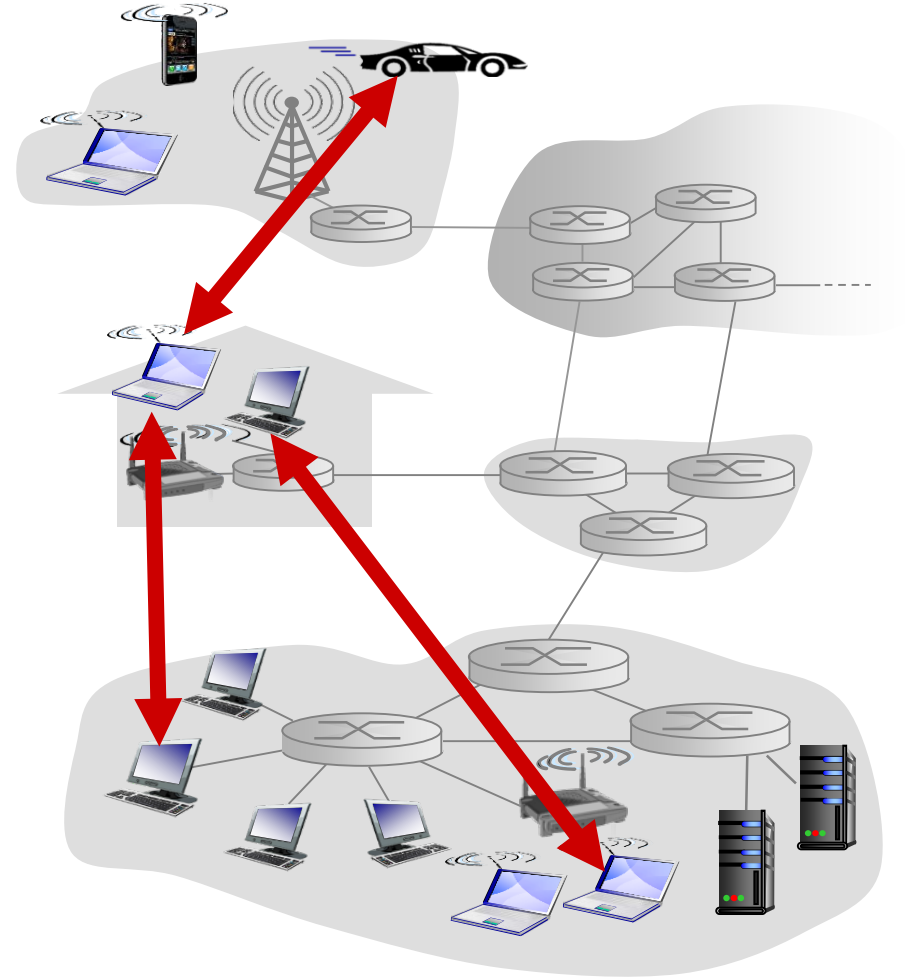
Pure P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers are intermittently connected and change IP addresses

examples:

- file distribution (BitTorrent)
- Streaming (Kankan)
- VoIP (Skype)

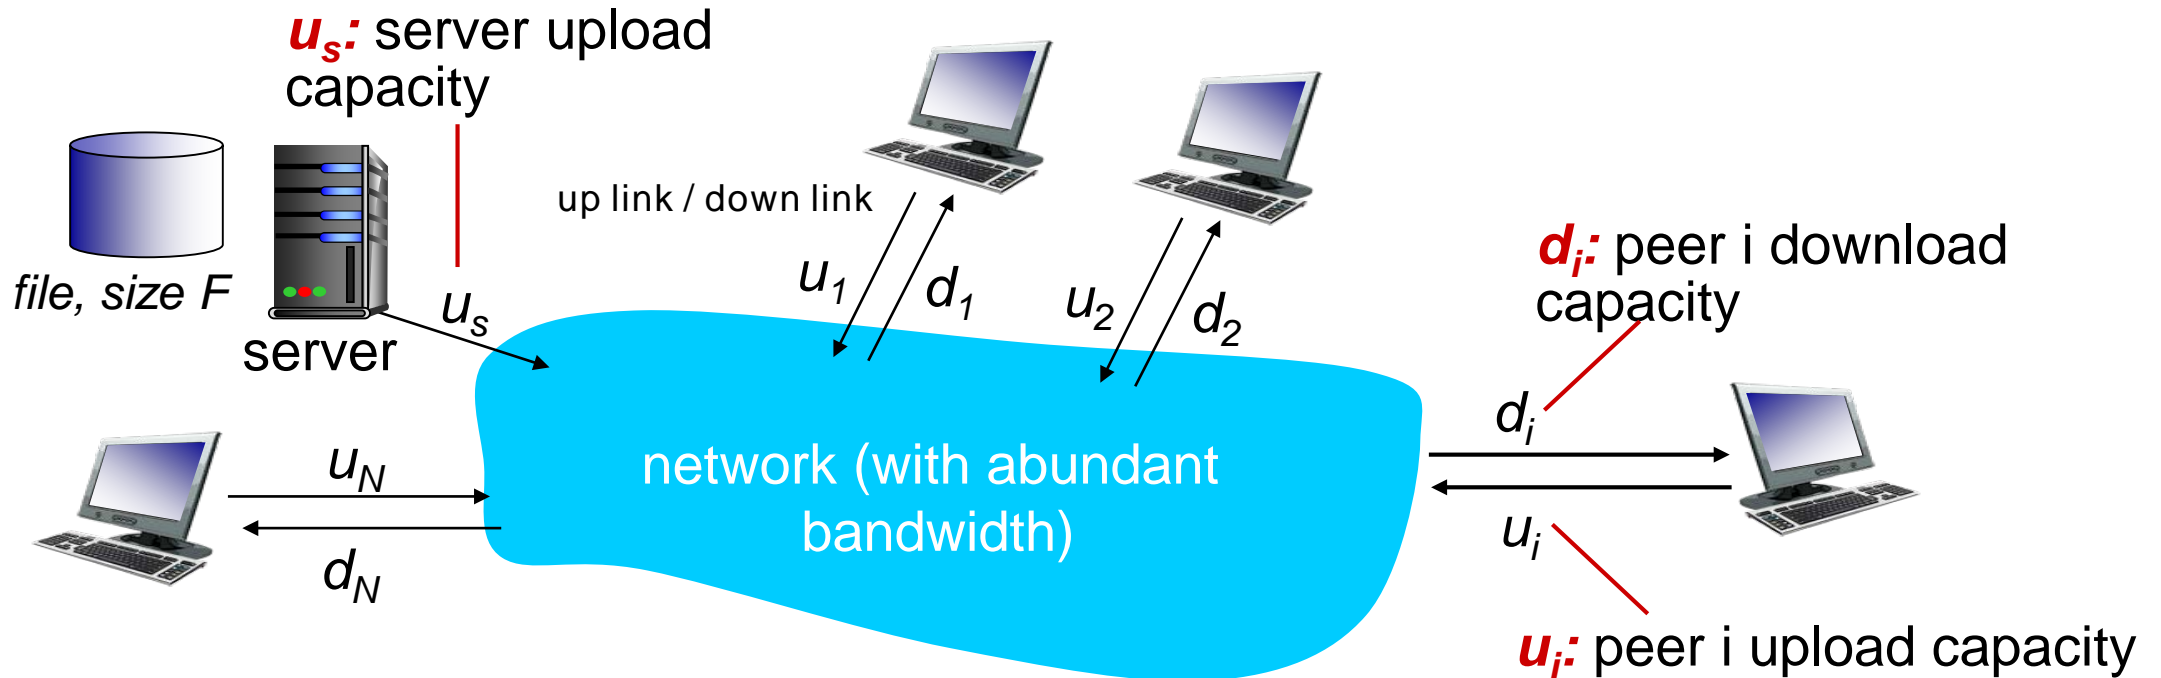
voice ip



File distribution: client-server vs P2P

Question: how much **time** to distribute file (size F) from one server to N peers?

- Server upload capacity & peer upload/download capacity are limited resources



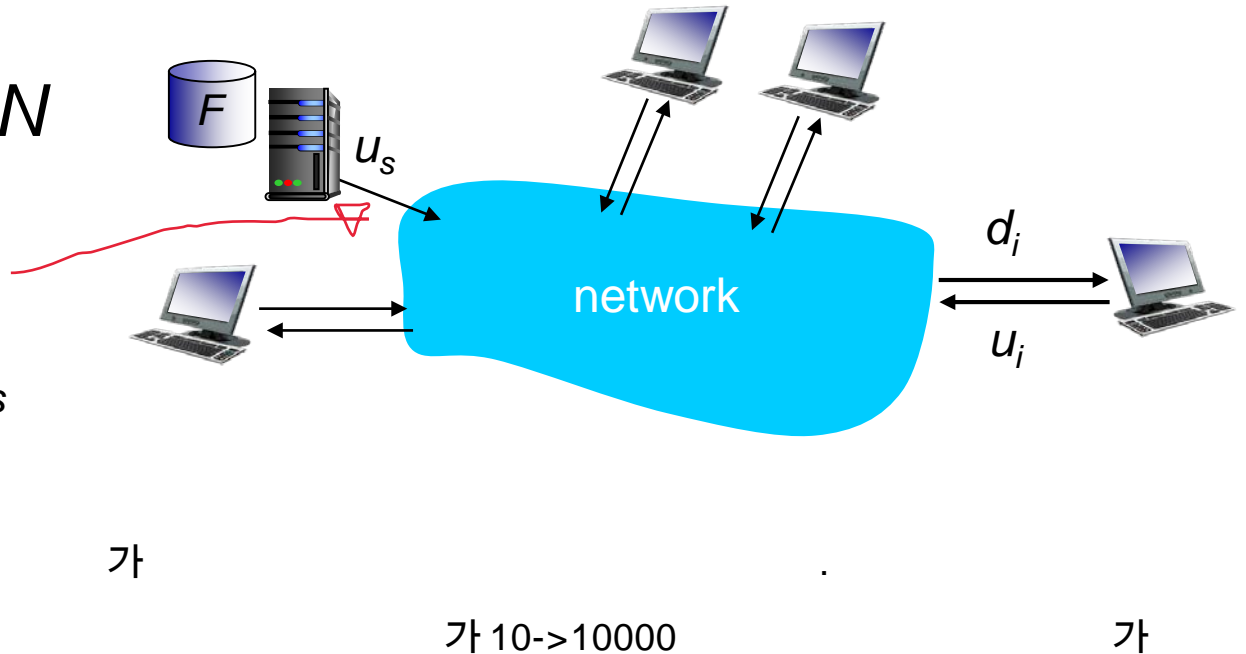
File distribution time: client-server

- ❖ **server transmission:** must sequentially send (upload) N file copies:

- time to send one copy: F/u_s
- time to send N copies: NF/u_s

- ❖ **client:** each client must download file copy

- d_{\min} = min client download rate
- min client download time: F/d_{\min}



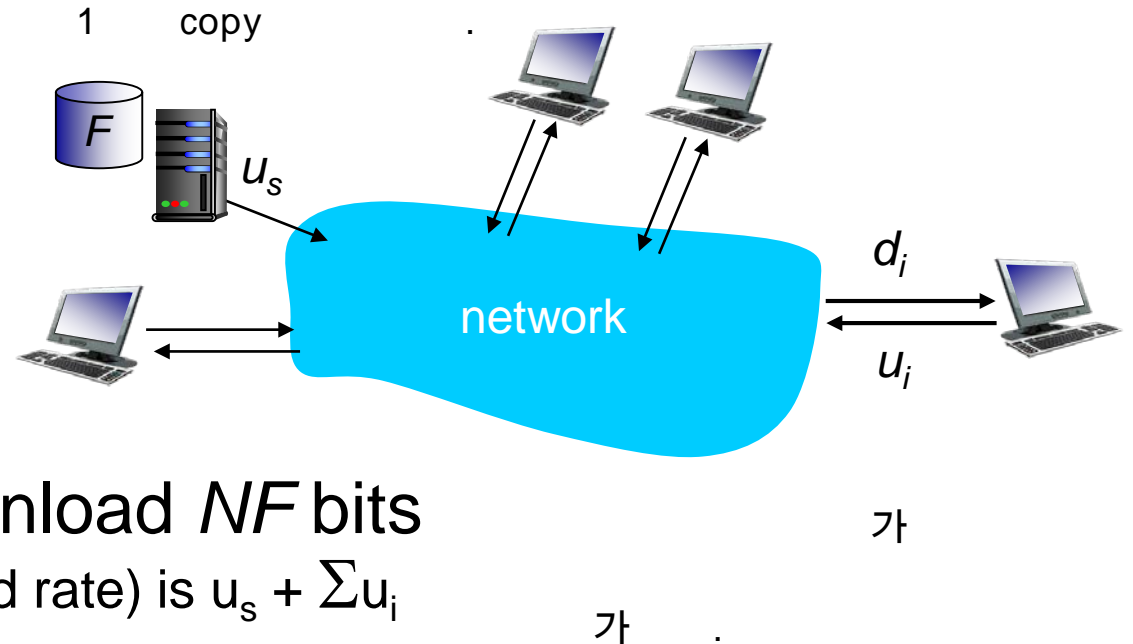
time to distribute F
to N clients using
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

File distribution time: P2P

- ❖ **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- ❖ **client:** each client must download file copy
 - min client download time: F/d_{\min}
- ❖ **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



time to distribute F
to N clients using
P2P approach

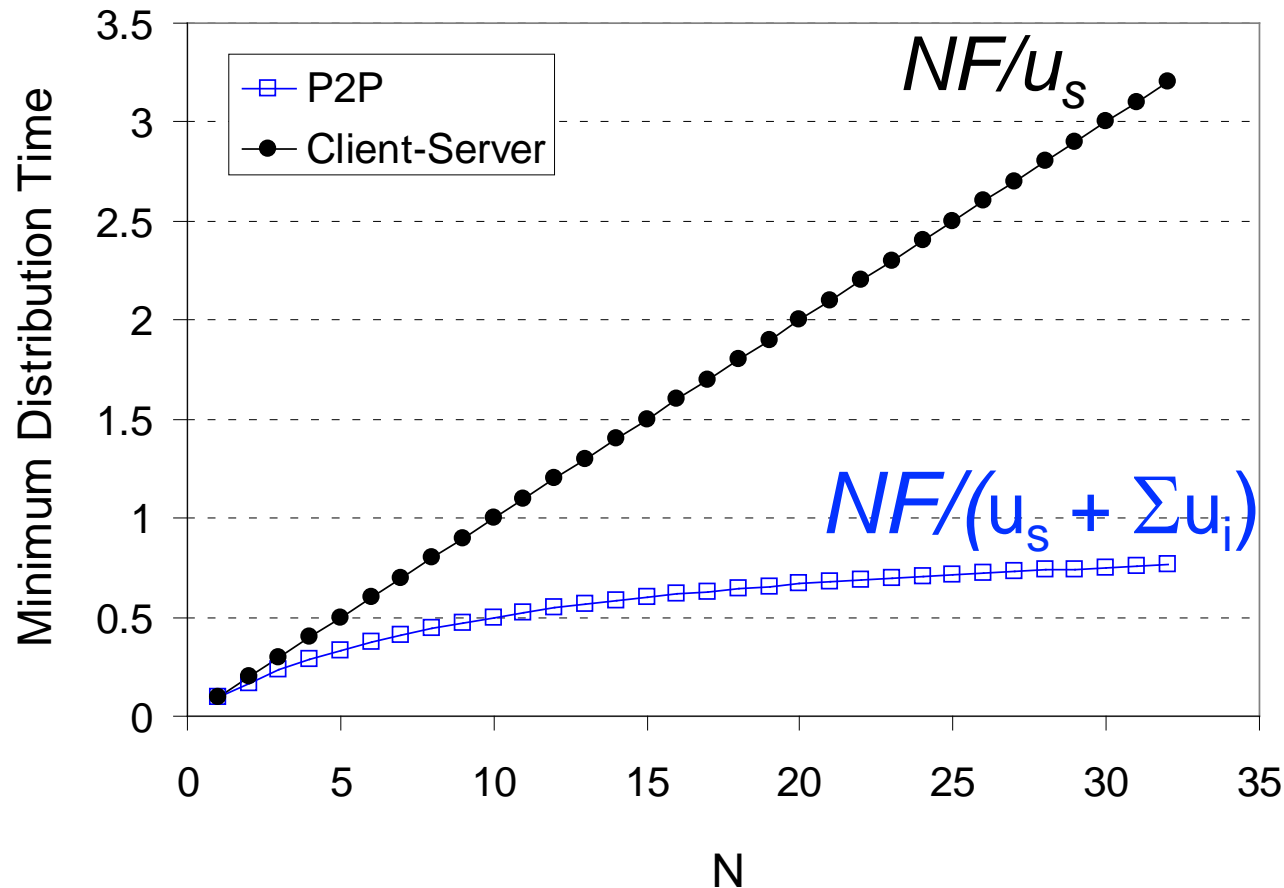
$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



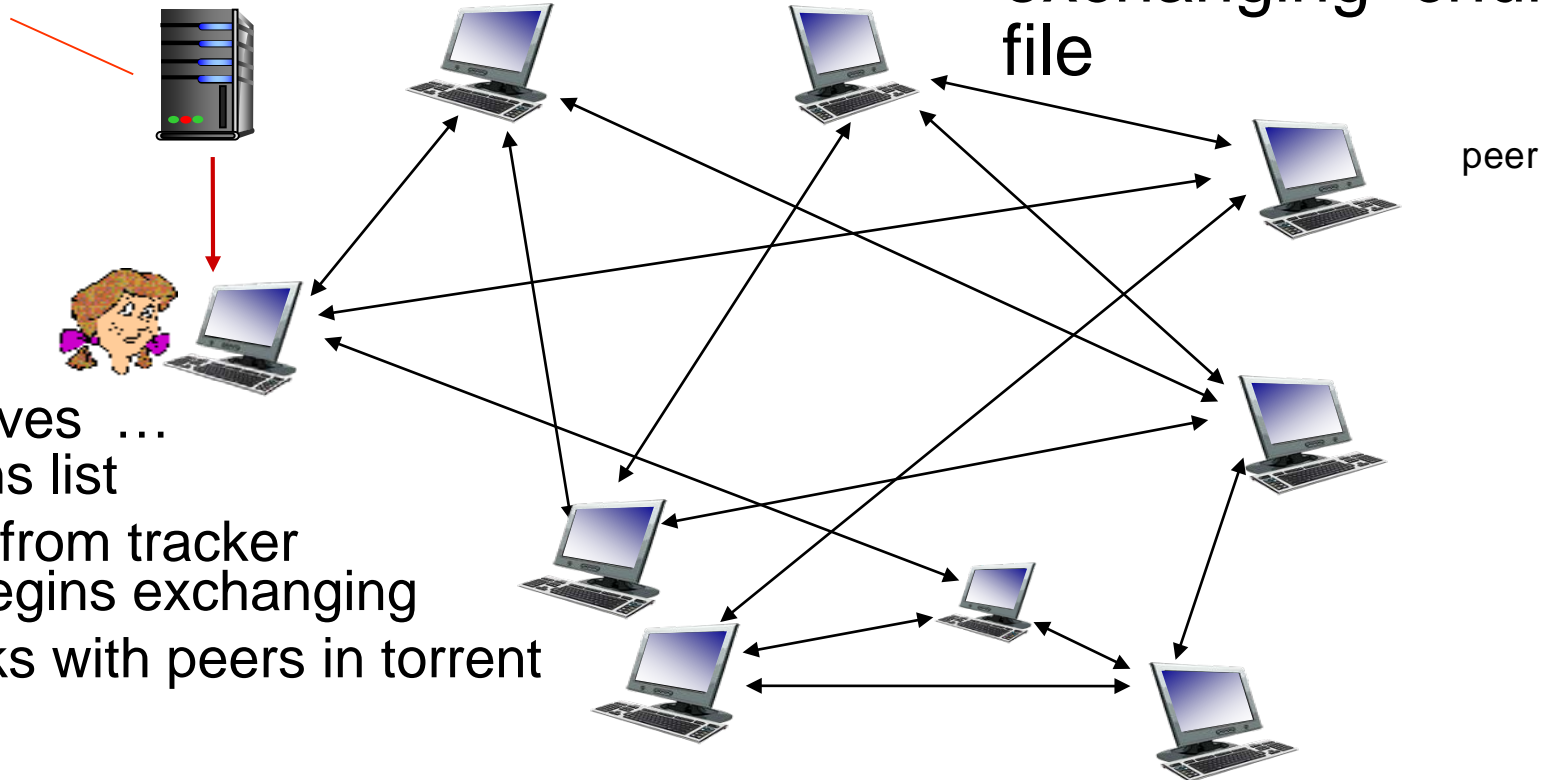
Self
Scalability

P2P file distribution: BitTorrent

- ❖ file divided into 256Kb chunks
- ❖ peers in torrent send/receive file chunks

tracker: tracks peers
participating in torrent

torrent: group of peers
exchanging chunks of a
file



- ① Alice arrives ...
... obtains list
- ② of peers from tracker
... and begins exchanging
- ③ file chunks with peers in torrent

가
가 .
가

P2P file distribution: BitTorrent

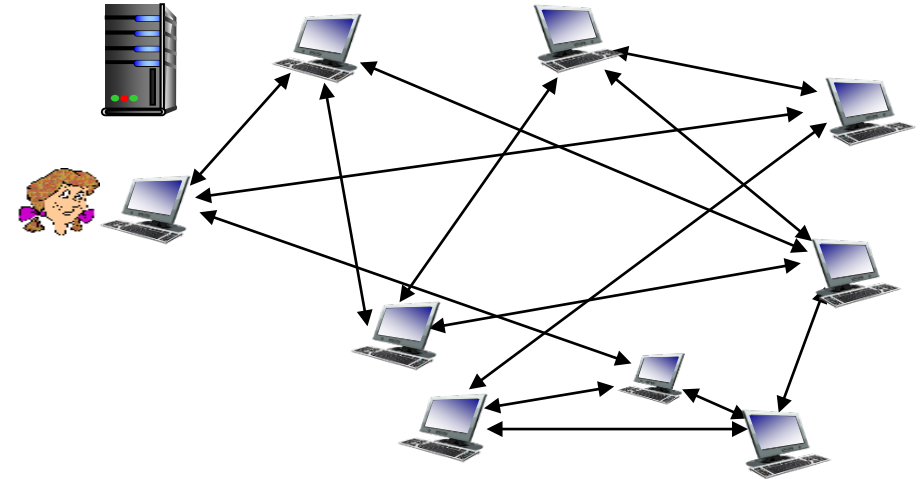


Read Chapter 2.5 (page 172-175) in textbook!

❖ peer joining torrent:

- has no chunks, but will accumulate them over time from other peers
- registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

peer
가
peer



- ❖ while downloading, peer uploads chunks to other peers
- ❖ peer may change peers with whom it exchanges chunks
- ❖ **churn**: peers may come and go
- ❖ once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: requesting, sending file chunks

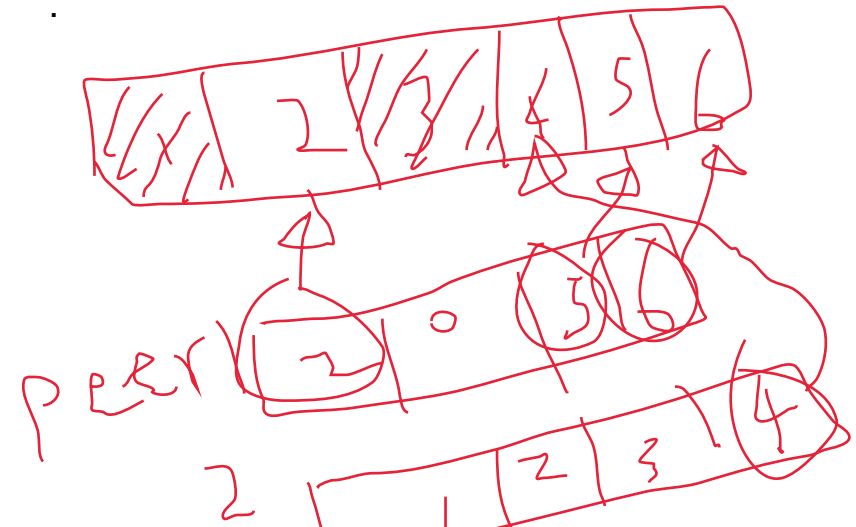
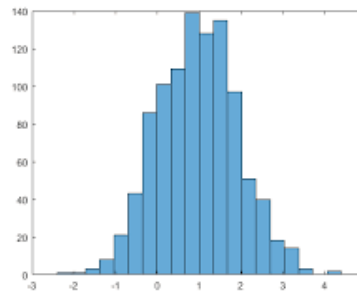
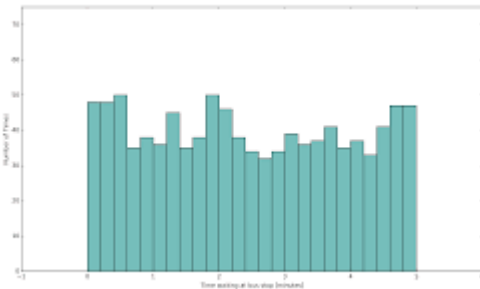
p2p

requesting chunks:

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for list of chunks that they have
- ❖ Alice requests missing chunks from peers, **rarest first**

peer가 10

- 1 chunk : 2
- 2 chunk : 5
- 3 chunk : 10



BitTorrent: requesting, sending file chunks

sending chunks: *tit-for-tat*

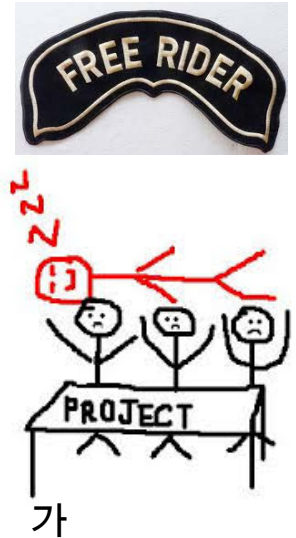
- ❖ Alice sends chunks to those four peers currently sending her chunks at highest rate
 - other peers are *choked* by Alice (do not receive chunks from her) = chunk
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “*optimistically unchoke*” this peer
 - newly chosen peer may join top 4

가 chunk

?

가

=

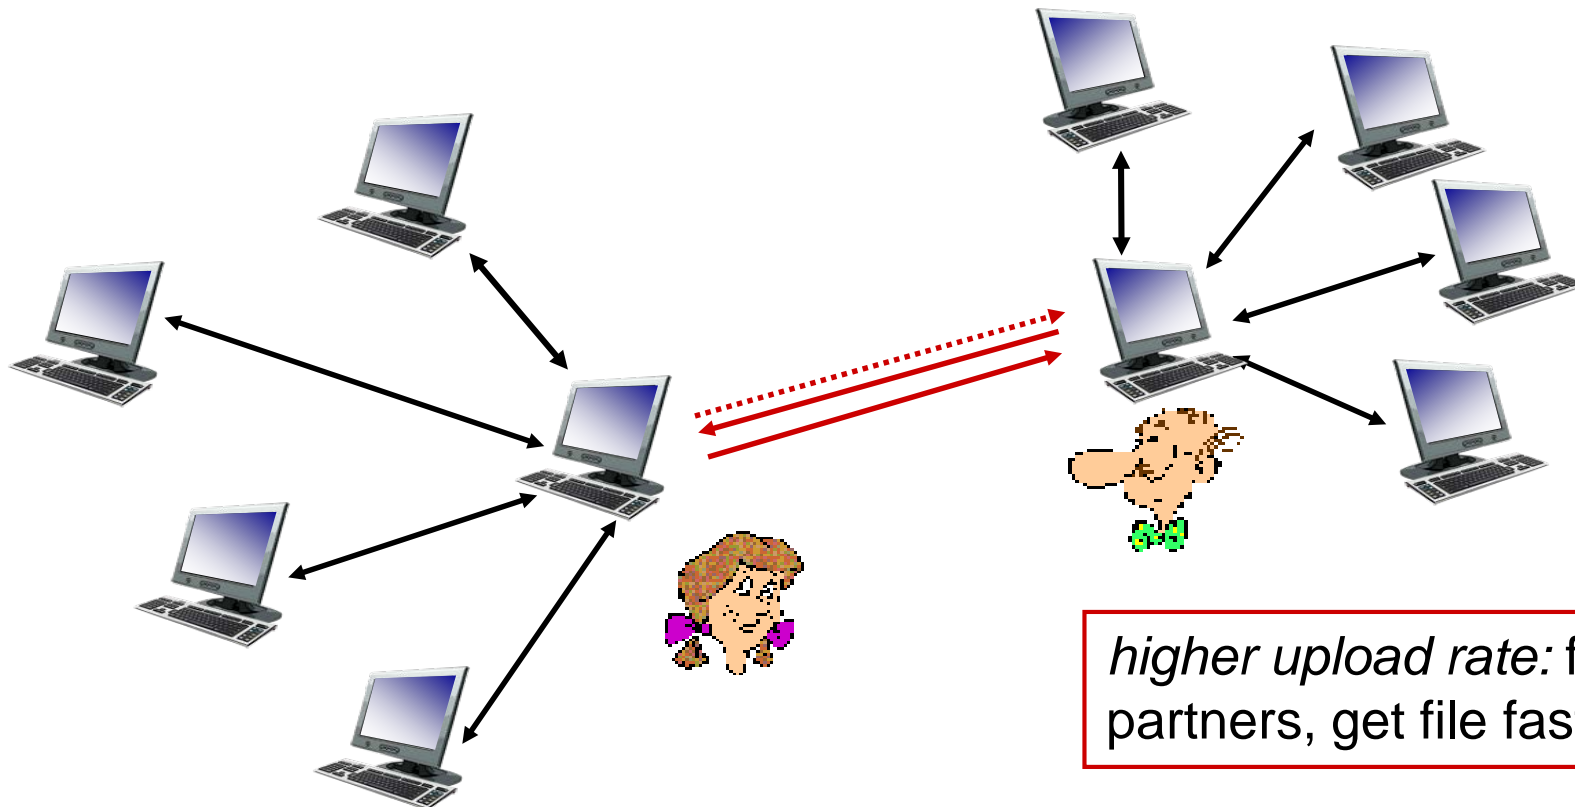


BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers

가

top four provider.



The rest?
Choke!

higher upload rate: find better trading partners, get file faster !

Chapter 2: outline

2.1 principles of
network
applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3,
IMAP

2.4 DNS


2.5 P2P applications

2.6 video streaming
and content
distribution
networks (CDNs)

traffic

video streaming

Video Streaming and CDNs: context

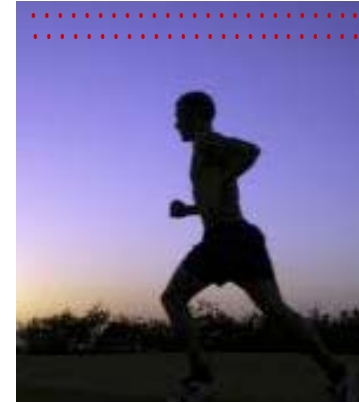
- video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B (10억) users?  scalability
 - single mega-video server won't work (why?)
- challenge: **heterogeneity**
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor) (,)
- **solution**: distributed, application-level infrastructure



Multimedia: video

- ❖ video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- ❖ digital image: array of pixels
 - each pixel represented by bits
- ❖ coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i



frame $i+1$

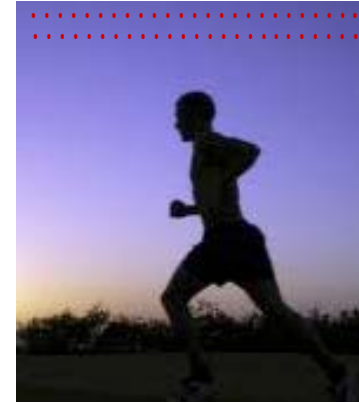
temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i

Multimedia: video

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)

(bandwidth)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

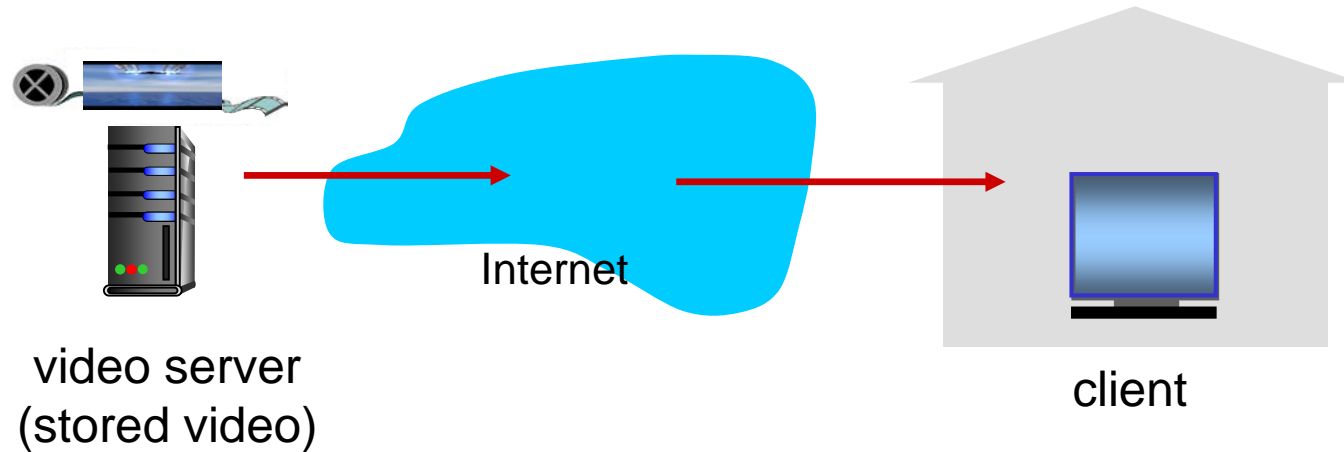
temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

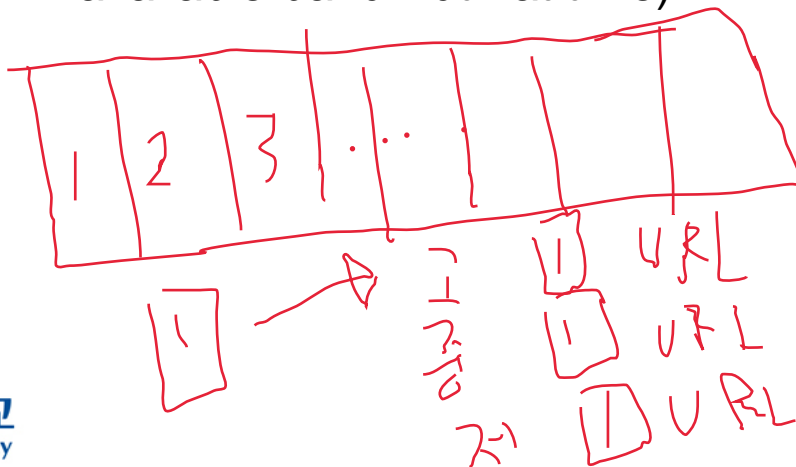
Streaming stored video:

simple scenario:



Streaming multimedia: DASH

- ❖ **DASH**: **D**ynamic, **A**daptive **S**treaming over **H**TTP
- ❖ **server**:
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - **manifest file**: provides URLs for different chunks
- ❖ **client**:
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)



manifest file url 가 .
 bandwidth .
 chunk .(3 4)
 bandwidth가
 encoding version .

Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ “intelligence” at client: client determines
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - *what encoding rate* to request (higher quality when more bandwidth available)
 - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

client

3

buffer

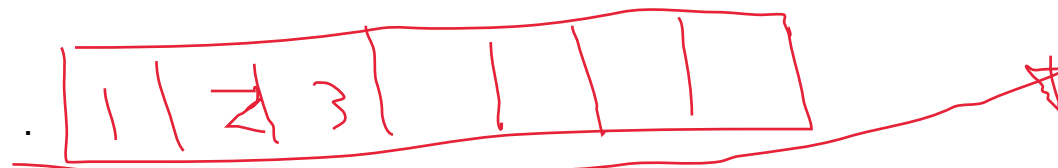
url

client

chunk
가

가

client가



buffer

Application Layer

- outgoing link가 bottle neck

....quite simply: this solution *doesn't scale*

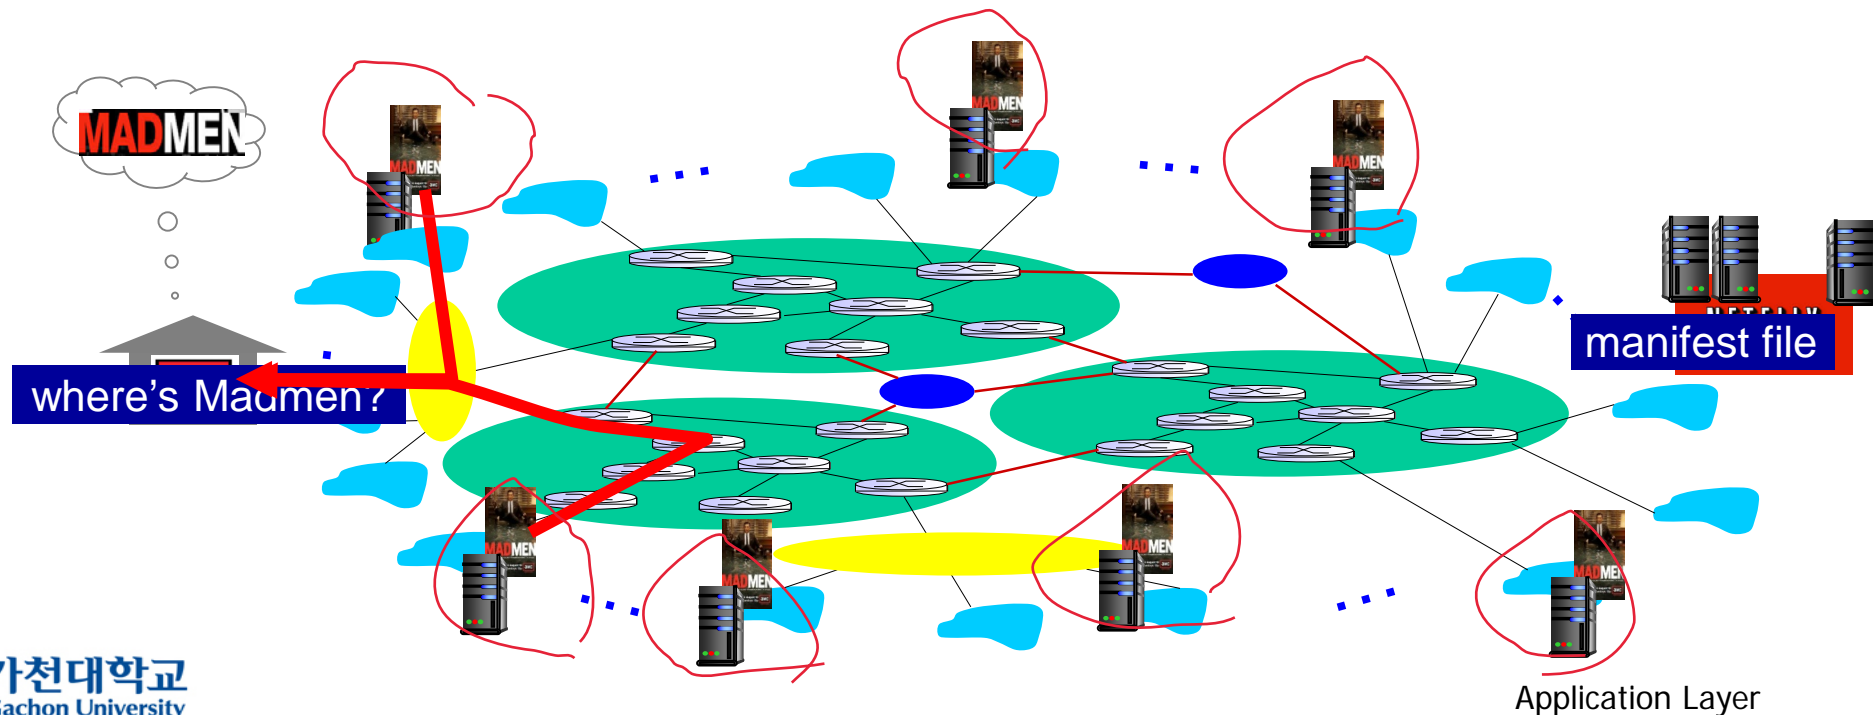
Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen

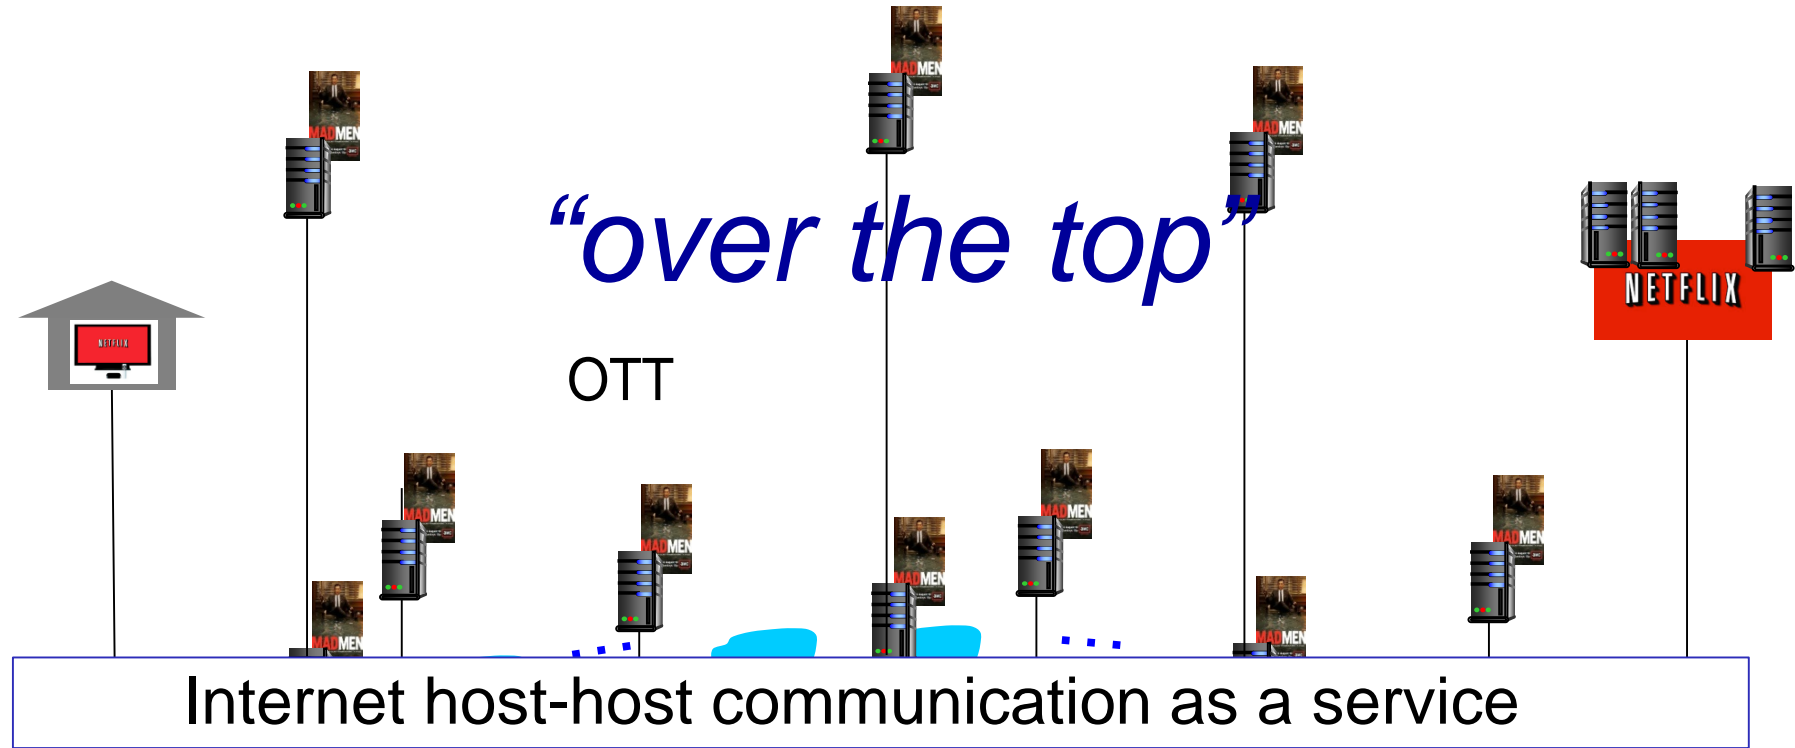
stream

- ① subscriber requests content from CDN
- ② directed to nearby copy, retrieves content
- ③ may choose different copy if network path congested

. ->



Content Distribution Networks (CDNs)



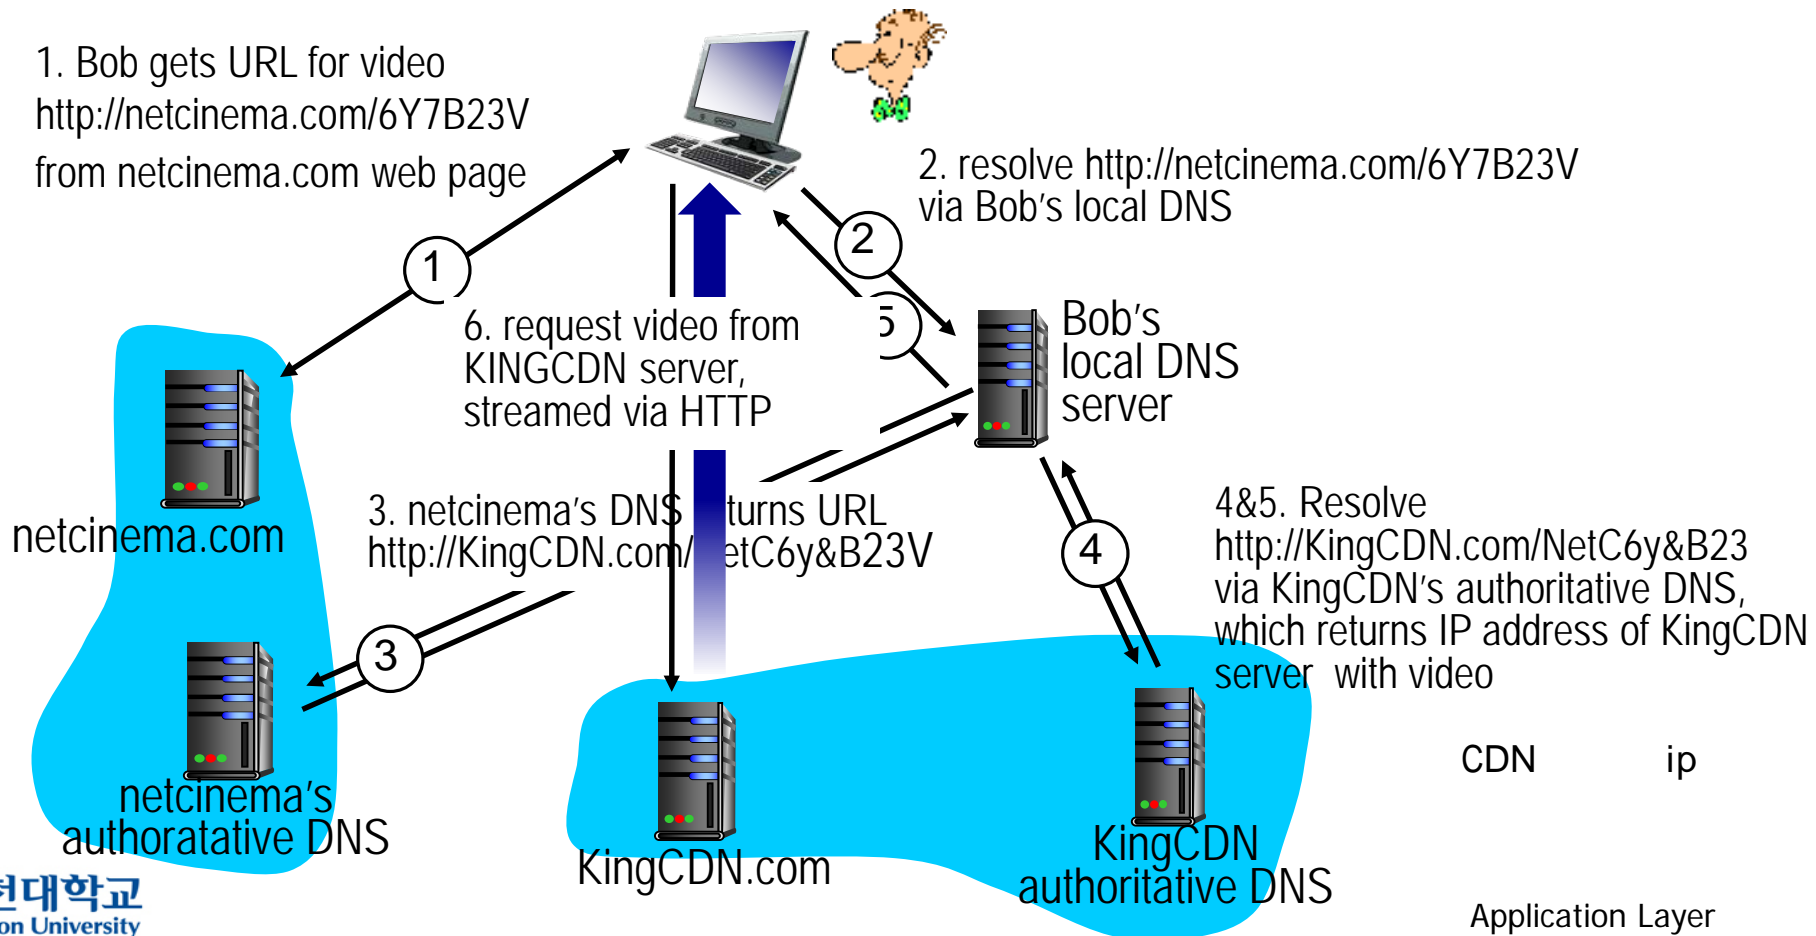
OTT challenges: coping with a congested Internet

- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



Case study: Netflix

