


Chapter 3

Transport Layer part 2

School of Computing
Gachon Univ.
Joohyung Lee

Most of slides from J.F Kurose and K.W. Ross. And, some slides from Prof. Joon Yoo



*Computer
Networking: A Top
Down Approach*

7th edition

Jim Kurose, Keith Ross
Pearson, 2017

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP


- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Principles of reliable data transfer

- ❖ important in application, transport, link layers
 - top-10 list of important networking topics!

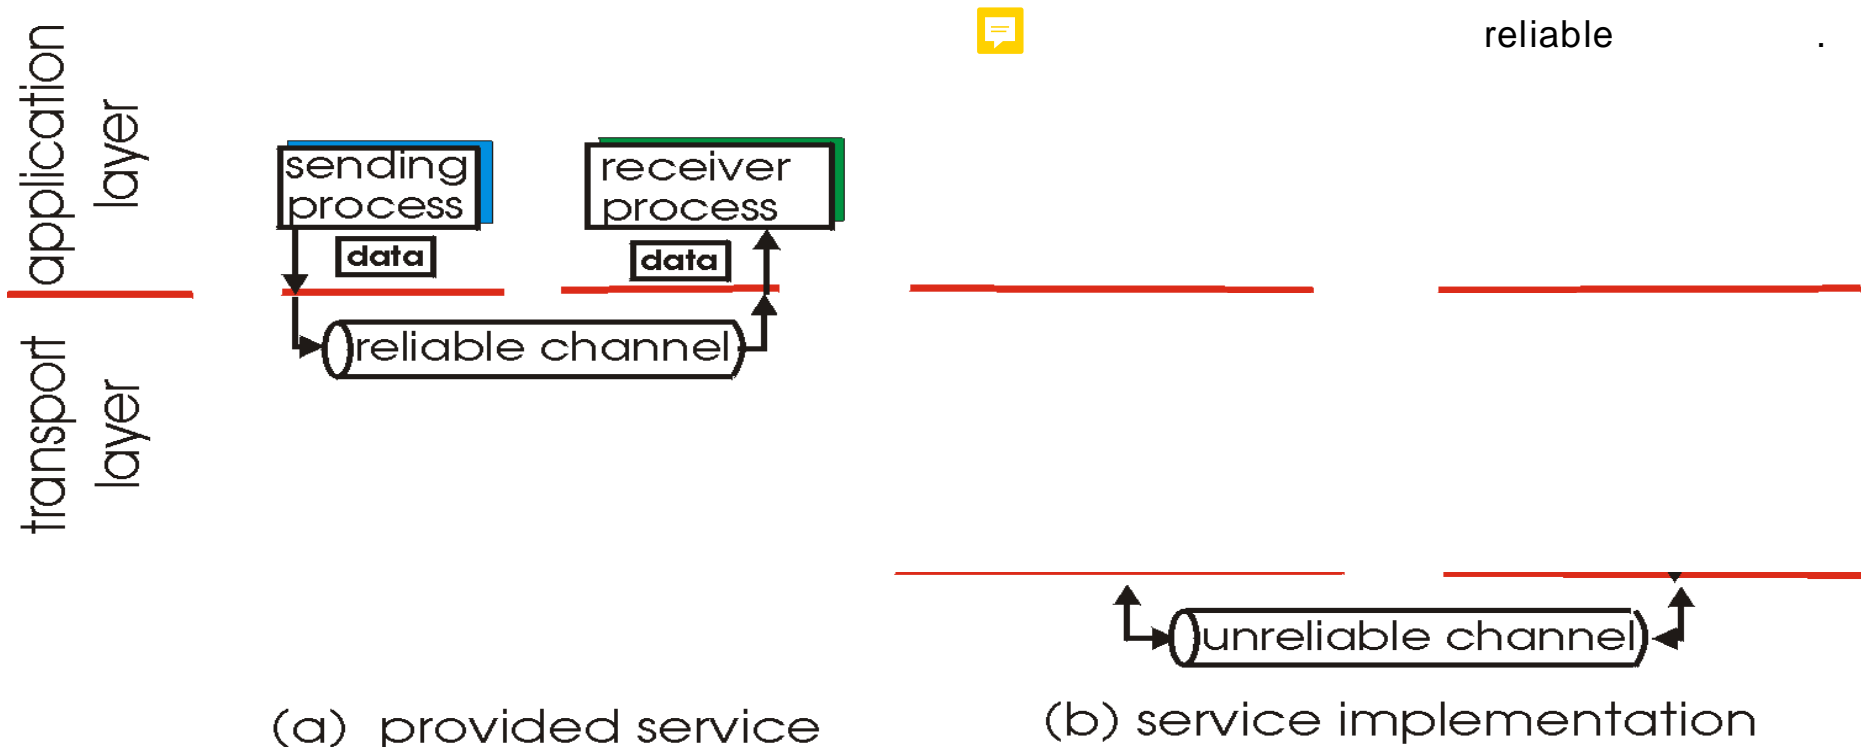


(a) provided service

- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of reliable data transfer

- ❖ important in application, transport, link layers
 - top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

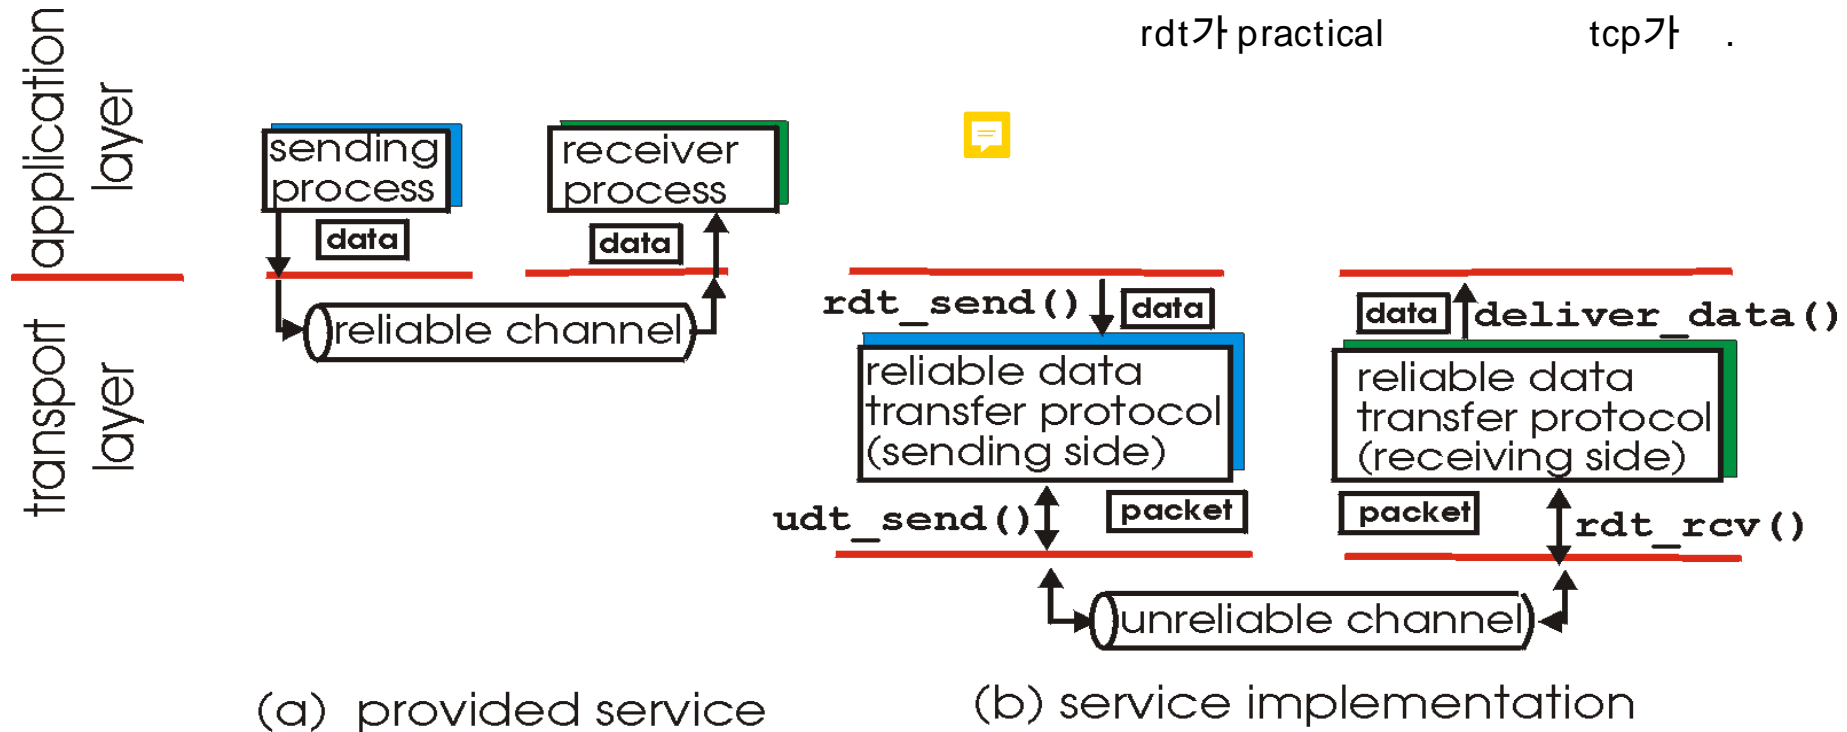
Principles of reliable data transfer

- ❖ important in application, transport, link layers
 - top-10 list of important networking topics!

Reliable data transfer protocols

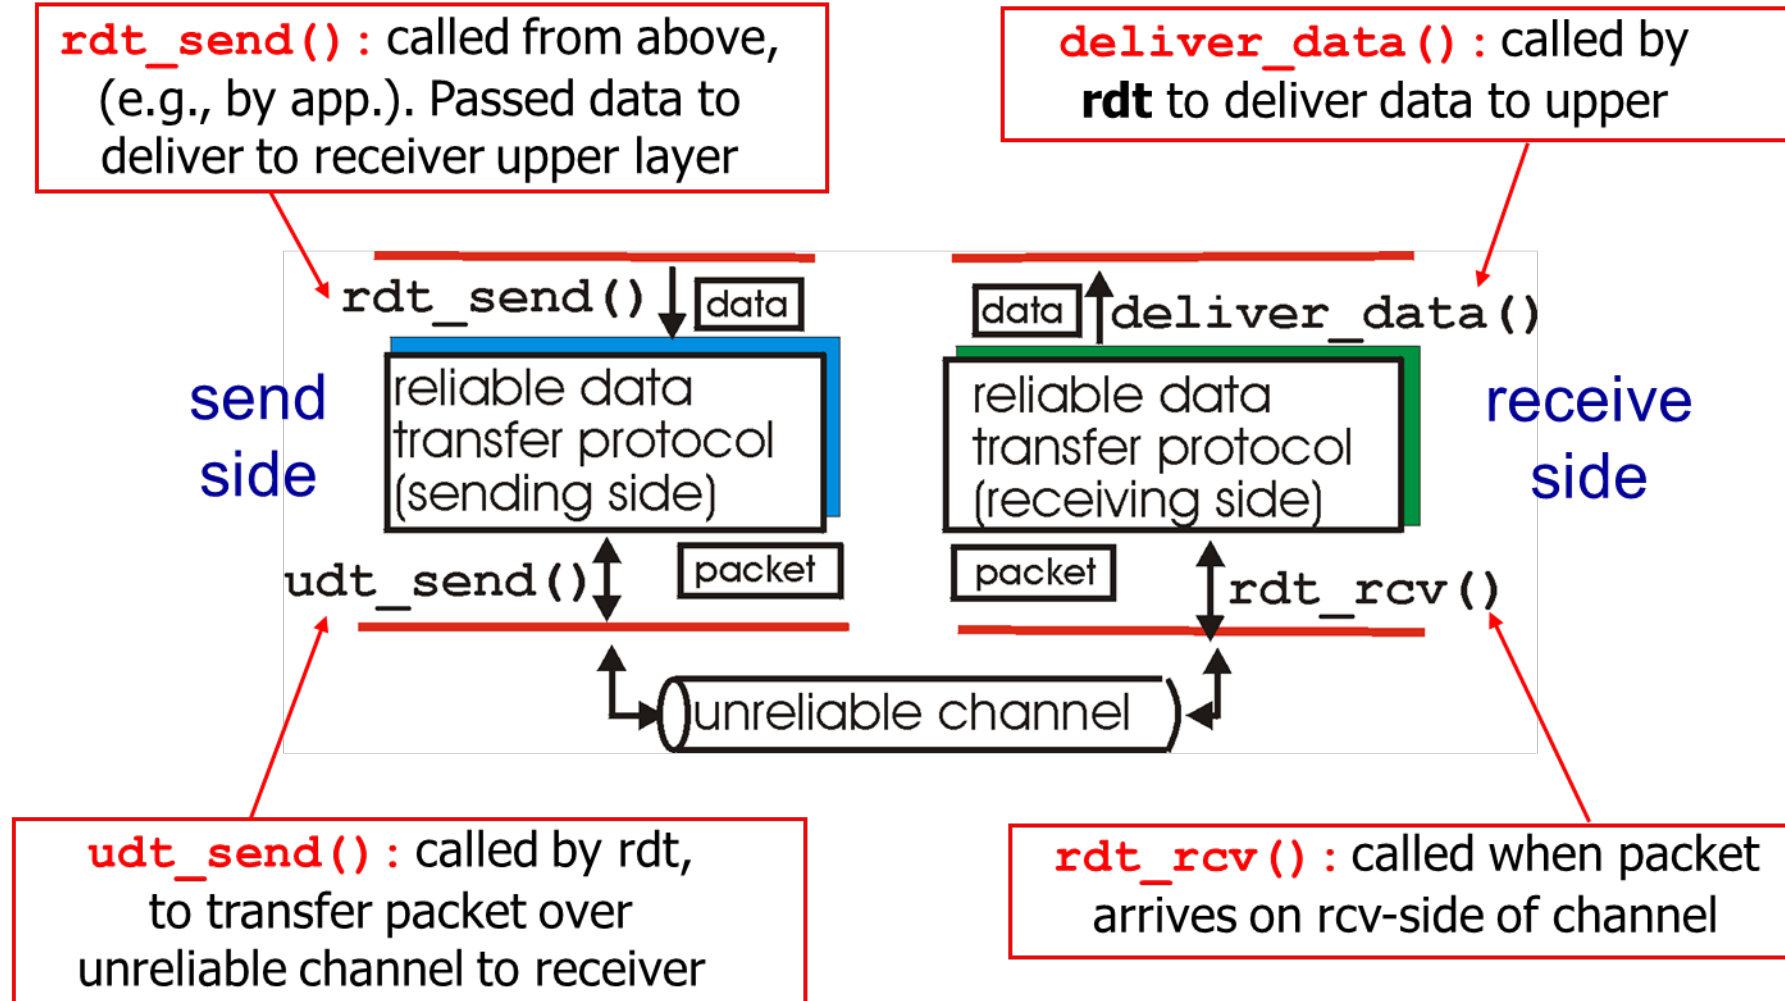
rdt가 practical

tcp가 .



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: getting started



Reliable data transfer: getting started

We'll:



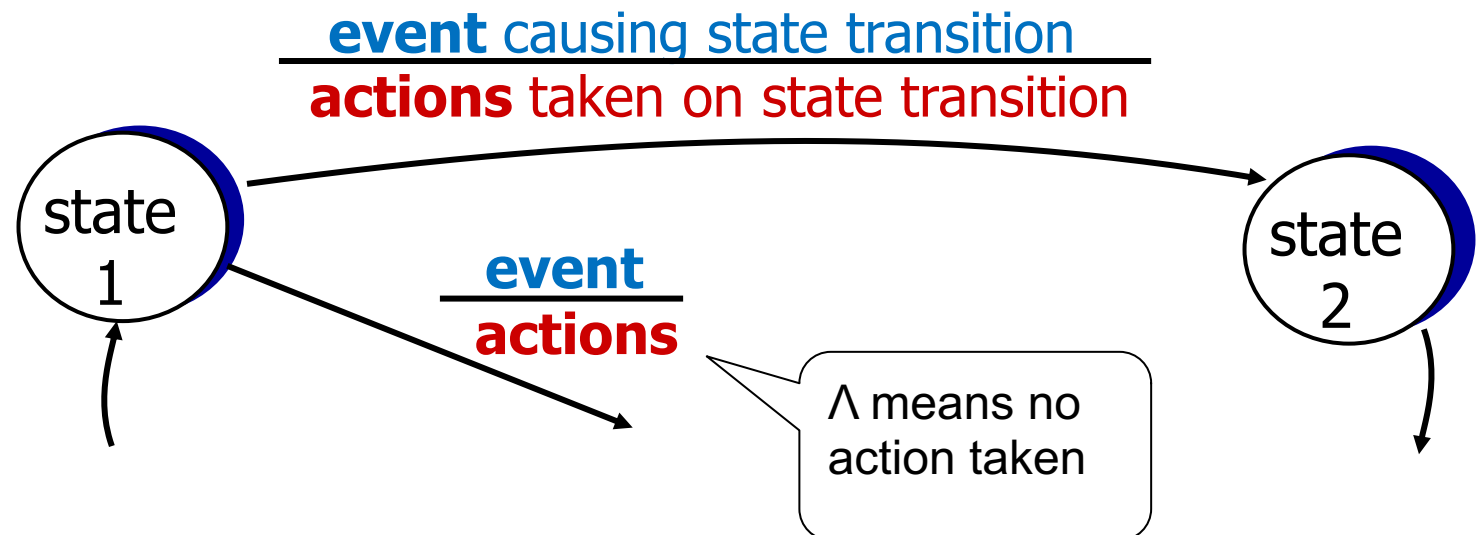
state machine

FSM

pseudo code,

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ use finite state machines (FSM) to specify sender, receiver

state: when in this “state”
next state is determined
by some **event**



Unreliable channel?

Potential Channel Errors



❖ bit errors (but no packet losses)

- How does receiver find bit errors?



❖ loss (drop) of packets

- How does sender find packet losses?



queuing loss
가

❖ reordering or duplication

- Arriving Pkts: 1, 2, 4, 5, 3, ... or 1, 2, 3, 3, ...




가 가

가

Characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt).

Different Conditions, Different Solutions

- ❖ The more difficult the condition, the smarter the solution to deal with the problem



rdt1.0: reliable transfer over a reliable channel

❖ underlying channel perfectly reliable

- no bit errors
- no loss of packets

❖ separate FSMs for sender, receiver:

- sender sends data into underlying channel
- receiver reads data from underlying channel

rdt 1.0

가 reliable

가

sender :

receiver :

application data 가

unreliable

receiver : high layer

. packet

packet

....?

(packet)

state(

)

rdt_rcv(packet)

extract (packet,data)

deliver_data(data)




recv from lower layer
send to upper layer

recv from upper layer
send to lower layer

sender

receiver

Initial state




rdt_send(data)

packet = make_pkt(data)

udt_send(packet)

data



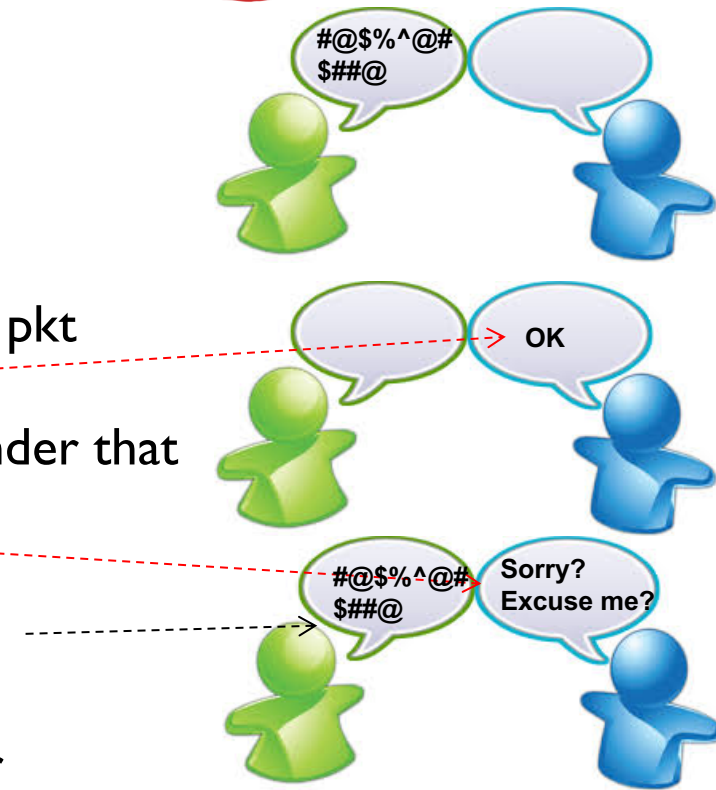
rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits ($0 \rightarrow 1$ or $1 \rightarrow 0$) in packet
 - checksum to detect bit errors
- ❖ *the question: how to recover from errors:*

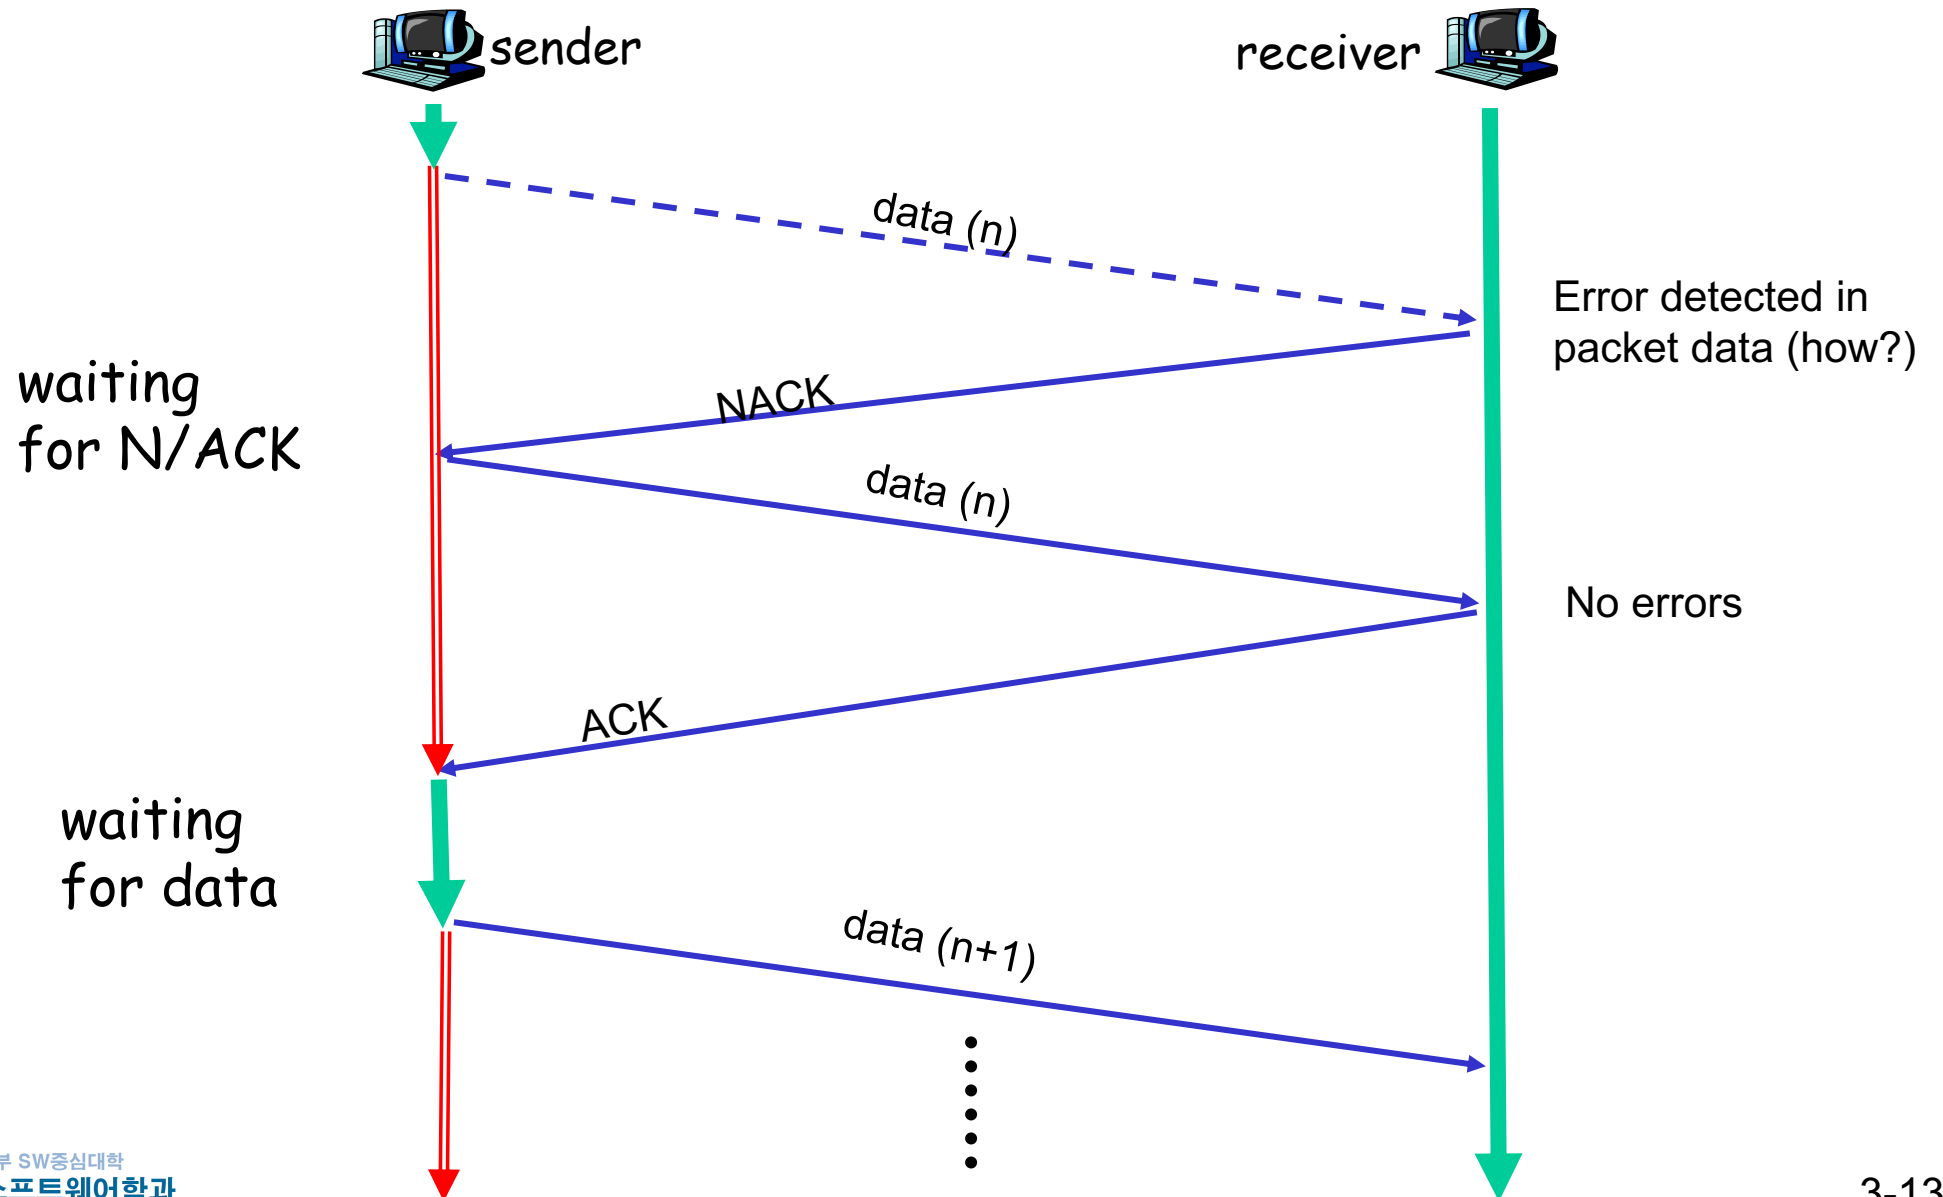
How do humans recover from “errors” during conversation?

rdt2.0: channel with bit errors

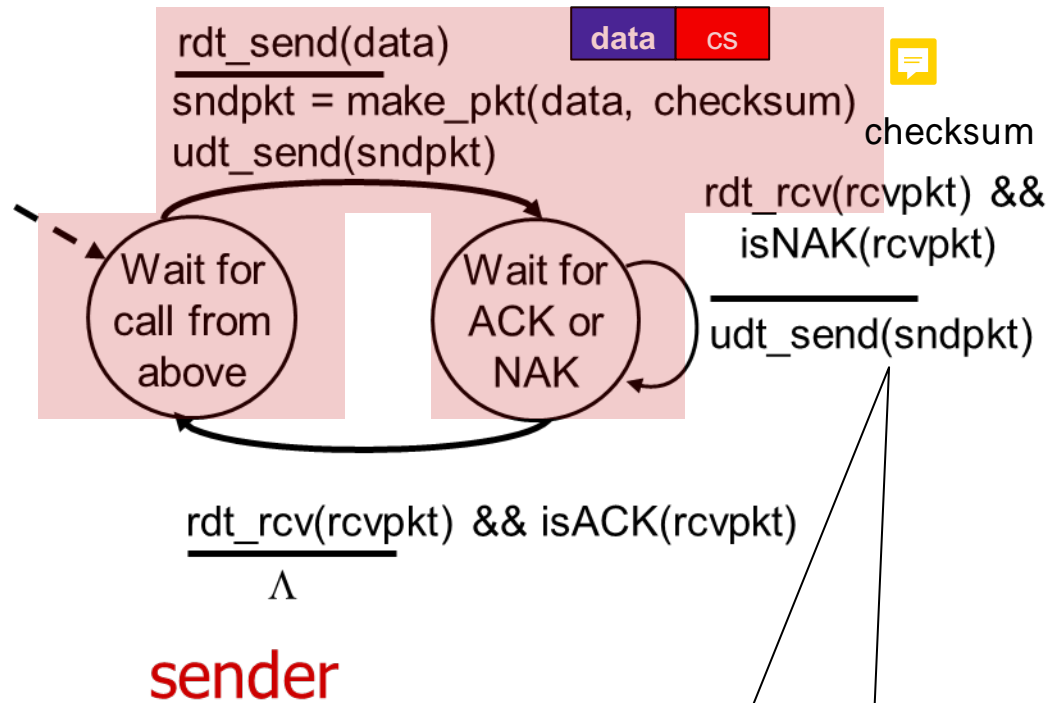
- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
 - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
 - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- ❖ new mechanisms in rdt2.0 (beyond rdt1.0):
 - error detection
 - feedback: control msgs (ACK,NAK) from receiver to sender
 - Retransmission




Big Picture of rdt2.0



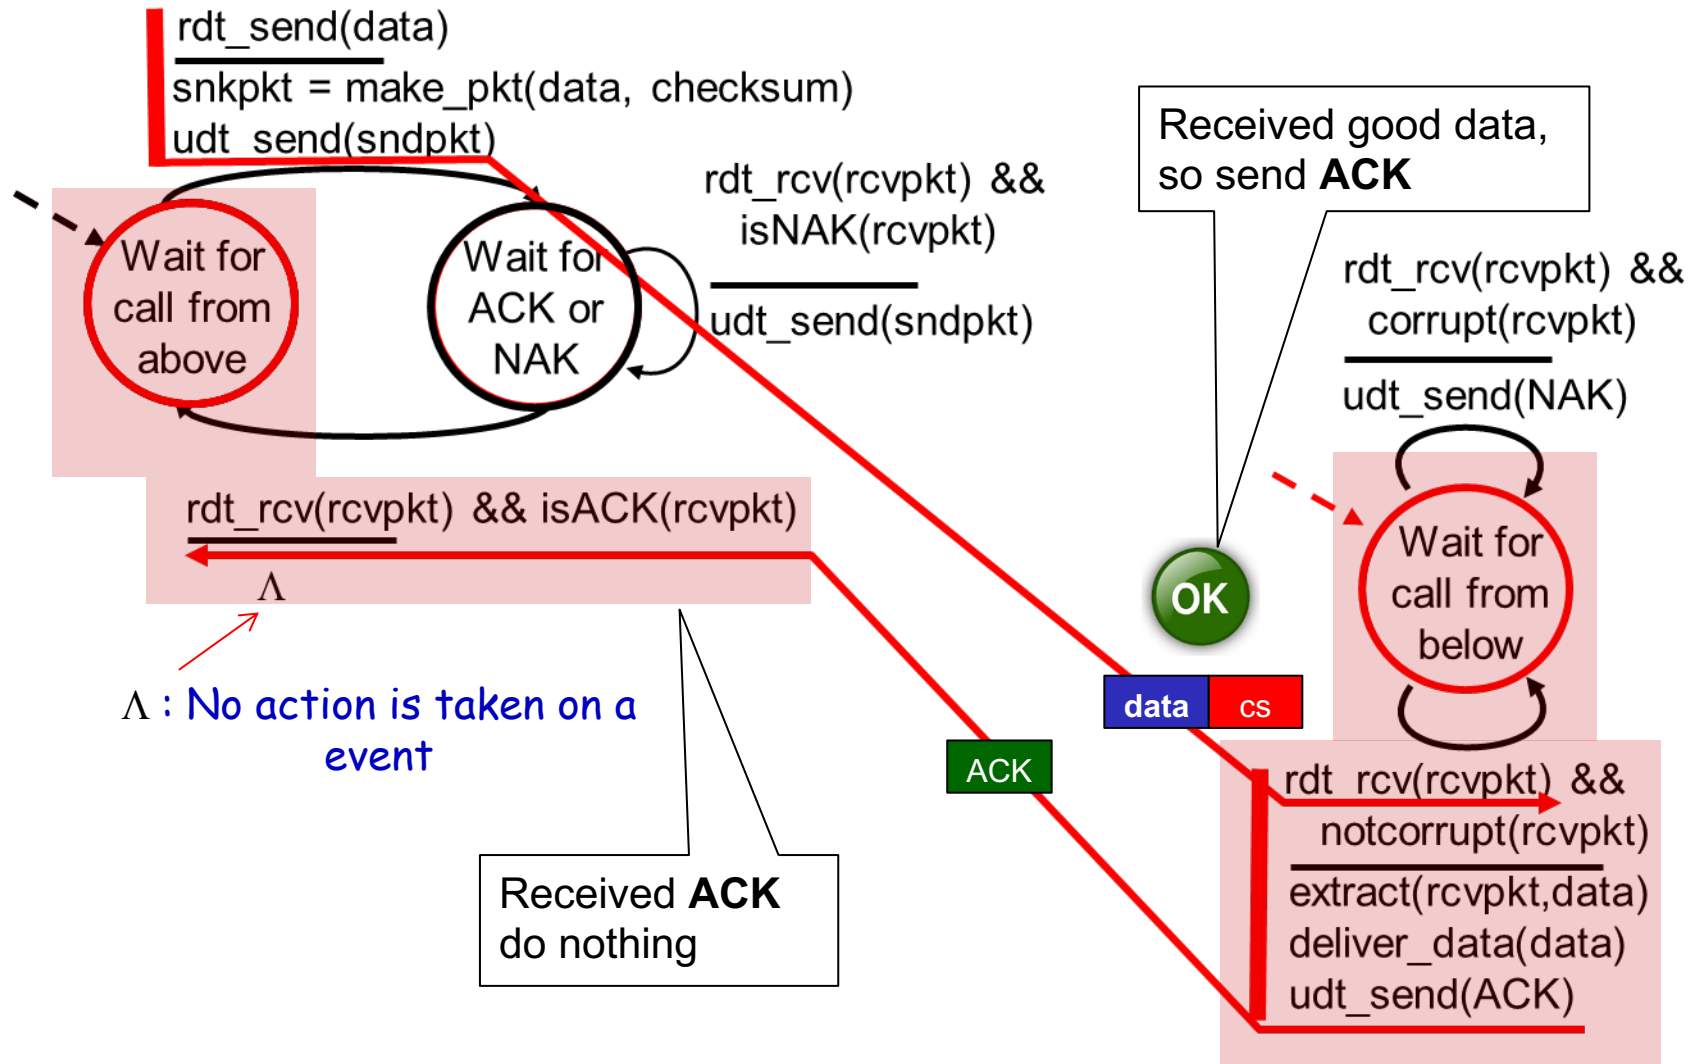
rdt2.0: FSM specification



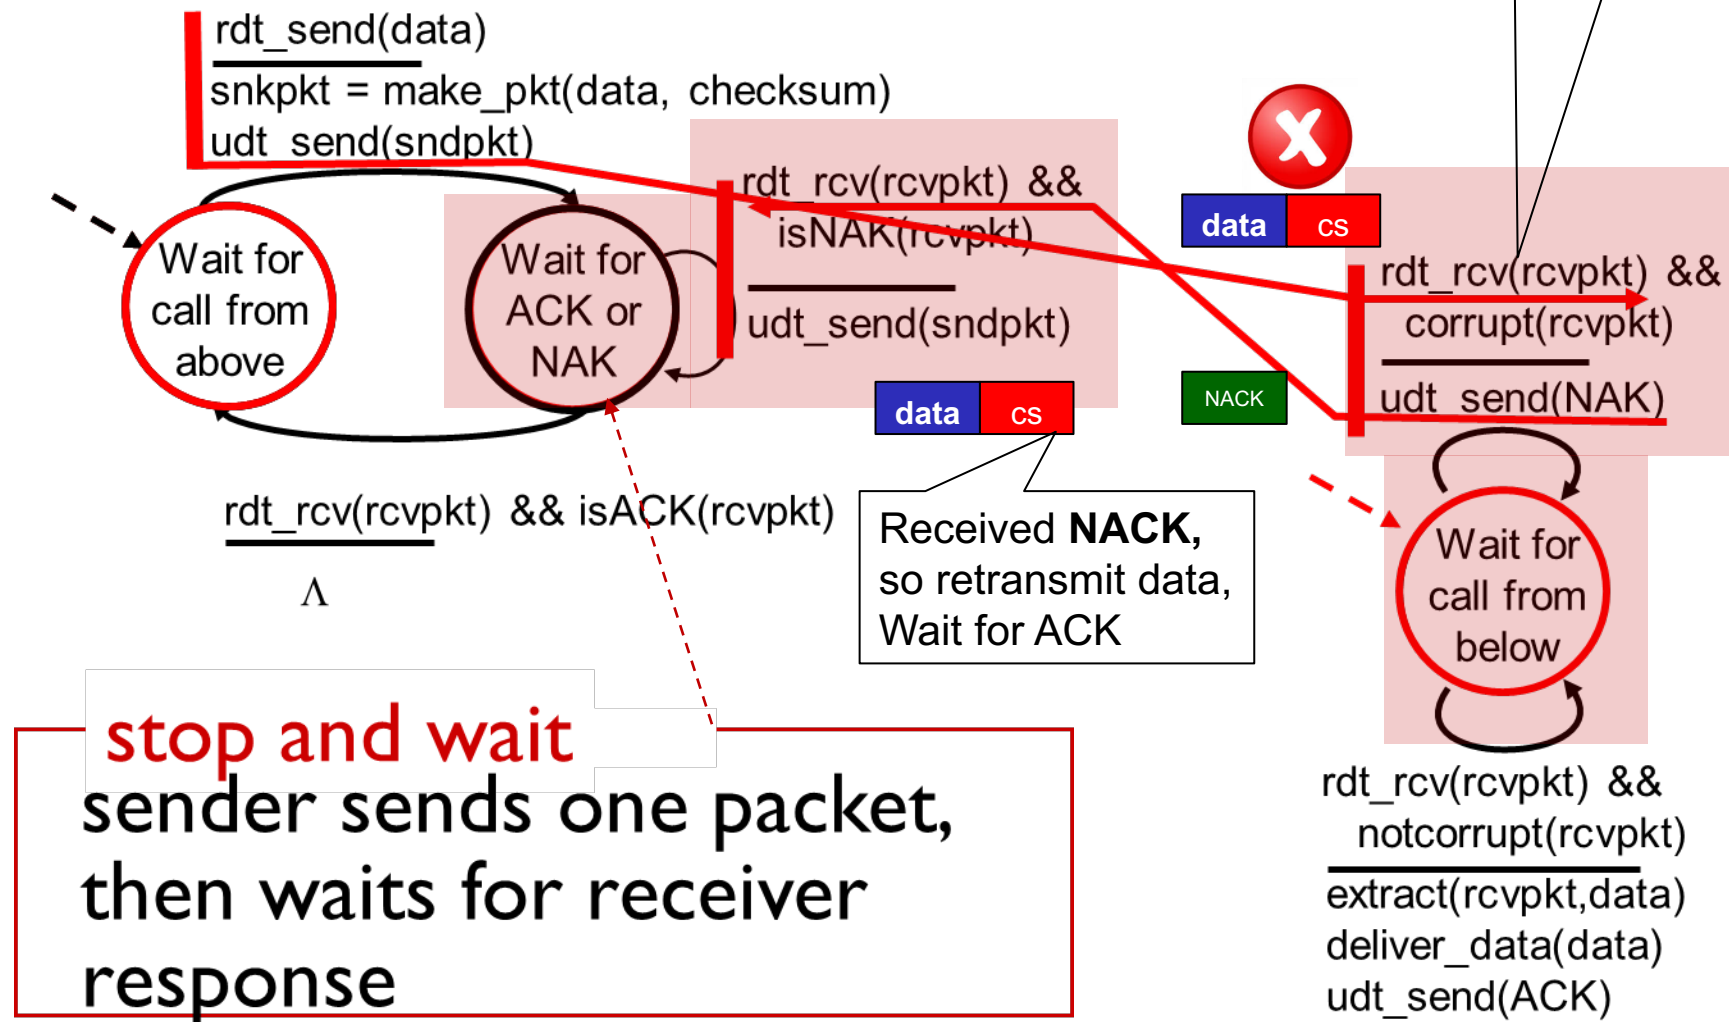
Send packet (with checksum)
and wait for ACK/NAK



rdt2.0: operation with no errors



rdt2.0: error scenario



rdt3.0: channels with errors and loss

new assumption:

underlying channel can also **lose** packets (data, ACKs) 

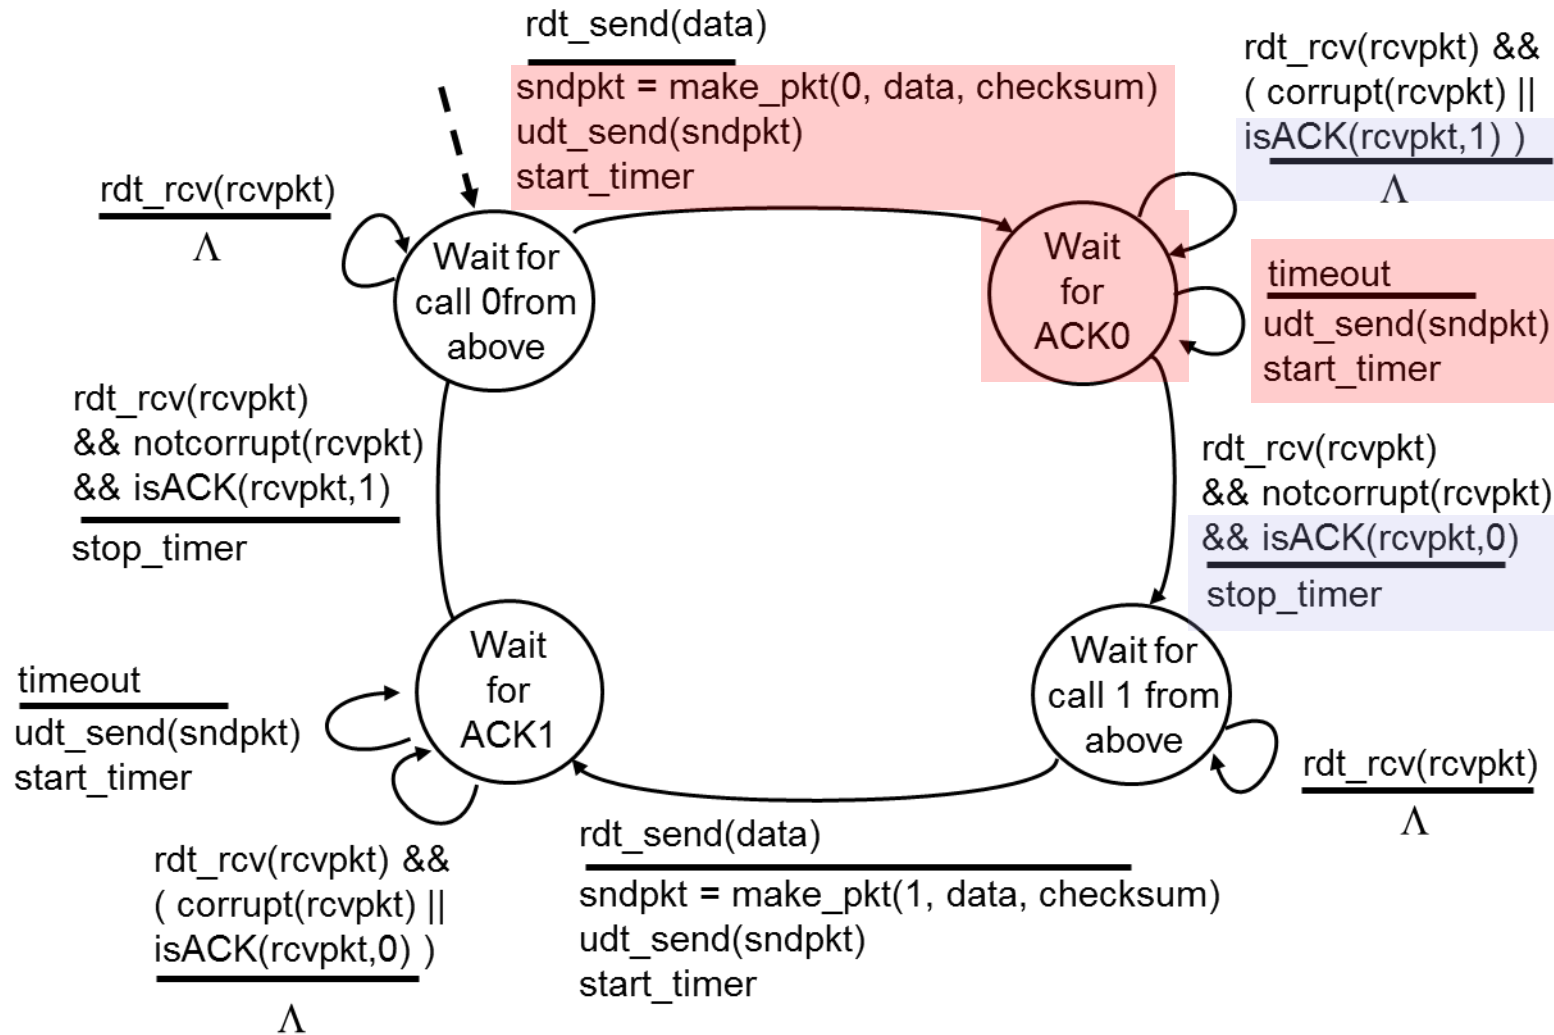
- Receiver will *not* receive any segments
- How will sender know that data packet (or ACK) has been lost?

approach: sender waits **reasonable** amount of time for ACK

- ❖ retransmits if no ACK received in this time
- ❖ requires countdown timer 

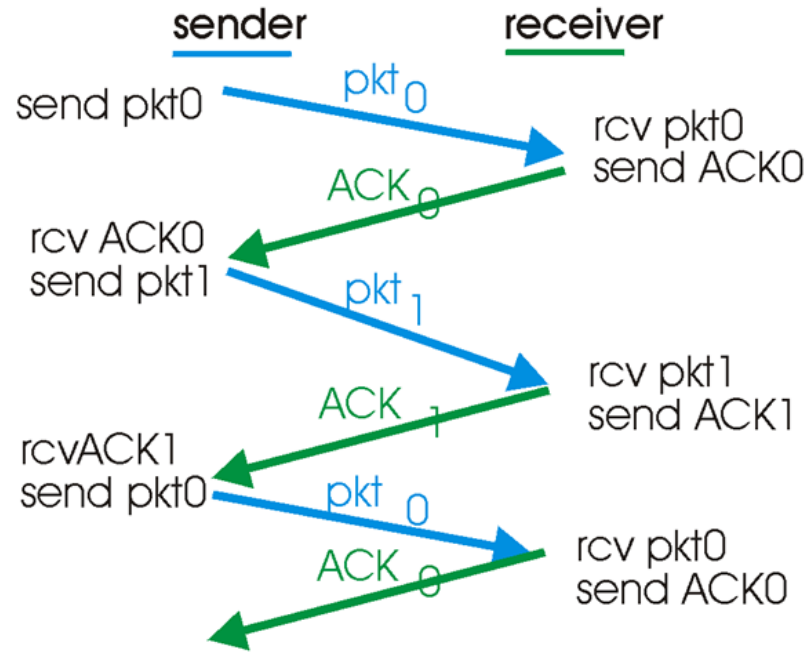


rdt3.0 sender



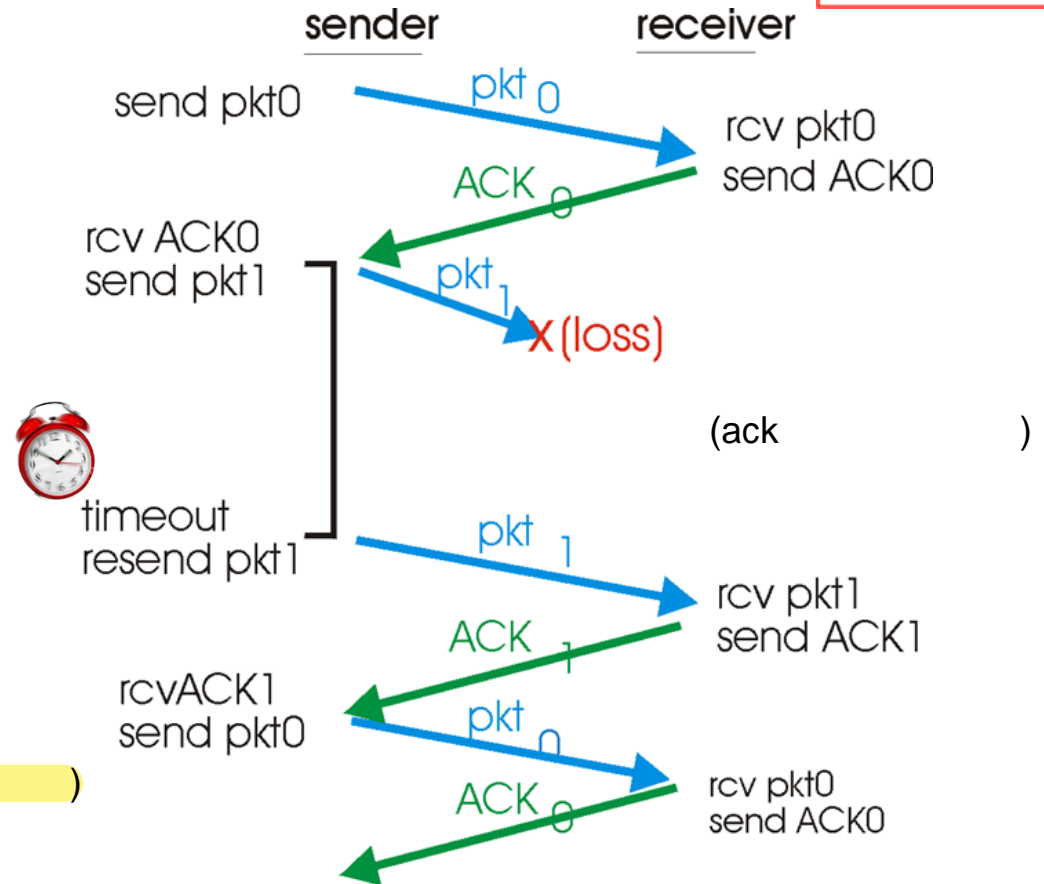
rdt3.0 in action

How many seq#?
Two sequence numbers are enough!




(a) operation with no loss (loss가 .)

0 1 가 .



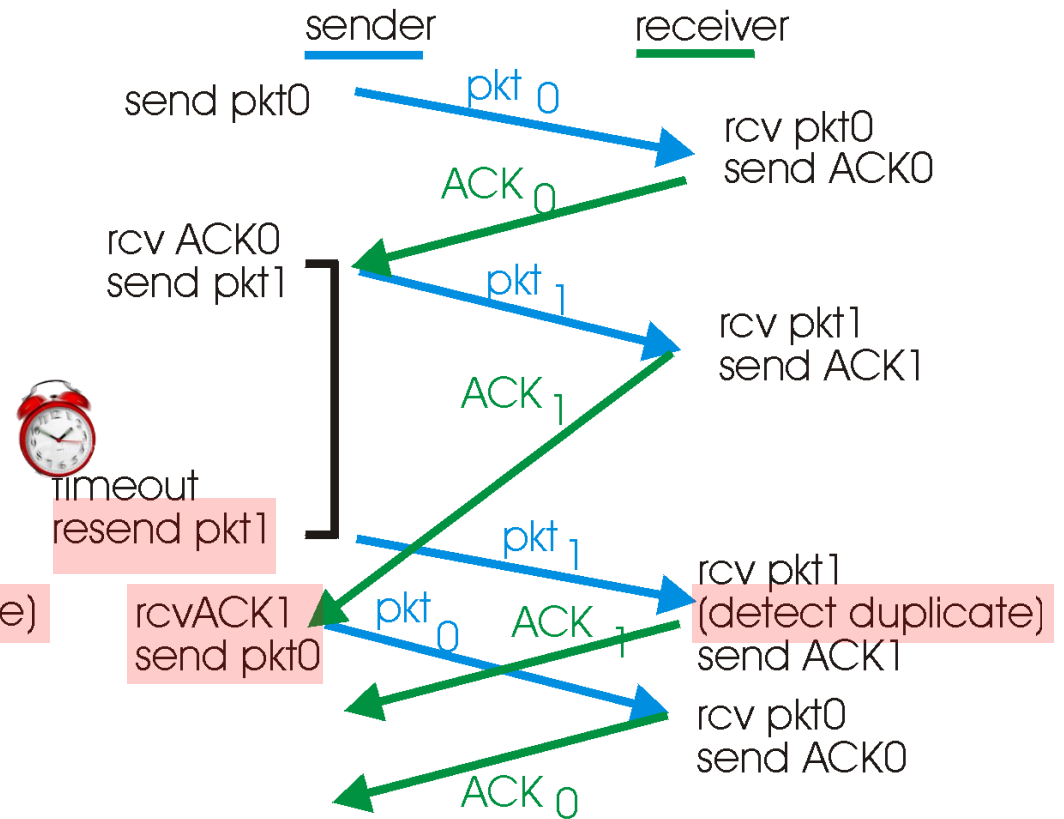
(b) lost packet

rdt3.0 in action



(c) lost ACK

sequence number(0 1)
pkt1



(d) premature timeout

timeout
ack1

loss
pkt0

Sequence Numbers

what happens if ACK is lost or arrives late?

- ❖ Sender will retransmit
- ❖ But, receiver has received the data packet
 - possible duplicate at receiver!
 - Handling duplicate at receiver
 - Is it a new data packet or retransmission?
- ❖ sender adds 1-bit **sequence number** to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkts

handling duplicates:

- ❖ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate: data seq. #'s handles this
- ❖ Also, ACK may be duplicate
 - pkt (or ACK) just delayed (not lost)
 - receiver must also specify seq # of pkt being ACKed


duplicated

Performance of rdt3.0


- ❖ rdt3.0 (stop-and-wait) is correct, but performance not good
- ❖ e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ us} (= 0.008 \text{ ms})$$

- U_{sender} : **utilization** – fraction of time sender busy sending
가



rdt3.0: stop-and-wait operation



- U_{sender} : **utilization** : 가

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027 \quad (\text{or just } 0.027\%)$$


- if $RTT = 30$ msec, 8000-bit (=1kB) pkt every 30 msec: 267kbps throughput over 1 Gbps link
- ❖ network protocol (rdt3.0) limits use of physical resources (1Gbps) !

Pipelined protocols

ack

pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased 0 1 range.
- buffering at sender and/or receiver . buffering




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

❖ two generic forms of pipelined protocols:
go-Back-N, selective repeat

Pipelining: increased utilization


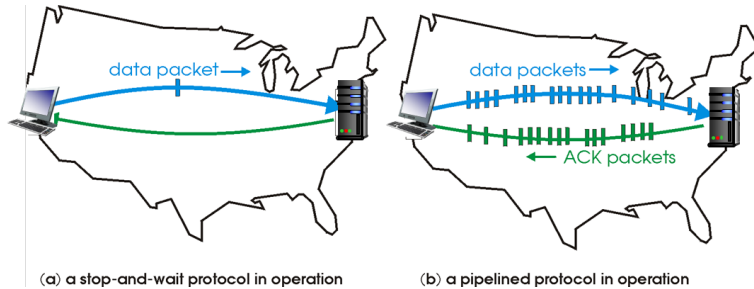


3-packet pipelining increases utilization by a factor of 3!

$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

$$= 0.0027N$$

Limit of N?



❖ Prev. example, $U = 0.0027N$

- Larger N means better utilization

❖ Can we increase N without any limit?

- Answer: No
- Flow Control and Congestion control (Next few chapters)

Pipelined protocols: Go-back-N (GBN)

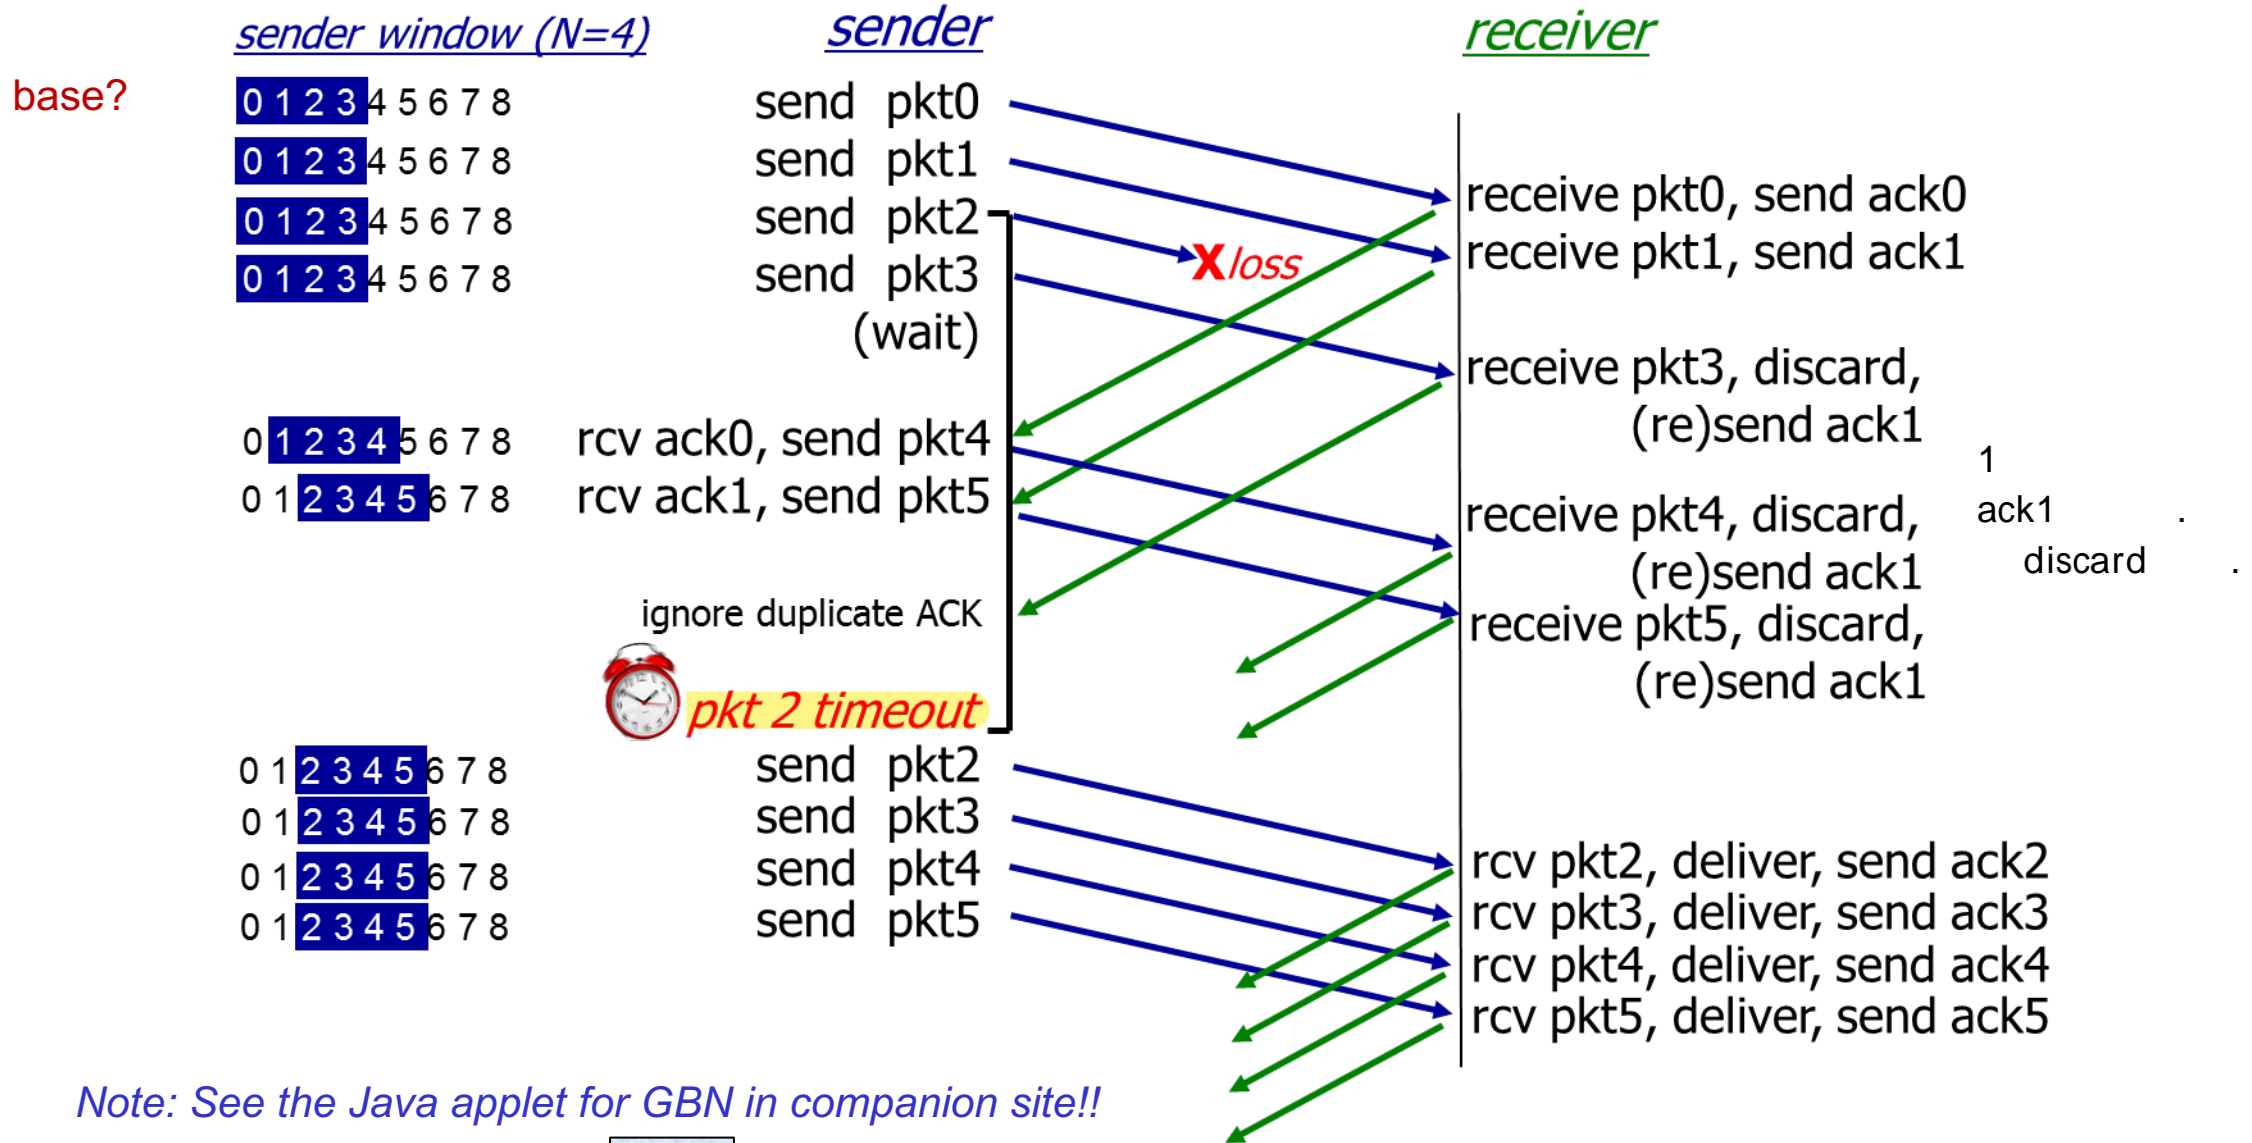
Go-back-N:

- ❖ Sender can send up to **N** (=window size) unacked packets
 - N: window size
 - k-bit seq # in pkt header (rtd3.0 had 1-bit)
- ❖ **base**: sender has timer for oldest unacked packet
- ❖ **Sliding-window protocol**

가 oldest unacked packet ack ...

timer 1 oldest unacked packet .

GBN in action




Note: See the Java applet for GBN in companion site!!



Go-Back-N: sender

- ❖ **window** of up to **N**, consecutive unack'ed pkts allowed




- ❖ **base**: seq# of oldest unacknowledged packet
- ❖ **nextseqnum**: smallest unused seq#
- ❖ Seq# uses k bits: seq# range is $[0, 2^k - 1]$.
 - ❖ TCP uses 32-bit seq#

Go-Back-N: sender

- ❖ receiver sends - “*cumulative ACK*”
 - ACK(n): all pkts *up to and including* n have been correctly received at receiver – “ n 까지 잘 받았음” (example: ACK loss)
 - Loss happens: receiver sends duplicate ACKs – “ n 까지 잘 받았는데 다음은 못 받았음”
- ❖ Timeout event
 - When timeout occurs, retransmit all unacked packets
 - Single timer for oldest in-flight pkt (=base)

GBN: sender extended FSM



rdt_send(data)

if (nextseqnum < base+N) { **Send up to N unacked packets**

 sndpkt[nextseqnum] = make_pkt(nextseqnum, data, chksum)

 udt_send(sndpkt[nextseqnum])

 if (base == nextseqnum)

 start_timer


 nextseqnum++

 } else
 refuse_data(data)

Set timer for oldest (base) packet

Λ
base=1
nextseqnum=1

rdt_rcv(rcvpkt)
&& corrupt(rcvpkt)



timeout



Retransmit all sent unacked packets

start_timer

udt_send(sndpkt[base])

udt_send(sndpkt[base+1])

...

udt_send(sndpkt[nextseqnum-1])

rdt_rcv(rcvpkt) &&
notcorrupt(rcvpkt)

base = getacknum(rcvpkt)+1

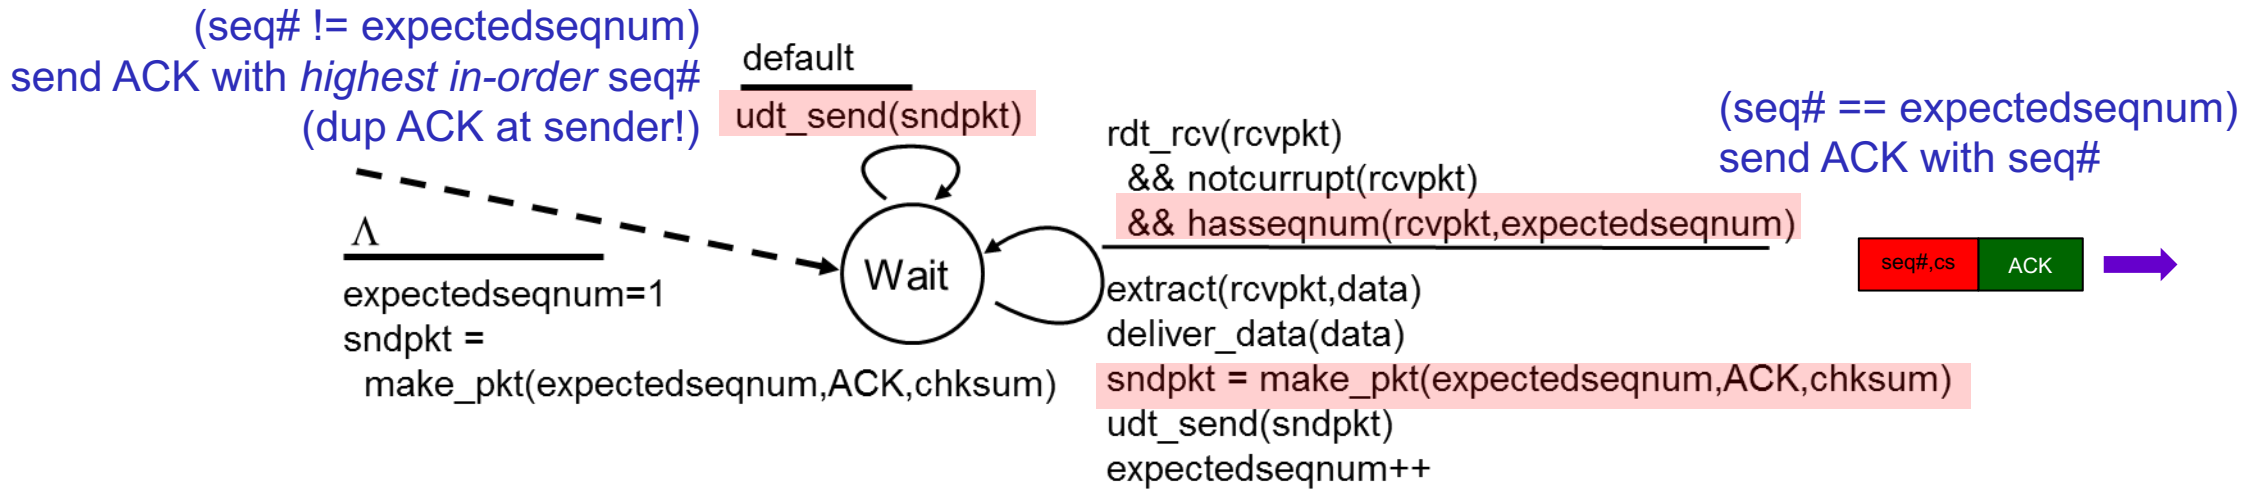
If (base == nextseqnum)

 stop_timer

else

 start_timer

GBN: receiver extended FSM



- ❖ If data packet n is received correctly, and is *in-order* (=packet $n-1$ has also been received), then send ACK for packet n
- ❖ Otherwise (if *out-of-order*), resend an ACK for *most recently received in-order packet*
 - may generate duplicate ACKs
 - need only remember **expectedseqnum** (= seq# of next in-order packet)
 - *Out-of-order* packets: discard (don't buffer): *no receiver buffering!*
 - Why not just store out-of-order pkts? Simpler receiver; no need to buffer out-of-order pkts, they will be retransmitted by GBN sender anyway.

Pipelined protocols: Selective Repeat (SR)

가 .

Selective Repeat:

- ❖ sender can have up to **N** unack'ed packets in pipeline (same as GBN)
- ❖ receiver sends *individual ack* for each packet
- ❖ sender maintains timer for each unacked packet
 - when timer expires, retransmit only that unacked packet


- cumulative ack .

Selective repeat

- ❖ receiver *individually* acknowledges all correctly received pkts
 - buffers pkts, as needed, for eventual in-order delivery to upper layer
- ❖ sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt
- ❖ sender window
 - N consecutive seq #'s
 - limits seq #'s of sent, unACKed pkts

Selective repeat in action

loss가
window



Note: See the Java applet for GBN in companion site!!



memory

가

Selective repeat

sender

data from above:

- ❖ if next available seq # in window, send pkt

timeout(n):

- ❖ resend only pkt n , restart timer (compare GBN sender)

ACK(n) in [sendbase, sendbase+N]:

- ❖ mark pkt n as received
- ❖ if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver


pkt n in [rcvbase, rcvbase+N-1]

- ❖ send ACK(n)
- ❖ out-of-order: **buffer**
- ❖ in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt


otherwise:

- ❖ ignore

Selective repeat: sender, receiver windows



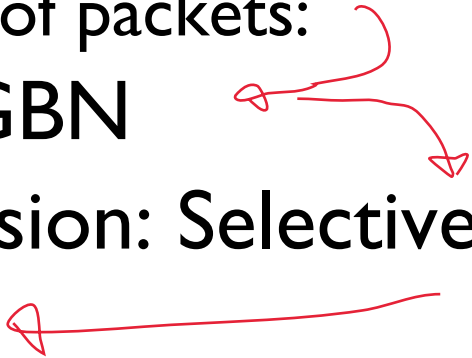
(a) sender view of sequence numbers



(b) receiver view of sequence numbers

Receiver more complex!
GBN receiver only needed to remember **expectedseqnum**

Reliable Data Transfer: Summary

- ❖ Stop-and-wait
 - rdt 2.0: bit errors (but no packet losses):
 - rdt 3.0: loss (drop) of packets:
 - ❖ Better utilization: GBN
 - ❖ Efficient retransmission: Selective Repeat
 - ❖ Next step: TCP!
- 
- Hand-drawn red arrows: one from 'rdt 3.0: loss (drop) of packets:' to 'Better utilization: GBN', and another from 'Next step: TCP!' to 'Efficient retransmission: Selective Repeat'.