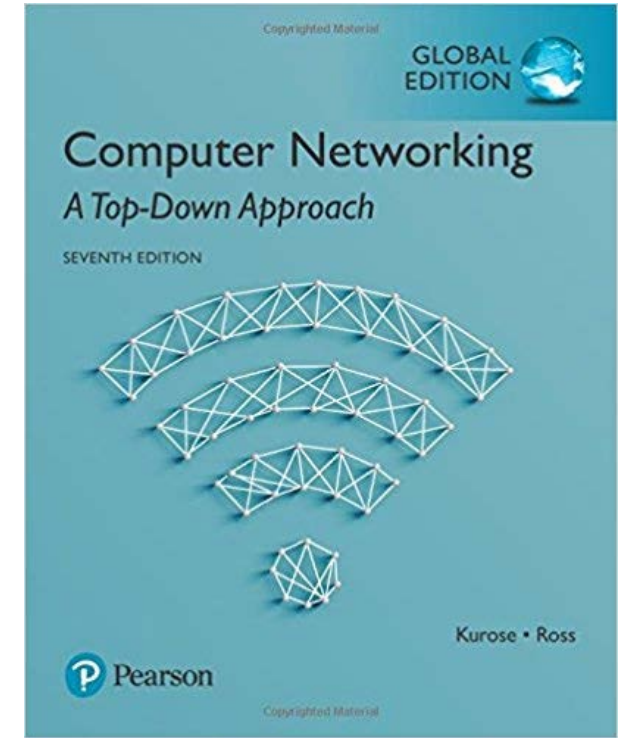


Chapter 3

Transport Layer part 4

School of Computing
Gachon Univ.
Joohyung Lee

Most of slides from J.F Kurose and K.W. Ross. And, some slides from Prof. Joon Yoo



*Computer
Networking: A Top
Down Approach*

7th edition

Jim Kurose, Keith Ross
Pearson, 2017

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

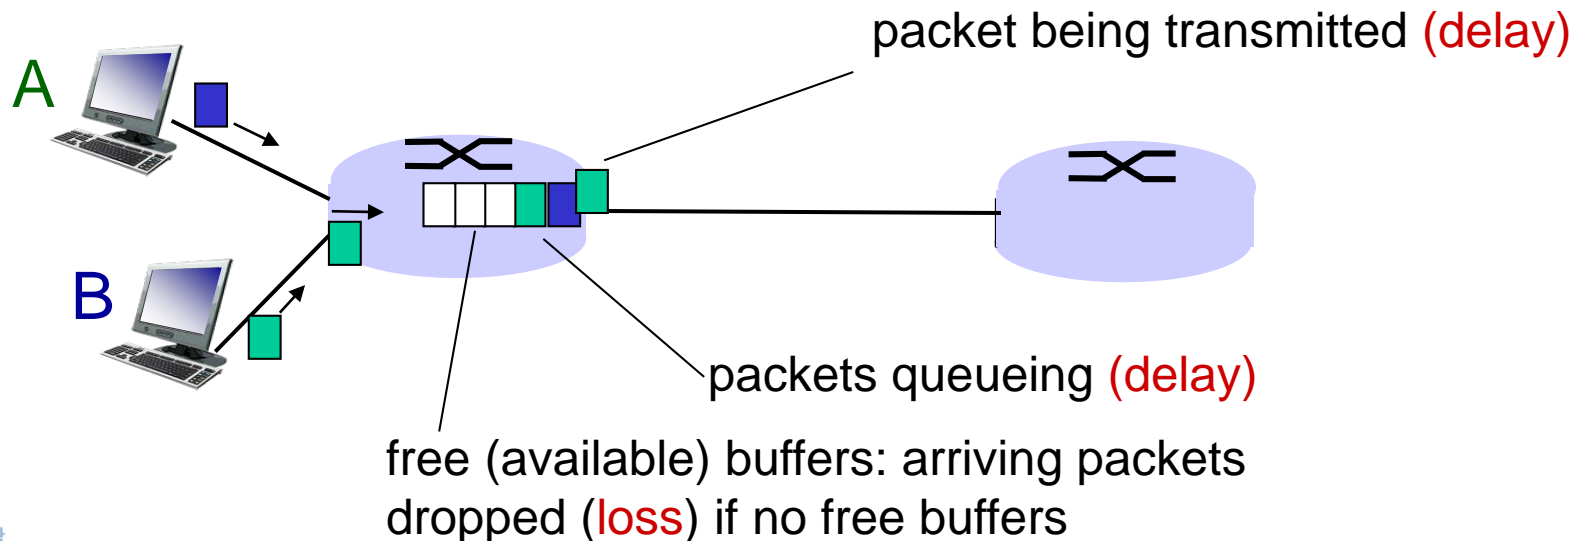
3.7 TCP congestion control

How do loss and delay occur?

packets *queue* in router buffers

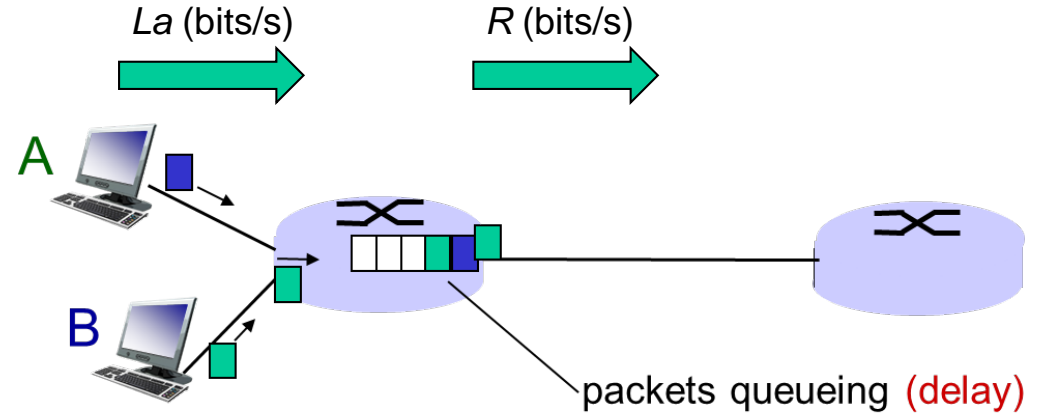
- ❖ packet arrival rate to link (temporarily) exceeds output link capacity
- ❖ packets queue, wait for turn
- ❖ Unlike other delays (d_{proc} , d_{trans} , and d_{prop}), queueing delay (d_{queue}) can vary from packet to packet

↓
queueing delay : variation
()



Queueing delay (Ch. 1.4.2)

- ❖ L : packet length (bits)
- ❖ a : average packet arrival rate (packets/s)
- ❖ R : link bandwidth (bps)



❖ $La/R = \rho$: **Traffic Intensity**

link capacity

queueing delay

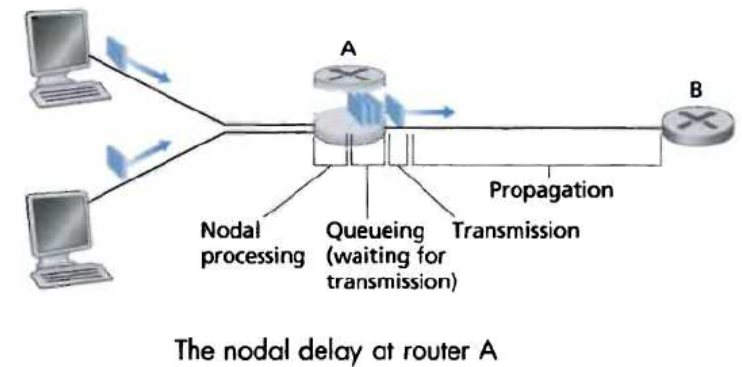
- ❖ $\rho > 1$: more “work” arriving than can be serviced; queue will tend to increase without bound; average delay infinite.
 - This should never happen!
- ❖ Design the system so that $\rho \leq 1$

queue가

Principles of congestion control

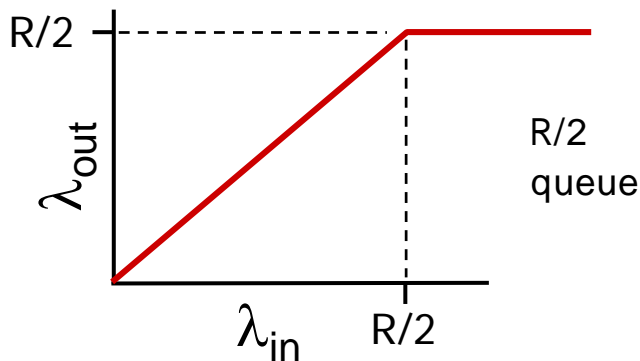
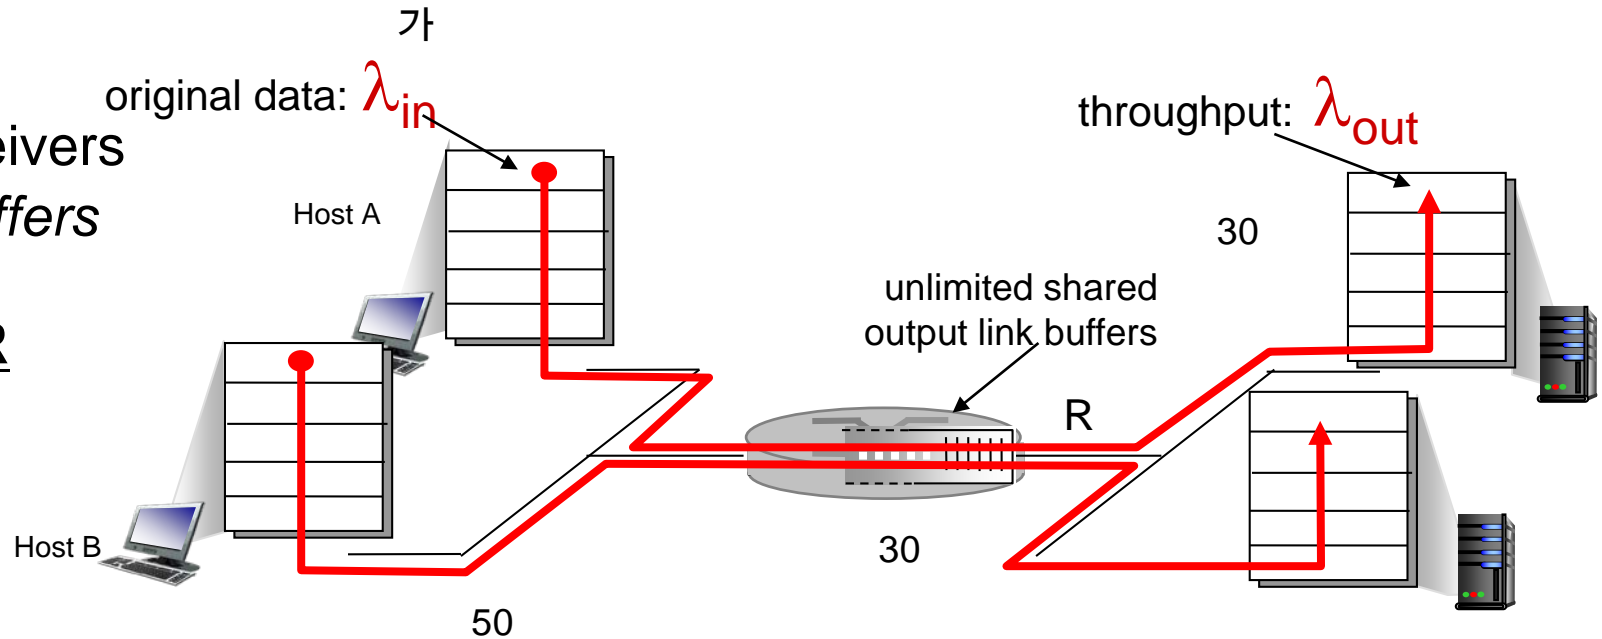
congestion:

- ❖ informally: “too many sources sending too much data too fast for network to handle”
- ❖ Symptoms (signs):
 - long delays (queueing in router buffers) = queueing delay
 - **lost packets** (buffer overflow at routers)
- ❖ **Congestion Control:**
 - ways to treat the cause of **network congestion**
 - different from **flow control**!
 - flow control considered the *receiver's* capacity

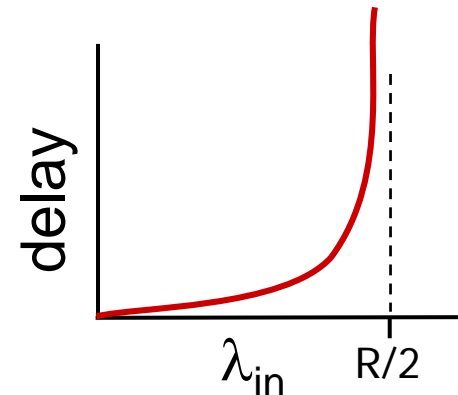


Causes/costs of congestion: scenario 1

- ❖ two senders, two receivers
- ❖ one router, *infinite buffers*
 - packet loss?
- ❖ output link capacity: R
 - each flow shares $R/2$
- ❖ no retransmission



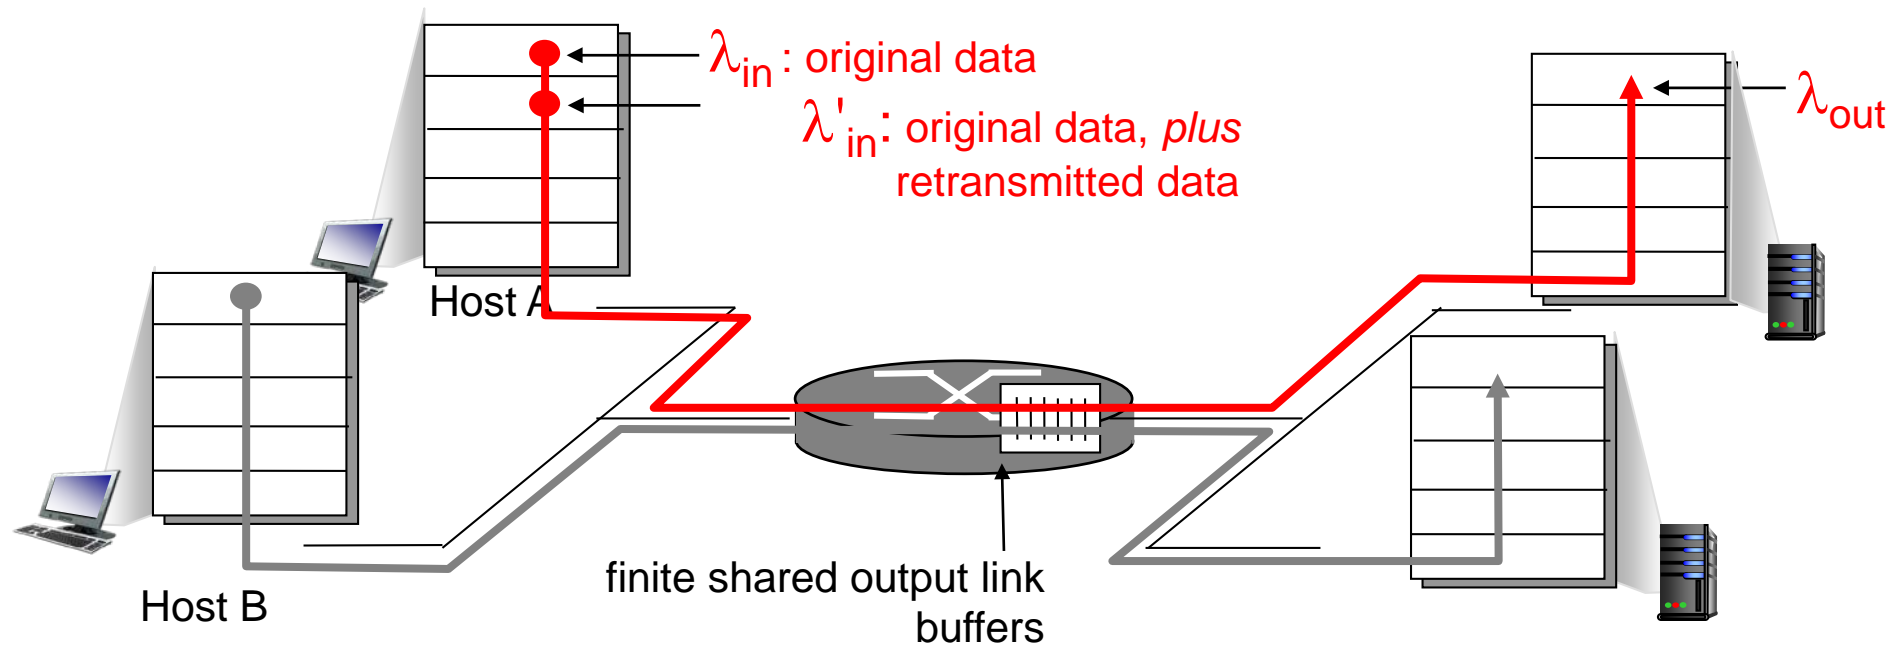
- ❖ maximum per-connection throughput: $R/2$



- ❖ **large queuing delays** as arrival rate, λ_{in} , approaches capacity

Causes/costs of congestion: scenario 2

- ❖ one router, *finite* buffers
 - packet loss?
- ❖ sender retransmission of timed-out packet
 - application-layer input = application-layer output: $\lambda_{in} = \lambda_{out}$
 - transport-layer input includes *retransmissions* : $\lambda'_{in} \geq \lambda_{in}$



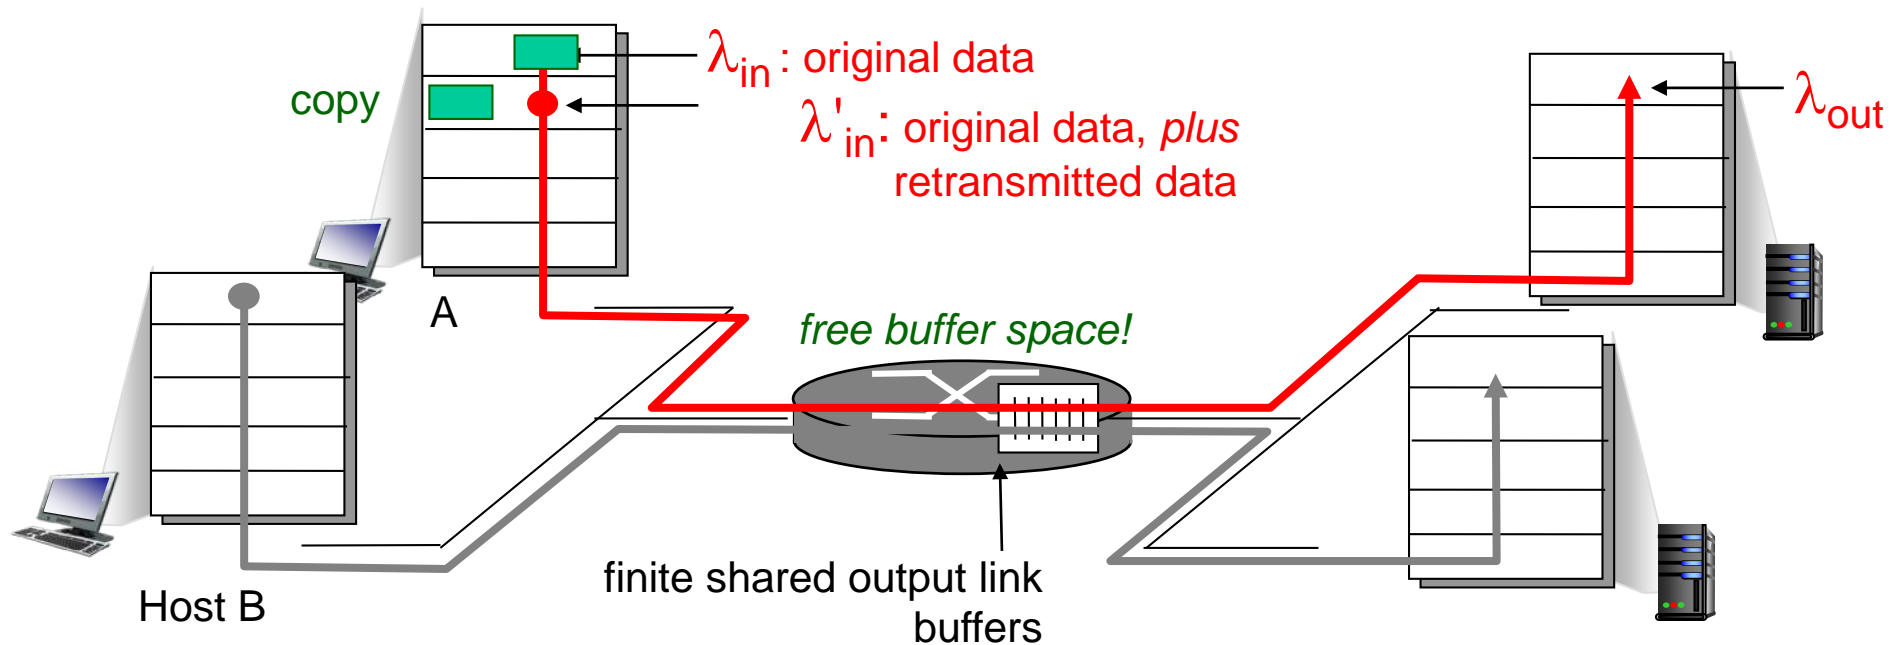
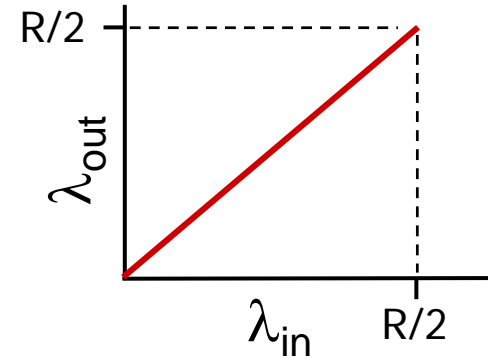
Causes/costs of congestion: scenario 2

idealization1: perfect knowledge

- ❖ sender sends only when router buffers available

- $\lambda'_{in} = \lambda_{in} = \lambda_{out}$

loss가



Causes/costs of congestion: scenario 2

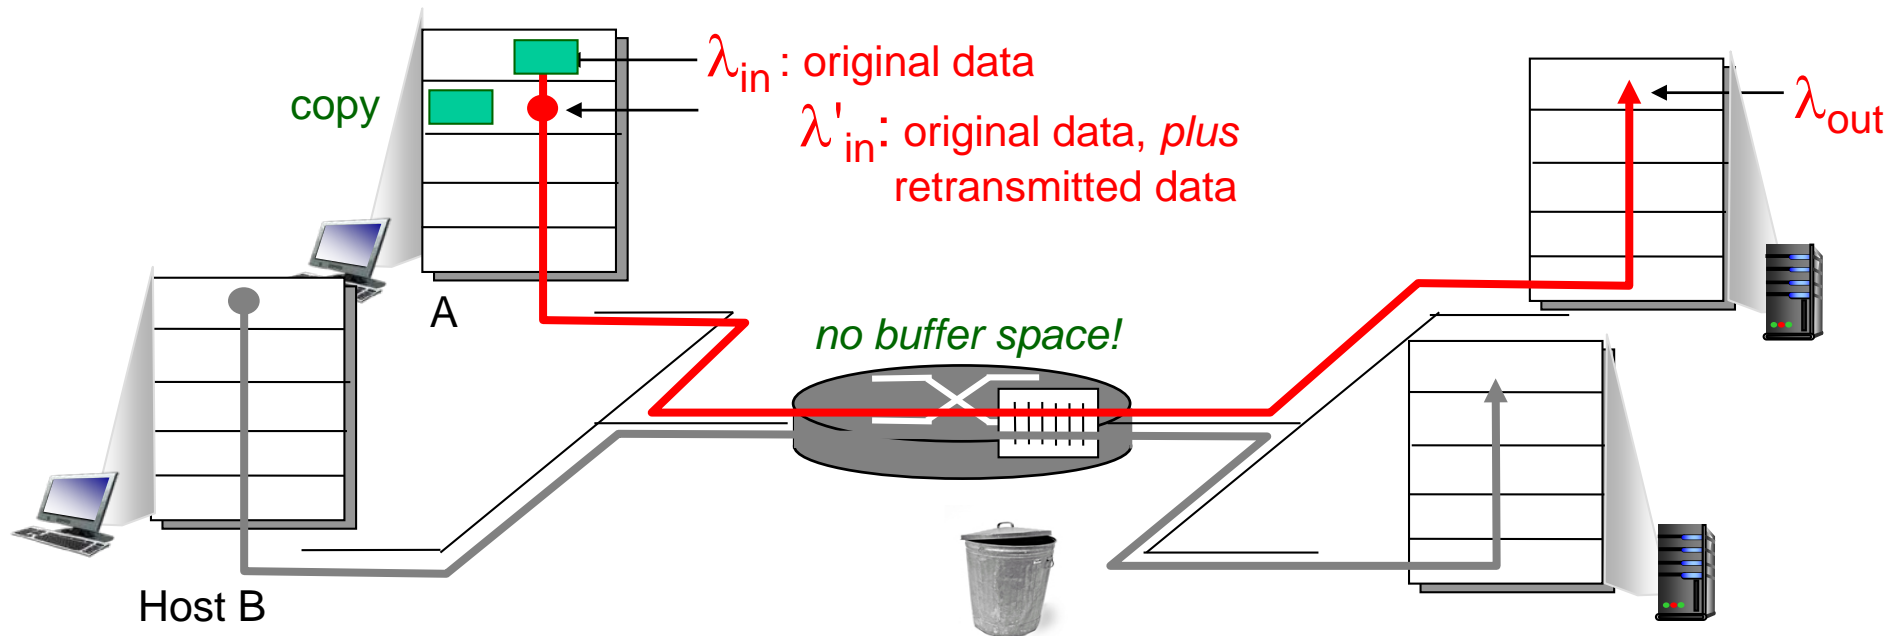
Idealization2: known loss

packets can be lost, dropped at router due to full buffers

가

- ❖ sender only resends if packet *known* to be lost

- $\lambda'_{in} \geq \lambda_{in}$ 가



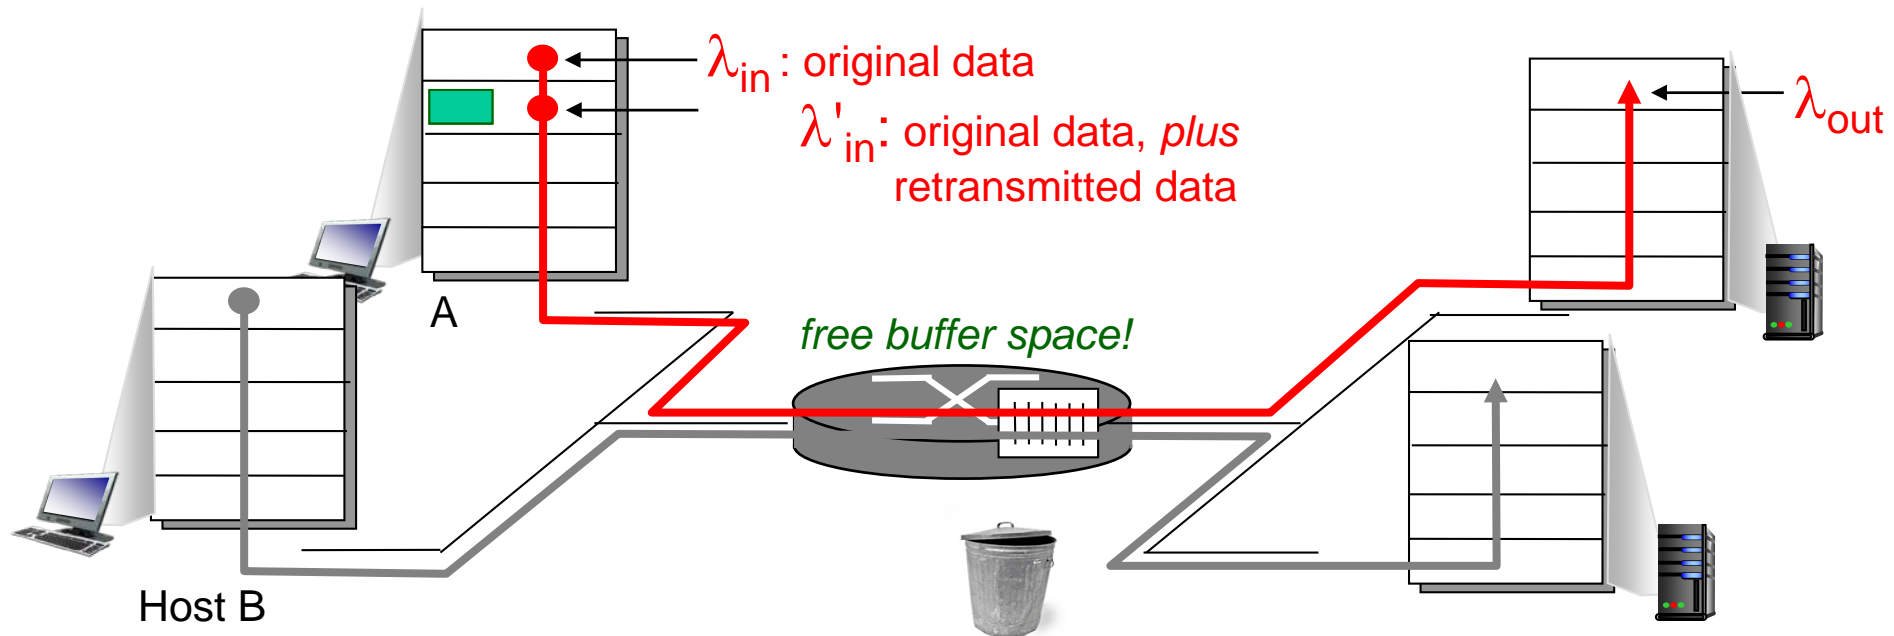
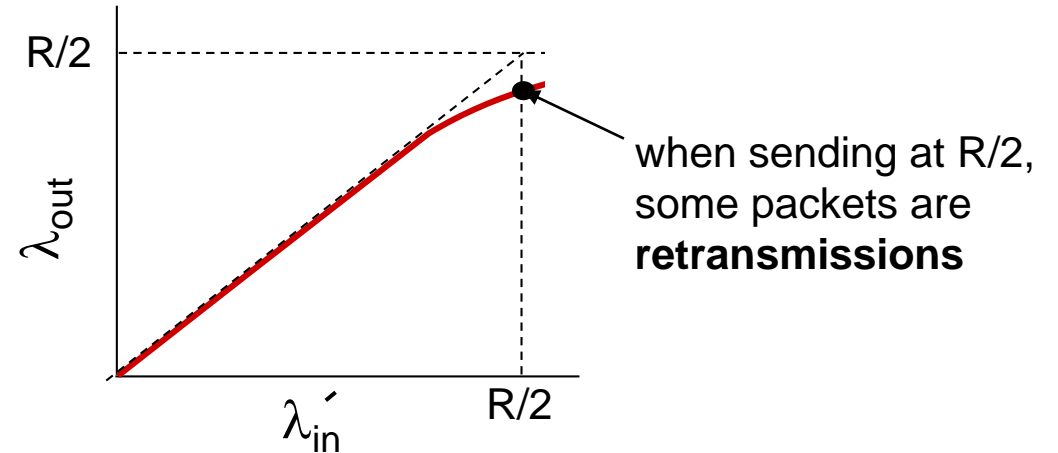
Causes/costs of congestion: scenario 2

Idealization2: known loss

packets can be lost, dropped at router due to full buffers

- ❖ sender only resends if packet *known* to be lost

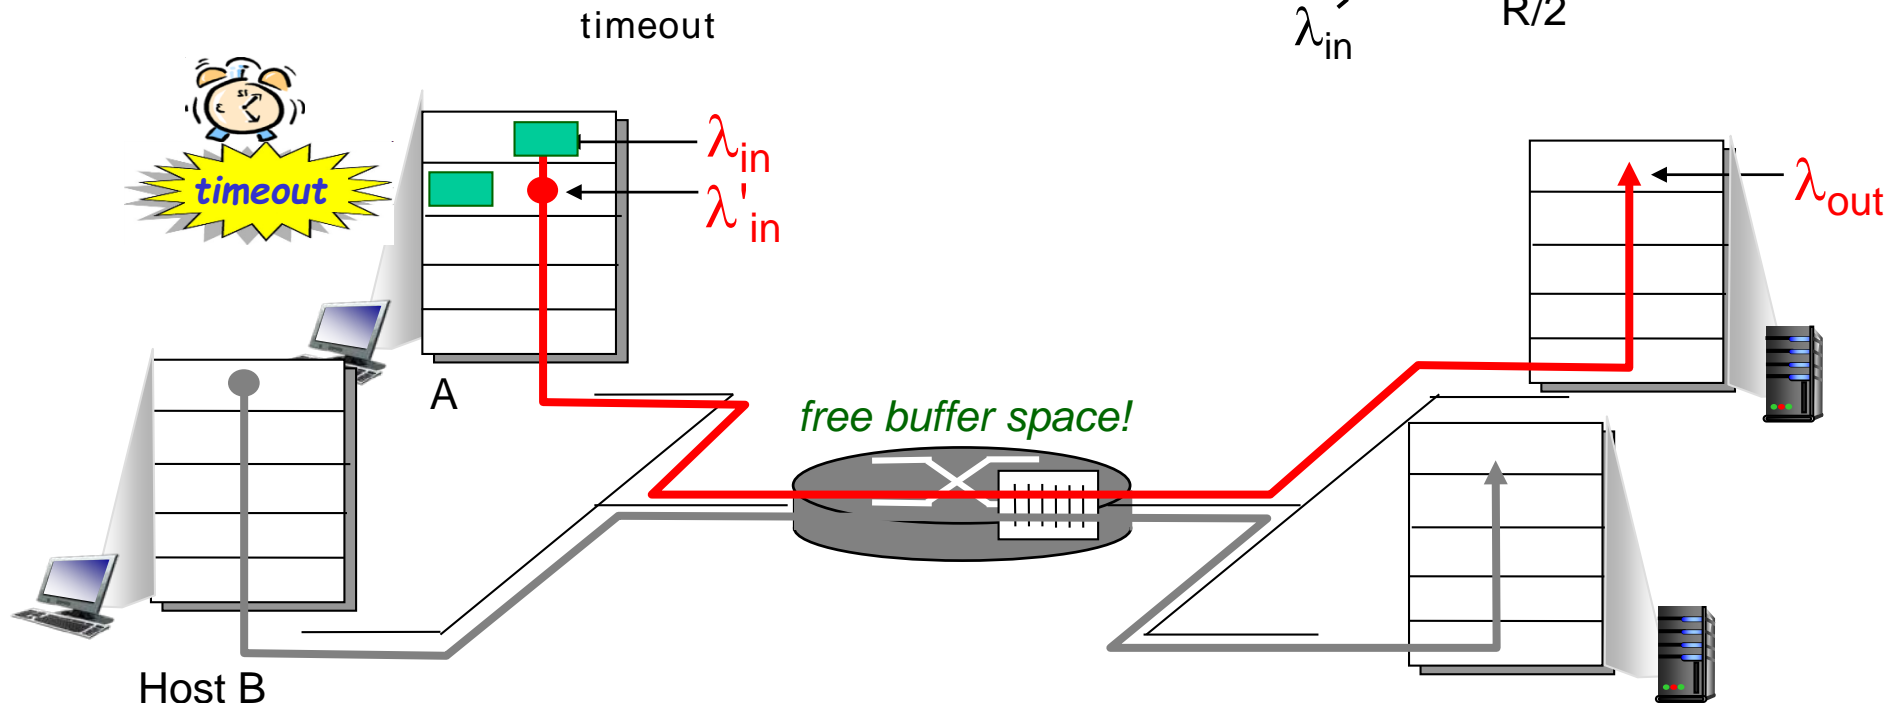
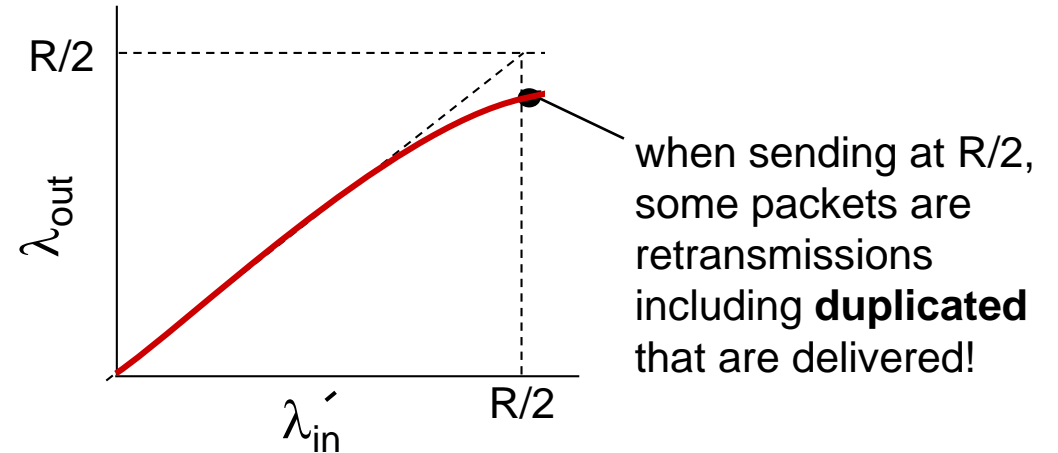
- $\lambda'_{in} \geq \lambda_{in}$



Causes/costs of congestion: scenario 2

Realistic: duplicates

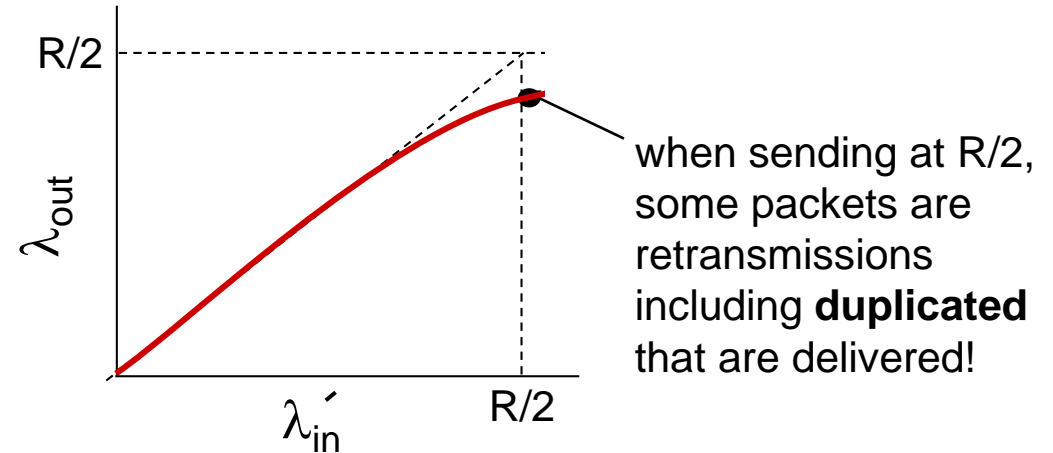
- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending *two* copies, both of which are delivered



Causes/costs of congestion: scenario 2

Realistic: duplicates

- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending *two* copies, both of which are delivered

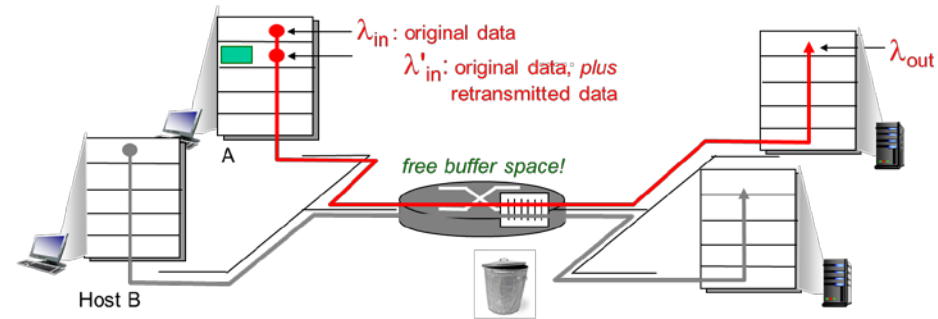


“costs” of congestion:

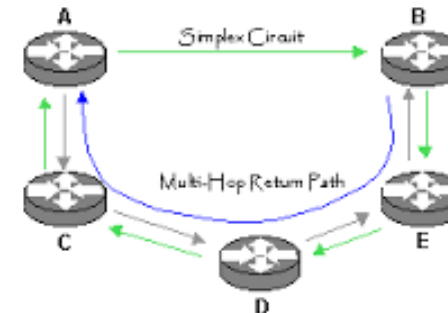
- ❖ large queueing delays
- ❖ more work (retrans) for given “goodput”
- ❖ unneeded retransmissions: link carries multiple copies of pkt
 - decreasing goodput

From single-hop to multi-hop path

- ❖ Until now (scenario 2), we only considered a **single-hop** path congestion



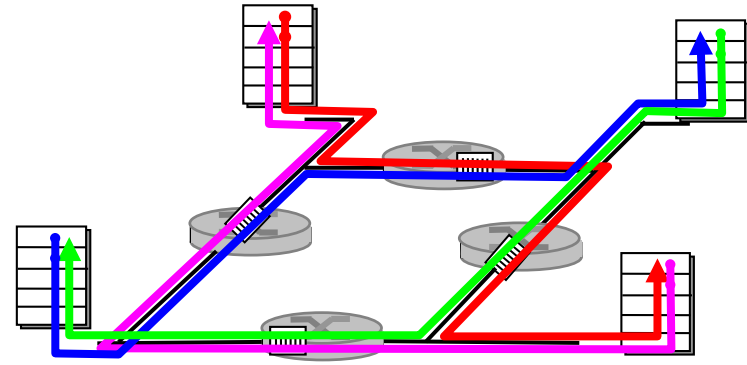
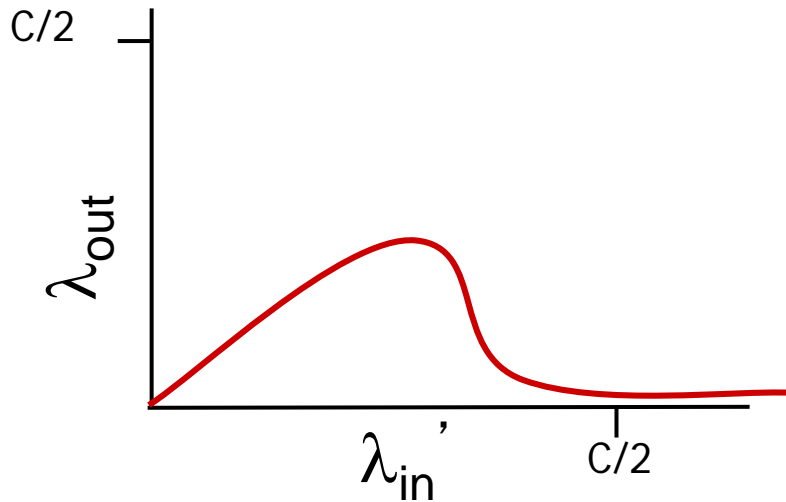
- ❖ But, generally in the Internet, the network path is **multi-hop** (Scenario 3)
- ❖ Congestion can happen at any link along the multi-hop path



congestion

goodput

Causes/costs of congestion: scenario 3



another “cost” of congestion:

- ❖ when packet dropped, any transmission capacity used along the path for that packet was wasted!

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

: congestion
()

TCP congestion control

- ❖ TCP/IP protocol stack was operational in early 1980s
- ❖ But, Internet was suffering from **congestion collapse**
 - hosts would send their packets into the Internet as fast as received window would allow (only flow control)
 - congestion would occur at some router and packet drop happens, and hosts would retransmit their packets, resulting in even more congestion
- ❖ TCP congestion control introduced by Van Jacobson in later 1980s
 - Each host (end systems) determines how much capacity is available in the network

link
capacity가



sending rate

TCP congestion control

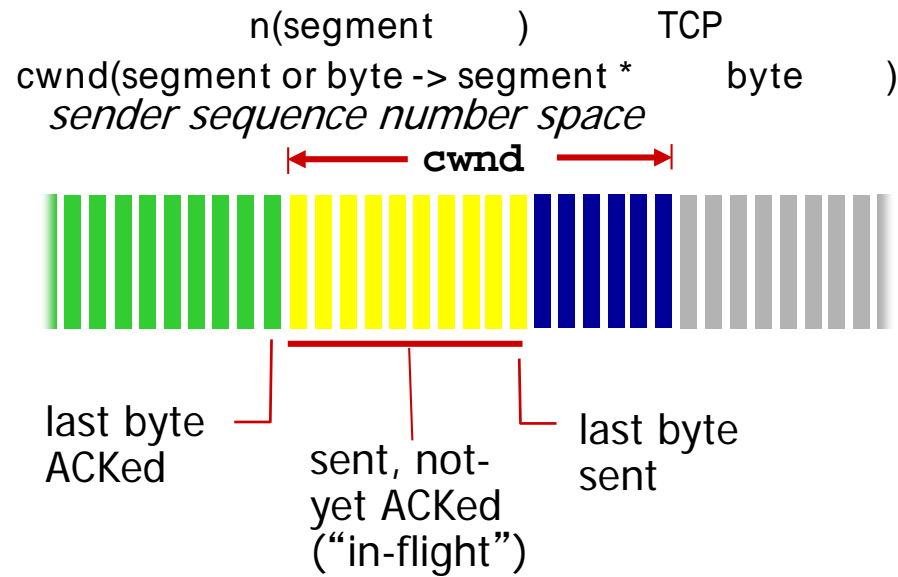
- ❖ End-to-end congestion control
 - router
 - client
 - link capacity
 - traffic intensity
- ❖ Sender needs to determine network capacity, which changes over time
 - n
 - $(n=1 \quad 1 \quad \text{ack} \quad)$
- ❖ Sender limits the sending rate as a function of available network capacity (i.e., network congestion)
 - e.g., little congestion – increase **send rate**
 - e.g., congestion along path – decrease **send rate**
- ❖ The only feedback to sender: ACK
 - ack 가 congestion
 - ACK arrival: a packet has arrived safely at receiver
 - ACK timeout: has not arrived -> congestion

TCP congestion control

❖ Questions

1. **How does the sender limit the rate?** ⁿ
2. How does the sender know there is congestion on path between itself and destination?
3. What algorithm should sender use to change send rate as function of congestions?

TCP Congestion Control: congestion window



- ❖ sender limits transmission:

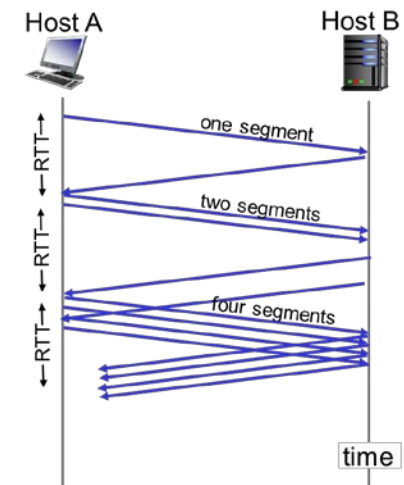
$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion

TCP sending rate:

- ❖ roughly: send **cwnd** bytes, wait **RTT** for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$



Transport Layer

TCP Congestion Control: congestion window

Congestion control(cwnd) and Flow control (advertised wnd) determine TCP sending rate

window

Send Window = MIN(flow control window, congestion window)

Note: advertised window is not so dynamic!

= advertised window

TCP congestion control

❖ Questions

1. How does the sender limit the rate?
- 2. How does the sender know there is congestion on path between itself and destination?**
3. What algorithm should sender use to change send rate as function of congestions?

Network Congestion

congestion

- ❖ How can the TCP sender know there is congestion?
 - Where does congestion actually happen?
 - However, network layer (e.g., router) does not send explicit message of congestion
- ❖ Loss event!
 - When there is excessive congestion, one (or more) router buffers along the path overflows, causing packet drop – **loss!**
 - Indication of loss
 - Timeout
 - Receipt of three duplicate ACKs

TCP congestion control

❖ Questions

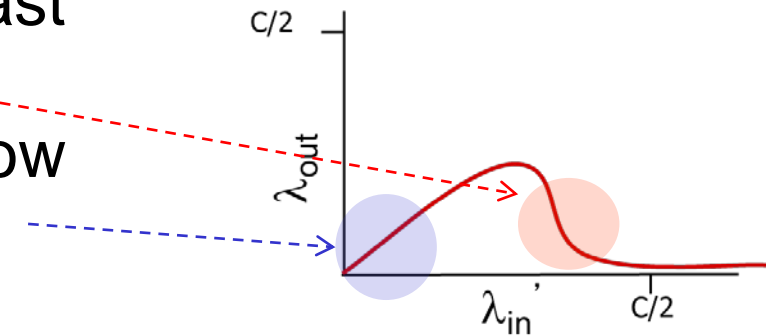
1. How does the sender limit the rate?
2. How does the sender know there is congestion on path between itself and destination?
3. **What algorithm should sender use to change send rate as function of congestions?**

TCP congestion control principles

❖ How should TCP sender determine send rate (or cwnd)?

- If TCP senders collectively send too fast
 - network congestion may happen
- If TCP senders cautiously send too slow
 - underutilize the network

utilization



❖ TCP principles

- A lost segment implies congestions, and hence, the TCP sender's rate should be decreased
- An acknowledged segment indicates successful delivery, and hence, the sender's rate can be increases when ACK arrives

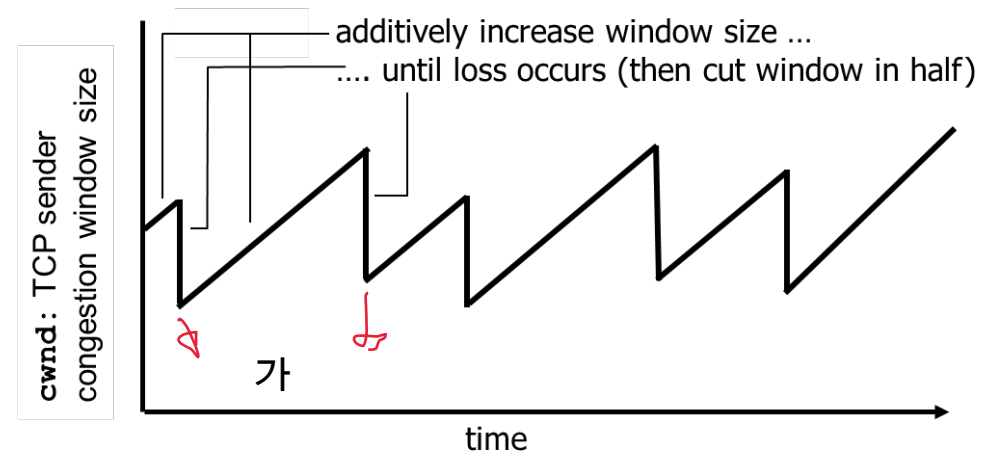
TCP congestion control principles

❖ Bandwidth probing

- TCP sender increases its rate in response to arriving ACKs until loss event occurs – then decrease rate
- TCP sender increases its rate to probe for rate at which congestion begins
- Then backs off from that rate, and then begins probing again to see if the congestion rate has changed.

❖ TCP Congestion Control Algorithm

- Slow start, congestion avoidance, fast recovery



TCP Congestion Control: Two modes

- ❖ TCP congestion control is governed by two parameters:

- **Congestion Window (cwnd)**
- **Slow-start threshold Value (ssthresh)**

(slow)

2

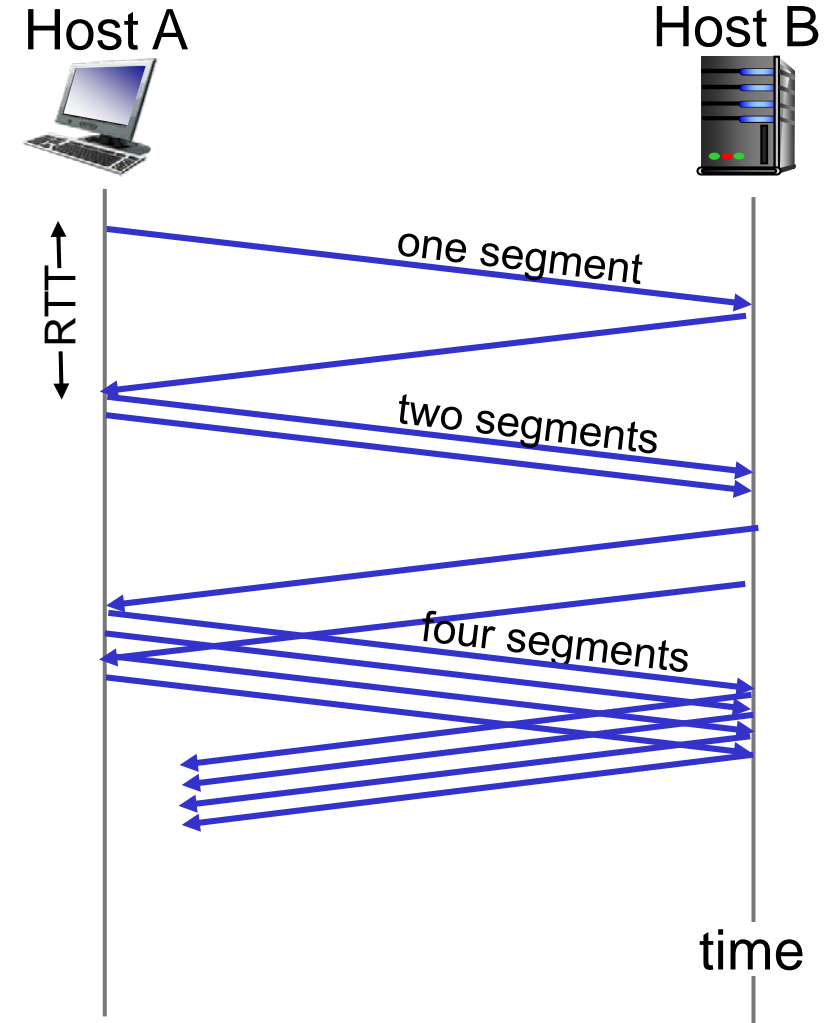
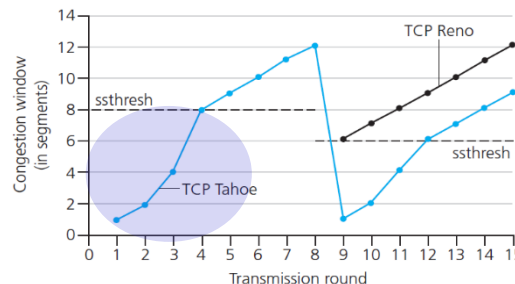
exponential *2 가
linear

- ❖ Congestion control works in two modes:

- **Slow Start** ($cwnd < ssthresh$)
- **Congestion Avoidance** ($cwnd \geq ssthresh$)
- **(Fast Recovery)**

TCP “Slow” Start

- ❖ when connection begins, increase rate exponentially until first loss event:
 - initially $cwnd = 1$ MSS
 - double $cwnd$ every RTT
 - done by incrementing $cwnd$ for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast



From Slow Start to Congestion Avoidance

TCP slows down the increase of `cwnd` if `cwnd` reaches the slow-start threshold value `ssthresh`

- If $cwnd \geq ssthresh$ then each time an ACK is received, increment `cwnd` as follows:

- $cwnd = cwnd + (1 \text{ MSS})/cwnd$
(linear increase per RTT)

1 가 4 -> 1/4 1/4 1/4 1/4
1/4
1

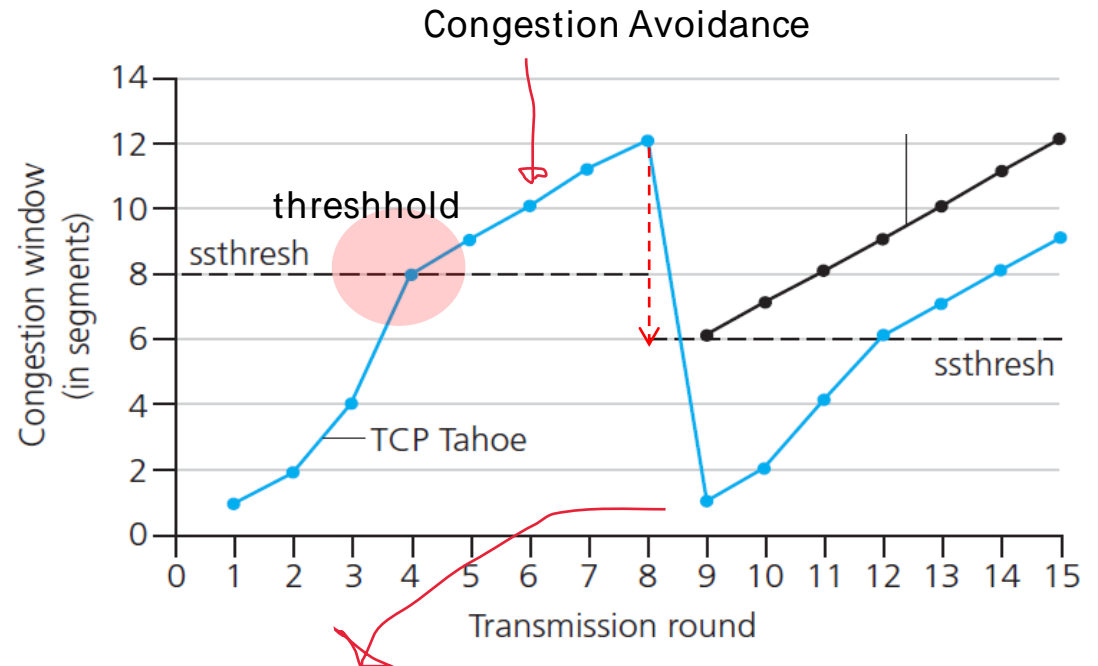
Setting of `ssthresh`:

- ❖ Initially, variable `ssthresh` is set to Advertised window size
- ❖ on loss event, `ssthresh` is set to 1/2 of `cwnd` just before loss event

TCP: switching from slow start to CA

Q: when should the exponential increase switch to linear?

A: when `cwnd` gets to 1/2 of its value before timeout. = *ssthresh*



1

1

TCP Tahoe Algorithm

Initially:

cwnd = 1 MSS;

ssthresh = advertised window size;

New Ack received:

if (cwnd < ssthresh)

/* Slow Start*/

cwnd = cwnd + 1 MSS; => 2

else

/* Congestion Avoidance */

cwnd = cwnd + (1 MSS)/cwnd; => 1 가

Timeout:

/* Multiplicative decrease */

ssthresh = cwnd/2;

cwnd = 1 MSS;

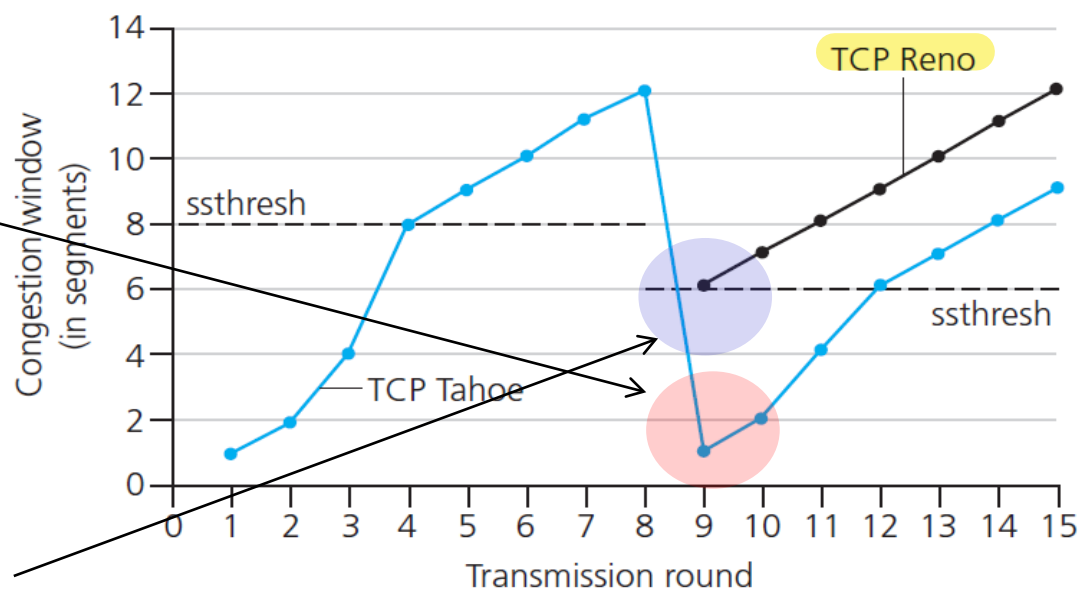
TCP: detecting, reacting to loss

❖ loss indicated by **timeout**:

- **cwnd** set to 1 MSS;
- window then grows exponentially (as in slow start) to threshold, then grows linearly

❖ **Fast Recovery (TCP Reno)**: loss indicated by 3 duplicate ACKs

- dup ACKs indicate network capable of delivering some segments (better than timeout)
- **cwnd** is cut in half window then grows linearly

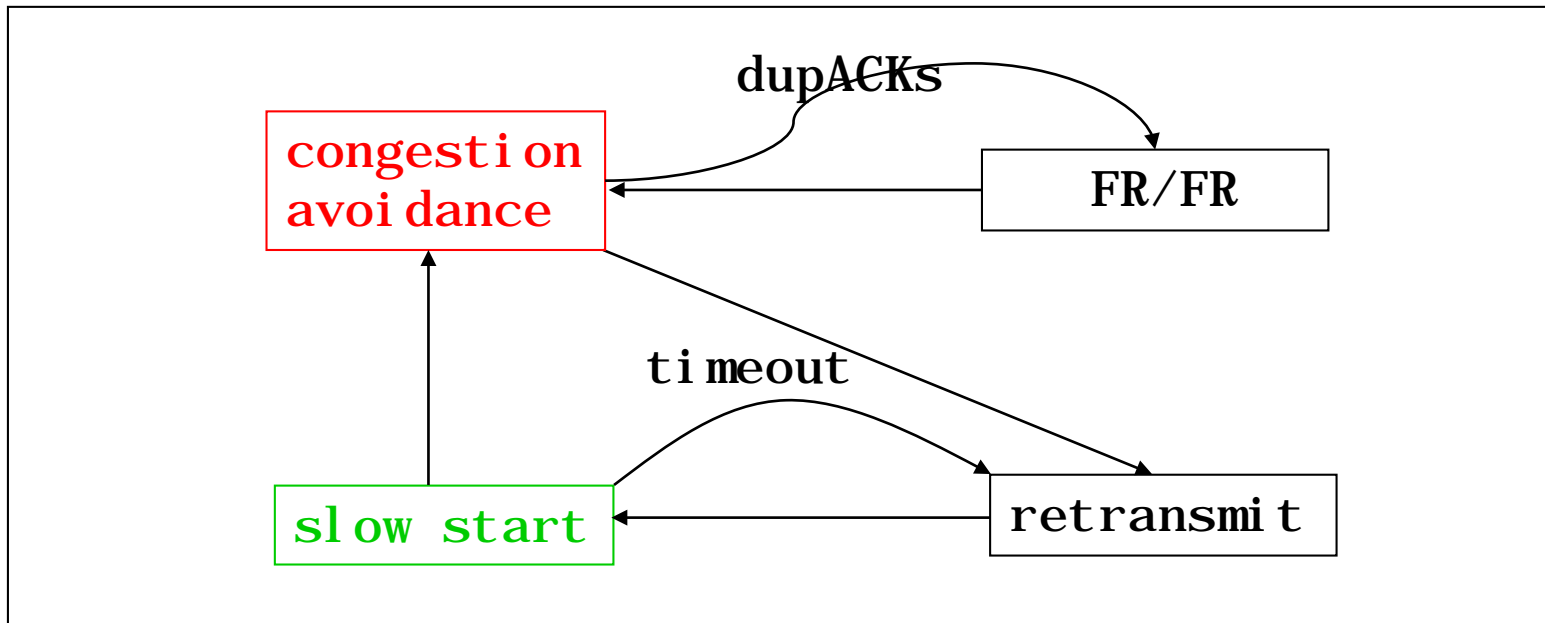


timeout loss threshold가

TCP Reno Summary

❖ Basic ideas

- Fast recovery avoids slow start
- dupACKs: fast retransmit + fast recovery
- Timeout: fast retransmit + slow start

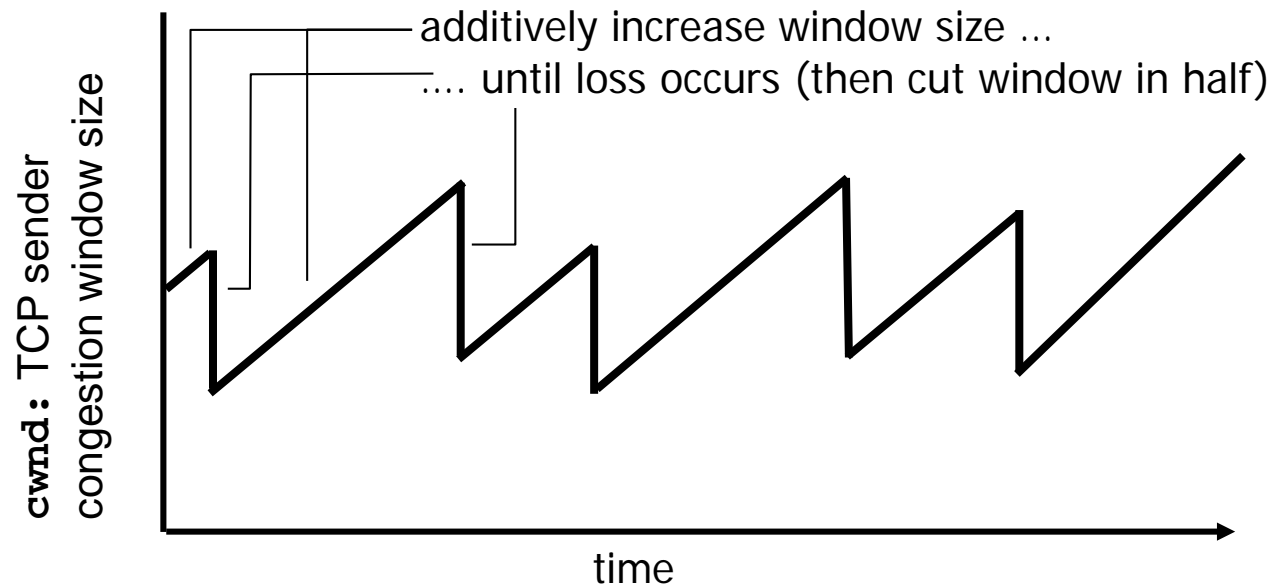


Additive Increase Multiplicative Decrease (AIMD)

- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase `cwnd` by 1 MSS every RTT until loss detected
 - *multiplicative decrease*: cut `cwnd` in half after loss (3-dup acks)



AIMD saw tooth behavior: probing for bandwidth



Chapter 3: summary

- ❖ principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- ❖ instantiation, implementation in the Internet
 - UDP
 - TCP